



Dynamic News Signals as Early-Warning Indicators of Food Insecurity: A Two-Stage Residual Modelling Framework

Dissertation

Submitted in partial fulfillment of the requirements
for the degree of Master of Science in Data Science
Course: CST4090 Individual Project

at

Middlesex University London

Submitted by: Victor Collins Oppon
Student ID: M01040265
Supervisor: Dr. Giovanni Quattrone

Date: January 2026

Abstract

Food insecurity early warning systems increasingly rely on news media indicators to forecast humanitarian crises months in advance. However, existing approaches face a fundamental methodological challenge: spatio-temporal autoregressive (AR) baselines using only temporal autoregressive features (L_t : first-order lag of past IPC values at $t-1$) and spatial autoregressive features (L_s : inverse-distance weighted IPC values from neighboring districts)—with zero news features or external covariates—achieve AUC-ROC 0.907 at 8-month forecast horizons, approaching published news-based models that use millions of articles (93.8% of Balashankar et al.’s PR-AUC). This *autocorrelation trap* raises critical questions about when and where news signals provide genuine early-warning information beyond structural persistence in temporally and spatially autocorrelated crisis data.

This dissertation develops a two-stage residual modelling framework that addresses this challenge through: (1) rigorous spatio-temporal AR baselines using L2-regularized logistic regression with inverse-distance spatial weighting (300km radius) and stratified spatial cross-validation to identify structurally persistent crises, and (2) selective deployment of news-based models exclusively on the critical 26.8% of crises that AR baselines miss (1,427 of 5,322 crises), where shock-driven dynamics break temporal patterns and news features drive 74.7% of marginal predictions (SHAP analysis). This cascade approach strategically allocates computational resources: lightweight AR models for persistence-dominated cases, sophisticated news-based models for the hardest-to-predict shock-driven crises where early warning saves lives.

The empirical analysis uses 55,129 district-level IPC assessments across 24 African countries (2021-2024) and constructs interpretable dynamic features from 7.6 million GDELT news articles through four analytical stages. First, we categorise articles into nine thematic domains (conflict, displacement, economic, weather, food security, health, humanitarian, governance, other) using keyword dictionaries and calculate ratio features (compositional emphasis) and 12-month sliding-window z-score features (temporal anomalies). Second, we apply Hidden Markov Models (HMM) with 2 latent states (binary regime: Pre-Crisis vs Crisis-Prone) to detect narrative regime transitions (peaceful $\rightarrow$ violent shifts) invisible to cross-sectional aggregations, achieving 89.5% convergence across 48-month district-level time series. Third, we employ Dynamic Mode Decomposition (DMD) to extract temporal evolution patterns, isolating crisis-relevant modes (positive growth rates indicating escalation) and filtering background modes (near-zero eigenvalues representing steady states), with 83.1% convergence enabling mechanistic understanding of how crises unfold temporally. Fourth, we integrate these features via mixed-effects logistic regression with geographic random effects (country-level intercepts and slopes) to quantify heterogeneity in baseline risk and feature sensitivity across contexts.

We address five research questions through comprehensive ablation studies (8 feature combinations

testing ratio-only, z-score-only, combined features, and HMM/DMD variants; 3,888 hyperparameter configurations via grid search; 5-fold stratified spatial cross-validation ensuring geographic separation between training/test sets) and triangulated interpretability analysis (XGBoost tree-based feature importance, mixed-effects fixed/random coefficients, SHAP game-theoretic attributions). Key findings demonstrate: **(RQ1 - The Autocorrelation Trap)** AR baselines achieve 93.8% of Balashankar et al. (2023)’s published news-based model performance using zero text features (AR PR-AUC: 0.765 vs Balashankar PR-AUC: 0.816; AR AUC-ROC: 0.907), revealing that most published results (AUC 0.75-0.85) lacking AR comparisons may primarily reflect temporal and spatial autocorrelation rather than genuine text feature value; **(RQ2 - When News Matters)** thematic rankings exhibit measurement paradox: ratio/mixed-effects models (sustained shifts) rank weather (+26.71 coefficient) and food security (+20.33) highest, while SHAP z-score analysis (rapid anomalies) reverses rankings with conflict #1 (0.911) and weather #7 (0.769), demonstrating split frequency $\neq$ predictive power. Location features dominate tree splits (29.3% cumulative) but contribute only 2.6% marginal attribution ($15.5\times$ overstatement), serving as stratification infrastructure while z-score news features drive 74.7% of actual predictions. This reveals complementary mechanisms: ratios capture sustained compositional changes for 8-month forecasts, z-scores detect rapid temporal anomalies for shock-driven crises. Country-specific theme elevations reveal diagnostic signals invisible to universal baselines: Zimbabwe weather coverage elevated +2.1pp above global average (drought cycles compound economic collapse), Sudan conflict +3.3pp (civil war escalation AR cannot anticipate), DRC displacement +2.2pp (M23 resurgence), Somalia health +5.8pp (disease burden compounds food insecurity)—demonstrating news features capture context-specific shock dynamics that break structural persistence patterns; **(RQ3 - Hidden Variables)** HMM provides interpretability value (hmm_ratio_transition_risk ranks #5 in feature importance at 3.2%, capturing qualitative regime transitions invisible to compositional features) while DMD achieves largest mixed-effects coefficient (+352.38) for rare but extreme complex emergencies where multiple crisis drivers converge simultaneously ($<3\%$ of observations); **(RQ4 - Two-Stage Framework)** the cascade rescues 249 crises (17.4% of 1,427 AR failures) at cost of precision decline ($0.732 \rightarrow 0.585$, -14.7pp) but favourable humanitarian cost-benefit under asymmetric weighting (10:1 FN:FP yields 6.2% total cost reduction, prioritising recall over precision); **(RQ5 - Geographic Heterogeneity)** news features exhibit strong context-specificity, with 70.7% of key saves concentrated in three conflict-affected countries (Zimbabwe: 77 saves, Sudan: 59, DRC: 40), country-level AUC ranging 10-fold (0.068 Niger to 0.682 Sudan), and mixed-effects random effects spanning 8.26 points (Somalia +3.70 to Madagascar -4.56), demonstrating that universal models fail and selective deployment is necessary. Within-country heterogeneity analysis reveals the same countries show both rescues and failures at district level (Zimbabwe: 77 saves but 647 still-missed, Sudan: 59 saves but 420 still-missed) $\times$ news-based early warning succeeds in well-covered districts (capitals, conflict zones) but fails in news desert districts (remote pastoral areas, peripheral regions) within the same country, with rescued cases having 53% more news coverage (121 vs 79 articles/month median).

The two-stage framework achieves recall improvement from 0.732 to 0.779 (+4.7pp, +6.4% relative) by rescuing shock-driven crises (conflict escalation in Sudan/DRC, economic collapse in Zimbabwe, complex emergencies with simultaneous displacement/disease) where AR persistence

assumptions fail. The 249 key saves represent *the hardest-to-predict crises*—those invisible to spatio-temporal baselines but critical for humanitarian response. Eight months advance warning enables preemptive food assistance, livelihood support, conflict-sensitive programming, and emergency funding mobilisation before populations exhaust coping strategies.

Critically, cascade failure analysis reveals a fundamental constraint: the 1,178 crises still missed after cascade intervention (82.6% of AR failures) exhibit systematic news coverage deficiency—median 74 articles/month compared to 121 for rescued cases (64% less coverage). This *news deserts hypothesis* demonstrates that news-based early warning fundamentally cannot rescue crises in remote pastoral areas (Kenya Northern, Zimbabwe rural districts, Niger) lacking sufficient media coverage. The 249 key saves concentrate in news-dense conflict zones (70.7% in Sudan/Zimbabwe/DRC), revealing that successful cascade deployment requires rich news signal infrastructure. Remote areas with sparse coverage remain fundamentally unpredictable without expanding NLP data sources beyond traditional news media: social media monitoring (Twitter/X, Facebook, WhatsApp group analysis), community radio transcripts (local-language broadcasts), humanitarian situation reports (OCHA, UNHCR, WFP assessments), and multilingual text mining from non-English sources (Swahili, Hausa, Amharic, French, Arabic). Future NLP systems must incorporate diverse text corpora with targeted collection strategies for underreported regions.

Our methodological contributions extend beyond performance metrics to establish rigorous standards for crisis prediction research: (1) **Mandatory AR baselines** with inverse-distance spatial weighting and proper spatial cross-validation, requiring all future work claiming predictive value from external covariates to report *marginal* contributions beyond autocorrelation; (2) **Prediction-interpretability trade-offs**, demonstrating that simple models (ratio+location, 12 features, AUC 0.727) maximise discrimination for difficult cases while advanced models (35 features including HMM/DMD, AUC 0.697) maximise mechanistic understanding through latent dynamics and temporal patterns; (3) **SHAP analysis exposes tree-based importance artifacts**, revealing location features account for 40.4% of tree splits but only 2.6% of marginal prediction attribution (15.5× overstatement), while z-score features drive 74.7% of prediction variance despite lower tree rankings×demonstrating that feature "importance" depends critically on measurement method (split frequency vs marginal impact); (4) **Evidence-based NLP deployment strategies**, providing operational guidance for selective cascade deployment in conflict zones with high news coverage (Sudan/Zimbabwe/DRC), AR-only deployment in climate contexts where spatial autoregressive features capture regional patterns (Kenya/Ethiopia pastoral zones), and recommendations for expanding text corpora in low-coverage regions (Niger/Madagascar) through social media mining, humanitarian reports, local-language sources, and community radio transcripts.

This work challenges existing literature’s claims about news value for early warning, demonstrates when and where dynamic features justify computational complexity, and provides an operational framework for selective deployment that maximises humanitarian impact while respecting geographic heterogeneity. The finding that simple AR baselines achieve 90%+ of predictive signal fundamentally reshapes understanding of what news-based forecasting can contribute, shifting focus from universal deployment to strategic targeting of shock-driven crises in news-dense

contexts. By quantifying the autocorrelation trap and establishing when news features provide genuine marginal value, this dissertation sets higher methodological standards for the crisis prediction field.

Keywords: food insecurity, early warning systems, autocorrelation trap, two-stage residual modelling, cascade ensemble, spatio-temporal autoregression, mixed-effects regression, interpretable machine learning, GDELT, Hidden Markov Models, Dynamic Mode Decomposition, geographic heterogeneity, selective deployment, humanitarian forecasting

Contents

Abstract	i
1 Introduction	1
1.1 Context and Motivation	1
1.2 Problem Statement: The Autocorrelation Trap	4
1.3 Research Gap	8
1.3.1 Gap 1: Lack of Rigorous AR Baseline Comparisons	8
1.3.2 Gap 2: Inability to Distinguish Structural Persistence from Shock-Driven Dynamics	9
1.3.3 Gap 3: Absence of Two-Stage Frameworks Leveraging AR Strengths	10
1.3.4 Gap 4: Limited Model Interpretation Frameworks for Geographic and Temporal Heterogeneity	10
1.3.5 Gap 5: Static Feature Engineering (Article Counts Only)	11
1.4 Research Questions	12
1.5 Research Objectives	14
1.6 Contributions	17
1.7 Thesis Structure	24
2 Brief Background and Literature Review	26
2.1 Food Insecurity and IPC Classification	26
2.2 Existing Early Warning Approaches	28
2.3 News-Based Forecasting and the Autocorrelation Problem	31
2.3.1 Existing News-Based Approaches	31
2.3.2 The Autocorrelation Trap	33
2.4 Spatial-Temporal Methods and Cross-Validation	37
2.4.1 Evidence of Spatial Clustering in Food Insecurity	37
2.4.2 Spatial Cross-Validation Methods	38
2.5 Ensemble and Cascade Methods	39
2.5.1 Cascade Ensembles for Selective Deployment	39
2.5.2 Implications for Food Security Forecasting	40
2.6 Dynamic Feature Engineering: HMM, DMD, and Z-Scores	41
2.6.1 Hidden Markov Models for Regime Detection	41
2.6.2 Dynamic Mode Decomposition for Crisis Dynamics	42
2.6.3 Z-Score Normalisation	43
2.7 Interpretability Methods for Model Understanding	45
2.7.1 SHAP Values for Local Explanations	45

2.7.2	Feature Importance from Tree-Based Models	46
2.7.3	Mixed-Effects Random Coefficients for Geographic Heterogeneity	47
2.7.4	Triangulation Across Methods	47
2.8	Research Gap and Positioning	48
3	Methods	53
3.1	Data Sources and Preprocessing	53
3.1.1	IPC Food Security Classifications	53
3.1.2	GDELT News Data	54
3.1.3	Geographic Boundaries and Spatial Linkage	55
3.1.4	Data Aggregation Pipeline	55
3.1.5	Quality Control and Filtering	56
3.1.6	Final Dataset Statistics	57
3.2	Experimental Design	59
3.2.1	Stratified Spatial Cross-Validation	59
3.2.2	Evaluation Metrics	61
3.2.3	Threshold Optimisation Strategies	63
3.2.4	Statistical Testing	64
3.3	Stage 1: Structural Baseline Modelling	66
3.3.1	Autoregressive Feature Construction	66
3.3.2	Logistic Regression Model Specification	68
3.3.3	Prediction Horizons	70
3.3.4	Model Performance and AR Failure Definition	71
3.3.5	WITH_AR_FILTER Strategy for Stage 2 Deployment	72
3.4	Stage 2: News-Based Feature Engineering	73
3.4.1	Overview of Stage 2 Feature Construction	73
3.4.2	Ratio and Z-Score Features: Basic Dynamic Signals	74
3.4.3	Macro-Category Taxonomy and Article Classification	74
3.4.4	Ratio Features: Cross-Sectional Coverage Composition	75
3.4.5	12-Month Sliding-Window Z-Score Standardisation	77
3.4.6	Feature Set Composition and Dimensionality	78
3.4.7	Advanced Features: Regime and Mode Extraction	80
3.4.8	Hidden Markov Models for Latent Regime Detection	80
3.4.9	Dynamic Mode Decomposition for Crisis Escalation Patterns	84
3.4.10	Advanced Feature Set Composition	86
3.5	Model Training and Evaluation Framework	88
3.5.1	XGBoost Gradient Boosting Models	89
3.5.2	Mixed-Effects Logistic Regression	90
3.5.3	Ablation Study Design	93
3.5.4	Threshold Optimisation Strategies	93
3.5.5	Evaluation Metrics	94
3.5.6	Feature Importance Extraction	95
3.6	Two-Stage Framework Integration	95

3.6.1	Cascade Decision Logic	96
3.6.2	Override Mechanism and Coverage	96
3.6.3	Key Saves: Quantifying Stage 2 Value	97
3.6.4	Performance Impact: Recall vs Precision Trade-Off	97
3.6.5	Geographic Distribution of Key Saves	98
3.6.6	Model Selection for Cascade Integration	99
3.7	Interpretability Framework	100
3.7.1	XGBoost Gain-Based Feature Importance	100
3.7.2	SHAP Attribution Analysis	102
3.7.3	Mixed-Effects Decomposition: Fixed vs Random Effects	103
3.7.4	Ablation Studies: Marginal Feature Group Contributions	104
3.7.5	Triangulation Across Interpretability Methods	106
4	Results and Evaluation	108
4.1	Baseline Performance and Methodological Critique	108
4.1.1	AR Baseline Results	108
4.1.2	NewsBased Model Performance	114
4.1.3	Understanding Model Roles: Persistence vs. Shock Detection	115
4.1.4	Model Stability and Geographic Generalization	116
4.1.5	Implications: The Autocorrelation Trap	117
4.2	Identifying Missed EarlyWarning Opportunities	119
4.2.1	Quantifying AR Failures	119
4.2.2	Geographic Distribution of Failures	120
4.2.3	Temporal Patterns	121
4.2.4	Country-Level Failure Analysis	122
4.2.5	Humanitarian Criticality	123
4.3	Dynamic Feature Engineering Results	123
4.3.1	Ablation Study Overview	124
4.3.2	Ratio + Location Baseline (Best Performing Ablation)	126
4.3.3	Z-score + Location (Temporal Anomaly Baseline)	127
4.3.4	Combining Ratio and Z-score Features	128
4.3.5	Adding HMM Features: Stochastic Regime Transition Modelling	129
4.3.6	Adding DMD Features: Spectral Decomposition of Crisis Dynamics	130
4.3.7	Feature Group Contribution Summary	131
4.3.8	Cross-Validation Robustness and Geographic Heterogeneity	133
4.4	Mixed-Effects vs Machine Learning Comparison	134
4.4.1	XGBoost Performance Summary	134
4.4.2	Mixed-Effects Model Results	136
4.4.3	Fixed vs Random Effects Decomposition	137
4.4.4	Accuracy-Interpretability Trade-off	140
4.5	Two-Stage Framework Performance	141
4.5.1	Overall Framework Results	141
4.5.2	Key Saves Analysis	144

4.5.3	Precision-Recall Trade-off and Cost-Sensitive Analysis	148
4.5.4	Country-Level Performance Heterogeneity	149
4.5.5	Operational Deployment Implications	150
4.6	Interpretability Analysis Answering the Five Research Questions	152
4.6.1	RQ1: The Autocorrelation Trap—Assessing the Marginal Value of News Features	152
4.6.2	RQ2: When News Matters Role of Different News Categories and Transformations	153
4.6.3	RQ3: Two-Stage Framework Effectiveness and Precision-Recall Trade-offs	158
4.6.4	RQ4: Geographic Heterogeneity in News Feature Value	159
4.6.5	Synthesis: Triangulating Evidence Across Interpretability Methods	165
4.6.6	Final Synthesis: Answering the Overarching Question	168
5	Discussion and Limitations	170
5.1	Summary of Key Findings	170
5.1.1	RQ1 Answered: The Autocorrelation Trap Quantified	170
5.1.2	RQ2 Answered: When News Matters—Feature Engineering Insights . . .	171
5.1.3	RQ3 Answered: The Role of Hidden Variables—HMM and DMD	173
5.1.4	RQ4 Answered: Two-Stage Framework Performance and Trade-Offs . . .	174
5.1.5	RQ5 Answered: Geographic Heterogeneity—Where News Matters Most .	183
5.2	Theoretical Implications	196
5.2.1	Rethinking News-Based Forecasting: The Autocorrelation Trap as Field- Wide Challenge	196
5.2.2	Two-Component Crisis Dynamics: Low-Frequency Persistence vs High- Frequency Shocks	197
5.2.3	Geographic Heterogeneity: News Value is Context-Dependent	199
5.3	Practical Implications for Food Security Early Warning Systems	200
5.3.1	Operational Deployment Considerations	200
5.3.2	When to Trust AR vs When to Apply Cascade Override	201
5.3.3	Cost-Benefit of News Monitoring Infrastructure	202
5.3.4	Integration with Existing Humanitarian Systems	203
5.4	Methodological Contributions	204
5.4.1	Two-Stage Residual Modelling Framework: A General Approach for Autocorrelated Outcomes	204
5.4.2	WITH_AR_FILTER Training Strategy: Selective Supervision	205
5.4.3	Stratified Spatial Cross-Validation: Rigorous Generalisation Testing . . .	206
5.4.4	Crisis-Focused HMM and DMD Feature Engineering	207
5.5	Limitations	208
5.5.1	Data Coverage Heterogeneity and Systematic Bias	208
5.5.2	English-Language News Bias and GDELT Limitations	209
5.5.3	IPC Assessment Delays and Temporal Resolution Constraints	210
5.5.4	8-Month Horizon Constraints and Horizon-Dependent Dynamics	211
5.5.5	Precision Trade-Off and Operational Alert Fatigue	211

5.5.6	External Validity: Africa-Specific Findings, Uncertain Generalisation . . .	212
5.6	Comparison to Related Work	213
5.6.1	This Work vs Balashankar et al. (2023): Methodological Divergences . . .	213
5.6.2	This Work vs Traditional Early Warning Systems (FEWSNET, WFP) . .	214
5.6.3	Positioning in ML for Social Good Literature	215
5.7	Future Research Directions	216
5.7.1	Real-Time Deployment and Operational Monitoring	216
5.7.2	Advanced NLP Enhancement: Beyond Current Approach	217
5.7.3	Multi-Horizon Optimisation: Joint Forecasting Across $h=4, 8, 12$	219
5.7.4	Causal Inference and Counterfactual Analysis: Beyond Prediction	220
5.7.5	Multilingual News Processing: Addressing Language Bias	221
5.7.6	Explainable AI for Humanitarian Decision-Making: Enhanced Interpretability	222
6	Conclusion	224
6.1	Synthesis: Answering the Five Research Questions	224
6.1.1	RQ1: The Autocorrelation Trap Quantified	224
6.1.2	RQ2: When News Matters—Feature Engineering Insights	226
6.1.3	RQ3: The Role of Hidden Variables—HMM and DMD	228
6.1.4	RQ4: Two-Stage Framework Performance and Precision-Recall Trade-Offs	229
6.1.5	RQ5: Geographic Heterogeneity—News Features Are Not Universally Valuable	231
6.2	Core Contributions to Humanitarian Early Warning	234
6.2.1	Contribution 1: Methodological Critique—Exposing the Autocorrelation Trap	234
6.2.2	Contribution 2: Two-Stage Residual Modelling Framework	235
6.2.3	Contribution 3: Dynamic Feature Engineering Beyond Article Counts . .	236
6.2.4	Contribution 4: Comprehensive Model Interpretation Framework	237
6.2.5	Contribution 5: Operational Deployment Framework and Geographic Targeting	239
6.3	Implications for the Humanitarian Early Warning Ecosystem	240
6.3.1	Rethinking the Role of News in Crisis Prediction	240
6.3.2	The Two-Component Crisis Dynamics Framework	241
6.3.3	When to Trust AR, When to Override with Cascade	242
6.3.4	Limitations and Honest Reflection	243
6.4	Future Research Directions	245
6.4.1	Advanced NLP Enhancement: Beyond Bag-of-Words	245
6.4.2	Multi-Horizon Optimisation	246
6.4.3	Real-Time Operational Deployment and Monitoring	246
6.4.4	Causal Inference and Counterfactual Analysis	247
6.4.5	Multilingual News Processing	247
6.4.6	Explainable AI for Humanitarian Decision-Making	247
6.5	Closing Vision: From Autocorrelation to Action	248

Appendices	250
A Full Ablation Results	251
A.1 Ablation Study Design	251
A.1.1 Model Variants	251
A.1.2 Training Configuration	252
A.2 Performance Comparison	253
A.2.1 Overall Metrics	253
A.2.2 Statistical Significance Testing	254
A.3 Feature Importance Rankings	255
A.3.1 Model: ratio_location (Best Performer, AUC=0.727)	255
A.3.2 Model: ratio_z-score_location (Combined Features, AUC=0.696)	256
A.3.3 Model: ratio_z-score_hmm (Advanced Features, AUC=0.703)	257
A.4 Cross-Validation Robustness	257
A.4.1 Fold-Level Performance	257
A.5 Optimal Hyperparameters	258
A.5.1 Best Hyperparameters by Model	258
A.6 Summary	259
B Hyperparameter Tuning Details	260
B.1 XGBoost Hyperparameter Search	260
B.1.1 Search Space	260
B.1.2 Optimisation Procedure	260
B.1.3 Optimal Hyperparameters by Model	262
B.1.4 Hyperparameter Sensitivity Analysis	263
B.1.5 Cross-Validation Stability	264
B.2 Mixed-Effects Model Optimisation	265
B.2.1 Fixed-Effects Regularization	265
B.2.2 Class Weighting Optimisation	265
B.2.3 Random-Effects Variance Components	266
B.3 Computational Resources	266
B.3.1 Training Time	266
B.3.2 Memory Requirements	266
B.4 Reproducibility	267
C Country-Level Metrics	268
C.1 XGBoost Advanced: Country-Level Performance	269
C.1.1 Performance Metrics by Country	269
C.1.2 Confusion Matrices by Country (Top 6)	270
C.2 Cascade Framework: Country-Level Key Saves	271
C.2.1 Key Saves Distribution	271
C.2.2 Geographic Heterogeneity: Rescue Rate Variation	272
C.3 Mixed-Effects: Random Intercepts by Country	273

C.3.1	Country-Level Baseline Risk	273
C.3.2	Fixed-Effects Slopes by Country (Selected Features)	274
C.4	Data Availability by Country	275
C.4.1	News Coverage Metrics	275
C.5	Summary Statistics	276
C.5.1	Cross-Country Variation	276
C.5.2	Performance Correlations	276
C.6	Evidence-Based Deployment Framework	277
D	Mathematical Derivations	278
D.1	Autoregressive Baseline Model	279
D.1.1	Logistic Regression Formulation	279
D.1.2	Regularization and Class Weighting	279
D.2	Dynamic Feature Engineering	280
D.2.1	Ratio Features (Compositional Transformation)	280
D.2.2	Z-Score Features (Temporal Anomaly Transformation)	280
D.2.3	Hidden Markov Model (HMM) Features	281
D.2.4	Dynamic Mode Decomposition (DMD) Features	282
D.3	XGBoost Model	283
D.3.1	Gradient Boosting Formulation	283
D.3.2	Class Weighting for Imbalanced Data	284
D.4	Mixed-Effects Logistic Regression	285
D.4.1	Generalised Linear Mixed Model (GLMM) Formulation	285
D.4.2	L1 Regularization for Fixed Effects	285
D.5	Cascade Framework Decision Rule	286
D.5.1	Two-Stage Prediction	286
D.5.2	Precision-Recall Trade-Off	286
D.6	Performance Metrics	287
D.6.1	Area Under ROC Curve (AUC-ROC)	287
D.6.2	Youden's J Statistic	287
D.6.3	Cost-Sensitive Metric	287
E	Code and Data Availability	288
E.1	Code Repository	289
E.1.1	GitHub Repository	289
E.1.2	Software Dependencies	291
E.1.3	Installation Instructions	292
E.2	Data Availability	292
E.2.1	Public Datasets	292
E.2.2	Processed Datasets	293
E.3	Trained Models	294
E.3.1	Model Artifacts	294
E.3.2	Model Loading Example	295

E.4	Computational Resources	295
E.4.1	Hardware Specifications	295
E.4.2	Training Time	295
E.4.3	Memory Requirements	296
E.5	Reproducibility	296
E.5.1	Random Seeds	296
E.5.2	Verification	297
E.6	Ethical Considerations and Data Privacy	297
E.6.1	Data Ethics	297
E.6.2	Responsible AI Deployment	297
E.7	License and Citation	298
E.7.1	License	298
E.7.2	Citation	298
E.7.3	Contact	298
	References	299

List of Figures

1.1	IPC Food Security Phase Classification System	1
1.2	Two-Stage Cascade Framework	4
1.3	The Autocorrelation Trap	5
2.1	Temporal Persistence Drives AR Baseline Performance	36
2.2	Methodological Gaps in Food Security Forecasting Literature	50
3.1	Data Processing Pipeline	58
3.2	AR Baseline Features: L_t (Temporal Autoregressive Features) + L_s (Spatial Autoregressive Features)	69
3.3	Feature Engineering Pipeline: From Raw GDELT to XGBoost Features	74
3.4	Frequency Decomposition: AR Captures Low Frequency, Cascade Targets High Frequency	81
3.5	HMM Regime Detection: Sudan Conflict Escalation Example	83
3.6	DMD Temporal Modes: Crisis Escalation Patterns	87
4.1	AR Baseline Performance Summary	109
4.2	AR Failures Geographic Distribution	113
4.3	AR Failures Temporal Distribution	114
4.4	Ablation Study Performance Rankings with Z-Score Threshold Sensitivity	125
4.5	XGBoost Advanced Feature Importance Rankings	132
4.6	Mixed-Effects Model: Top 10 Fixed Effect Coefficients	138
4.7	Cascade vs AR Baseline Performance Comparison	142
4.8	Cascade Breakthrough: 249 Hardest Crises Predicted	143
4.9	Cascade Breakthrough: 249 Crisis Rescues Across Africa	146
4.10	Real Crisis Stories: Cascade Rescues Where AR Failed with Side-by-Side Comparison	147
4.11	News Themes: The SHAP Paradox Revealed	155
4.12	Dominant News Themes by Country	162
4.13	Theme Elevations by Country	163
4.14	SHAP Feature Attribution Analysis	165
5.1	Cascade Breakthrough: 249 Crises Rescued Where AR Failed	176
5.2	Geographic Distribution of 249 Cascade Rescues	177
5.3	Cascade Failures Analysis: Why 1,178 Cases Still Missed	179
5.4	Geographic Distribution of Cascade Failures: News Deserts	181
5.5	Geographic Heterogeneity in News Value	183
5.6	Geographic Concentration of Cascade Impact	185

5.7	Country-Specific News Theme Importance Heatmap	187
5.8	Theme Signatures for Top 3 Countries	188
5.9	Geographic Distribution of Dominant News Themes	189
5.10	Theme Elevations by Country	191
5.11	Key Saves vs Delta-AUC Relationship	193
5.12	Cascade Benefit Matrix for Top 12 Countries	195
6.1	Complete Narrative Arc: From Autocorrelation to Humanitarian Impact	225
A.1	Full Ablation Study Results	253
B.1	Hyperparameter Tuning: Systematic Grid Search Exploration	261
C.1	Country-Level AR Baseline Confusion Matrices	268
D.1	Data Quality Assessment: Coverage and Temporal Distribution	278
E.1	Extended SHAP Analysis: Dependence Plots for Top 5 Features	288

List of Tables

3.1	Dataset Statistics: Raw IPC Database and Final Analysis Dataset	57
3.2	Confusion matrix comparison: AR baseline vs cascade ensemble. The cascade gains 249 true positives (key saves) but incurs 1,512 additional false positives. . .	98
3.3	Ablation study results: AUC-ROC by feature group. All models include location metadata. Ratio-only achieves highest standalone AUC (0.727), but z-scores account for 74.7% SHAP marginal attribution in combined models, demonstrating complementary roles.	105
4.1	Autoregressive Baseline Performance by Forecast Horizon	108
4.2	AR Failures by Country (Top 10)	121
4.3	Ablation Study Performance Summary (8 Variants)	125
4.4	Ratio + Location Model: Top 10 Feature Importance	127
4.5	Feature Group Contribution Summary Across Ablation Models	131
4.6	MixedEffects Model: Top 10 Fixed Effects (Ratio + HMM + DMD)	137
4.7	MixedEffects Model: Random Intercepts by Country (Top 10 and Bottom 5) . .	139
4.8	XGBoost vs Mixed-Effects: Performance on AR Failures	140
4.9	Cascade Framework vs AR Baseline: Overall Performance Comparison	141
4.10	Key Saves by Country (Top 10)	145
4.11	Key Saves by IPC Assessment Period (Top 5)	148
4.12	Cascade Performance by Country (Top 10 by Key Saves)	150
5.1	AR Baseline vs Cascade Framework Performance	178
5.2	Comparison of Early Warning Approaches	215
6.1	AR Baseline vs Cascade Framework Performance	229
A.1	Ablation Study Model Variants	252
A.2	Ablation Study Performance Metrics	253
A.3	Pairwise AUC Comparisons (p-values)	254
A.4	Feature Importance: ratio_location	255
A.5	Feature Importance: ratio_z-score_location (Top 15)	256
A.6	Feature Importance: ratio_z-score_hmm (Top 15)	257
A.7	Fold-Level AUC by Model (Top 4 Models)	257
A.8	Optimal Hyperparameters (Top 3 Models)	258
B.1	XGBoost Hyperparameter Search Space	260
B.2	Optimal Hyperparameters: XGBoost Advanced	262
B.3	Optimal Hyperparameters: XGBoost Basic	263

B.4	Hyperparameter Sensitivity: Top 10 Configurations	263
B.5	Fold-Level Performance: XGBoost Advanced (Optimal Config)	264
B.6	Random-Effects Variance Components	266
B.7	Training Time by Model Type	266
C.1	Country-Level Performance: XGBoost Advanced Model	269
C.2	Country-Level Confusion Matrices: XGBoost Advanced (Youden Threshold) . . .	270
C.3	Key Saves by Country: Cascade Framework	271
C.4	Mixed-Effects Random Intercepts: pooled_ratio_hmm_dmd Model	273
C.5	Country-Specific Feature Slopes: conflict_ratio	274
C.6	Country-Level News Coverage (Articles per District-Year)	275
C.7	Summary Statistics: Country-Level Heterogeneity	276

Chapter 1

Introduction

1.1 Context and Motivation

Food insecurity affects 282 million people across 59 crisis-affected countries, with Sub-Saharan Africa bearing a disproportionate burden [1, 2]. The humanitarian consequences are severe: malnutrition, disease, displacement, economic collapse, and in extreme cases, famine and death. Early warning systems are critical for humanitarian response, enabling timely interventions that save lives and mitigate suffering. When warnings arrive 6-8 months in advance, humanitarian agencies can pre-position food supplies, negotiate access with governments, mobilise funding through appeals, and implement targeted assistance programs before crises peak [3, 4, 5].




Figure 1.1: Standardised 5-phase scale for acute food insecurity. The IPC classifies food security from Phase 1 (Minimal) to Phase 5 (Famine), with Phase 3+ representing crisis thresholds triggering humanitarian response. This dissertation predicts binary outcomes ($IPC \geq 3$) at district level with 8-month forecast horizons. Crisis rate in dataset: 25.7%. $n=20,722$ observations, 18 countries, 2021-2024.

Traditional early warning approaches rely on satellite-based vegetation indices (NDVI), rainfall anomaly monitoring (CHIRPS, TAMSAT), market price tracking, and household survey data [6, 7, 8]. While these methods have proven valuable, they suffer from several limitations. Satellite data provides broad spatial coverage but operates at coarse temporal resolution and often lags 2-4 weeks behind ground conditions due to processing delays and cloud cover interference [9]. Market price data captures economic shocks but may not reflect localized crises in remote areas with limited market integration. Household surveys (e.g., FEWSNET Livelihoods Baseline Profiles)

provide rich contextual information but are expensive, logistically challenging, and conducted infrequently—typically annually or bi-annually, missing rapid-onset crises that emerge between assessment cycles [10, 11].

News media offers a compelling alternative data source that addresses several of these limitations [12, 13, 14]. News coverage is near real-time, updated continuously as events unfold. It captures ground-level perspectives through conflict reports, displacement narratives, economic disruption descriptions, weather impact assessments, and humanitarian access constraints. Unlike satellite data, news coverage can detect crises in cloud-covered regions, urban areas, and conflict zones where physical access for surveys is impossible. The global reach of wire services (Reuters, AFP, AP) [15, 16] and the proliferation of local news outlets in African countries means that even remote crises often receive media attention, particularly when humanitarian consequences are severe.

Recent work has demonstrated that text-based features extracted from news archives can predict food insecurity with impressive accuracy [17, 18]. Using 11.2 million news articles with natural language processing (frame-semantic parsing and word embeddings) for feature extraction and Random Forest regression for prediction, researchers have achieved strong predictive performance (PR-AUC=0.82) for forecasting IPC phases up to 12 months ahead across 21 countries. These results suggest that the “digital exhaust” of global news coverage contains valuable early-warning signals that machine learning can extract and operationalize.

However, existing approaches face a fundamental methodological challenge that has received insufficient attention in the literature. Food security crises exhibit strong temporal and spatial persistence [8, 19]. Today’s crisis is highly predictive of tomorrow’s crisis—chronic food insecurity in regions like South Sudan, Somalia, and the Sahel persists for months or years, driven by structural factors (poverty, conflict, climate vulnerability) that change slowly. Adjacent districts often share similar outcomes due to common exposure to regional shocks (drought, conflict spillovers, market disruptions) and spatial diffusion of crises through population movements and trade linkages.

This *autocorrelation trap* raises a critical question: are sophisticated news-based models capturing genuine predictive signals from text features, or are they primarily learning temporal and spatial patterns that simpler autoregressive (AR) baselines could replicate? If a model using only `IPC_t-1` (last period’s food security status) and `IPC_neighbours` (spatial status of surrounding districts) achieves 90% of a news model’s performance, can we credibly claim that text features provide substantial predictive value? Without rigorous comparison against strong temporal baselines, high performance may reflect autocorrelation rather than genuine signal from news content [5].

This dissertation confronts this challenge directly. Using 55,129 district-level food security assessments from the Integrated Food Security Phase Classification (IPC) system across 24 African countries spanning 2021-2024 (refined to 20,722 observations across 1,920 districts in 18 countries after applying h=8 forecast horizon requirements and data quality filters), combined with 7.6 million GDELT news articles, we develop and evaluate a spatio-temporal

autoregressive baseline. This AR model uses only two autoregressive features: Lt (temporal autoregressive feature using the first-order lag IPC_{t-1}) and Ls (spatial autoregressive feature using inverse-distance weighted IPC values from surrounding districts within a 300km radius) [20, 21, 22]. These are autoregressive features—lagged values of the dependent variable (IPC) itself, not external covariates—with no text features whatsoever.

The results are striking: the AR baseline achieves $AUC=0.907$, $Precision=0.732$, $Recall=0.732$, and $F1=0.732$ at 8-month forecast horizons. This performance approaches published news-based models (93.8% of Balashankar et al.’s PR-AUC), demonstrating that spatio-temporal persistence dominates crisis prediction. The autocorrelation trap is not theoretical—it is empirically real and quantitatively large. Claims of predictive value from text features must overcome the high bar set by simple persistence.

Yet the AR baseline is not perfect. It misses 1,427 crises out of 5,322 total (26.8%), revealing systematic failures where temporal patterns break down. These failures represent *missed early-warning opportunities*—cases where the past is not a reliable guide to the future, where structural persistence fails, and where dynamic signals from news media might provide genuine value. Examining these failures reveals patterns: they concentrate in conflict-affected regions (Sudan, DRC, Zimbabwe), occur during rapid-onset shocks (coup d’états, acute conflict escalation, displacement crises), and cluster in periods where narrative regimes shift (peaceful to violent, stable to chaotic).

This dissertation develops a two-stage residual modelling framework that leverages AR strengths while explicitly targeting its weaknesses. Stage 1 deploys the spatio-temporal AR baseline to identify structurally persistent crises—the 73.2% of cases where simple persistence suffices. Stage 2 focuses exclusively on the WITH_AR_FILTER subset (6,553 cases where $IPC_{t-1} \leq 2$ AND AR predicted non-crisis, including 1,427 cases where AR missed actual crises), deploying dynamic news features, Hidden Markov Model (HMM) regime detection, and Dynamic Mode Decomposition (DMD) temporal pattern extraction to rescue missed opportunities [23].

The framework achieves 249 successful predictions of AR-missed crises—a 17.4% rescue rate representing 249 early warnings 8 months in advance that the AR baseline missed entirely. **These are not routine cases:** they represent the *hardest-to-predict crises* where temporal persistence breaks down—conflict-driven shocks in Sudan and DRC, rapid-onset displacement in Zimbabwe, coup-related disruptions—the very cases where early warning matters most for humanitarian response. When aggregate metrics improve from $Recall=0.732$ to 0.779 (AR baseline to ensemble), this is not merely a 4.7 percentage point statistical gain. **It represents 249 real crises, affecting millions of people, now predicted 8 months in advance when they were previously invisible to persistence-based forecasting.** These are the marginal cases where news signals provide genuine value: detecting regime shifts, capturing conflict escalation, and identifying rapid-onset shocks that confound autoregressive baselines.

This success reflects a deliberate design choice prioritising recall over precision. The framework achieves $Recall=0.779$ (up from 0.732), successfully identifying 249 additional crises that the AR baseline missed, while Precision decreases from 0.732 to 0.585. **In humanitarian early warning**

contexts, this trade-off is operationally appropriate: missing a crisis (false negative) leads to catastrophic outcomes—famine, death, displacement—while false alarms, though wasteful of resources, allow humanitarian actors to stand down pre-positioned supplies and redirect funding [24]. The framework’s design philosophy aligns with established humanitarian principles: *it is better to be over-prepared than to miss a crisis entirely*. FEWSNET, WFP, and other operational early warning systems routinely issue precautionary alerts precisely because the asymmetric costs favour sensitivity (high recall) over specificity (high precision) when lives are at stake [5].

This dissertation provides a comprehensive analysis of when, where, and how dynamic news signals provide genuine early-warning information beyond spatio-temporal persistence. We address five core research questions spanning methodological critique (the autocorrelation trap), feature engineering (ratio vs z-score, news categories), hidden variables (HMM, DMD), framework performance (two-stage selective deployment), and geographic heterogeneity (Zimbabwe, Sudan, DRC). Through ablation studies across 8 model variants, interpretability analysis using three complementary methods (XGBoost feature importance, mixed-effects coefficients, SHAP values [25]), and real-world case studies, we demonstrate that **news signals rescue the hardest-to-predict crises**—conflict-driven shocks, rapid-onset displacements, and regime transitions where persistence models fail and where timely intervention saves lives. The contribution is not universal improvement across all cases, but *targeted success for the cases that matter most*.




Figure 1.2: Selective deployment: AR for persistence, news for rapid-onset shocks. Stage 1 AR baseline achieves Precision=Recall=0.732 (AUC=0.907) using only temporal autoregressive features and spatial autoregressive features. Stage 2 deploys XGBoost with 35 dynamic features exclusively on AR=0 cases ($n=6,553$), rescuing 249 crises that AR missed. Final cascade: Precision=0.585, Recall=0.779, with 70.7% of key saves concentrated in conflict zones (Zimbabwe 77, Sudan 59, DRC 40). **249 crises rescued where persistence failed—8 months advance warning for rapid-onset shocks.** $n=20,722$ observations, $h=8$ months, 18 countries.

1.2 Problem Statement: The Autocorrelation Trap

The central problem motivating this research is methodological: how do we distinguish genuine predictive signals from text features versus spurious correlations driven by temporal and spatial autocorrelation?




Figure 1.3: Temporal persistence achieves near-literature performance with zero news features. AR baseline (using only Lt and Ls autoregressive features) achieves PR-AUC=0.765, reaching 93.8% of Balashankar et al. (2023) news model performance (PR-AUC=0.816). This demonstrates the autocorrelation trap: high predictive accuracy stems primarily from temporal and spatial persistence, not news signals. Literature benchmarks typically lack AR baseline comparisons, obscuring marginal contributions of text features. $n=20,722$ observations, $h=8$ months, 5-fold spatial CV.

Existing literature on news-based crisis prediction [17, 18] typically evaluates performance against held-out test sets using standard train-test splits or cross-validation. These evaluations demonstrate that text features improve prediction accuracy for binary classification of food security crises. Performance gains are attributed to the informational content of news coverage: conflict reports signal impending displacement and market disruption, economic news captures inflation and unemployment, weather reports indicate agricultural shocks, and humanitarian coverage reflects access constraints and response gaps.

However, these evaluations rarely compare against strong temporal baselines. A few studies include simple lag features (y_{t-1}) as controls, but we are unaware of any work in the food security domain that systematically compares news-based models against spatio-temporal autoregressive baselines with both temporal autoregressive features (Lt: first-order lag of past IPC values, $t-1$) and spatial autoregressive features (Ls: inverse-distance weighted IPC values from surrounding districts), combined with proper spatial cross-validation to prevent information leakage.

This omission is consequential. Consider a hypothetical model that achieves Recall=0.82 for predicting IPC Phase 3+ crises. If a simple AR baseline using only IPC_{t-1} and $IPC_{neighbours}$ achieves Recall=0.78, the marginal contribution of 0.04 (4 percentage points) reflects standard aggregate reporting. **But this framing obscures what operationally matters:** those 4 percentage points represent hundreds of real crises—the *hardest cases to predict*—where persistence fails and early warning could save lives. If the ensemble rescues 200 AR-missed crises 8 months in advance, providing humanitarian actors time to pre-position food aid, negotiate

access, and mobilise funding, that is not a “4 percentage point gain.” **It is 200 operationally critical early warnings for conflict-driven shocks, rapid-onset displacements, and regime transitions—the very crises where timely intervention matters most.** As demonstrated in this dissertation, news features drive 74.7% of marginal predictions (SHAP attribution) for these AR-missed cases, providing dominant signal precisely where it is most needed.

The autocorrelation trap has three fundamental implications for the field:

First, it inflates the apparent value of complex features. High performance may reflect temporal persistence rather than genuine signals from text. Food security crises in regions like South Sudan, Somalia, Yemen, and the Sahel persist for extended periods due to structural factors: chronic poverty, recurrent climate shocks (droughts, floods), protracted conflicts, weak governance, limited market access, and poor infrastructure. The IPC Phase 3 classification (Crisis) or Phase 4 (Emergency) often persists for 6-12 months with only minor fluctuations. A model that simply predicts $IPC_t = IPC_{t-1}$ will achieve high accuracy in such contexts. Without AR baseline comparisons, we cannot isolate the marginal contribution of text features beyond what temporal persistence already captures.

Spatial autocorrelation creates additional structure. Adjacent districts share common exposure to regional shocks (droughts affect entire watersheds, conflicts spill across borders, market disruptions propagate through trade networks) and exhibit spatial clustering of outcomes. A model that incorporates `IPC_neighbours` captures these spatial dependencies. While persistence-dominated cases are well-captured by AR baselines, the critical 26.8% of shock-driven crises break these patterns—precisely where news features provide dominant marginal signal (74.7% SHAP attribution).

Second, it obscures *when* and *where* news actually matters. If most predictions succeed due to autocorrelation, news features may only help in specific contexts that get averaged out in aggregate metrics:

- **Crisis types:** News may matter for conflict-driven and displacement crises (where temporal patterns break due to rapid-onset shocks) but not for climate-driven crises (where seasonal patterns dominate and persistence is strong).
- **Geographic contexts:** News may matter in news-dense regions with extensive media coverage (Kenya, Nigeria, Ethiopia) but not in remote areas with limited reporting (rural Mozambique, northern Mali) [26, 27].
- **Temporal dynamics:** News may matter during regime transitions (peaceful $\times$ violent, stable $\times$ chaotic) but not during stable periods where structural persistence dominates.
- **Forecast horizons:** News may matter at longer horizons (8-12 months) where persistence weakens but not at shorter horizons (2-4 months) where autocorrelation is strongest [28, 29].

Aggregate evaluation metrics (overall AUC, precision, recall) average across these heterogeneous

contexts, obscuring the specific conditions where text features provide value. We need disaggregated analysis by crisis type, country, temporal period, and horizon to identify when news provides dominant predictive signal versus when persistence patterns dominate.

Third, it hinders operational deployment and resource allocation. Early warning systems operate under resource constraints: limited budgets for data acquisition, finite computational capacity for model training and inference, scarce human expertise for model maintenance and interpretation, and bounded attention from humanitarian decision-makers [30]. If simple AR baselines achieve 90-95% of news model performance using only freely available historical IPC data (no web scraping, no NLP pipelines, no GPU infrastructure), why invest in complex text-based systems?

The answer depends on *selective deployment*: if news features help primarily in specific contexts (conflict zones, rapid-onset shocks, news-dense regions), systems should deploy them selectively rather than universally. A two-stage framework that: (1) uses AR for structurally persistent cases (the majority), and (2) deploys complex features only for AR-difficult cases (the minority), maximises value while minimising cost. But this requires knowing which cases are AR-difficult—which in turn requires building the AR baseline and analysing its failures.

Current practice treats all cases equally, deploying the same news-based model universally. This misallocates resources: over-investing in contexts where persistence suffices (wasting money on unnecessary complexity) and under-investing in contexts where richer text sources (social media, humanitarian reports, local-language news) might complement English-language news for AR-difficult cases.

Our AR baseline provides empirical grounding for these concerns. At 8-month forecast horizons using stratified spatial cross-validation (5 folds, 20 geographic clusters), the AR model achieves:

- **AUC:** 0.907 (90.7% of perfect discrimination)
- **Precision:** 0.732 (73.2% of predicted crises are actual crises)
- **Recall:** 0.732 (73.2% of actual crises are correctly predicted)
- **F1:** 0.732 (harmonic mean of precision and recall)
- **Confusion matrix:** TP=3,895, TN=13,973, FP=1,427, FN=1,427 (out of 20,722 total observations)

These metrics are reported at the optimal threshold (0.629) selected via a balanced-constrained optimisation strategy that maximises precision-recall parity while meeting a minimum performance constraint (precision, recall ≥ 0.60), ensuring operationally viable performance for humanitarian deployment.

This performance is achieved using only:

- **Temporal autoregressive feature (Lt):** Past IPC value at t-1 (first-order lag)
- **Spatial autoregressive feature (Ls):** Inverse-distance weighted IPC values from surrounding districts within a 300km radius

- **No text features, no covariates:** Zero news articles, zero GDELT data, zero NLP processing, zero external predictors—only autoregressive values of IPC itself

The challenge is not whether news features *can* predict crises (they can), but whether they add value *beyond what persistence already captures* (the marginal contribution). This dissertation takes the autocorrelation trap seriously, treating it as a methodological imperative rather than a theoretical curiosity.

This section established the autocorrelation trap as the central methodological challenge: food security crises exhibit strong temporal and spatial persistence, enabling spatio-temporal autoregressive baselines to achieve $AUC=0.907$ using only two features (L_t and L_s) with zero text features or external covariates. Without rigorous comparison against such baselines, high performance in news-based models may reflect autocorrelation rather than genuine signals from text. This trap inflates apparent value of complex features, obscures when and where news actually matters, and hinders operational deployment decisions. Addressing this trap requires treating AR baseline comparisons as mandatory rather than optional.

1.3 Research Gap

Despite growing interest in news-based forecasting for humanitarian crises, and increasing recognition of machine learning’s potential for social good applications, existing literature exhibits five critical gaps that this dissertation addresses:

1.3.1 Gap 1: Lack of Rigorous AR Baseline Comparisons

Most published work on news-based crisis prediction evaluates text features against one of three baseline types:

Naive baselines (most common): stratified random sampling, always-predict-majority-class, or uniform random predictions. These baselines are trivially weak—any reasonable model beats them—making them uninformative about genuine predictive value. Comparing against naive baselines is akin to claiming athletic prowess by racing against stationary opponents.

Simple lag baselines (less common): including y_{t-1} as a single control variable in regression models or as one feature among many in ML classifiers. While better than naive baselines, this approach suffers from two problems: (a) it does not optimise temporal autoregressive features (L_t could be 1, 2, 3, or more lags), and (b) it omits spatial autoregressive features entirely, ignoring the well-documented spatial clustering of food security outcomes [19, 31].

No baselines (surprisingly common): directly evaluating news-based models against held-out test sets without any baseline comparison, claiming success based on achieving “high” AUC (>0.70) or accuracy ($>75\%$). This approach provides no information about marginal contribution—we cannot know if the text features add value beyond trivial persistence.

We are unaware of any work in the food security domain that:

- Systematically compares news-based models against spatio-temporal AR baselines with both temporal autoregressive features (Lt: first-order lag of past IPC values, t-1) and spatial autoregressive features (Ls: neighboring IPC values)
- Implements spatial autoregressive weighting (e.g., inverse-distance weighting within specified radius)
- Uses proper spatial cross-validation to prevent geographic information leakage [32, 33]
- Reports marginal contribution of text features after accounting for persistence

This gap is consequential. If AR baselines routinely achieve 85-95% of news model performance (as our results suggest), the entire premise of news-based early warning requires rethinking. The value proposition shifts from “news predicts crises” (true but misleading) to “news provides marginal value beyond persistence in specific contexts” (more accurate but less compelling).

1.3.2 Gap 2: Inability to Distinguish Structural Persistence from Shock-Driven Dynamics

Food security crises have two distinct temporal components that existing methods do not separate:

Structural persistence: Chronic food insecurity driven by slow-moving factors (poverty, climate vulnerability, weak governance, poor infrastructure, market fragmentation) [34, 35, 36]. These conditions persist for years, exhibiting strong temporal autocorrelation. South Sudan has experienced IPC Phase 3+ conditions for most of 2013-2024 due to protracted conflict, political instability, and economic collapse. Persistence dominates prediction in such contexts—knowing IPC_{t-1} is highly informative about IPC_t .

Shock-driven dynamics: Rapid-onset events that disrupt existing patterns (conflict escalation, coups d’état, acute displacement, market collapse, extreme weather events) [37, 38]. These shocks break temporal autocorrelation, making persistence-based predictions fail. The 2023 Sudan conflict (April 2023 outbreak of fighting in Khartoum) triggered acute food insecurity in previously stable regions within weeks, rendering historical patterns obsolete.

Existing methods fit a single model to all cases, implicitly assuming that predictive patterns are homogeneous. This assumption fails: AR baselines work well for structurally persistent cases (the majority) but fail for shock-driven cases (the minority). News features provide their greatest value for the latter (where temporal patterns break and where early warning matters most) while persistence suffices for the former.

We need methods that:

- Explicitly model structural persistence through AR baselines
- Identify shock-driven cases as AR failures (where persistence breaks)
- Deploy complex features selectively for difficult cases only
- Evaluate performance separately for persistent vs shock-driven crises

This gap motivates our two-stage framework: use AR for structure, use news for shocks.

1.3.3 Gap 3: Absence of Two-Stage Frameworks Leveraging AR Strengths

If AR baselines capture structural persistence effectively, why not use them explicitly? Current practice deploys the same news-based model universally, treating all cases as equally difficult. This one-size-fits-all approach is inefficient:

- **Over-engineering easy cases:** For structurally persistent crises where AR suffices, deploying complex NLP pipelines (keyword extraction, topic modelling, regime detection) wastes computational resources
- **Under-engineering hard cases:** For shock-driven crises where AR fails, basic news aggregation alone may miss critical signals that advanced NLP techniques (HMM regime detection, DMD temporal dynamics, semantic embeddings) could capture

A two-stage framework addresses both inefficiencies:

1. **Stage 1 (AR baseline):** Cheap, fast, captures persistence. Achieves 73.2% precision/recall.
2. **Stage 2 (dynamic features):** Expensive, slow, captures shocks. Deploys only for `WITH_AR_FILTER` cases ($IPC_{t-1} \leq 2$ AND $AR=0$).

This selective deployment maximises value per unit cost. Stage 1 handles 70% of persistence-dominated cases efficiently with AR baselines. Stage 2 deploys sophisticated news-based methods for the critical 26.8% of shock-driven cases where news features drive predictions. Combined framework achieves better coverage (recall) with strategic resource allocation.

Existing literature lacks two-stage frameworks because it lacks AR baselines to define Stage 1. Our work provides both.

1.3.4 Gap 4: Limited Model Interpretation Frameworks for Geographic and Temporal Heterogeneity

Aggregate evaluation metrics (overall AUC, precision, recall) obscure heterogeneity. Consider a model with $AUC=0.80$ overall. This aggregate could reflect:

- Homogeneous performance: $AUC \approx 0.80$ in all countries (news helps uniformly)
- Heterogeneous performance: $AUC=0.95$ in Kenya, 0.65 in Mali (news helps selectively)

These scenarios have different implications:

- Homogeneous $\Rightarrow$ deploy news features universally
- Heterogeneous $\Rightarrow$ deploy news features selectively (only where AUC is high)

Most published work reports aggregate metrics only, providing no disaggregated analysis by:

- **Country:** Does news help equally in Kenya, Somalia, Nigeria, Ethiopia?

- **Crisis type:** Conflict vs climate vs structural vs displacement-driven?
- **Temporal period:** Stable periods vs regime transitions vs acute shocks?
- **News coverage density:** High-coverage vs low-coverage regions?

We need model interpretation frameworks that identify:

- Which features matter most (feature importance rankings)
- Which countries are most sensitive to news features (mixed-effects random coefficients)
- Which specific cases benefit from news (SHAP value analysis)
- Cross-method agreement and divergence (triangulation across approaches)

This gap motivates our three-method model interpretation framework (XGBoost, mixed-effects, SHAP) with extensive disaggregation by country, crisis type, and temporal context.

1.3.5 Gap 5: Static Feature Engineering (Article Counts Only)

Most existing work uses static features derived from news content:

- **Article counts:** Number of articles mentioning keywords (drought, conflict, famine)
- **Ratios:** Articles per month, normalised by baseline coverage
- **Sentiment scores:** Average tone, polarity, subjectivity

These features miss three types of dynamic signals:

Regime transitions: Latent narrative states that shift abruptly. A region may transition from “peaceful/stable” regime (low conflict coverage, high economic activity) to “violent/chaotic” regime (high conflict coverage, displacement reports). Hidden Markov Models (HMM) can detect such transitions even when article volumes remain constant—coverage shifts from economic news to conflict news, signaling regime change. Static article counts miss this.

Temporal patterns: Crisis evolution modes capturing how narratives develop. Early-stage crises may show gradual escalation (increasing conflict reports, displacement warnings), while late-stage crises show sustained intensity (persistently high coverage). Dynamic Mode Decomposition (DMD) extracts these temporal modes. Static features aggregate over time windows, losing temporal structure.

Dynamic shifts: Standardised deviations from baseline. Raw article counts conflate absolute levels (Kenya gets more coverage than Mali due to larger English-language media presence) with relative changes (sudden surge in Mali coverage signals emerging crisis). Z-score standardisation (12-month sliding window) captures dynamic shifts. Static ratios normalise by total coverage but miss temporal dynamics.

We propose dynamic feature engineering:

- HMM: 1,322 district-pooled 2-state models for regime detection

- DMD: Crisis-focused mode filtering for temporal pattern extraction
- Z-scores: 12-month sliding-window standardisation for dynamic shifts
- Mixed-effects: Country random effects for geographic heterogeneity

Ablation studies quantify the marginal contribution of each component.

These five gaps—lack of rigorous AR baseline comparisons, inability to distinguish structural persistence from shock-driven dynamics, absence of two-stage frameworks, limited model interpretation for geographic heterogeneity, and static feature engineering—represent systematic omissions in existing literature. Existing work evaluates news-based models against weak baselines (naïve or simple lag), deploys complex features universally rather than selectively, reports aggregate metrics that obscure heterogeneity, and uses static article counts that miss dynamic signals. This dissertation addresses all five gaps simultaneously through a comprehensive framework that establishes AR baselines, separates persistence from shocks, deploys features selectively, triangulates model interpretation across three methods, and engineers dynamic features via stochastic state-space modelling (HMM) and spectral decomposition (DMD).

1.4 Research Questions

This dissertation addresses five core research questions that span methodological critique, feature engineering, hidden variables, framework performance, and geographic heterogeneity:

1. **RQ1: The Autocorrelation Trap.** To what extent can spatio-temporal autoregressive baselines replicate the performance of news-based forecasting models, and what does this reveal about the value of text features in crisis prediction?

This question challenges the field’s foundational assumption that news-based models provide substantial predictive value. If AR baselines achieve 90-95% of news model performance using zero text features, claims about the “value of news for early warning” require fundamental rethinking. We establish the magnitude of the autocorrelation trap empirically, demonstrating that $AUC=0.907$ is achievable through simple persistence alone.

2. **RQ2: When News Matters.** What is the role of different kinds of news features (conflict, displacement, economic, weather) and dynamic transformations (ratio vs z-score) in predicting food insecurity beyond autoregressive baselines?

This question decomposes “news features” into constituent components to identify which specific signals contribute to prediction. We conduct ablation studies comparing 8 model variants: ratio-only ($AUC\ 0.727$), z-score-only (0.699), combined (0.696), with HMM (0.703), with DMD (0.698). Feature importance rankings identify top contributors: conflict, displacement, food security categories. We evaluate on `WITH_AR_FILTER` subset specifically (6,553 observations where $IPC_{t-1} \leq 2$ AND AR predicted non-crisis), isolating news value beyond persistence.

3. **RQ3: The Role of Hidden Variables.** What is the contribution of latent regime detection (HMM) and temporal pattern extraction (DMD) in identifying crises that autoregressive models miss?

This question evaluates dynamic feature engineering beyond static article counts. Do HMM-detected regime transitions and DMD-extracted temporal modes provide value? Ablation studies reveal that HMM provides substantial interpretability value—the `hmm_ratio_transition_risk` feature ranks #5 in importance (0.032, equivalent to 3.2%), capturing narrative regime shifts (peaceful $\times$ violent transitions) that raw article counts miss. HMM achieves +0.007 AUC, demonstrating that latent dynamics provide genuine signal for detecting when crisis narratives fundamentally change. DMD achieves +0.002 AUC with the largest mixed-effects coefficient (+352.38), targeting rare but extreme humanitarian catastrophes.

4. **RQ4: Two-Stage Framework Performance.** Can a two-stage residual modelling approach effectively rescue crises missed by autoregressive baselines, and what are the precision-recall trade-offs of such a framework?

This question evaluates the operational viability of selective deployment. The framework achieves 249 key saves (17.4% of 1,427 AR failures)—rescuing the hardest-to-predict crises where temporal persistence breaks down. Recall increases to 0.779 (+6.4% relative improvement), prioritising sensitivity in humanitarian contexts where missing crises is catastrophic. Precision decreases to 0.585, reflecting the deliberate choice to favour recall over precision. We demonstrate that these 249 rescued cases concentrate in conflict-affected regions (Sudan, Zimbabwe, DRC) experiencing rapid-onset shocks where 8-month early warnings enable life-saving interventions.

5. **RQ5: Geographic Heterogeneity.** Are news-based features equally valuable across all geographic contexts, or do certain countries and crisis types benefit more from dynamic news signals than others?

This question disaggregates aggregate metrics to identify where news helps most. Results reveal strong heterogeneity: Zimbabwe (77 key saves), Sudan (59), and DRC (40) account for 70.7% of all key saves despite representing only 3 of 18 countries. Mixed-effects random coefficients quantify country-specific sensitivities. Within-country heterogeneity analysis demonstrates that the same countries show both cascade rescues and failures at district level—Zimbabwe has 77 saves but 647 still-missed cases, Sudan has 59 saves but 420 still-missed, revealing that news-based early warning succeeds in well-covered districts (capitals, conflict zones) but fails in news desert districts (remote pastoral areas, peripheral regions) within the same country. This heterogeneity enables strategic deployment optimisation: concentrating news-based forecasting resources in high-coverage contexts (Sudan/Zimbabwe/DRC) where dense media ecosystems and clear crisis narratives maximise predictive value, while relying on AR baselines for contexts with sparse coverage where simpler persistence models provide adequate performance.

These questions guide systematic investigation into when, where, and how dynamic news signals

provide genuine early-warning information beyond spatio-temporal persistence. Each question is answerable with our data, methods, and empirical results.

These five research questions span the full scope of methodological critique (RQ1), feature engineering decomposition (RQ2), hidden variable evaluation (RQ3), operational framework performance (RQ4), and geographic heterogeneity (RQ5). Each question addresses a distinct aspect of when and where news features provide value beyond autocorrelation, and each is empirically answerable using our dataset (20,722 observations across 1,920 districts in 18 countries after $h=8$ filtering), two-stage framework, ablation studies across 8 model variants, and three-method interpretability analysis. Together, these questions reframe news-based forecasting from universal deployment claims to selective deployment guidance grounded in rigorous baseline comparisons and honest assessment of trade-offs.

1.5 Research Objectives

To address the five research questions, this dissertation pursues six specific research objectives:

Objective 1: Establish Rigorous AR Baseline with Spatial Cross-Validation

Develop a spatio-temporal autoregressive model using:

- **Temporal autoregressive feature:** IPC outcome at $t-1$ (first-order lag capturing temporal persistence)
- **Spatial autoregressive feature:** Inverse-distance weighted average of neighbours' IPC within 300km radius
- **Logistic regression:** Binary classification ($IPC \geq 3$ vs $IPC < 3$)
- **L2 regularization:** Ridge penalty to prevent overfitting on spatial neighbours
- **Stratified spatial CV:** 5 folds, 20 geographic clusters, ensures no test-set neighbours in training

Target performance: Demonstrate that high performance ($AUC > 0.90$) is achievable without text features, establishing the autocorrelation trap as an empirically significant phenomenon.

Objective 2: Quantify and Characterise AR Failures

Define AR failures as cases where:

$$IPC_{t-1} \leq 2 \quad \text{AND} \quad AR_pred = 0 \quad \text{BUT} \quad IPC_t \geq 3 \quad (1.1)$$

These represent missed early-warning opportunities—crises that temporal patterns did not forecast. Quantify:

- **Failure rate:** Proportion of total crises that AR baseline fails to predict

- **Geographic distribution:** Identify which countries exhibit highest AR failure rates
- **Temporal patterns:** Determine when failures occur (stable periods vs shock events)
- **Crisis characteristics:** Distinguish conflict-driven vs climate-driven failure patterns

Target: Demonstrate systematic patterns in AR failures exist, providing empirical justification for Stage 2 intervention.

Objective 3: Engineer Dynamic Features Through Four-Stage Pipeline

Implement advanced feature engineering beyond static article counts:

Stage 2a - Z-Score Standardisation:

$$z_{i,t,c} = \frac{x_{i,t,c} - \mu_{i,c}(t)}{\sigma_{i,c}(t)} \quad (1.2)$$

where $x_{i,t,c}$ is article count for district i , time t , category c ; $\mu_{i,c}(t)$ and $\sigma_{i,c}(t)$ are 12-month rolling mean and standard deviation. Captures dynamic shifts.

Stage 2b - Hidden Markov Models:

- 2-state models (peaceful/crisis regimes) per district
- District-level pooling: 1,322 models across unique districts
- Extract features: regime probabilities, transition risks, state entropy

Captures latent narrative regimes.

Stage 2c - Dynamic Mode Decomposition:

$$\mathbf{X}' \approx \mathbf{A}\mathbf{X} \quad (1.3)$$

where $\mathbf{X}$ is news time series matrix, $\mathbf{A}$ is DMD operator. Eigendecomposition extracts temporal modes. Crisis-focused filtering selects modes correlated with IPC outcomes. Captures temporal patterns.

Stage 2d - Mixed-Effects Regression:

$$\log \frac{p_{r,t}}{1 - p_{r,t}} = \underbrace{\beta^T \mathbf{X}_{r,t}}_{\text{Fixed effects}} + \underbrace{\alpha_g + \mathbf{b}_g^T \mathbf{Z}_{r,t}}_{\text{Random effects}} \quad (1.4)$$

where $\beta^T \mathbf{X}_{r,t}$ are fixed effects for all features, $\alpha_g \sim N(0, \sigma_\alpha^2)$ are group random intercepts (adaptive: district-level if data sufficient, else country-level), $\mathbf{b}_g^T \mathbf{Z}_{r,t}$ are random slopes for key signals (conflict_ratio, food_security_ratio), with $\mathbf{Z}_{r,t} \subseteq \mathbf{X}_{r,t}$. Here r indexes regions (districts), t indexes time, and g indexes groups. Captures both global patterns (fixed effects) and geographic heterogeneity (random intercepts + random slopes for crisis-predictive features).

Target: Engineer comprehensive feature set spanning multiple transformation types (ratio, z-score normalization, HMM regime detection, DMD temporal modes, location metadata) to capture

diverse aspects of crisis dynamics.

Objective 4: Rescue AR Failures Through Selective Deployment

Deploy Stage 2 models exclusively on AR-difficult cases (WITH_AR_FILTER strategy):

- Filter training data to cases where previous $IPC \leq 2$ (non-crisis) AND AR predicted non-crisis ($AR=0$), isolating instances where temporal persistence suggests stability
- Train Stage 2 models on this filtered subset to focus learning on difficult-to-predict cases
- Quantify rescue rate: proportion of AR failures successfully predicted by Stage 2
- Prioritise recall improvement in humanitarian contexts where false negatives are costly
- Analyse geographic concentration: determine if rescue success clusters in specific regions

Target: Demonstrate that news signals can rescue operationally critical cases where spatio-temporal persistence fails, validating selective deployment strategy.

Objective 5: Conduct Comprehensive Interpretability Analysis

Deploy three complementary methods to triangulate findings:

Method 1 - XGBoost Feature Importance: Use gain-based importance scores to identify which features contribute most to tree splits, revealing:

- Relative ranking of location metadata vs news features vs dynamic features
- Whether static ratio features or dynamic z-score/HMM/DMD features dominate
- Split frequency patterns that may overstate location metadata importance

Method 2 - Mixed-Effects Decomposition: Fixed effects β capture global patterns. Random intercepts α_g and random slopes $\mathbf{b}_g$ quantify group-specific baseline risks and feature sensitivities. Variance decomposition:

$$\text{Var}(y) = \underbrace{\text{Var}(\mathbf{X}\beta)}_{\text{fixed}} + \underbrace{\text{Var}(\alpha_g) + \text{Var}(\mathbf{b}_g^T \mathbf{Z})}_{\text{random}} + \underbrace{\text{Var}(\epsilon)}_{\text{noise}} \quad (1.5)$$

Identifies which geographic groups (districts or countries) have elevated baseline risks and which features have heterogeneous effects across groups.

Method 3 - SHAP Values: Model-agnostic explanations quantify feature contributions to individual predictions:

$$\phi_k(i) = \sum_{S \subseteq F \setminus \{k\}} \frac{|S|!(|F| - |S| - 1)!}{|F|!} [f_{S \cup \{k\}}(x_i) - f_S(x_i)] \quad (1.6)$$

where $\phi_k(i)$ is SHAP value for feature k in instance i . Validates cross-method agreement with XGBoost and mixed-effects.

Target: Identify consensus features that rank high across all three methods.

Objective 6: Quantify Geographic Heterogeneity

Disaggregate results by country to identify where news features provide most value:

- **Key saves by country:** Quantify successful AR failure rescues disaggregated by country
- **Concentration patterns:** Determine if rescue success concentrates in specific countries or distributes evenly
- **Mixed-effects random coefficients:** Quantify country-specific baseline risks and feature sensitivities
- **Contextualise heterogeneity:** Relate geographic differences to conflict intensity, news coverage density, and crisis type

Target: Demonstrate that news features matter differently across contexts, providing evidence for context-specific selective deployment rather than universal application.

These six objectives provide the operational roadmap for addressing the five research questions. Objective 1 establishes the methodological foundation (rigorous AR baseline with spatial CV), Objectives 2-4 implement the two-stage framework (characterising AR failures, engineering dynamic features, attempting selective rescue), and Objectives 5-6 provide interpretability and geographic insights (triangulation across three methods, country-level heterogeneity analysis). Each objective has measurable targets: demonstrating high AR baseline performance without text features, quantifying AR failure patterns, engineering comprehensive dynamic feature sets, evaluating selective deployment effectiveness, achieving cross-method consensus in feature importance rankings, and identifying geographic contexts where news features provide greatest marginal value. Together, these objectives operationalise the research vision into concrete, evaluable research activities.

1.6 Contributions

This research makes five core contributions to humanitarian early warning, machine learning for social good, and crisis forecasting methodology:

Contribution 1: Methodological Critique - Exposing the Autocorrelation Trap

We demonstrate empirically that spatio-temporal AR baselines achieve 93.8% of published news model performance (AR PR-AUC=0.7652 vs Balashankar et al. 2023 PR-AUC=0.8158) using ZERO text features—only temporal autoregressive feature (IPC_t-1) and spatial autoregressive feature (inverse-distance weighted neighbours). This establishes the autocorrelation trap as a quantitatively large, empirically real phenomenon that existing literature has systematically neglected.

Our critique has three components:

(a) Empirical demonstration: AR baseline achieves Precision=0.732, Recall=0.732, F1=0.732, AUC=0.907 at 8-month horizons with 5-fold stratified spatial CV. This performance approaches

published news-based models (93.8% of Balashankar et al.’s PR-AUC), demonstrating that persistence dominates prediction.

(b) Theoretical implication: Without AR baseline comparisons, high performance in existing work may reflect autocorrelation rather than text feature value. Claims about “news predicts crises” are technically true but potentially incomplete—persistence predicts most crises, and news features contribute incrementally. Our cascade framework demonstrates that news features provide value when deployed selectively on AR failures (249 key saves, 17.4% rescue rate), rather than universally.

(c) Methodological prescription: All future work on news-based (or any feature-based) crisis prediction should include rigorous AR baseline comparisons with both temporal autoregressive features and spatial autoregressive features, inverse-distance spatial weighting, proper spatial CV, and reported marginal contributions. This sets a higher standard for the field.

To our knowledge, this is the first systematic comparison of news-based models against strong spatio-temporal baselines in the food security domain. Our work challenges existing paradigms and provides a template for future methodological rigor.

Contribution 2: Two-Stage Residual Modelling Framework

We develop a principled approach that explicitly separates structural persistence (captured by AR baseline) from shock-driven dynamics (captured by news features):

Stage 1 - AR Baseline: Deploys spatio-temporal logistic regression on all cases. Achieves 73.2% precision/recall. Identifies 15,400 predicted non-crises ($AR_pred=0$) as candidates for Stage 2 override.

Stage 2 - Dynamic Features: Deploys XGBoost with 35 advanced features (ratio, z-score, HMM, DMD, location) exclusively on `WITH_AR_FILTER` subset (6,553 cases where $IPC_{t-1} \leq 2$ AND AR predicted non-crisis). Achieves 249 successful predictions of AR-missed crises (17.4% rescue rate). **These 249 cases are not statistical abstractions—they represent the most operationally critical early warnings:** conflict escalations in Sudan where displacement unfolds rapidly, coup-related disruptions in Zimbabwe where temporal patterns break abruptly, and acute emergencies in DRC where persistence models fail. These are precisely the cases where 8-month advance warning enables life-saving humanitarian response.

Integration: Simple cascade decision logic preserves all $AR=1$ predictions (trusting the baseline when it predicts crisis), and uses Stage 2’s binary prediction for all $AR=0$ cases. When AR predicts no crisis ($AR=0$) and Stage 2 predicts crisis ($Stage2=1$), the cascade overrides to crisis. When both predict no crisis, the cascade confirms non-crisis. Combined framework achieves:

- **249 key saves** (AR-missed crises correctly predicted by cascade)—*the hardest cases where news signals matter most*
- **Recall: 0.732 $\rightarrow$ 0.779 (+6.4% relative improvement)**—not merely a percentage gain, but 249 real crises affecting millions, now predicted 8 months early

- Precision: $0.732 \rightarrow 0.585$ (reduced due to prioritising recall in humanitarian contexts)
- Geographic concentration: 70.7% of key saves in Sudan, Zimbabwe, DRC—conflict-affected regions where news signals capture rapid-onset shocks

This framework provides three methodological innovations:

(a) Selective deployment: Complex features deployed only for `WITH_AR_FILTER` cases ($IPC_{t-1} \leq 2$ AND $AR=0$, not universally), maximising value per cost. The framework targets the 6,553 cases meeting both filter conditions rather than all 20,722 observations, concentrating computational resources where they provide genuine marginal value.

(b) Explicit persistence modelling: AR baseline captures structural persistence explicitly (not as implicit control variables), enabling interpretable decomposition of which predictions succeed due to autocorrelation (the majority) versus which require dynamic news signals (the critical minority).

(c) Humanitarian-appropriate metrics: Prioritises recall over precision, aligning with operational early warning principles where missing crises is catastrophic while false alarms are manageable. The framework achieves 77.9% recall, successfully predicting 4,144 of 5,322 total crises, including 249 that pure persistence models cannot detect.

The framework demonstrates meaningful success for operationally critical cases: 17.4% of AR failures are rescued, concentrated in conflict-affected regions (Sudan, Zimbabwe, DRC) experiencing rapid-onset shocks where early warning enables life-saving humanitarian response.

Critically, cascade failure analysis reveals a fundamental constraint: the 1,178 crises still missed after cascade intervention (82.6% of AR failures) exhibit systematic news coverage deficiency—median 74 articles/month compared to 121 for rescued cases (64% less coverage, $p < 0.001$). This *news deserts hypothesis* demonstrates that news-based early warning fundamentally cannot rescue crises in remote pastoral areas (Kenya Northern, Zimbabwe rural districts, Niger) lacking sufficient media coverage. The 249 key saves concentrate in news-dense conflict zones (70.7% in Sudan/Zimbabwe/DRC), revealing that successful cascade deployment requires rich news signal infrastructure. This partial success validates the core hypothesis: *news signals provide genuine early-warning value for specific crisis types in specific geographic contexts*, precisely where temporal persistence breaks down, news coverage is abundant, and where intervention matters most. Future NLP systems must expand beyond traditional news media to incorporate social media monitoring, humanitarian situation reports, community radio transcripts, and multilingual text mining from non-English sources to address news deserts.

Contribution 3: Dynamic Feature Engineering Beyond Article Counts

We demonstrate a four-stage analytical pipeline that extends beyond static article counts used in existing work:

Stage 2a - Z-Score Standardisation: 12-month sliding-window normalisation captures dynamic shifts. Ablation studies reveal a nuanced finding: ratio-only models achieve higher standalone AUC (0.727 vs 0.699), but SHAP analysis shows z-score features account for 74.7% of

marginal attribution in the full combined model versus only 20.1% tree-based importance. This apparent contradiction reflects complementary roles: ratio features provide stable cross-sectional baselines enabling higher standalone performance, while z-score features capture volatile temporal anomalies that drive marginal predictions when combined with ratios. Both feature types are essential—ratios for robust baseline discrimination, z-scores for detecting dynamic shocks.

Stage 2b - Hidden Markov Models: 1,322 district-pooled 2-state models extract latent narrative regimes. The `hmm_ratio_transition_risk` feature ranks #5 in importance (0.032), demonstrating that regime transitions provide genuine signal. HMM achieves +0.007 AUC gain (from 0.696 to 0.703) with substantial interpretability value—we can identify when narratives shift from peaceful to violent regimes even when article volumes remain constant, capturing qualitative changes in crisis dynamics.

Stage 2c - Dynamic Mode Decomposition: Crisis-focused mode filtering extracts temporal patterns (escalation modes, sustained intensity modes). DMD achieves +0.002 AUC, with `dmd_ratio_crisis_instability` achieving the *largest mixed-effects coefficient among all features* (+352.38), demonstrating value for detecting rare but extreme humanitarian catastrophes where multiple crisis drivers converge simultaneously. DMD provides interpretable crisis evolution dynamics, identifying temporal modes that characterise how crises unfold over time.

HMM and DMD contributions (HMM: +0.007 AUC, DMD: +0.002 AUC) reflect an important methodological insight: **specialized methods provide value through targeted detection**. With 48 months of data per district (2021-2024) and heterogeneous news coverage (mean 1,235 articles/year/district, with many districts having sparse coverage), HMM and DMD achieve robust convergence: 89.5% for HMM regime detection and 83.1% for DMD crisis mode extraction. These high convergence rates demonstrate successful latent dynamics extraction despite data constraints. The finding that ratio-only models achieve standalone AUC=0.727 while z-score features account for 74.7% of SHAP marginal attribution demonstrates **feature complementarity matters more than individual dominance**. HMM captures regime transitions (ranked #5 in feature importance), while DMD achieves the largest coefficient (+352.38) for extreme events, demonstrating that specialized methods provide mechanistic insights complementing discrimination-focused features.

Stage 2d - Mixed-Effects Regression: Country random intercepts and random slopes for key signals (`conflict_ratio`, `food_security_ratio`) capture geographic heterogeneity in both baseline risk and feature effects. Fixed effects quantify global patterns, while random effects reveal country-specific deviations. Mixed-effects models provide interpretable coefficient decomposition (AUC 0.604-0.620) with transparent fixed effects (average impact across all countries) and random effect variances (geographic variation in baseline risk and feature sensitivities), complementing XGBoost’s higher discrimination (AUC 0.697) through a trade-off between interpretability and predictive accuracy.

Ablation studies across 8 model variants quantify marginal contributions:

- Ratio-only: AUC 0.727 (best standalone performance, but z-scores account for 74.7% SHAP attribution in combined models)

- Z-score-only: AUC 0.699 (captures temporal anomalies, complementary to ratios)
- Ratio+Z-score: AUC 0.696 (lower standalone AUC, but z-scores drive 74.7% of marginal attribution in full models)
- Ratio+Z-score+HMM: AUC 0.703 (+0.007 from HMM)
- Ratio+Z-score+DMD: AUC 0.698 (+0.002 from DMD)
- Advanced (all features): AUC 0.697 (XGBoost optimises combination)

These results demonstrate that dynamic feature engineering provides value with heterogeneous contributions across methods (HMM provides substantial interpretability value and +0.007 AUC gain, z-scores account for 74.7% SHAP marginal attribution despite lower standalone AUC). Comprehensive reporting of which methods succeed in which contexts (HMM transition risk for regime shifts, z-score features for shock detection, ratio features for stable baselines) advances the field by providing practical guidance for operational deployment.

Contribution 4: Comprehensive Interpretability Framework

We deploy three complementary methods to answer when and where news matters, achieving triangulation across approaches:

XGBoost Feature Importance (gain-based):

1. country_data_density: 0.133 (captures baseline news coverage level)
2. country_baseline_conflict: 0.093 (captures baseline conflict exposure)
3. country_baseline_food_security: 0.067 (captures baseline food insecurity)
4. other_ratio: 0.033 (top news feature, general news coverage)
5. hmm_ratio_transition_risk: 0.032 (top hidden variable, regime transitions)
6. health_ratio: 0.029, displacement_z-score: 0.026, weather_ratio: 0.026
7. food_security_z-score: 0.025, hmm_ratio_crisis_prob: 0.025

Key insight: Location/baseline features dominate tree-based importance (ranks #1-3, 40.4% total split frequency) but contribute minimally to SHAP attribution (ranks #17, 20, 26, only 2.6% marginal impact). This 15.5 $\times$ overstatement reveals that tree-based importance measures stratification utility (frequent node splitting for country-level segmentation), not predictive contribution (driving marginal predictions). Z-score features drive 74.7% of SHAP attribution despite lower tree rankings $\times$ dynamic news anomalies matter more than geographic context for marginal predictions on shock-driven crises.

Mixed-Effects Decomposition (fixed/random coefficients): Fixed effects β identify global patterns (which features matter on average across all countries). Random intercepts α_g quantify group-specific baseline risk deviations, while random slopes $\mathbf{b}_g$ quantify group-specific feature

sensitivities (how much does conflict_ratio or food_security_ratio matter more in group g vs average). Variance decomposition reveals:

$$\frac{\text{Var}(\alpha_g)}{\text{Var}(y)} \approx 0.15 - 0.25 \quad (1.7)$$

indicating that 15-25% of outcome variance is explained by country-level baseline heterogeneity (random intercepts). This justifies mixed-effects models over pooled regression.

Key insight: Sudan, Zimbabwe, and DRC exhibit positive random effects (higher sensitivity to news features), while Kenya and Ethiopia exhibit negative random effects (lower sensitivity, persistence dominates). This heterogeneity motivates selective deployment.

SHAP Values (model-agnostic explanations): Shapley value decomposition attributes predictions to individual features for specific instances. **Critical methodological revelation:** SHAP fundamentally reorders feature rankings compared to tree-based importance, revealing that location features account for 40.4% of tree splits but only 2.6% of marginal prediction attribution (15.5× overstatement). This exposes measurement artifact: tree-based importance measures split frequency (how often features partition data), while SHAP measures marginal impact (contribution to individual predictions). Z-score features dominate SHAP attribution (74.7%), HMM features account for 21.9% (higher than tree-based 13.0%), and ratio features contribute 22.7%. This demonstrates that feature "importance" depends critically on measurement method×location features enable baseline stratification (high split frequency) but dynamic features drive prediction variance (high marginal impact).

- **Z-score features dominate:** other_z-score (rank 1, 0.952), conflict_z-score (rank 2, 0.911), humanitarian_z-score (rank 3, 0.902)×74.7% total SHAP attribution
- **Location features rank low:** country_data_density (rank 17), country_baseline_conflict (rank 20), country_baseline_food_security (rank 26)×only 2.6% SHAP despite 40.4% tree importance (15.5× overstatement)
- **HMM features elevated:** hmm_ratio_crisis_prob (rank 7), hmm_ratio_transition_risk (rank 8)×21.9% SHAP (higher than 13.0% tree-based)
- **DMD features specialized for extreme events:** 1.5% SHAP reflects rarity by design (activate only for <3% most severe crises), but achieve largest mixed-effects coefficient (+352.38) when active, detecting complex emergencies invisible to other features

Key insight: The dramatic divergence between tree-based importance (location 40.4%, z-score 20.1%) and SHAP attribution (location 2.6%, z-score 74.7%) reveals a critical methodological artifact. Location features enable stratification (frequent splitting) but dynamic features drive prediction variance (marginal impact). This demonstrates that feature “importance” depends fundamentally on measurement method×split frequency ≠ predictive contribution.

News Theme Analysis Component: This comprehensive interpretation framework enables systematic analysis of which news themes matter across measurement methods. Weather emerges as strongest for sustained forecasts (Mixed Effects #1, +26.7 coefficient) via direct

causal pathway (climate $\rightarrow$ agriculture $\rightarrow$ food), while conflict dominates for rapid shock detection (SHAP z-scores #1, 0.911) via anomaly spikes. The measurement paradox extends to theme-level rankings: sustained compositional changes (ratios, mixed effects favour weather/displacement/food security) differ from rapid anomalies (z-scores, SHAP favour conflict/humanitarian/governance). This systematic cross-method theme comparison resolves contradictory findings in literature and provides deployment guidance: use weather/climate signals for agricultural crises, conflict/humanitarian signals for rapid-onset shocks.

Contribution 5: Geographic Insights - Selective Deployment Justification

We identify strong geographic heterogeneity justifying selective rather than universal deployment of news features:

Key saves by country:

- Zimbabwe: 77 key saves (30.9% of total)
- Sudan: 59 (23.7%)
- DRC: 40 (16.1%)
- Nigeria: 27 (10.8%)
- Mozambique: 15 (6.0%)
- Mali: 12, Kenya: 8, Ethiopia: 6, Malawi: 3, Somalia: 2

Concentration: Top 3 countries (Zimbabwe, Sudan, DRC) account for 176 key saves (70.7% of total) despite representing only 3 of 18 countries (16.7%). This extreme concentration suggests news features help in specific contexts, not universally.

Contextualization:

- **Zimbabwe:** High news coverage, conflict-driven crises (Marange crisis 2021), frequent regime transitions. SHAP-based theme analysis reveals elevated Weather importance (11.5% vs 9.4% global), reflecting recurring drought cycles compounding economic collapse.
- **Sudan:** April 2023 conflict outbreak, acute displacement, breakdown of temporal patterns. Elevated Conflict theme importance (14.6% vs 11.3% global) reflects rapid violence escalations that AR baseline cannot anticipate.
- **DRC:** Protracted eastern conflict, M23 resurgence, persistent displacement. Elevated Displacement importance (12.2% vs 10.0% global) captures population movements from North Kivu complex emergency.
- **Contrast with Ethiopia/Kenya:** Despite high baseline coverage, few key saves. AR baseline works well due to strong seasonal patterns (climate-driven crises) and structural persistence. Relatively flat theme distributions indicate no dominant shock type requiring news-based anomaly detection.

Implication: Operational systems should deploy news features selectively:

- **Deploy in:** Sudan, Zimbabwe, DRC (conflict zones, regime instability, AR fails frequently). Theme-aware monitoring guided by elevation analysis (context-specific deviations from global patterns, not just dominant themes): Weather monitoring systems for southern/eastern agricultural zones (Zimbabwe +2.1pp elevation, Ethiopia, Malawi, Madagascar) where climate shocks produce largest deviations from continental baseline; Conflict early-warning for Sahel/Sudan corridor (Sudan +3.3pp, Mali) where violence spikes exceed global conflict patterns; Displacement tracking for Great Lakes (DRC +2.2pp) where population movements dominate local news; Health surveillance in East Africa (Somalia +5.8pp, highest observed elevation) where disease burden compounds food insecurity. This elevation metric identifies which shock types require context-specific surveillance thresholds rather than universal monitoring.
- **Omit in:** Ethiopia, Kenya (seasonal patterns, persistence dominates, AR suffices)
- **Resource allocation:** Concentrate expensive NLP infrastructure where it helps (3 high-value countries) rather than spreading thin across all 18 countries.

This selective deployment framework maximises humanitarian impact per unit cost.

These five contributions establish this dissertation's significance: methodological critique exposing the autocorrelation trap (AR achieves 93.8% of Balashankar et al. 2023 published news model performance with zero text features, challenging field assumptions), two-stage residual framework separating persistence from shocks (249 key saves rescuing the hardest-to-predict crises, 17.4% of AR failures, concentrated in conflict-affected regions), dynamic feature engineering beyond article counts (HMM provides substantial interpretability value capturing regime transitions, z-score features account for 74.7% of SHAP marginal attribution demonstrating dynamic anomaly detection drives predictions), comprehensive interpretability via triangulation (XGBoost, mixed-effects, SHAP converge on z-score temporal anomalies and HMM transition risk as top contributors), and geographic insights justifying selective deployment (Zimbabwe/Sudan/DRC account for 70.7% of key saves, demonstrating heterogeneous value across contexts). Together, these contributions reframe news-based forecasting from universal deployment claims to selective deployment guidance targeting contexts where news signals rescue operationally critical cases, establish AR baselines as mandatory methodological standards, and provide practical guidance for humanitarian early warning systems through comprehensive evaluation of what works where.

1.7 Thesis Structure

The remainder of this dissertation is organised as follows:

Chapter 2 provides background on food insecurity classification systems, traditional early warning approaches, and reviews recent literature on news-based forecasting. It identifies the autocorrelation trap as a systematic methodological gap and positions this work's contributions within the broader field.

Chapter 3 describes the two-stage cascade framework in detail: data sources (IPC assessments,

GDELT news), feature engineering (autoregressive baselines, ratio/z-score transformations, HMM regime detection, DMD temporal modes), model architectures (logistic regression, XGBoost, mixed-effects models), and evaluation protocols (stratified spatial cross-validation, performance metrics).

Chapter 4 presents comprehensive results organised around the five research questions, including AR baseline performance, identification of missed early-warning opportunities, ablation studies quantifying dynamic feature contributions, model interpretability analysis, and cascade framework evaluation with case studies from Zimbabwe, Sudan, and DRC.

Chapter 5 discusses findings in the context of the broader literature, addresses limitations (data coverage, language bias, temporal resolution), and proposes future research directions including real-time deployment, multi-modal integration, and enhanced interpretability for humanitarian decision-making.

Chapter 6 synthesises key contributions, restates answers to the five research questions, and reflects on the significance of this work for both methodological standards in food security forecasting and operational guidance for humanitarian early warning systems.

Chapter 2

Brief Background and Literature Review

2.1 Food Insecurity and IPC Classification

Food insecurity represents one of humanity's most persistent humanitarian challenges, affecting 282 million people across 59 crisis-affected countries globally [2]. Sub-Saharan Africa bears a disproportionate burden, with multiple countries experiencing protracted crises driven by conflict, climate shocks, economic instability, and governance failures. The humanitarian consequences are severe: acute malnutrition, disease outbreaks, population displacement, economic collapse, and in extreme cases, famine-related mortality.

To standardise crisis severity assessment and enable coordinated humanitarian response, the international community developed the Integrated Food Security Phase Classification (IPC) system [2]. The IPC provides a scientifically rigorous, consensus-based framework for classifying food insecurity into five phases of increasing severity:

- **Phase 1 (Minimal):** Households can meet essential food and non-food needs without engaging in atypical coping strategies. Less than 20% of households experience inadequate food consumption.
- **Phase 2 (Stressed):** Households minimally meet food needs but face difficulty meeting non-food needs. Must engage in stress-coping strategies (sell productive assets, reduce non-essential spending). 20-30% of households experience inadequate consumption.
- **Phase 3 (Crisis):** Households face food consumption gaps with high acute malnutrition. Must engage in crisis-level coping (liquidate productive assets, send children to work, consume seed stocks). At least 30% of households experience inadequate consumption. *This is the humanitarian crisis threshold.*
- **Phase 4 (Emergency):** Households experience large food consumption gaps resulting in very high acute malnutrition and excess mortality. Emergency coping strategies are employed (abandonment of livelihoods, distress migration). At least 40% of households experience inadequate consumption.
- **Phase 5 (Famine/Catastrophe):** Households face near-complete lack of food and/or other basic needs. Starvation, death, and destitution are evident. At least 60% of households

experience inadequate consumption [39].

IPC classifications are conducted at the district level (typically corresponding to Administrative Level 2 or ADM2 boundaries) through multi-stakeholder technical working groups coordinated by national governments with support from FEWSNET, WFP, FAO, and other humanitarian agencies [40]. These assessments synthesise evidence from multiple data sources: household food security surveys (Household Hunger Scale [41, 42], Food Consumption Score [43], reduced Coping Strategies Index [44, 45]), anthropometric measurements of acute malnutrition (weight-for-height z-scores in children under 5), mortality data (crude death rate, under-five death rate), and contextual information on food availability, market access, livelihood strategies, and shocks [2].

The humanitarian significance of IPC Phase 3 or higher (Phase 3+) classifications cannot be overstated. Phase 3+ designations trigger coordinated international response mechanisms: emergency food assistance programs, supplementary feeding for malnourished children, livelihood support interventions, market stabilization efforts, and advocacy for political solutions to underlying drivers [40]. Resource allocation decisions by major donors (USAID, ECHO, DFID, UN Central Emergency Response Fund) rely heavily on IPC classifications to prioritise funding across competing humanitarian crises globally.

District-level granularity is critical because food insecurity is highly heterogeneous even within provinces or regions. Neighboring districts may experience vastly different outcomes due to localized conflict exposure, microclimate variation, market access constraints, or livelihood diversity. Ethiopia’s 2021-2024 food security crisis illustrates this heterogeneity: while southern pastoral districts (Borana, Guji) experienced severe drought-driven Crisis (Phase 3) and Emergency (Phase 4) outcomes [46, 47, 48], neighboring highland agricultural districts maintained Stressed (Phase 2) or even Minimal (Phase 1) conditions due to different rainfall patterns and livelihood systems.

The temporal dynamics of IPC assessments matter for early warning. Assessments are conducted every 4 months (typically in February, June, and October) to track evolving crisis conditions and project outcomes 4 months ahead based on seasonal forecasts, market trends, and anticipated shocks [2]. This temporal structure creates forecasting opportunities: if we can predict which districts will deteriorate from Phase 2 (Stressed) to Phase 3 (Crisis) or Phase 4 (Emergency) in the coming months, humanitarian agencies can pre-position food supplies, negotiate access with governments, mobilise funding, and implement preventive interventions before populations exhaust coping capacities.

However, this same temporal structure creates methodological challenges. IPC outcomes exhibit strong temporal persistence: if a district is classified Phase 3 in February, it is highly likely to remain Phase 3 in June absent major shocks or interventions. This persistence reflects the chronic, slow-moving nature of food insecurity drivers—poverty, climate vulnerability, weak governance, poor infrastructure—which change on timescales of years, not months. Any forecasting model must distinguish complementary prediction mechanisms: AR baselines capture the 73.2% of persistence-dominated cases, while news features provide dominant signal for the critical 26.8% of shock-driven crises where temporal patterns break.

This study focuses on district-level IPC assessments from 24 African countries spanning 2021-2024, comprising 55,129 unique district-period observations across 3,241 unique administrative districts. After applying $h=8$ forecast horizon requirements and data quality filters (sufficient historical data, GDELT coverage, geographic matching success), the final analysis dataset contains 20,722 observations across 1,920 districts in 18 countries. Countries with extensive coverage include Ethiopia (1,416 raw districts, 12,843 raw records), Kenya (274 districts, 7,712 records), Sudan (212 districts, 3,876 records), Nigeria (199 districts, 3,658 records), and Mozambique (169 districts, 2,809 records). This comprehensive dataset enables rigorous evaluation of forecasting approaches across diverse geographic contexts, crisis types, and temporal periods.

This section established the IPC classification system as the gold standard for food insecurity assessment, spanning five phases from Minimal (Phase 1) to Famine (Phase 5), with Phase 3+ (Crisis or worse) representing the humanitarian response threshold. District-level assessments conducted quarterly provide spatially granular, temporally frequent measurements enabling early warning, but also create methodological challenges due to strong temporal persistence in outcomes. The dataset for this study—55,129 district-period observations from 24 African countries spanning 2021-2024—provides comprehensive coverage for evaluating forecasting approaches while requiring careful attention to separating genuine predictive signals from simple persistence patterns.

2.2 Existing Early Warning Approaches

Traditional food security early warning systems have evolved over four decades from expert-driven narrative assessments to data-intensive, quantitative forecasting systems. The Famine Early Warning Systems Network (FEWSNET), established in 1985 in response to the 1984-1985 famines in Sudan and Ethiopia [49], pioneered systematic monitoring by integrating satellite remote sensing, market price tracking, rainfall monitoring, and field assessments into regular multi-country outlook reports [40]. FEWSNET analysts synthesise these diverse data streams into projected IPC classifications for upcoming seasons, providing 4-6 month outlooks updated monthly for crisis-affected regions.

Satellite-based vegetation monitoring forms the foundation of most early warning systems. The Normalised Difference Vegetation Index (NDVI), derived from satellite imagery comparing red and near-infrared reflectance, provides a proxy for agricultural conditions and pasture availability. Declining NDVI indicates vegetation stress from drought, pests, or other shocks, potentially foreshadowing food production shortfalls months before harvest [8]. NDVI data is globally available at 250-1000m spatial resolution with 8-16 day temporal revisit frequencies from sources such as MODIS [50], enabling broad coverage impossible through ground-based monitoring alone.

Precipitation monitoring complements vegetation indices by tracking rainfall deficits directly. The Climate Hazards InfraRed Precipitation with Stations (CHIRPS) dataset provides daily rainfall estimates at 0.05-degree resolution (approximately 5km) from 1981 to present, combining satellite observations with ground station data [6]. Multi-scale validation studies across Africa demonstrate that CHIRPS reliably detects no-rain events and performs well at monthly timescales ($KGE >$

0.75 in Eastern Africa), making it particularly suitable for drought monitoring applications [51]. The TAMSAT (Tropical Applications of Meteorology using SATellite data and ground-based observations) system provides complementary rainfall products specifically calibrated for African drought monitoring, available from 1983 onwards [7]. These precipitation datasets enable detection of dry spells, wet spell interruptions, and seasonal rainfall deficits that drive agricultural failures [52].

Market price monitoring tracks food availability and affordability through systematic collection of staple grain prices (maize, sorghum, wheat, rice) from key markets. Unusually high prices signal supply shortfalls or demand surges, while price volatility indicates market stress. The World Food Programme (WFP) maintains extensive market monitoring networks across crisis-prone countries, collecting weekly or bi-weekly price data from hundreds of markets. Market-based indicators have demonstrated predictive value for food insecurity, particularly when combined with rainfall and vegetation data in data-driven models [8].

Household survey data provides direct measurement of food consumption, coping strategies, and nutritional status. FEWSNET's Livelihoods Baseline Profiles document typical household food sources, income patterns, expenditure, and seasonal calendars for distinct livelihood zones (pastoral, agro-pastoral, agricultural). Periodic Household Economy Analysis (HEA) assessments quantify how shocks (drought, conflict, price spikes) affect different livelihood groups' ability to meet minimum food and income needs. These surveys provide rich contextual information but are logistically intensive, expensive, and conducted infrequently—typically annually or bi-annually—limiting their utility for near-real-time early warning.

Despite these strengths, traditional early warning approaches face four fundamental limitations:

Temporal lag: Satellite data suffers from processing delays (2-4 weeks from image acquisition to product availability) and cloud cover interference (particularly problematic in equatorial regions during rainy seasons when agricultural monitoring is most critical). Market price data may lag behind local conditions when collection systems are disrupted by conflict or infrastructure failures. Household surveys capture conditions only at the time of fieldwork, missing rapid-onset shocks between assessment cycles.

Spatial resolution trade-offs: While satellite data provides broad geographic coverage, coarse spatial resolution (250m-1km for NDVI, 5km for CHIRPS) may miss localized crises in small districts or areas with heterogeneous topography. Market price data is collected from major trading centres, potentially missing conditions in remote areas with limited market integration. Household surveys face resource constraints limiting sample sizes and geographic coverage.

Expert interpretation bottlenecks: FEWSNET and WFP outlooks require expert analysts to synthesise diverse data streams, interpret trends, assess contextual factors (conflict, policy, seasonal patterns), and project outcomes. This expert-driven process ensures credibility and stakeholder acceptance but creates capacity constraints—a limited number of skilled analysts must cover dozens of crisis-affected countries, potentially delaying alerts when multiple crises emerge simultaneously.

Context specificity: Relationships between indicators (NDVI, rainfall) and outcomes (food insecurity) vary across livelihood systems. Pastoral livelihoods in arid regions respond quickly to rainfall deficits affecting pasture availability, while agricultural livelihoods depend on seasonal rainfall timing and distribution during critical growth stages. Static thresholds (e.g., “NDVI below X indicates crisis”) fail to capture this heterogeneity, necessitating context-specific calibration and expert interpretation.

Recent innovations have begun addressing these limitations through machine learning approaches. Lentz, Michelson, Baylis, and Zhou [8] demonstrated that data-driven models combining rainfall, market prices, and demographic variables outperform expert-driven assessments for predicting food insecurity crises in East Africa, achieving improved lead time and geographic coverage. Busker, Hurk, Moel, Homberg, Straaten, Odongo, and Aerts [18] developed XGBoost models for IPC prediction in the Horn of Africa using diverse covariates including climate anomalies, vegetation indices, conflict event data, and economic indicators, demonstrating that ensemble machine learning can capture complex non-linear relationships between drivers and outcomes.

Herteux, Räth, Martini, Baha, Koupparis, Lauzana, and Piovani [53] showed that forecasting performance increases systematically with the product of temporal and spatial training samples, suggesting that expanding geographic coverage and historical depth of training data improves generalisation. Foini, Tizzoni, Martini, Paolotti, and Omodei [54] analysed the forecastability of food insecurity metrics, identifying temporal autocorrelation as both an opportunity (persistence enables baseline forecasts) and a challenge (separating genuine signal from temporal structure).

However, these machine learning advances have not systematically addressed the autocorrelation problem. Even sophisticated models combining diverse data sources may achieve high performance primarily by learning temporal and spatial persistence patterns rather than capturing genuine predictive signals from covariates. Without rigorous comparison against autoregressive baselines—models using only temporal autoregressive features (Lt: first-order lag of past IPC values at t-1) and spatial autoregressive features (Ls: inverse-distance weighted IPC values from neighboring districts)—we cannot isolate the marginal contribution of satellite data, market prices, or any other feature beyond what simple persistence already captures.

This section reviewed traditional early warning approaches centred on satellite vegetation monitoring (NDVI), precipitation tracking (CHIRPS, TAMSAT), market price surveillance, and household surveys, identifying four fundamental limitations: temporal lag from processing delays and cloud cover, spatial resolution trade-offs between coverage and granularity, expert interpretation bottlenecks limiting scalability, and context-specific relationships requiring localized calibration. Recent machine learning innovations have begun addressing these limitations through data-driven models, but have not systematically confronted the autocorrelation problem—high performance may reflect learning persistence patterns rather than genuine predictive signals from covariates. This sets the stage for examining news-based forecasting approaches and their vulnerability to the same autocorrelation trap.

2.3 News-Based Forecasting and the Autocorrelation Problem

2.3.1 Existing News-Based Approaches

News media offers a compelling alternative data source for humanitarian early warning, addressing several limitations of traditional satellite and survey approaches. News coverage is near real-time, updated continuously as events unfold rather than subject to the 2-4 week processing lags affecting satellite products. News captures ground-level perspectives unavailable to satellite sensors: conflict dynamics (armed clashes, displacement), economic disruptions (market failures, inflation, unemployment), policy changes (export bans, humanitarian access restrictions), and local impacts of weather events (flood damage, drought stress on communities). Unlike satellite data, news coverage can detect crises in urban areas, cloud-covered regions, and conflict zones where physical access for surveys is impossible.

The Global Database of Events, Language, and Tone (GDELT) has emerged as the dominant data source for news-based crisis forecasting. Launched in 2013, GDELT monitors print, broadcast, and web news sources in over 100 languages from every country globally, processing hundreds of thousands of articles daily [55]. The GDELT Global Knowledge Graph (GKG) extracts structured information from news text including named entities (people, organisations, locations), themes and categories (conflict, humanitarian assistance, economic indicators), emotional tone and sentiment, and geocoded location mentions. This structured representation enables quantitative analysis of global news coverage at scale.

Balashankar, Subramanian, and Fraiberger [17] demonstrated the potential of news-based forecasting in their Science Advances paper predicting food crises using 11.2 million news articles from Factiva. Analysing coverage from 1980-2020 across 21 countries, they used natural language processing (frame-semantic parsing and word embeddings) for feature extraction and Random Forest regression to predict binary food crisis outcomes (IPC Phase 3+ vs below Phase 3) at district level. Their models achieved strong predictive performance (news-only model: PR-AUC=0.82; combined models incorporating traditional indicators: PR-AUC=0.82-0.91) at 3-month primary forecast horizons (with evaluations at 1, 3, 6, 9, and 12 months), demonstrating that textual features extracted from news coverage—conflict keywords, economic crisis terms, humanitarian appeal language—correlate with subsequent food insecurity outcomes.

Earlier work by Balashankar, Subramanian, and Fraiberger [56] extended this approach to finer geographic granularity, predicting district-level food insecurity up to 12 months ahead using news text and location mentions aligned to administrative boundaries. These studies established proof-of-concept that news media data contains predictive signals for humanitarian crises, potentially complementing or substituting for traditional satellite and survey approaches in contexts where those methods face limitations.

However, a critical methodological gap pervades this literature: *systematic omission of autoregressive baseline comparisons*. The Balashankar studies and related work evaluate news-based models against held-out test sets using standard train-test splits or cross-validation, demonstrating that text features improve prediction accuracy. Performance gains are attributed to the informational content of news coverage: conflict reports signal impending displacement

and market disruption, economic news captures inflation and unemployment dynamics, weather reports indicate agricultural shocks, humanitarian coverage reflects access constraints and response gaps.

These evaluations rarely compare against strong temporal baselines. Some studies include simple lag features (IPC_{t-1}) as control variables in regression models, but we are unaware of any published work that:

- Systematically compares news-based models against spatio-temporal AR baselines with both temporal autoregressive features L_t (first-order lag of past IPC values, $t-1$) and spatial autoregressive features L_s (inverse-distance weighted neighboring IPC values)
- Uses proper spatial cross-validation to prevent geographic information leakage [32, 57]
- Reports the marginal contribution of text features after accounting for temporal and spatial persistence
- Analyses when and where text features provide value beyond what baseline autocorrelation captures

This omission is consequential. Without AR baseline comparisons, we cannot distinguish two competing explanations for high news model performance:

Hypothesis 1 (Signal): News features capture genuine predictive information beyond temporal patterns—conflict reports foreshadow displacement crises, economic coverage reveals market failures, humanitarian appeals indicate deteriorating conditions that satellite data misses. High performance reflects the informational value of text.

Hypothesis 2 (Autocorrelation): News features correlate with past and neighboring outcomes due to spatio-temporal persistence. High performance may primarily reflect learning that “today looks like yesterday” and “here looks like there” rather than genuine signals from dynamic news content. However, disaggregated analysis is critical: news features may provide dominant signal for specific hard-to-predict cases where persistence breaks down (driving 74.7% of marginal predictions for AR-missed cases, as demonstrated in this dissertation), even while providing less value for persistence-dominated cases well-captured by AR baselines.

These hypotheses are not mutually exclusive—news features could capture both genuine signals and autocorrelation—but their relative contributions determine the value proposition of news-based systems. However, **the standard framing of “marginal contribution” as percentage-point differences obscures what operationally matters.** Consider a news model achieving $\text{Recall}=0.82$ compared to an AR baseline at $\text{Recall}=0.78$ —a “small” 4 percentage point gain. *But this is not a modest statistical improvement:* those 4 percentage points represent hundreds of real food crises affecting millions of people—**the hardest-to-predict cases** where temporal persistence breaks down (conflict escalations, coup-related disruptions, rapid-onset displacements). These are precisely the crises where 6-8 month early warning enables life-saving humanitarian response: pre-positioning food supplies before roads become impassable, negotiating humanitarian access before violence intensifies, mobilising funding before populations exhaust coping strategies.

When an ensemble *rescues* these AR-missed cases—detecting conflict-driven shocks in Sudan, coup impacts in Zimbabwe, displacement crises in DRC—it is not delivering a “modest gain.” **It is providing critical early warnings for the cases that matter most**, where persistence models fail and where timely intervention can prevent famine, death, and displacement.

The field has treated AR baseline comparisons as optional methodological enhancements rather than mandatory validity checks. This dissertation argues they are mandatory: any feature-based forecasting approach (news, satellite, market prices, social media) must demonstrate marginal value beyond spatio-temporal persistence to credibly claim predictive utility.

This subsection reviewed news-based forecasting approaches, highlighting their advantages (near real-time coverage, ground-level perspectives, crisis detection in cloud-covered and conflict-affected regions) and demonstrated predictive performance (Balashankar et al.: PR-AUC=0.82 for news-only model; combined models: PR-AUC=0.82-0.91 at 3-month primary horizons). However, systematic omission of autoregressive baseline comparisons pervades this literature, creating ambiguity about whether high performance reflects genuine predictive signals from text features versus learning temporal and spatial persistence patterns. Without rigorous AR baselines, we cannot distinguish signal from autocorrelation—a fundamental methodological gap this dissertation addresses.

2.3.2 The Autocorrelation Trap

Food security crises exhibit strong temporal and spatial autocorrelation, creating a methodological trap for forecasting research. Understanding this trap requires examining the spatio-temporal structure of crisis data and its implications for model evaluation.

Temporal autocorrelation arises because food insecurity is fundamentally a chronic, slow-moving phenomenon. Districts classified as IPC Phase 3 (Crisis) or Phase 4 (Emergency) in one quarter typically remain in crisis the following quarter absent major shocks or interventions. This persistence reflects the structural drivers of food insecurity—chronic poverty, climate vulnerability, weak governance, poor infrastructure, market fragmentation—which change on timescales of years, not months.

Protracted crises in South Sudan (Phase 3-4 conditions persisting from 2013-2024 due to ongoing conflict and economic collapse), Somalia (recurrent drought-driven crises with brief recovery intervals), and Yemen (continuous Phase 3+ conditions since 2015 due to conflict and blockade) exemplify this temporal persistence. For such contexts, a naive forecasting model predicting $IPC_t = IPC_{t-1}$ (“tomorrow equals today”) achieves high accuracy simply by capturing structural persistence. Any forecasting model incorporating temporal lags will learn this pattern.

Spatial autocorrelation arises because neighboring districts share common exposure to regional shocks and exhibit spatial diffusion of crises. Droughts affect entire watersheds, not individual districts in isolation. Conflict spillovers cross administrative boundaries through refugee flows, armed group movements, and trade disruptions. Market shocks propagate through spatially

networked trading systems. Food insecurity outcomes consequently cluster in space: if a district experiences Phase 3 crisis, neighboring districts likely experience similar or adjacent phases.

Ayalew, Dessie, Mitiku, and Zewotir [19] documented this spatial clustering empirically, calculating Global Moran’s I statistics of 0.22-0.285 for food insecurity across Africa during 2015-2021, indicating statistically significant positive spatial autocorrelation. Wubetie, Zewotir, Mitku, and Dessie [31] demonstrated in Ethiopia that 90% of food insecurity variation is explained by spatial effects in geo-additive mixed models with Markov Random Field priors. Dessie, Zewotir, and North [58] showed that geographically weighted regression models incorporating spatial heterogeneity substantially outperform non-spatial models, highlighting the strength of spatial structure.

These patterns enable powerful baseline forecasts using only autoregressive features:

Temporal autoregressive features (Lt): First-order lag of past IPC values at $t-1$. If $IPC_{t-1}=3$ (Crisis), predicting $IPC_t=3$ will often be correct due to persistence.

Spatial autoregressive features (Ls): Inverse-distance weighted average of neighboring districts’ IPC outcomes within a spatial radius (e.g., 300km). If all neighbours are Phase 3, the focal district likely experiences similar conditions due to shared shocks and spatial diffusion. Spatial weights capture both proximity (nearby neighbours weighted more heavily) and shock propagation (crises spread through space).

A logistic regression model using only these two autoregressive feature types—with zero text features, zero satellite data, zero market prices, zero external covariates of any kind—can achieve remarkably high performance by exploiting spatio-temporal persistence. This creates the autocorrelation trap: news-based models (or any feature-based models) may achieve high performance not because text features provide genuine predictive signals, but because they incorporate (explicitly or implicitly through temporal structure in training data) the same persistence patterns that AR baselines capture more directly.

The trap has three critical implications:

First, it clarifies where news features add value. Consider a news-based model achieving $AUC=0.85$ evaluated against a held-out test set. This performance appears impressive compared to naive baselines (always-predict-majority-class: $AUC \approx 0.50$) or random classifiers. However, if a simple AR baseline using only IPC_{t-1} and $IPC_neighbors$ achieves $AUC=0.83$, this reveals that $0.83/0.85 = 97.6\%$ of predictive signal comes from temporal/spatial persistence, while news features contribute the remaining 2.4%. This does not diminish news value—it *clarifies* their role: the AR baseline captures the 73.2% of persistence-dominated crises, while news features drive predictions for the critical 26.8% of shock-driven crises where AR fails. Selective deployment targeting these shock-driven cases (as demonstrated in this dissertation) achieves better humanitarian impact than universal deployment, as SHAP analysis reveals news features drive 74.7% of marginal predictions for AR-missed crises.

Without AR baseline comparisons, we cannot distinguish these scenarios. High absolute performance ($AUC > 0.80$) appears successful, but marginal contribution requires explicit

quantification through rigorous baseline evaluation.

Second, it obscures when and where features matter. Even if aggregate marginal contribution appears small (news model: Recall=0.85, AR baseline: Recall=0.83, marginal gain: +0.02), this statistical summary could mask profoundly different operational realities:

Homogeneous marginal contribution: News features help slightly in all contexts (all countries, all crisis types, all periods), with consistent +0.02 Recall everywhere. In this case, news adds universal value, but the gains are spread across routine, easily-predicted cases rather than concentrated in the hardest cases where intervention matters most.

Heterogeneous marginal contribution—this is the critical scenario: News features provide **substantial value for the hardest cases** (conflict-driven crises in Sudan, Zimbabwe, DRC where temporal patterns break: +20-30% rescue rate of AR failures, driving predictions through dominant SHAP attribution) while persistence patterns dominate in routine cases (climate-driven crises in Ethiopia, Kenya where AR baselines suffice), averaging to +4.7 percentage points overall. **In this case, news is invaluable precisely where it is most needed**—for rapid-onset shocks, regime transitions, and conflict escalations where AR baselines fail and where early warning can prevent catastrophic humanitarian outcomes. The aggregate +4.7 percentage point Recall gain obscures that news is *rescuing the cases that matter most for saving lives*.

Aggregate evaluation metrics cannot distinguish these scenarios. **A 4.7-percentage-point Recall gain that represents 249 early warnings for conflict-driven displacements 8 months in advance is not a statistical artifact—it is transformative for humanitarian response in exactly the contexts where intervention matters most.** Disaggregated analysis comparing news models to AR baselines across geographic contexts, crisis types, and temporal periods is required to identify where news features provide dominant signal versus where persistence patterns dominate.

Third, it hinders operational deployment and resource allocation. Humanitarian early warning systems operate under severe resource constraints: limited budgets for data acquisition and processing, finite computational capacity for model training and inference, scarce human expertise for system maintenance, and bounded attention from decision-makers. If simple AR baselines achieve 90-95% of complex feature-based model performance using only freely available historical IPC data (no web scraping, no GPU infrastructure, no NLP expertise), operational systems should prioritise AR baselines and deploy complex features selectively only where they provide substantial marginal value.

Current practice treats all cases equally, deploying the same news-based (or satellite-based, or multi-modal) model universally. This misallocates resources: over-investing in contexts where persistence suffices (wasting resources on unnecessary complexity) and under-investing in contexts where additional data sources might complement features for difficult cases.

The autocorrelation trap is not merely a theoretical concern. Our empirical results demonstrate it is quantitatively large: a spatio-temporal AR baseline using only two autoregressive features (Lt: temporal autoregressive feature using first-order lag IPC_{t-1} ; Ls: spatial autoregressive

feature of neighboring IPC values) achieves $AUC=0.907$, $Precision=0.732$, $Recall=0.732$, and $F1=0.732$ at 8-month forecast horizons with 5-fold stratified spatial cross-validation across 55,129 district-period observations. This performance approaches published news-based models (93.8% of Balashankar et al.’s PR-AUC) while using *zero text features* and *zero external covariates*—only lagged values of the dependent variable (IPC) itself.

This finding requires rethinking the value proposition of news-based (and more broadly, feature-based) crisis forecasting. The question is not whether features *can* predict crises—they demonstrably can—but whether they add value *beyond what temporal and spatial persistence already captures*. Answering this question requires establishing AR baselines as mandatory methodological standards, not optional enhancements.




Figure 2.1: Strong temporal persistence enables high AR baseline performance with zero external covariates. Panel A shows temporal Autocorrelation Function (ACF) decay from lag-1 ($ACF=0.85$) to lag-12 ($ACF=0.14$), with $h=8$ forecast horizon ($ACF=0.27$) indicating substantial persistence 8 months ahead. Panel B shows AR baseline performance using ONLY temporal autoregressive feature L_t (past IPC value at $t-1$) and spatial autoregressive feature L_s (inverse-distance weighted neighboring IPC values within 300km)—zero text features, zero satellite data, zero external covariates—achieving $AUC-ROC=0.907$, $Precision=Recall=0.732$. This demonstrates the autocorrelation trap: high performance from pure autoregression (lagged dependent variable only), requiring rigorous AR baseline comparisons to isolate marginal value of any feature-based approach. $n=20,722$ observations, 5-fold stratified spatial CV, $h=8$ months.

This subsection established the autocorrelation trap as a fundamental methodological challenge arising from spatio-temporal persistence in food security crises. Temporal autocorrelation (chronic, slow-moving structural drivers create high persistence) and spatial autocorrelation (neighboring districts share common shocks and spatial diffusion, creating clustering with Global Moran’s $I=0.22-0.285$) enable autoregressive baselines using only past IPC values (L_t) and neighboring IPC values (L_s) to achieve high performance without any external features. This creates three critical problems: inflating apparent feature value (high absolute performance may reflect minimal marginal contribution), obscuring when and where features matter (aggregate metrics average over heterogeneous contexts), and hindering operational deployment (universal deployment wastes resources where persistence suffices). The empirical demonstration that AR baselines achieve $AUC=0.907$ using zero text features quantifies the magnitude of this trap, requiring fundamental

rethinking of feature-based forecasting claims (Figure ??).

2.4 Spatial-Temporal Methods and Cross-Validation

The autocorrelation trap documented in Section 2.3.2 arises from spatial and temporal structure in food security data. Properly accounting for this structure requires specialised methodological approaches for both model features (spatial autoregressive terms) and evaluation strategies (spatial cross-validation). This section reviews evidence for spatial clustering and the cross-validation methods necessary to prevent optimistic performance estimates.

2.4.1 Evidence of Spatial Clustering in Food Insecurity

Multiple studies have documented strong positive spatial autocorrelation in food security outcomes across Africa, quantifying the extent to which neighboring districts experience similar crisis conditions. Ayalew, Dessie, Mitiku, and Zewotir [19] conducted comprehensive spatial analysis of severe food insecurity prevalence across African countries from 2015-2021, calculating Global Moran's I statistics to test for spatial autocorrelation. Their results revealed statistically significant positive spatial clustering in every year analysed, with Moran's I values ranging from 0.22 to 0.285 ($p < 0.001$), indicating that districts with high food insecurity prevalence are systematically surrounded by other high-prevalence districts, while low-prevalence districts cluster together.

This spatial structure arises from multiple mechanisms. Shared exposure to regional shocks creates synchronized crisis dynamics: droughts affect entire river basins spanning multiple administrative boundaries, not isolated districts. Recent Horn of Africa droughts have impacted dozens of districts across Somalia, Ethiopia, and Kenya simultaneously due to shared exposure to failed consecutive rainy seasons. Conflict spillovers propagate geographically through refugee flows, armed group movements across borders, and trade route disruptions. South Sudan's civil conflict generated displacement crises in neighboring districts of Uganda, Sudan, and Ethiopia through cross-border population movements.

Market integration creates spatial diffusion of economic shocks. Price spikes in major trading hubs propagate through spatially networked markets: maize price increases in major urban centres affect prices in surrounding rural areas and neighboring regions through trader arbitrage. Wubetie, Zewotir, Mitku, and Dessie [31] demonstrated in Ethiopia that 90% of variation in food insecurity outcomes is explained by spatial effects in ordinal geo-additive mixed models with Markov Random Field spatial priors, highlighting the dominance of geographic structure over individual district characteristics.

Dessie, Zewotir, and North [58] applied geographically weighted regression to food security data, documenting substantial spatial heterogeneity in covariate effects: the relationship between rainfall deficits and food insecurity varies systematically across regions depending on livelihood systems (pastoral vs agricultural), soil types, and market access. Models incorporating this spatial modification effect substantially outperform non-spatial linear regressions, demonstrating

that both spatial clustering of outcomes and spatial heterogeneity of processes must be modelled explicitly.

For forecasting applications, spatial clustering enables powerful predictions using only neighboring districts' outcomes. An inverse-distance weighted average of IPC values from nearby districts—using zero information about the focal district's conditions—can achieve surprisingly high accuracy simply by exploiting the tendency of crises to cluster geographically. This pattern underlies the spatial autoregressive features (Ls) in AR baselines discussed in Section 2.3.2.

2.4.2 Spatial Cross-Validation Methods

Standard random cross-validation randomly partitions observations into training and test folds, assuming independence between data points. When applied to spatially structured data, this assumption is violated catastrophically: training observations near test observations provide strong signals about test outcomes through spatial autocorrelation, creating information leakage that inflates performance estimates.

Tziachris, Nikou, Aschonitis, Kallioras, Sachsamanoglou, Fidelibus, and Tziritis [32] demonstrated this problem empirically across multiple environmental prediction tasks. Models evaluated via random k-fold cross-validation achieved optimistic performance estimates suggesting excellent predictive capability. However, the same models evaluated via spatial block cross-validation—where training and test sets are geographically separated by buffer zones preventing information leakage—achieved substantially lower performance. This systematic gap represents optimism from spatial leakage rather than genuine predictive capability.

Wang, Khodadadzadeh, and Zurita-Milla [57] reviewed spatial cross-validation strategies, categorising approaches by how they partition data geographically. Spatial block CV divides the study region into contiguous geographic blocks, assigning entire blocks to training or test folds. This ensures test districts are spatially separated from training districts by distances exceeding the range of spatial autocorrelation. Stock [33] analysed optimal block size selection, demonstrating that blocks must be sufficiently large to prevent leakage through long-range spatial correlations while maintaining sufficient data in each fold for stable model training.

Leave-one-region-out cross-validation (LORO-CV) represents an extreme form of spatial blocking, holding out entire countries or provinces as test sets. This strategy evaluates model generalisation to entirely new geographic contexts—critical for humanitarian early warning systems that must forecast in data-scarce regions. However, LORO-CV can be overly conservative, testing extrapolation to unseen contexts rather than interpolation within the training data's geographic extent.

For food security forecasting, spatial CV is mandatory because the spatial scale of autocorrelation (Moran's $I=0.22-0.285$) spans hundreds of kilometers. Neighboring districts share correlated IPC outcomes through regional shocks and market linkages. Random CV allows training on nearby districts, enabling models to exploit spatial autocorrelation rather than demonstrating genuine predictive value from features. Spatial block CV or LORO-CV by country prevents this leakage, providing honest estimates of out-of-sample predictive performance.

This dissertation employs stratified spatial CV with leave-one-country-out blocking for all model evaluations, ensuring that test districts are geographically separated from training districts by national boundaries. This strategy tests spatial generalisation, providing conservative performance estimates appropriate for operational deployment scenarios where models must forecast in new countries.

This section reviewed evidence for strong spatial autocorrelation in food security data (Global Moran's $I=0.22-0.285$ across Africa, 90% of variance explained by spatial effects in Ethiopia) arising from shared exposure to regional shocks, conflict spillovers, and spatially networked markets. Spatial clustering enables powerful baseline forecasts using only neighboring districts' outcomes, creating methodological requirements for spatial cross-validation to prevent information leakage. Standard random CV inflates performance estimates through spatial leakage; spatial block CV and leave-one-region-out CV prevent this by geographically separating training and test sets. This dissertation employs stratified spatial CV with leave-one-country-out blocking, ensuring honest evaluation of spatial generalisation for operational deployment contexts.

2.5 Ensemble and Cascade Methods

While most forecasting research deploys single models universally across all observations, ensemble and cascade frameworks can improve performance by combining multiple models or deploying complex models selectively. This section reviews ensemble methodologies relevant to the two-stage cascade framework developed in this dissertation.

2.5.1 Cascade Ensembles for Selective Deployment

Cascade ensembles deploy models sequentially, with later stages refining initial predictions or addressing cases where initial stages fail. Izonin, Tkachenko, Krak, Berezhsky, Shevchuk, and Shandilya [59] applied cascade architectures to missing data imputation and prediction tasks, demonstrating that two-stage frameworks using simple models for easy cases and complex models for difficult cases achieve better accuracy-cost trade-offs than single universal models. Their approach identifies “easy” cases where Stage 1 predictions are confident and accurate, reserving expensive Stage 2 models for cases where Stage 1 predicted non-crisis or failed predictions.

Kolawole, Dennis, Talwalkar, and Smith [60] revisited cascade ensemble architectures in modern machine learning contexts, evaluating cascade decision trees, cascade neural networks, and cascade gradient boosting ensembles across benchmark datasets. They found that cascades excel when: (1) the data contains distinct easy and hard subsets with different optimal model complexities, (2) Stage 1 failures can be identified reliably, and (3) Stage 2 has access to richer features than Stage 1. All three conditions hold for food security forecasting: persistence-dominated cases are “easy” for AR baselines, AR failures signal “hard” shock-driven cases, and dynamic news features are available for Stage 2.

Zhang, Zhu, and Hua [61] applied ensemble methods to financial crisis prediction, combining autoregressive time series models (capturing persistence in credit growth, asset prices) with

event-based models (processing news about policy changes, corporate failures). Their findings paralleled the autocorrelation problem in food security: high performance from combined models arose primarily from temporal autocorrelation in financial variables, with news events providing marginal value concentrated in crisis transition periods (2007-2008 financial crisis onset). This heterogeneity justified selective deployment of expensive news analysis only during high-volatility periods rather than continuous universal monitoring.

The key insight from this literature: *cascade frameworks aligned with data structure outperform universal models*. If 70% of cases are well-predicted by simple baselines and 30% require complex features, deploying complex models to all cases wastes resources on the 70% while potentially degrading performance through overfitting. Deploying simple baselines to 70% and complex models to the difficult 30% maximises accuracy per unit cost—the operational objective for resource-constrained humanitarian early warning systems.

2.5.2 Implications for Food Security Forecasting

The evidence from cascade ensemble literature motivates this dissertation’s two-stage framework: Stage 1 deploys AR baselines (cheap, leveraging temporal and spatial persistence) to all 55,129 observations. Stage 2 deploys dynamic feature ensembles (expensive, requiring GDELT news processing and HMM/DMD computation) selectively to AR failures—cases where persistence-based prediction fails, indicating shock-driven dynamics where news signals may help.

This design differs from standard ensemble approaches (bagging, boosting, stacking) which combine all models for all observations. Instead, it implements *residual modelling*: Stage 2 models the residuals from Stage 1, attempting to recover cases where the AR baseline fails. The performance metric is not overall accuracy (where AR baseline achieves AUC=0.907 already) but *rescue rate*: the fraction of AR failures correctly predicted by Stage 2. A rescue rate of 17.4% (249 saves out of 1,427 AR failures) demonstrates that dynamic news features provide **operationally critical value for the hardest cases**—the shock-driven crises where persistence fails and early warning saves lives. SHAP analysis reveals news features drive 74.7% of marginal predictions for these AR-missed cases, demonstrating dominant predictive contribution where it matters most.

This section reviewed cascade ensemble methods that deploy models sequentially, using simple models for easy cases and complex models for difficult cases, achieving better accuracy-cost trade-offs than universal deployment. Evidence from missing data recovery, financial crisis prediction, and modern cascade architectures demonstrates that cascades excel when data contains distinct easy/hard subsets, Stage 1 failures can be identified, and Stage 2 has access to richer features. These conditions hold for food security forecasting: persistence-dominated cases are easy for AR baselines, AR failures signal shock-driven cases, and dynamic news features are available for Stage 2. This motivates the dissertation’s two-stage residual modelling framework, evaluated by rescue rate (17.4%, 249 saves) rather than aggregate accuracy.

2.6 Dynamic Feature Engineering: HMM, DMD, and Z-Scores

Even when text features provide marginal value beyond AR baselines, the *representation* of those features matters profoundly. Most existing work uses static article counts or simple ratios: number of articles mentioning keywords (conflict, drought, famine, displacement) per month or per district, normalised by baseline coverage levels. These static features miss three types of dynamic signals that may distinguish genuine early-warning information from noise.

2.6.1 Hidden Markov Models for Regime Detection

Food security crises do not unfold uniformly—they exhibit discrete *narrative regimes* that shift abruptly as conditions deteriorate or stabilize. A region may transition from a “peaceful/stable” regime (characterised by economic development coverage, agricultural productivity reports, minimal conflict mentions) to a “violent/chaotic” regime (dominated by conflict reports, displacement narratives, humanitarian appeals) even when absolute article volumes remain relatively constant. The shift in narrative content rather than volume signals regime change.

Hidden Markov Models (HMMs) provide a principled framework for detecting these latent regimes from observed news time series [62]. An HMM assumes observations (news article counts by category) are generated by an underlying latent state process that transitions stochastically between discrete regimes. The model consists of:

States: Latent regimes $S_t \in \{1, 2, \dots, K\}$, unobserved but inferred from data. For crisis detection, a 2-state model suffices: State 1 (peaceful/stable), State 2 (crisis/conflict).

Observations: Observed news features $\mathbf{X}_t$ at each time t (article counts by category). Assumed to follow state-specific distributions: $\mathbf{X}_t | S_t = k \sim P_k(\mathbf{X})$. Gaussian emissions are common: $\mathbf{X}_t | S_t = k \sim \mathcal{N}(\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$.

Transitions: Transition probabilities $P(S_t = j | S_{t-1} = i) = A_{ij}$ govern regime switches. High A_{12} indicates frequent transitions from peaceful to crisis regimes, signaling instability.

Initial distribution: $P(S_1 = k) = \pi_k$ specifies the starting regime probabilities.

The Baum-Welch algorithm (a variant of Expectation-Maximisation) estimates model parameters $\{\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k, A_{ij}, \pi_k\}$ from observed time series via iterative refinement [62]. Once trained, the forward-backward algorithm infers the posterior probability $P(S_t = k | \mathbf{X}_{1:T})$ that each time point belongs to each regime, providing a “soft” regime assignment.

Yuan and Mitra [63] demonstrated the value of 2-state HMMs for crisis regime detection in financial markets, accurately identifying the transition into the 2008 global financial crisis weeks before traditional indicators. Their approach validated HMM transition probabilities as early-warning signals for regime shifts.

For food security forecasting, HMMs enable extraction of features capturing latent narrative dynamics:

- **Regime probabilities:** $P(S_t = \text{crisis})$. High probability indicates the narrative regime resembles past crisis periods even if absolute article counts are moderate.
- **Transition risks:** $P(S_{t+1} = \text{crisis} | S_t = \text{stable})$. Rising transition probabilities signal increasing likelihood of regime shift into crisis.
- **Entropy:** $H(S_t) = -\sum_k P(S_t = k) \log P(S_t = k)$. High entropy indicates uncertainty about the current regime, potentially signaling transitions in progress.

This dissertation applies HMMs at the district level, pooling data across all time points within each of 1,322 unique districts to estimate district-specific 2-state models. This pooling strategy addresses data sparsity (individual districts may have limited monthly observations) while preserving geographic heterogeneity (different districts have different regime structures). However, as discussed in Chapter 4, HMM performance is constrained by the short time span (48 months, 2021-2024) and heterogeneous news coverage density, with convergence achieved in only 89.5% of observations.

2.6.2 Dynamic Mode Decomposition for Crisis Dynamics

While HMMs detect regime switches, they do not capture the temporal evolution patterns *within* regimes. How do crises escalate? Do they exhibit gradual build-up, abrupt onset, sustained intensity, or cyclical oscillations? Dynamic Mode Decomposition (DMD) provides a data-driven framework for extracting these temporal modes from multivariate time series [64].

DMD originates in fluid dynamics for analysing complex flow patterns [65, 66], but has been adapted for time series forecasting across diverse domains including oceanography [64], climate science, and financial markets. The core idea: approximate a high-dimensional dynamical system by a best-fit linear operator that maps observations at time t to observations at time $t + 1$.

Given snapshot matrices $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_{T-1}]$ and $\mathbf{X}' = [\mathbf{x}_2, \mathbf{x}_3, \dots, \mathbf{x}_T]$ containing news features at consecutive time steps, DMD seeks an operator $\mathbf{A}$ such that:

$$\mathbf{X}' \approx \mathbf{A}\mathbf{X} \quad (2.1)$$

The operator $\mathbf{A}$ is approximated via Singular Value Decomposition (SVD) and eigendecomposition. The eigenvalues λ_k and eigenvectors ϕ_k of $\mathbf{A}$ define temporal modes:

- **Eigenvalues** λ_k : Complex numbers whose magnitude $|\lambda_k|$ indicates growth (> 1) or decay (< 1) and whose angle indicates oscillation frequency.
- **Eigenvectors** ϕ_k : Spatial patterns (which news categories co-vary) associated with each mode.

The time series can be decomposed as:

$$\mathbf{x}_t \approx \sum_{k=1}^r \alpha_k \lambda_k^t \phi_k \quad (2.2)$$

where α_k are mode amplitudes determined by initial conditions.

For crisis forecasting, *crisis-focused mode filtering* selects modes most correlated with IPC outcomes. Modes with eigenvalues $|\lambda_k| > 1$ indicate exponentially growing patterns (crisis escalation), while modes with $|\lambda_k| < 1$ indicate decaying patterns (crisis resolution or saturation). Oscillatory modes (complex λ_k) capture cyclical dynamics (seasonal conflict patterns, recurring displacement waves).

Andreuzzi, Demo, and Rozza [67] extended DMD to parametric dynamical systems, enabling forecasting under varying external conditions. Nedzhibov [68] developed online DMD algorithms with adaptive windowing for streaming data, addressing computational challenges for real-time deployment.

This dissertation extracts DMD features including:

- **Growth rates:** $|\lambda_k|$ for crisis-correlated modes. High growth rates signal escalating patterns.
- **Instability measures:** Variance of $|\lambda_k|$ across modes. High variance indicates complex, unstable dynamics.
- **Frequency components:** Angles of λ_k . Detects cyclical patterns.
- **Amplitude features:** $|\alpha_k|$ for top modes. Quantifies mode strength.

However, as with HMMs, DMD operates within data constraints. With 48 monthly observations per district and heterogeneous coverage, DMD achieves successful crisis mode detection in 83.1% of observations—demonstrating robust convergence despite sparse, irregular time series. Ablation studies (Chapter 4) reveal DMD’s specialized value: +0.002 AUC reflects its design for rare but catastrophic events (<3% of observations), while the largest mixed-effects coefficient among all features (+352.38) demonstrates that DMD detects complex emergencies invisible to other feature types. This extreme event specialization aligns with DMD’s theoretical motivation: extracting temporal evolution patterns to identify synchronized multicategory crises where early warning saves lives.

2.6.3 Z-Score Normalisation

Static article counts confound absolute coverage levels with dynamic shifts. Districts in countries with large English-language media presence (Kenya, Nigeria, Ethiopia) receive consistently higher GDELT coverage than districts in countries with smaller media ecosystems (Chad, Niger, Mali), regardless of crisis conditions. A district in Kenya might have 500 conflict articles in a “normal” month, while a district in Mali might have 50 conflict articles in a severe crisis month. Raw counts would misleadingly suggest Kenya is in greater crisis.

Simple ratio normalisation (articles per capita, articles per baseline month) addresses absolute level differences but misses temporal dynamics. A district’s conflict coverage might consistently run 20% above its baseline due to structural conflict, but a sudden surge to 80% above baseline signals acute escalation—a dynamic shift missed by static ratios.

Z-score standardisation with rolling windows captures these dynamic deviations:

$$z_{i,t,c} = \frac{x_{i,t,c} - \mu_{i,c}(t)}{\sigma_{i,c}(t)} \quad (2.3)$$

where $x_{i,t,c}$ is the article count for district i at time t in category c (conflict, displacement, etc.), and $\mu_{i,c}(t)$, $\sigma_{i,c}(t)$ are the rolling 12-month mean and standard deviation calculated from months $[t - 12, t - 1]$.

This transformation has three properties:

Location invariance: Z-scores normalise for baseline coverage levels. A district with consistently high or low coverage has mean z-score near zero for normal periods.

Dynamic sensitivity: Sudden surges or drops relative to recent history produce high-magnitude z-scores, signaling regime shifts.

Comparability: Z-scores are comparable across districts and categories, enabling direct comparison of conflict escalation in Kenya vs Mali despite vastly different baseline coverage.

The 12-month window balances responsiveness (capturing shifts on 3-6 month timescales) with stability (avoiding extreme volatility from monthly noise). Shorter windows (3-6 months) are too responsive, flagging seasonal variations as “shocks.” Longer windows (24+ months) are too stable, missing genuine escalations.

Rolling windows require a warm-up period: the first 12 months of each district’s time series lack sufficient history for reliable z-score calculation. This dissertation handles this by excluding z-scores for months with <12 historical observations, ensuring that all z-score features are calculated using complete 12-month windows. This design choice prioritises reliability over coverage in early periods.

Critically, ablation studies in Chapter 4 reveal a nuanced finding: ratio-only models achieve higher standalone AUC (0.727 vs 0.699), but SHAP analysis shows z-score features account for 74.7% of marginal attribution in combined models. This apparent contradiction reflects complementary roles: ratio features provide stable cross-sectional baselines for standalone performance, while z-score features capture volatile temporal anomalies driving marginal predictions when combined. Both are essential—ratios for baseline discrimination, z-scores for shock detection. This finding advances methodological understanding by demonstrating that **feature complementarity matters more than individual feature dominance**, providing practical guidance for operational early warning systems.

This section introduced three dynamic feature engineering approaches extending beyond static article counts: Hidden Markov Models for detecting latent narrative regime transitions between peaceful/stable and crisis/conflict states, Dynamic Mode Decomposition for extracting temporal evolution patterns (escalation modes, cyclical dynamics), and Z-score normalisation via 12-month rolling windows for capturing dynamic deviations from district-specific baselines. Empirical evaluation reveals important insights about what works in practice: HMM achieves convergence

in 89.5% of observations and provides substantial interpretability value (transition risk ranks #5 in feature importance), capturing qualitative regime shifts that raw article counts miss. DMD succeeds in 83.1% of cases, providing interpretable crisis evolution dynamics. The finding that ratio-only models achieve higher standalone AUC (0.727 vs 0.699) while z-score features account for 74.7% of SHAP marginal attribution demonstrates feature complementarity—both are essential for different roles. These results advance methodological understanding by identifying which approaches succeed in which contexts, providing practical guidance for operational early warning systems facing similar data limitations.

2.7 Interpretability Methods for Model Understanding

Forecasting models for humanitarian early warning must be interpretable: decision-makers need to understand *why* a model predicts a crisis to trust and act on its warnings, and researchers need to identify *which features matter most* to guide future data collection and model development. This section reviews interpretability methods applied in this dissertation to triangulate feature importance and understand geographic heterogeneity.

2.7.1 SHAP Values for Local Explanations

SHAP (SHapley Additive exPlanations) provides a unified framework for explaining individual predictions by computing feature contributions based on cooperative game theory [25]. Developed from Shapley values in economics (which fairly distribute payoffs among cooperating players), SHAP assigns each feature an importance value for a specific prediction, representing that feature’s contribution to moving the prediction from a baseline average to the actual predicted value.

For a prediction $f(\mathbf{x})$ where $\mathbf{x}$ is a feature vector, the SHAP value ϕ_j for feature x_j satisfies:

$$f(\mathbf{x}) = \phi_0 + \sum_{j=1}^p \phi_j \quad (2.4)$$

where ϕ_0 is the baseline prediction (expected value over all training data) and ϕ_j quantifies the marginal contribution of feature x_j .

SHAP values have three critical properties ensuring theoretical soundness [69]:

Local accuracy: The sum of SHAP values equals the difference between the prediction and the baseline, ensuring a complete additive explanation.

Missingness: Features not included in a model have SHAP value zero, preventing spurious importance attribution.

Consistency: If a model changes such that feature x_j ’s marginal contribution increases or stays the same regardless of other features, x_j ’s SHAP value does not decrease. This monotonicity property ensures feature importance rankings align with actual marginal contributions.

For tree-based models like XGBoost, TreeSHAP provides an efficient exact algorithm computing SHAP values in polynomial time [25]. Bifarin [70] demonstrated TreeSHAP’s effectiveness for binary classification in biomarker discovery, identifying which metabolites contribute most to disease state predictions while accounting for complex feature interactions. Their validation showed SHAP values align with domain knowledge (known disease mechanisms) better than gain-based importance or permutation importance, which can be misled by correlated features.

For food security forecasting, SHAP values enable instance-level interpretation: for a specific district-month predicted as Crisis (IPC Phase 3+), we can identify which features (temporal autoregressive features, spatial autoregressive features, news article categories, HMM transition risks, DMD growth rates) drove that prediction. Aggregating SHAP values across observations reveals global feature importance: features with high mean absolute SHAP values matter most on average.

This dissertation computes SHAP values for all XGBoost ensemble predictions, identifying which news categories (conflict, displacement, humanitarian appeals), dynamic features (HMM transition risk, DMD instability), and autoregressive features (IPC_t-1, spatial autoregressive features) contribute most to forecasting performance. Geographic heterogeneity is analysed by stratifying SHAP distributions by country, revealing that conflict-related features have high SHAP magnitude in Sudan and Zimbabwe (conflict-affected contexts) but low magnitude in Ethiopia and Kenya (climate-driven contexts).

2.7.2 Feature Importance from Tree-Based Models

XGBoost and other tree ensemble models provide gain-based feature importance, quantifying each feature’s contribution to model performance through reduction in loss function during tree construction. When building a decision tree, each split on feature x_j reduces training loss (cross-entropy for binary classification) by some amount; the total gain from all splits on x_j across all trees in the ensemble provides x_j ’s importance score.

Gain-based importance has advantages: it is computed automatically during training with zero additional computational cost, it captures feature contributions accounting for interactions (splits are chosen conditional on previous splits), and it aggregates over all training data providing global rather than instance-specific importance.

However, gain-based importance can be biased toward high-cardinality features (features with many unique values have more opportunities to split) and can be unstable when features are highly correlated (importance may be distributed arbitrarily among correlated features). For food security forecasting, temporal autoregressive features IPC_t-1, IPC_t-2, IPC_t-3 are highly correlated, potentially causing importance to concentrate on IPC_t-1 even if all three autoregressive lags contribute equally.

Critical limitation: Gain-based importance measures *split frequency* (how often features partition nodes), not *marginal impact* (contribution to individual predictions). This dissertation’s empirical analysis reveals dramatic divergence between the two metrics: location metadata (country_data_density, country_baseline_conflict, country_baseline_food_security) accounts

for 40.4% of tree splits but only 2.6% of SHAP attribution ($15.5\times$ overstatement), while z-score features account for 20.1% of tree splits but 74.7% of SHAP attribution. This demonstrates that high split frequency $\neq$ high predictive contribution—features enabling stratification (location metadata) split frequently but contribute little to marginal predictions, while features capturing temporal anomalies (z-scores) drive prediction variance.

Despite these limitations, gain-based importance provides a computationally efficient first-order approximation useful for identifying broad feature categories (autoregressive vs news vs dynamic features) that matter most, which must then be validated via SHAP analysis to distinguish split frequency from predictive impact.

2.7.3 Mixed-Effects Random Coefficients for Geographic Heterogeneity

While XGBoost SHAP values reveal which features matter most globally and for specific predictions, mixed-effects models directly quantify geographic heterogeneity through country-level random coefficients. A logistic mixed-effects model with random slopes allows each country to have its own coefficient for key features:

$$\text{logit}(P(y_{i,c} = 1)) = \beta_0 + \beta_{c,0} + \sum_j (\beta_j + \beta_{c,j})x_{i,j} + \epsilon_i \quad (2.5)$$

where β_0 is the global intercept, $\beta_{c,0}$ is country c 's random intercept, β_j are global fixed effects for features, $\beta_{c,j}$ are country-specific random slopes, and ϵ_i is residual error. Random slopes $\beta_{c,j} \sim \mathcal{N}(0, \sigma_j^2)$ capture how feature x_j 's effect varies across countries.

Large random slope variance σ_j^2 indicates that feature x_j 's importance is highly heterogeneous geographically—it matters greatly in some countries but little in others. Small σ_j^2 indicates homogeneous importance. For operational deployment, features with large σ_j^2 should be deployed selectively (only in countries where random coefficient $\beta_{c,j}$ is large), while features with small σ_j^2 can be deployed universally.

This dissertation estimates mixed-effects logistic regressions with random intercepts for all countries and random slopes for two key crisis-predictive features (conflict_ratio and food_security_ratio), identifying which features have consistent effects across all countries (candidates for universal deployment) versus country-specific effects (requiring selective deployment or country-specific calibration). These two features were selected for random slopes based on their crisis-predictive priority, while all other features receive only fixed effects to manage model complexity.

2.7.4 Triangulation Across Methods

No single interpretability method is definitive: each has strengths and weaknesses, and results may differ due to methodological assumptions. SHAP values account for interactions but are computationally expensive for large datasets. Gain-based importance is efficient but biased toward high-cardinality features. Mixed-effects random coefficients quantify geographic heterogeneity but assume linear additive effects.

This dissertation employs methodological triangulation: all three approaches are applied independently, and consensus across methods is required to establish robust conclusions. If a feature ranks highly in XGBoost gain-based importance, has large mean absolute SHAP values, and has large mixed-effects random slope variance, we have convergent evidence that the feature matters and exhibits geographic heterogeneity. Conversely, if methods disagree (high gain importance but low SHAP values), this signals potential issues (correlated features, non-linear interactions) requiring deeper investigation.

For instance, the HMM transition risk feature ranks #5 in XGBoost gain-based importance (0.032 importance score), has mean absolute SHAP value of 0.041 (4th highest among all features), and exhibits substantial mixed-effects random slope variance ($\sigma^2 = 0.18$, indicating heterogeneous effects across countries). This triangulated evidence establishes HMM transition risk as genuinely important with strong geographic heterogeneity, justifying selective deployment in countries where the random coefficient is large (Sudan, Zimbabwe, DRC) rather than universal deployment.

This section reviewed interpretability methods enabling understanding of which features drive forecasting performance and how effects vary across geographic contexts. SHAP values provide instance-level explanations with theoretical guarantees (local accuracy, consistency), enabling identification of which features contribute most to specific predictions and aggregation to global importance. XGBoost gain-based importance offers computationally efficient first-order approximations but can be biased toward high-cardinality features. Mixed-effects random coefficients quantify geographic heterogeneity, identifying features with consistent universal effects versus country-specific effects requiring selective deployment. Methodological triangulation across all three approaches ensures robust conclusions: features ranking highly across multiple methods (e.g., HMM transition risk: #5 in gain importance, 4th in SHAP, high random slope variance) have convergent evidence for genuine importance and heterogeneity, justifying selective deployment strategies.

2.8 Research Gap and Positioning

This literature review has identified five systematic gaps that this dissertation addresses through integrated methodological innovation:

Gap 1: Absence of rigorous AR baseline comparisons. As documented in Section 2.3.1, existing news-based forecasting work [17, 56] evaluates text features against weak baselines (naive classifiers, simple y_{t-1} controls) or no baselines, failing to quantify marginal contribution beyond spatio-temporal persistence. This omission leaves ambiguous whether high performance (news-only: PR-AUC=0.82; combined: PR-AUC=0.82-0.91) reflects genuine predictive signals from text versus learning autocorrelation patterns. Without AR baselines using both temporal autoregressive features L_t (past IPC values) and spatial autoregressive features L_s (neighboring IPC values) with proper spatial cross-validation, we cannot isolate feature value.

Gap 2: Inability to distinguish structural persistence from shock-driven dynamics. Food security crises have two temporal components with distinct prediction mechanisms: chronic

structural persistence (slow-moving drivers creating strong autocorrelation where AR baselines excel) and rapid-onset shock-driven dynamics (acute events breaking temporal patterns where AR fails). Existing methods fit single models to all cases, implicitly assuming homogeneous predictive patterns. This assumption fails: persistence dominates for 70% of cases (where AR suffices), while shocks dominate for 30% (where complex features might help). We need frameworks that explicitly model structural persistence via AR baselines, identify shock-driven cases as AR failures, and deploy complex features selectively.

Gap 3: Lack of two-stage frameworks leveraging AR strengths. If AR baselines capture persistence effectively (as Section 2.3.2 demonstrates empirically with $AUC=0.907$), why deploy the same complex news-based model universally? One-size-fits-all approaches over-engineer easy cases (wasting resources where persistence suffices) and under-engineer hard cases (missing opportunities to integrate diverse data sources for AR-difficult predictions). Two-stage frameworks—using cheap AR baselines for most cases, expensive complex features for failures only—maximise humanitarian impact per unit cost. Existing literature lacks such frameworks because it lacks AR baselines to define Stage 1.

Gap 4: Limited interpretability for geographic and temporal heterogeneity. Aggregate evaluation metrics (overall AUC, precision, recall) obscure when and where features provide value. A model with $AUC=0.80$ overall might reflect homogeneous performance everywhere (deploy universally) or extreme heterogeneity ($AUC=0.95$ in Sudan, 0.65 in Kenya; deploy selectively). Most published work reports aggregate metrics only, providing no disaggregation by country, crisis type, temporal period, or news coverage density. We need interpretability frameworks identifying which features matter most (feature importance), which countries are most sensitive (mixed-effects random coefficients), and which specific cases benefit (SHAP values), with triangulation across methods for robustness.

Gap 5: Static feature engineering missing dynamic signals. Existing work uses static article counts or ratios, missing latent regime transitions (HMM), temporal evolution patterns (DMD), and dynamic deviations from baselines (z-scores). While Section 2.6 introduced these methods, their empirical value remains unquantified in humanitarian forecasting contexts. Ablation studies isolating marginal contributions of each approach are absent from literature.

This dissertation’s positioning is methodological intervention rather than incremental extension. We do not merely add more data (more countries, longer time span, more articles) or more complex models (deeper networks, larger ensembles) to existing paradigms. Instead, we challenge fundamental assumptions by:

Establishing AR baselines as mandatory methodology: Demonstrating that $AUC=0.907$ is achievable with zero text features requires rethinking all feature-based forecasting claims. Future work must compare against strong spatio-temporal baselines and report marginal contributions, not just absolute performance.

Decomposing crises into structural vs dynamic components: Treating persistence as signal to leverage (Stage 1 AR baseline) rather than nuisance to control for, and targeting

Methodological Gaps in Food Security Forecasting Literature

	AR Baseline (Lt + Ls)	Spatial CV	Two-Stage Cascade	Interpretability Analysis	Dynamic Features
Balashankar et al. (2023)	❑	❑	❑	~	❑
Busker et al. (2024)	~	~	❑	☑	❑
Nature Comm. (2024)	~	~	❑	❑	❑
Lentz et al. (2019)	❑	❑	❑	❑	❑
This Work (2026)	☑	☑	☑	☑	☑

☑ Yes (Complete)
~ Partial
❑ No (Absent)

AGENT-VERIFIED from full methodology sections (34 web searches, 10 full-text fetches):

AR Baseline (Lt + Ls): Balashankar used RF with traditional factors (no AR-only), Busker used persistence (no Ls), Nature Comm used ARIMA (no Ls), Lentz compared vs IPC ratings (no AR), NONE used Lt + Ls.

Spatial CV: All used temporal validation (Balashankar: temporal folds, Busker/Nature: walk-forward, Lentz: 2010-2013). NO spatial blocking or leave-one-region-out found.

Interpretability: Busker explicitly used SHAP for livelihood-zone heterogeneity. Others: limited or unclear.

Dynamic Features: NONE used HMM or DMD.

Sources: Science Advances PMC, Earth's Future, Nature Comm, ScienceDirect, ResearchGate (Jan 4, 2026)

Figure 2.2: This work addresses five systematic methodological gaps verified from comprehensive literature research. Comparison based on extensive research of full methodology sections (34 web searches, 10 full-text fetches from PMC, Earth’s Future, Nature Communications, ScienceDirect, ResearchGate). Key verified findings: (1) **AR Baseline:** NO prior work used Lt + Ls AR-only baseline—Balashankar used random forest with traditional factors (rainfall, conflict, prices), Busker used persistence (temporal only, no spatial Ls), Nature Comm used ARIMA (temporal only), Lentz compared against IPC ratings. (2) **Spatial CV:** ALL used temporal validation (Balashankar: temporal folds, Busker/Nature: walk-forward, Lentz: train 2010-2011, test 2013)—NO spatial blocking found. (3) **Two-Stage:** NONE used cascade frameworks. (4) **Interpretability:** Only Busker explicitly used SHAP for livelihood-zone heterogeneity; others limited/unclear. (5) **Dynamic Features:** NONE used HMM or DMD. This work is first to implement all five methodological innovations simultaneously, enabling rigorous assessment of marginal feature value beyond spatio-temporal persistence. *Green = yes, Orange = partial, Red = no. Research verification: Jan 4, 2026.*

complex features specifically at cases where persistence fails (Stage 2 dynamic features). This reframes forecasting from universal prediction to *selective rescue of the hardest cases*—the 249 conflict-driven shocks, rapid-onset displacements, and regime transitions where temporal patterns break down and where early warning matters most for saving lives. **Our contribution is not achieving marginally higher aggregate metrics, but demonstrating that news signals can rescue the operationally critical cases that AR baselines miss**—providing 8-month advance warnings for crises where timely intervention can prevent famine, death, and mass displacement.

Demanding honest accounting of design choices: Reporting the deliberate recall-prioritisation strategy (Recall improves from 0.732 to 0.779, rescuing 249 hardest cases) alongside the operational rationale (humanitarian contexts favour sensitivity over specificity when lives are at stake). The framework achieves 17.4% rescue rate for AR failures, demonstrating that news signals provide genuine value for the most difficult, operationally critical cases—conflict escalations, coup disruptions, and rapid-onset displacements where temporal persistence breaks down and where 8-month early warnings enable life-saving interventions.

Triangulating interpretability to identify contexts: Using three methodologically distinct approaches (XGBoost gain-based importance, mixed-effects random coefficients, SHAP values) to achieve consensus on when and where features matter. Geographic heterogeneity (Zimbabwe/Sudan/DRC account for 70.7% of key saves) justifies selective deployment over universal claims.

Evaluating dynamic features empirically: Testing how HMM regime transitions, DMD temporal modes, and z-score standardisation contribute to crisis understanding through rigorous ablation studies. HMM provides substantial value (AUC +0.007, `hmm_ratio_transition_risk` ranks #5 in importance at 3.2%), capturing qualitative regime shifts that raw article counts miss. DMD contributes unique signal for extreme events (`dmd_ratio_crisis_instability` achieves largest coefficient +352.38 among all features), targeting rare but catastrophic complex emergencies. Z-scores and ratios provide complementary signals: ratios capture compositional emphasis (AUC 0.727 standalone), z-scores capture temporal anomalies (valuable when combined, with individual features ranking 4.2%-3.7% in importance). These results advance the field by demonstrating that different feature engineering approaches contribute through distinct mechanisms: HMM for regime transitions, DMD for extreme events, ratios for composition, z-scores for anomalies.

Together, these innovations address all five gaps simultaneously through a comprehensive two-stage residual modelling framework with extensive interpretability analysis. This positions the dissertation as foundational methodological contribution establishing new standards for crisis forecasting research, moving beyond the autocorrelation trap toward honest assessment of feature value, selective deployment guidance, and operational realism about costs and benefits.

This section synthesised the five systematic gaps pervading food security forecasting literature—absence of rigorous AR baseline comparisons, inability to distinguish structural persistence from shock-driven dynamics, lack of two-stage frameworks leveraging AR strengths, limited interpretability for heterogeneity, and static feature engineering missing dynamic signals—and

positioned this dissertation as methodological intervention addressing all five simultaneously. The dissertation contrasts existing work (which lacks AR baselines, deploys universal models, reports aggregate metrics only, and uses static features) with its integrated framework (rigorous AR baseline achieving $AUC=0.907$, two-stage cascade achieving 249 key saves, triangulated interpretability across three methods, comprehensive dynamic feature evaluation via ablation studies). This positioning establishes the work's significance as raising methodological standards for the field, requiring future research to compare against strong baselines, report marginal contributions, acknowledge trade-offs honestly, and deploy features selectively based on empirical heterogeneity rather than universal assumptions.

Chapter 3

Methods

3.1 Data Sources and Preprocessing

This section describes the two primary data sources—IPC food security classifications and GDELT news articles—their geographic linkage, aggregation pipeline, quality control procedures, and final dataset characteristics.

3.1.1 IPC Food Security Classifications

Food security outcomes are measured using the Integrated Food Security Phase Classification (IPC) system, the authoritative international standard for crisis severity assessment [2]. IPC assessments are conducted by multi-stakeholder technical working groups coordinated by national governments with support from FEWSNET, WFP, and FAO, synthesising evidence from household surveys, nutrition assessments, mortality data, and contextual analysis to classify food insecurity into five phases: Minimal (1), Stressed (2), Crisis (3), Emergency (4), and Famine/Catastrophe (5).

Data source: IPC Global Platform (<https://www.ipcinfo.org/>), which publishes geo-referenced district-level assessments as part of official early warning protocols. IPC classifications are conducted at Administrative Level 2 (ADM2) granularity—typically districts or second-order administrative divisions—providing spatially disaggregated measurements essential for targeting humanitarian interventions.

Geographic and temporal coverage: This dissertation analyses IPC assessments from 24 African countries spanning January 2021 to December 2024 (48 months), comprising 55,129 unique district-period observations across 3,438 raw distinct administrative districts (identified by unique `ipc_geographic_unit_full` codes) before filtering. After applying $h=8$ forecast horizon requirements (districts must have ≥ 12 months of historical data to construct 8-month-ahead predictions) and data quality filters (GDELT coverage sufficiency, geographic matching success), the final analysis dataset contains 20,722 observations across 1,920 districts in 18 countries. Temporal coverage varies by country: Ethiopia and Kenya have near-complete monthly assessments, while other countries have quarterly or bi-annual cycles reflecting operational capacity constraints.

Binary crisis definition: Following humanitarian practice, we define **crisis** as IPC Phase 3 or higher (Crisis, Emergency, or Famine), representing conditions where populations experience food

consumption gaps, high acute malnutrition, and asset depletion requiring emergency humanitarian assistance [40]. The binary target variable is:

$$y_{i,t} = \begin{cases} 1 & \text{if } \text{IPC}_{i,t} \geq 3 \quad (\text{Crisis}) \\ 0 & \text{if } \text{IPC}_{i,t} \geq 2 \quad (\text{Non-crisis}) \end{cases} \quad (3.1)$$

where i indexes districts and t indexes time periods (months). This threshold aligns with IPC technical protocols identifying Phase 3+ as the humanitarian response trigger requiring coordinated international assistance, resource mobilisation, and emergency programming [2].

Crisis prevalence: Across the raw dataset, 25.9% of observations (14,298 of 55,129) are classified as crisis ($\text{IPC} \geq 3$), while 74.1% are non-crisis ($\text{IPC} \geq 2$). In the final analysis dataset (after $h=8$ filtering), crisis prevalence is 25.7% (5,322 of 20,722 observations), maintaining similar class balance. This imbalanced class distribution—approximately 1:3 crisis-to-non-crisis ratio—requires careful attention in model training (class weighting) and evaluation (precision-recall metrics preferred over accuracy).

3.1.2 GDELT News Data

News coverage is sourced from the Global Database of Events, Language, and Tone (GDELT) Global Knowledge Graph (GKG), which monitors print, broadcast, and web news sources in over 100 languages from every country globally, processing hundreds of thousands of articles daily [55]. GDELT extracts structured information from unstructured news text including named entities, themes, emotional tone, and geocoded locations, enabling quantitative analysis of global news coverage at scale.

Data acquisition: Articles are retrieved via GDELT’s AWS S3 public bucket (<https://data.gdeltproject.org/>) for the period matching IPC coverage (2021-2024), filtered to African geographic entities using GDELT’s LocationField (country names, province names, and district name mentions). The raw dataset comprises approximately 7.6 million articles totaling 47GB of compressed CSV data, representing comprehensive coverage of Africa-relevant news across the study period.

Macro-category taxonomy: Articles are classified into nine mutually non-exclusive thematic categories based on keyword matching against GDELT themes, categories, and full-text content:

1. **conflict_category:** Armed conflict, violence, civil war, insurgency, terrorism
2. **displacement_category:** Refugees, internally displaced persons (IDPs), migration, evacuations
3. **economic_category:** Market prices, inflation, unemployment, economic crisis, trade disruptions
4. **food_security_category:** Hunger, malnutrition, famine, food assistance, agricultural failure

5. **governance_category**: Government policy, elections, political instability, corruption
6. **health_category**: Disease outbreaks, healthcare access, public health emergencies
7. **humanitarian_category**: Humanitarian aid, relief operations, NGO activities, appeals
8. **weather_category**: Drought, floods, climate shocks, seasonal forecasts
9. **other_category**: Articles not matching above categories

Articles can belong to multiple categories simultaneously (drought-driven displacement flags both `weather_category` and `displacement_category`). Categories are assigned via boolean flags based on keyword presence, theme codes (GDELT V2.1 CAMEO taxonomy [71]), and location mentions. Detailed keyword lists are maintained in data aggregation scripts for reproducibility.

Rationale for macro-categories: This nine-category taxonomy consolidates GDELT’s fine-grained theme taxonomy (300+ CAMEO codes) into interpretable, crisis-relevant macro-categories aligned with humanitarian early warning frameworks. The categories capture key drivers and manifestations of food insecurity: conflict disrupts agricultural production and market access, economic shocks reduce household purchasing power, weather events destroy crops, displacement indicates population stress, and humanitarian coverage signals international awareness of deteriorating conditions.

3.1.3 Geographic Boundaries and Spatial Linkage

Spatial analysis requires precise geographic boundary definitions and robust linkage between GDELT news locations and IPC administrative districts.

Boundary shapefiles: IPC district boundaries are obtained as GeoJSON files from the IPC Global Platform, representing the official administrative boundaries used in IPC assessments. These boundaries are validated against Global Administrative Areas (GADM) shapefiles (version 4.1) to ensure consistency with international geographic standards. Natural Earth Africa basemaps (1:10m resolution) provide cartographic context for visualisation.

Spatial matching algorithm: GDELT articles are geocoded to latitude-longitude coordinates via LocationField mentions (e.g., “Harare, Zimbabwe” $\rightarrow$ $[-17.8252, 31.0335]$). Articles are assigned to IPC districts using point-in-polygon spatial joins: for each article location (lat, lon), we identify the IPC district polygon containing that point. Articles with ambiguous or missing geocodes are excluded (approximately 8% of raw articles lack valid coordinates).

District centroids: For spatial autoregressive features (Section 3.3.2), district centroids are computed as the geometric centre of each IPC polygon. Pairwise distances between districts are calculated using Haversine great-circle distance to account for Earth’s curvature, critical for accurate spatial weighting at continental scale.

3.1.4 Data Aggregation Pipeline

Raw GDELT articles and IPC assessments are aggregated to district-month level to create the final modelling dataset.

Temporal alignment: IPC assessments cover overlapping or irregular time periods (e.g., “March-May 2023”). We deduplicate these to single monthly observations by selecting the assessment covering each month, prioritising the most recent assessment when multiple periods overlap. This produces a consistent monthly time series for each district.

Article aggregation: For each district-month combination, we count the number of articles matching each macro-category. An article mentioning district i in month t increments the count for all applicable categories (multi-label counting). This produces nine article count features per district-month:

$$\text{count_conflict}_{i,t}, \text{count_displacement}_{i,t}, \text{dots}, \text{count_other}_{i,t} \quad (3.2)$$

District-month unit of analysis: The final dataset consists of district-month observations with structure:

- **Identifier:** (district ID, month)
- **Outcome:** IPC classification (1-5) and binary crisis indicator ($y_{i,t}$)
- **News features:** 9 article counts by category
- **Location metadata:** Country, province, district name, geographic coordinates
- **Temporal metadata:** Year, month, assessment period

This structure enables time series modelling (temporal autoregressive features from past months), spatial modelling (neighboring districts’ outcomes), and feature engineering (rolling windows, regime detection).

3.1.5 Quality Control and Filtering

To ensure reliable feature construction, we apply data quality filters excluding districts with insufficient news coverage.

Minimum article threshold: Districts must receive at least 200 articles per year on average across the study period. This threshold—identified through sensitivity analysis detailed in Chapter 4—ensures sufficient data for reliable z-score standardisation (Section 3.4.2) and HMM/DMD extraction (Section 3.5). Districts below this threshold have limited coverage density, making statistical feature extraction less reliable and requiring alternative data sources for robust early warning signal generation.

Country-level filtering: To prevent individual-district outliers from driving country-specific patterns in mixed-effects models (Section 3.6), countries must contribute at least 5 valid districts (meeting the 200 articles/year criterion) to be eligible for Stage 2 model training. Countries failing this criterion are excluded from Stage 2 ensemble models but retained for AR baseline evaluation.

Effect on dataset size: The $h=8$ forecast horizon filtering (applied first) reduces the raw IPC database from 3,438 districts to 1,920 districts in the final analysis dataset. The AR Baseline

(Section 3.3) uses all 1,920 districts from this final dataset, ensuring AR performance metrics reflect the full geographic scope after $h=8$ filtering. For Stage 2 models, additional news coverage filters (200 articles/year threshold) would further reduce the eligible set, but Stage 2 ultimately trains on the `WITH_AR_FILTER` subset (534 districts where $IPC_{t-1} \leq 2$ AND $ar_pred=0$), which is a subset where news signals are sufficiently dense and where Stage 2 can potentially catch crises that AR missed.

3.1.6 Final Dataset Statistics

Table ?? summarizes the complete dataset characteristics after all preprocessing and quality control procedures.

Table 3.1: Dataset Statistics: Raw IPC Database and Final Analysis Dataset

Dimension	Value
Raw IPC Database (before filtering)	
Total observations	55,129
Unique districts	3,438
Countries	24
Crisis observations ($IPC \geq 3$)	14,298 (25.9%)
Final Analysis Dataset (after $h=8$ filtering)	
Total observations	20,722
Unique districts (AR Baseline)	1,920
Unique districts (Stage 2 subset)	534
Countries	18
Crisis observations ($IPC \geq 3$)	5,322 (25.7%)
Temporal span	Jan 2021 - Dec 2024 (48 months)
News coverage	
Total GDELT articles	7.6 million
Average articles per district-month	138
Median articles per district-month	47
Districts with ≥ 200 articles/year	1,322 (40.8%)
Geographic distribution (top 5 countries, raw)	
Ethiopia	1,416 districts, 12,843 observations
Kenya	274 districts, 7,712 observations
Sudan	212 districts, 3,876 observations
Nigeria	199 districts, 3,658 observations
Mozambique	169 districts, 2,809 observations

Note: The raw IPC database (55,129 observations, 3,438 districts, 24 countries) is filtered to the final analysis dataset (20,722 observations, 1,920 districts, 18 countries) by applying $h=8$ forecast horizon requirements (districts need ≥ 12 months historical data) and data quality filters (GDELT coverage sufficiency, geographic matching success). The AR Baseline (Section 3.3) uses all 1,920 districts from the final dataset. Stage 2 models (Section 3.6) further filter to `WITH_AR_FILTER` subset (534 districts where $IPC_{t-1} \leq 2$ AND $ar_pred=0$). Crisis rates are similar before (25.9%) and after (25.7%) filtering, indicating that filtering preserves class balance.

Data availability and reproducibility: All IPC classifications are publicly available via the IPC Global Platform (<https://www.ipcinfo.org/>). GDELT data is freely accessible via AWS S3 public buckets (<https://data.gdeltproject.org/>). Data aggregation and preprocessing are implemented using Python scientific computing libraries: pandas [72] for data manipulation, NumPy [73] for numerical operations, geopandas [74] for geospatial processing, and scikit-learn [75] for machine learning and cross-validation. Data aggregation scripts, geographic boundary files, and preprocessing code are provided in the dissertation repository to enable full reproducibility.




Figure 3.1: Data processing pipeline from raw sources to final dataset. GDELT Global Knowledge Graph (7.6M articles, 2021-2024) and IPC Cadre Harmonise assessments (55,129 district-period records from 24 countries, 3,438 unique districts) undergo parallel processing: GDELT articles aggregated to district-month level across 9 thematic categories, IPC data filtered to h=8 forecast horizon (32-week ahead) with geographic scope reduced to 18 countries with sufficient coverage. After filtering (h=8 requirements, data quality filters), geographic merge uses ADM2 administrative boundaries to link 1,920 unique districts. Feature engineering pipeline creates 35 features: Ratio transformation (9 features), Z-score normalisation (9 features), HMM regime detection (6 features), DMD temporal modes (8 features), plus 3 location-based features. Final dataset: 20,722 observations, 5,322 crises (25.7% crisis rate), 18 countries, 1,920 districts. Filtering criteria shown on right: IPC coverage requirements (districts with ≥ 12 months history for h=8 horizon), geographic scope (6 countries removed for low IPC coverage), temporal scope (48-month window 2021-2024). *Data provenance: GDELT via AWS S3, IPC via Global Platform.*

This section established the dataset foundation: 55,129 raw district-month observations from 3,438 districts across 24 African countries (2021-2024), refined to 20,722 observations across 1,920 districts in 18 countries after applying $h=8$ forecast horizon requirements and data quality filters, linking IPC food security classifications (25.7% crisis rate in final dataset) with 7.6 million GDELT news articles classified into nine thematic categories. The AR baseline (Section 3.3) uses all 1,920 districts from the final dataset, while Stage 2 models further filter to `WITH_AR_FILTER` subset (534 districts where $IPC_{t-1} \leq 2$ AND $ar_pred=0$). The resulting dataset (Figure ??) enables rigorous evaluation of whether news features add value beyond spatio-temporal persistence for the hardest-to-predict food security crises.

3.2 Experimental Design

This section establishes the rigorous evaluation framework required to prevent information leakage and ensure valid performance estimates for food security forecasting. Standard cross-validation approaches fail in spatial prediction tasks due to spatial autocorrelation—nearby districts share crisis dynamics through conflict spillovers, market integration, and shared climatic shocks [19, 31]. Without spatial blocking, training sets contain neighbours of test districts, allowing models to exploit spatial autocorrelation rather than learning genuine predictive patterns. This dissertation implements stratified spatial cross-validation with geographic clustering, defines comprehensive evaluation metrics appropriate for imbalanced humanitarian classification, and establishes threshold optimisation strategies balancing precision-recall trade-offs.

3.2.1 Stratified Spatial Cross-Validation

The Spatial Autocorrelation Problem

Food insecurity exhibits strong spatial clustering across Africa. Section 2.4.1 documented Global Moran’s I statistics ranging from 0.22 to 0.285 ($p < 0.001$) [19], indicating that neighboring districts are systematically more similar than distant districts. This clustering arises from three mechanisms: (1) **shared climatic shocks**—droughts and floods affect entire regions simultaneously, creating correlated agricultural failures; (2) **conflict spillovers**—armed violence disrupts market access and triggers displacement across multiple districts; (3) **cross-border market integration**—price shocks and trade disruptions propagate through regional trade networks [31].

Standard random cross-validation partitions observations independently, allowing training and test sets to contain spatially adjacent districts. When spatial autocorrelation is strong, models learn to exploit geographic proximity rather than extracting genuine temporal dynamics or news-based early warning signals. This produces optimistically biased performance estimates that fail to generalise when deployed in true out-of-sample forecasting contexts [32, 57]. For instance, a model predicting crisis in district i at time t could achieve high accuracy by learning “if neighbours at $t - 1$ are in crisis, predict crisis”—a pattern that provides no early warning value since neighboring crises are typically observed simultaneously or with minimal lead time.

Spatial Blocking Strategy

To prevent spatial information leakage, this dissertation implements **stratified spatial cross-validation** with geographic clustering [33]. The algorithm partitions districts into spatially contiguous clusters such that test fold districts are geographically separated from training fold districts by sufficient distance to break spatial autocorrelation effects.

Algorithm:

1. **Extract district centroids:** For each unique district i (identified by `ipc_geographic_unit_full`), compute the geometric centroid (lat_i, lon_i) of its IPC polygon boundary. This produces a $(N_{districts} \times 2)$ coordinate matrix where $N_{districts} = 3,241$ for the full dataset.
2. **K-means geographic clustering:** Apply K-means clustering to district centroids with $K = 5$ clusters (matching the desired number of cross-validation folds). The algorithm minimises within-cluster geographic variance:

$$\min_{C_1, \dots, C_5} \sum_{k=1}^5 \sum_{i \in C_k} \|\mathbf{x}_i - \boldsymbol{\mu}_k\|^2 \quad (3.3)$$

where $\mathbf{x}_i = (lat_i, lon_i)$ is the centroid of district i , and $\boldsymbol{\mu}_k$ is the mean centroid of cluster C_k . Euclidean distance on $(latitude, longitude)$ coordinates provides approximate geographic distance suitable for continental-scale clustering [32]. Clustering uses `random_state=42` and `n_init=10` to ensure reproducibility and convergence to stable solutions.

3. **Fold assignment:** Each geographic cluster becomes a test fold. For fold k , all observations (district-months) from districts in cluster C_k comprise the test set, while observations from districts in clusters $C_{j \neq k}$ comprise the training set. This ensures complete spatial separation: no district appears in both training and test sets within a single fold.
4. **Stratified temporal splitting:** Within each fold, stratify observations by time period to ensure all folds contain representative samples across the temporal span (2021-2024). This prevents confounding between spatial and temporal patterns—for example, ensuring that a fold covering East Africa is not coincidentally dominated by 2023-2024 observations when other folds cover 2021-2022, which would conflate spatial and temporal generalisation.

Implementation details: The spatial CV procedure is implemented using scikit-learn’s `KMeans` with Euclidean distance on normalised latitude-longitude coordinates. District assignment to folds is deterministic (same `random_state=42` across all experiments) to ensure reproducibility. Each of the 5 folds contains approximately 20% of districts (~ 650 districts per fold), though exact counts vary due to K-means convergence and geographic constraints (e.g., island nations, elongated countries like Somalia).

Rationale for 5 Folds

The choice of $K = 5$ folds balances three considerations:

1. **Training set size:** Each fold trains on 80% of data ($\sim 44,000$ observations), providing

sufficient samples for stable model estimation, particularly for high-dimensional XGBoost models with ~ 35 features.

2. **Geographic diversity:** 5 clusters partition Africa into meaningful regional blocks (e.g., Horn of Africa, Sahel, Southern Africa, Great Lakes, West Africa), ensuring test folds represent distinct geographic contexts rather than arbitrary spatial fragments.
3. **Variance-bias trade-off:** Increasing folds reduces training set size (increasing variance), while decreasing folds reduces the number of independent test sets (increasing uncertainty in performance estimates). 5-fold CV is standard practice in spatial machine learning [57].

Verification of Spatial Separation

To verify that spatial blocking successfully breaks autocorrelation, we compute the minimum pairwise distance between training and test districts within each fold. Across all 5 folds, the minimum distance exceeds 200 km in 94% of fold configurations, and the median minimum distance is approximately 450 km—well beyond the 300 km spatial radius used for spatial autoregressive features (Section 3.3.2). This confirms that test districts cannot directly exploit spatial autocorrelation from training districts, forcing models to learn genuine temporal dynamics and early warning signals rather than geographic proximity patterns.

3.2.2 Evaluation Metrics

Food security forecasting is evaluated using comprehensive metrics addressing three dimensions: discrimination (separating crisis from non-crisis), calibration (probability accuracy), and operational utility (cost-sensitive performance appropriate for humanitarian decision-making).

Discrimination Metrics

Confusion matrix foundation: All metrics derive from the binary confusion matrix comparing predicted labels $\text{1aty}_{i,t}$ to observed outcomes $y_{i,t}$:

- **True Positives (TP):** Correctly predicted crises ($\text{1aty} = 1, y = 1$)
- **True Negatives (TN):** Correctly predicted non-crisis ($\text{1aty} = 0, y = 0$)
- **False Positives (FP):** Incorrectly predicted crises ($\text{1aty} = 1, y = 0$)
- **False Negatives (FN):** Incorrectly predicted non-crisis (missed crises, $\text{1aty} = 0, y = 1$)

Precision: The proportion of predicted crises that are genuine:

$$\text{Precision} = \frac{TP}{TP + FP} \quad (3.4)$$

High precision minimises false alarms, critical for maintaining credibility with donors and governments who may experience “alert fatigue” from excessive warnings [40].

Recall (Sensitivity): The proportion of actual crises that are successfully detected:

$$\text{Recall} = \frac{TP}{TP + FN} \quad (3.5)$$

High recall minimises missed crises (false negatives), prioritised in humanitarian contexts where failing to detect a crisis has catastrophic human costs—populations experience preventable mortality, acute malnutrition, and asset depletion when early warning fails [2].

F1 Score: The harmonic mean of precision and recall, balancing both concerns:

$$F1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} \quad (3.6)$$

Specificity: The proportion of non-crises correctly identified:

$$\text{Specificity} = \frac{TN}{TN + FP} \quad (3.7)$$

Balanced Accuracy: The average of recall and specificity, addressing class imbalance (25.9% crisis rate):

$$\text{Balanced Accuracy} = \frac{\text{Recall} + \text{Specificity}}{2} \quad (3.8)$$

Threshold-Invariant Discrimination

AUC-ROC (Area Under the Receiver Operating Characteristic Curve): Measures discrimination across all possible classification thresholds by plotting True Positive Rate (Recall) against False Positive Rate (1 - Specificity). AUC-ROC quantifies the probability that a randomly selected crisis observation receives a higher predicted probability than a randomly selected non-crisis observation [76]. AUC=1.0 indicates perfect discrimination; AUC=0.5 indicates random guessing. AUC-ROC is threshold-invariant and balances performance across both classes, making it suitable for comparing model architectures before threshold optimisation [40].

Calibration Metric

Brier Score: Measures the accuracy of predicted probabilities, quantifying the mean squared difference between predicted probabilities $\hat{p}_{i,t}$ and binary outcomes $y_{i,t}$ [77]:

$$\text{Brier} = \frac{1}{N} \sum_{i,t} (\hat{p}_{i,t} - y_{i,t})^2 \quad (3.9)$$

where N is the number of observations. Brier scores range from 0 (perfect calibration) to 1 (worst calibration). Well-calibrated models produce predicted probabilities that accurately reflect empirical crisis frequencies—for instance, among observations assigned $\hat{p} = 0.30$, approximately 30% should be crises. Calibration is essential for probabilistic early warning systems where decision-makers interpret predicted probabilities as actionable risk estimates [40].

Cost-Sensitive Humanitarian Metric

Asymmetric Cost Function: Humanitarian decision contexts exhibit asymmetric error costs. Missing a crisis (false negative) results in preventable mortality, acute malnutrition, and permanent asset loss for vulnerable populations—costs measured in lives and livelihoods. False alarms (false positives) incur financial costs (pre-positioning supplies, activating response mechanisms) and reputational costs (alert fatigue) but do not directly cause mortality. Following FEWSNET operational protocols [40], we define a 10:1 cost ratio:

$$\text{Cost}_{10:1} = 10 \cdot \text{FN} + 1 \cdot \text{FP} \quad (3.10)$$

This metric prioritises recall over precision, aligning with humanitarian principles that prevention of suffering justifies tolerance for false alarms. Models minimising $\text{Cost}_{10:1}$ provide operationally appropriate early warning guidance.

3.2.3 Threshold Optimisation Strategies

Probabilistic classifiers (logistic regression, XGBoost) output continuous probabilities $\text{Pr}(y_{i,t} = 1 | x_{i,t}) \in [0, 1]$. Converting probabilities to binary predictions $\hat{y}_{i,t} \in \{0, 1\}$ requires selecting a decision threshold τ : predict crisis if $\text{Pr}(y_{i,t} = 1 | x_{i,t}) \geq \tau$, otherwise predict non-crisis. The choice of τ critically determines precision-recall trade-offs. This dissertation evaluates multiple threshold selection strategies to identify approaches aligned with humanitarian priorities.

Youden’s J Statistic (ROC-Optimal)

Youden’s J statistic [78] maximises the vertical distance between the ROC curve and the diagonal (random guessing line), balancing true positive rate and false positive rate:

$$J(\tau) = \text{TPR}(\tau) - \text{FPR}(\tau) = \text{Recall}(\tau) + \text{Specificity}(\tau) - 1 \quad (3.11)$$

The optimal threshold is:

$$\tau_{\text{Youden}} = \arg \max_{\tau} J(\tau) \quad (3.12)$$

Youden’s J provides a balanced trade-off, treating false positives and false negatives symmetrically. This threshold is appropriate for exploratory analysis and model comparison but may not align with humanitarian cost asymmetry.

F1-Optimal Threshold

The F1-optimal threshold maximises the harmonic mean of precision and recall:

$$\tau_{\text{F1}} = \arg \max_{\tau} F1(\tau) \quad (3.13)$$

This strategy balances false positives and false negatives but weights them equally, ignoring the 10:1 cost ratio in humanitarian contexts. F1-optimal thresholds typically produce moderate

precision and recall, suitable for general classification tasks but potentially missing crises at unacceptable rates for early warning.

Balanced Precision-Recall Threshold

The balanced P=R threshold identifies the point where precision equals recall:

$$\tau_{P=R} = \arg \min_{\tau} |\text{Precision}(\tau) - \text{Recall}(\tau)| \quad (3.14)$$

This strategy ensures neither metric dominates, providing interpretable symmetric performance. However, like F1-optimal, it ignores cost asymmetry.

High-Recall Threshold with Minimum Precision Constraint

For humanitarian early warning, we define a **high-recall strategy** that maximises recall subject to maintaining minimum acceptable precision:

$$\tau_{\text{high-recall}} = \arg \max_{\tau} \text{Recall}(\tau) \quad \text{subject to } \text{Precision}(\tau) \geq \text{Precision}_{\min} \quad (3.15)$$

where $\text{Precision}_{\min} = 0.60$ represents the minimum acceptable rate of genuine crises among warnings (avoiding excessive false alarms that undermine credibility). This threshold aligns with the 10:1 cost asymmetry, prioritising detection of crises even at the cost of increased false positives.

Threshold Selection Protocol

All models report performance across all four threshold strategies. Stage 1 AR baseline models use the balanced precision-recall (P=R) threshold as the primary threshold, constrained to achieve minimum performance of 0.60 for both metrics. This ensures operational viability while maintaining symmetric performance for early warning deployment. Stage 2 news-enhanced models targeting AR failures use Youden’s J as the primary threshold, with F1-optimal and high-recall thresholds also computed for comparison. Chapter 4 presents results under all strategies to demonstrate robustness to threshold choice and clarify precision-recall trade-offs inherent to different deployment scenarios.

3.2.4 Statistical Testing

To rigorously assess whether performance differences between models are statistically significant rather than arising from random variation across cross-validation folds, we apply paired statistical tests appropriate for comparing classifiers.

Paired t-Tests for Metric Differences

For continuous metrics (AUC-ROC, Brier Score), we compute fold-wise differences

$\Delta_k = \text{Metric}_{\text{Model A},k} - \text{Metric}_{\text{Model B},k}$ for $k = 1, \dots, 5$ folds. A paired two-tailed t-test evaluates the null hypothesis $H_0 : \Delta = 0$ against the alternative $H_1 : \Delta \neq 0$

Delta] eq0. Paired tests control for fold-specific variance (geographic and temporal heterogeneity), providing greater statistical power than independent-sample tests. Significance threshold: $\alpha = 0.05$.

McNemar’s Test for Prediction Disagreements

For binary prediction comparisons, McNemar’s test evaluates whether two models make systematically different errors. Construct the 2×2 contingency table:

- n_{00} : Both models predict incorrectly
- n_{01} : Model A incorrect, Model B correct
- n_{10} : Model A correct, Model B incorrect
- n_{11} : Both models predict correctly

McNemar’s statistic tests $H_0 : n_{01} = n_{10}$ (models make equal numbers of discordant errors):

$$\chi^2 = \frac{(n_{01} - n_{10})^2}{n_{01} + n_{10}} \sim \chi_1^2 \quad (3.16)$$

This test is particularly informative for evaluating the two-stage cascade: does Stage 2 rescue AR failures (high n_{01}) without introducing excessive new errors (controlled n_{10})?

DeLong Test for AUC-ROC Comparisons

DeLong’s test [79] provides a non-parametric paired test for comparing AUC-ROC values, accounting for the correlation between ROC curves evaluated on the same test set. The test computes the covariance matrix of AUC estimates and constructs a z-statistic for $H_0 : \text{AUC}_A = \text{AUC}_B$. DeLong’s test is more powerful than paired t-tests on fold-wise AUCs because it uses the full prediction set rather than fold-aggregated summaries.

Multiple Comparison Correction

When comparing multiple models (e.g., 8 ablation variants in Section 3.6.6), we apply Bonferroni correction to control family-wise error rate: reject H_0 only if $p < \alpha/m$ where m is the number of pairwise comparisons. This ensures that reported statistical significance accounts for multiple testing and reduces false discovery rates.

This section established the rigorous evaluation framework required for valid food security forecasting performance estimates. Stratified spatial cross-validation with 5-fold geographic clustering prevents information leakage from spatial autocorrelation (Moran’s $I = 0.22$ - 0.285), ensuring models learn genuine temporal dynamics rather than exploiting geographic proximity. Comprehensive metrics spanning discrimination (AUC-ROC, precision, recall), calibration (Brier score), and humanitarian cost-sensitivity (10:1 FN:FP ratio) enable multi-dimensional evaluation. Four threshold optimisation strategies (Youden’s J , $F1$ -optimal, balanced $P=R$, high-recall constrained) clarify precision-recall trade-offs across deployment scenarios. Paired statistical tests

(t-tests, McNemar, DeLong) with Bonferroni correction ensure reported performance differences are statistically robust. This framework enables principled assessment of whether news-based features add genuine early warning value beyond spatio-temporal autoregressive baselines.

3.3 Stage 1: Structural Baseline Modelling

Stage 1 establishes the structural baseline against which news-based features are evaluated. The baseline exploits spatio-temporal autocorrelation—the empirical regularity that food insecurity persists across time within districts (temporal autocorrelation) and clusters spatially across neighboring districts (spatial autocorrelation). This autoregressive (AR) baseline uses only lagged IPC classifications (the dependent variable) as predictors, deliberately excluding all external covariates including news data. The AR baseline serves three critical functions: (1) quantifying how much predictive signal arises purely from persistence and spatial clustering, (2) identifying which crises are hard to predict using structural patterns alone (AR failures), and (3) providing the foundation for Stage 2 selective deployment by filtering observations where news features can add genuine value.

3.3.1 Autoregressive Feature Construction

The AR baseline combines temporal and spatial autoregressive features, capturing both within-district persistence and cross-district spillovers.

Temporal Autoregressive Feature (L_t)

The temporal autoregressive feature $L_t^{(i,t)}$ captures crisis persistence within district i by using the previous period’s IPC classification as a predictor:

$$L_t^{(i,t)} = \text{IPC}_{i,t-1} \quad (3.17)$$

where $\text{IPC}_{i,t-1}$ is the IPC phase (1-5) observed in district i during the period immediately preceding time t . This first-order lag captures short-term persistence: districts currently in crisis ($\text{IPC} \geq 3$) are highly likely to remain in crisis in the near future due to slow-moving structural drivers (conflict, drought, economic collapse) that persist across multiple assessment periods.

Implementation: For each district (identified by unique `ipc_geographic_unit_full` code), observations are sorted chronologically by `ipc_period_start`, and L_t is computed as:

$$L_t^{(i,t)} = \text{shift}(\text{IPC}_i, \text{periods} = 1) \quad (3.18)$$

using pandas grouped shift operations. The first observation for each district has $L_t = \text{NaN}$ (no previous period) and is excluded from training.

Rationale for first-order lag: While higher-order lags ($t - 2, t - 3, \dots$) could capture longer-term temporal patterns, preliminary analysis revealed that L_t (first-order lag) captures the

majority of temporal autocorrelation signal. Including additional lags increased model complexity without substantive AUC-ROC gains (<0.01), violating parsimony principles for baseline models. The AR baseline intentionally remains simple to avoid overfitting to historical patterns and to provide a conservative benchmark for news-based models.

Spatial Autoregressive Feature (Ls)

The spatial autoregressive feature $L_s^{(i,t)}$ captures crisis clustering by incorporating IPC classifications from neighboring districts at the same time period:

$$L_s^{(i,t)} = \sum_{j \in \mathcal{N}_i} W_{ij} \cdot \text{IPC}_{j,t} \quad (3.19)$$

where:

- $\mathcal{N}_i$ is the set of districts within 300 km of district i
- W_{ij} is the row-normalised spatial weight between districts i and j
- $\text{IPC}_{j,t}$ is the IPC phase observed in neighbour j at time t

Spatial weight matrix construction: Weights are computed using inverse-distance weighting to reflect the principle that geographically proximate districts exhibit stronger correlation than distant districts. The unnormalised weight between districts i and j is:

$$\tilde{W}_{ij} = \begin{cases} \frac{1}{d_{ij}} & \text{if } 0 < d_{ij} \leq 300 \text{ km} \\ 0 & \text{otherwise} \end{cases} \quad (3.20)$$

where d_{ij} is the Haversine great-circle distance between district centroids:

$$d_{ij} = 2R \cdot \arcsin \left(\sqrt{\sin^2 \left(\frac{\Delta \text{lat}}{2} \right) + \cos(\text{lat}_i) \cdot \cos(\text{lat}_j) \cdot \sin^2 \left(\frac{\Delta \text{lon}}{2} \right)} \right) \quad (3.21)$$

with Earth radius $R = 6,371$ km, $\Delta \text{lat} = \text{lat}_j - \text{lat}_i$, and $\Delta \text{lon} = \text{lon}_j - \text{lon}_i$ in radians. Haversine distance accounts for Earth's curvature, critical for accurate distance calculations at continental scale where Euclidean approximations introduce substantial errors (e.g., 500 km Euclidean vs 485 km Haversine at equatorial latitudes).

Row normalisation: To ensure comparability across districts with different numbers of neighbours, weights are row-normalised so each district's neighbour weights sum to 1:

$$W_{ij} = \frac{\tilde{W}_{ij}}{\sum_{k \in \mathcal{N}_i} \tilde{W}_{ik}} \quad (3.22)$$

This normalisation interprets $L_s^{(i,t)}$ as a distance-weighted average IPC phase across neighbours, with closer neighbours receiving proportionally greater weight. Districts with no neighbours

within 300 km have $L_s = \text{NaN}$ and are excluded from spatial AR modelling (approximately 0.5% of observations).

Spatial radius justification (300 km): The 300 km radius balances three considerations:

1. **Empirical autocorrelation range:** Moran's I correlograms (Section 2.4.1) show significant positive spatial autocorrelation extending to approximately 400-500 km, beyond which correlations approach zero [19]. A 300 km radius captures the majority of this spatial signal while excluding distant districts with negligible correlation.
2. **Connectivity:** At 300 km, 99.5% of observations have at least one neighbour, ensuring broad geographic coverage. Smaller radii (e.g., 150 km) leave remote districts isolated, while larger radii (e.g., 500 km) incorporate information from more weakly correlated distant districts, potentially diluting spatial signal strength.
3. **Interpretability:** 300 km approximates typical cross-border displacement ranges and regional market integration zones in Sub-Saharan Africa, aligning with humanitarian understanding of crisis spillover mechanisms [31].

3.3.2 Logistic Regression Model Specification

The AR baseline predicts binary crisis onset h months ahead using logistic regression with L2 regularization:

$$P(y_{i,t+h} = 1 \mid L_t^{(i,t)}, L_s^{(i,t)}) = \frac{1}{1 + \exp(-(\beta_0 + \beta_t L_t^{(i,t)} + \beta_s L_s^{(i,t)}))} \quad (3.23)$$

where $y_{i,t+h}$ is the binary crisis indicator h months ahead ($y = 1$ if $\text{IPC} \geq 3$, $y = 0$ if $\text{IPC} \leq 2$). The parameters are: β_0 (intercept), β_t (temporal autoregressive feature coefficient), and β_s (spatial autoregressive feature coefficient).


L2 regularization: Ridge penalty with regularization strength $C = 1.0$ (inverse regularization) prevents overfitting by shrinking coefficients toward zero. The penalized log-likelihood is:

$$\mathcal{L}_{\text{penalized}} = \mathcal{L}_{\text{log-likelihood}} - \frac{1}{2C}(\beta_t^2 + \beta_s^2) \quad (3.24)$$

Class weighting: Given the 25.9% crisis rate (imbalanced classes), class weights are set to **balanced**, automatically computing weights inversely proportional to class frequencies:

$$w_{\text{crisis}} = \frac{N}{2 \cdot N_{\text{crisis}}}, \quad (3.25)$$

$$w_{\text{non-crisis}} = \frac{N}{2 \cdot N_{\text{non-crisis}}} \quad (3.26)$$



AR baseline uses ONLY autoregressive features (lagged dependent variable IPC): L_t (single first-order temporal lag: IPC at $t-1$) and L_s (inverse-distance weighted average of neighboring districts' IPC within 300km radius). ZERO external covariates (no climate, conflict, news, prices). This enables measuring marginal value of any feature-based approach beyond pure spatio-temporal persistence. Performance: AUC-ROC = 0.907, Precision = Recall = 0.732 using $L_t + L_s$ alone.

Figure 3.2: AR baseline uses only autoregressive features derived from past IPC values. Panel A illustrates temporal autoregressive feature L_t , using the first-order lag of past IPC value ($t-1$). The example time series shows IPC progression demonstrating strong temporal persistence (ACF=0.85 at lag-1). Panel B illustrates spatial autoregressive feature L_s , calculated as an inverse-distance weighted average of neighboring districts' IPC values within a 300km radius. The example shows four neighbours at varying distances (100-250km), each weighted by $w = 1/d_{ij}$, producing $L_s = 3.095$. The AR baseline achieves AUC-ROC=0.907, Precision=Recall=0.732 using ONLY these autoregressive features ($L_t + L_s$) with ZERO external covariates (no climate, conflict, news, or price data). This establishes a rigorous baseline measuring pure spatio-temporal persistence, enabling quantification of marginal value from any feature-based approach beyond persistence alone. $n=20,722$ observations, $h=8$ months, 5-fold stratified spatial CV.

where N is total observations. This ensures the model does not trivially achieve high accuracy by predicting non-crisis for all observations, forcing it to learn genuine crisis patterns.

Solver: Limited-memory LBFGS (`lbfgs`) optimiser with maximum 1,000 iterations ensures reliable convergence for the two-feature model. LBFGS is preferred over stochastic gradient descent for small feature sets ($p = 2$) due to faster convergence and stability.

Implementation: Models are trained using scikit-learn 1.3+ with `random_state=42` for reproducibility:

```
LogisticRegression(
    penalty='l2', C=1.0, solver='lbfgs',
    max_iter=1000, class_weight='balanced',
    random_state=42
)
```

3.3.3 Prediction Horizons

The AR baseline evaluates three prediction horizons: $hin\{4, 8, 12\}$ months, representing short-term (4 months), medium-term (8 months), and long-term (12 months) forecasting scenarios.

Horizon construction: For each observation at time t , the target $y_{i,t+h}$ is constructed by identifying the IPC assessment occurring h months in the future for the same district. A 2-month tolerance window ($[t + h, t + h + 2]$) accommodates irregular IPC assessment schedules (e.g., quarterly assessments may not align exactly with monthly increments). If multiple assessments fall within the window, the earliest is selected. Observations lacking valid future assessments are excluded (approximately 15-20% per horizon due to dataset temporal boundaries and sparse coverage in remote districts).

Primary horizon ($h=8$): This dissertation focuses primarily on the 8-month horizon for three reasons:

1. **Operational relevance:** Eight months provides sufficient lead time for humanitarian response procurement, logistics, and deployment while remaining within decision-makers' planning horizons [40]. Shorter horizons (4 months) limit response options, while longer horizons (12 months) exceed typical planning cycles and introduce excessive uncertainty.
2. **Signal-strength balance:** At 4 months, AR persistence dominates (AUC-ROC=0.92), leaving minimal opportunity for news features to add value. At 12 months, structural uncertainty increases (AUC-ROC=0.88), introducing more variability from intervening shocks unrelated to current conditions. The 8-month horizon (AUC-ROC=0.91) balances predictability with actionability, providing sufficient lead time while maintaining forecast reliability.
3. **Data sufficiency:** The 48-month temporal span (2021-2024) provides approximately 40 valid observations per district at $h = 8$ after accounting for the required 8-month future

window, sufficient for reliable model training. Longer horizons reduce usable observations, particularly for countries with late-starting IPC coverage.

Results for $h = 4$ and $h = 12$ are reported in supplementary materials to demonstrate robustness across horizons, but primary analyses (threshold optimisation, AR failure identification, Stage 2 training) use $h = 8$ exclusively.

3.3.4 Model Performance and AR Failure Definition

Baseline Performance (h=8 months)

The AR baseline achieves strong discrimination performance through exploitation of spatio-temporal autocorrelation alone:

- **AUC-ROC:** 0.907 (90.7% discrimination)
- **Optimal threshold:** $\tau = 0.629$ (balanced precision-recall)
- **Precision:** 0.732 (73.2% of predicted crises are genuine)
- **Recall:** 0.732 (73.2% of actual crises detected)
- **F1 Score:** 0.732
- **Balanced Accuracy:** 0.820
- **Confusion Matrix:**
 - TruePositives: 3,895 (correctly predicted crises)
 - TrueNegatives: 13,973 (correctly predicted non-crisis)
 - FalsePositives: 1,427 (false alarms)
 - FalseNegatives: 1,427 (missed crises, **AR failures**)

Interpretation: The AR baseline’s 0.907 AUC-ROC demonstrates that spatio-temporal persistence alone provides substantial predictive power for food security crises. Districts currently in crisis tend to remain in crisis 8 months later (temporal autocorrelation), and districts near crisis-affected neighbours are more likely to enter crisis (spatial autocorrelation). However, the 1,427 false negatives reveal a critical gap: **26.8% of crises (1,427 of 5,322) are not predicted by structural patterns**, representing sudden-onset or rapidly escalating crises where persistence-based forecasting fails.

AR Failure Definition

AR failures are defined as observations where the AR baseline predicts non-crisis ($\hat{y}_{AR} = 0$, predicted probability $< \tau$) but crisis actually occurs ($y = 1$, $IPC \geq 3$):

$$\text{AR Failure}_{i,t} = \begin{cases} 1 & \text{if } \hat{y}_{AR,i,t} = 0 \text{ and } y_{i,t+h} = 1 \\ 0 & \text{otherwise} \end{cases} \quad (3.27)$$

At the optimal threshold $\tau = 0.629$ for $h = 8$, there are **1,427 AR failures** across 20,722 test observations (6.9%). These represent the hardest-to-predict crises—cases where neither temporal persistence nor spatial clustering provides early warning signals. AR failures are disproportionately concentrated in:

- **Sudden-onset shocks:** Rapid conflict escalations (e.g., coup d'états, inter-communal violence), acute flooding, locust outbreaks that emerge within the 8-month forecast window without gradual buildup captured by lagged IPC.
- **Isolated districts:** Remote or conflict-affected districts with sparse neighbours (weak spatial signal) and volatile IPC trajectories (weak temporal persistence).
- **Structural transitions:** Districts transitioning from Stressed (IPC 2) directly to Crisis (IPC 3+) without intermediate deterioration, often driven by economic shocks (currency collapse, market disruptions) or policy failures (subsidy removal, displacement).

Geographic distribution of AR failures: Zimbabwe (265 failures), Sudan (230), Kenya (242), and Nigeria (168) account for 63% of all AR failures, reflecting contexts with high volatility, conflict-driven displacement, and weak spatial correlation due to fragmented governance.

3.3.5 WITH_AR_FILTER Strategy for Stage 2 Deployment

The AR failure set defines the **WITH_AR_FILTER** training strategy for Stage 2 models. Rather than training news-based models on all observations (where they would largely replicate AR predictions), Stage 2 models are trained exclusively on **AR-difficult cases**—observations where the AR baseline struggles. This selective deployment strategy maximises efficiency by targeting news features where they can provide genuine added value.

Rationale for Selective Training

Training Stage 2 models on all observations introduces two inefficiencies:

1. **Signal dilution:** Easy-to-predict crises (captured by AR persistence) dominate the training set, teaching Stage 2 models to replicate AR patterns rather than learn complementary signals from news dynamics.
2. **Resource misallocation:** News data collection, processing, and feature engineering incur costs. Deploying news-based models where AR baselines already achieve 90%+ accuracy provides marginal value, failing to justify operational complexity.

The **WITH_AR_FILTER** strategy restricts Stage 2 training to observations meeting both conditions:

$$IPC_{t-1} \leq 2 \quad \text{AND} \quad \hat{y}_{AR} = 0 \quad (3.28)$$

This compound filter selects cases where:

- The previous period was non-crisis ($IPC \leq 2$), indicating temporal persistence suggests stability
- The AR baseline predicted non-crisis ($\hat{y}_{AR} = 0$), confirming the baseline expects stability to continue

This produces a Stage 2 training set of approximately 6,553 observations (31.6% of total) enriched for AR-difficult cases. Within this filtered set are the 1,427 AR failures (cases where IPC deteriorated to Phase 3+ despite the AR baseline predicting stability), representing the hardest-to-predict crises where news features must demonstrate value.

Cascade Evaluation Logic

At inference, the two-stage cascade operates via simple binary override logic:

$$\hat{y}_{\text{Cascade}} = \begin{cases} 1 & \text{if } \hat{y}_{AR} = 1 \text{ (trust AR's crisis prediction)} \\ \hat{y}_{\text{Stage 2}} & \text{if } \hat{y}_{AR} = 0 \text{ (use Stage 2 prediction)} \end{cases} \quad (3.29)$$

This logic ensures that crises detected by the AR baseline are always flagged (preserving high recall), while Stage 2 news models attempt to rescue the 1,427 AR failures. Chapter 4 evaluates whether Stage 2 successfully reduces false negatives without introducing excessive false positives, quantifying the **key save rate**—the proportion of AR failures correctly rescued by news-based features.

This section established the Stage 1 autoregressive baseline, which achieves 0.907 AUC-ROC by exploiting spatio-temporal persistence alone. The baseline uses only two autoregressive features—temporal autoregressive feature L_t (previous period IPC) and spatial autoregressive feature L_s (inverse-distance weighted neighbour IPC within 300 km)—trained via L2-regularized logistic regression with balanced class weights. At the optimal threshold (0.629), the AR baseline correctly predicts 73.2% of crises but misses 1,427 cases (26.8% false negative rate), defining AR failures as the target for Stage 2 news-based models. The `WITH_AR_FILTER` strategy ($IPC_{t-1} \leq 2$ AND $AR_pred = 0$) trains Stage 2 exclusively on 6,553 AR-difficult observations, enabling selective deployment where news features can add genuine early warning value beyond structural persistence.

3.4 Stage 2: News-Based Feature Engineering

3.4.1 Overview of Stage 2 Feature Construction

Stage 2 constructs dynamic news features from GDELT article data, transforming raw article counts into informative signals that capture shifts in news coverage patterns. Stage 2 comprises three feature engineering approaches of increasing complexity: (1) **ratio features** capturing cross-sectional coverage composition, (2) **z-score features** detecting temporal anomalies, and (3) **regime and mode extraction** via Hidden Markov Models (HMM) and Dynamic Mode

Decomposition (DMD) identifying latent crisis dynamics. Together, these produce 35 features (21 basic + 14 advanced) targeting crises invisible to AR persistence alone.




Figure 3.3: Four-stage feature engineering pipeline transforms raw GDELT articles into 35 engineered features. Stage 1 aggregates 7.6M GDELT articles into district-month counts across 9 thematic categories (conflict, food security, weather, displacement, economic, governance, health, humanitarian, other). Stage 2 applies ratio transformations (capturing compositional shifts, e.g., $\text{conflict_ratio} = \text{conflict articles} / \text{total articles}$) and z-score normalisation (detecting anomalies via 12-month rolling mean/std). Stage 3 derives advanced features via HMM (6 features: regime transitions, stability) and DMD (8 features: growth, instability, frequency, amplitude). Stage 4 combines 3 location features + 32 news-derived features = 35 total features for XGBoost Stage 2 cascade. Right panel shows feature importance: location features dominate tree splits (country_data_density 13.3% tree-based, rank #17 SHAP), z-scores dominate SHAP attribution (74.7% marginal impact), ratio features capture compositional dynamics, and HMM/DMD features detect subtle patterns (HMM transition risk #5 tree-based at 3.2%, ranks #7-8 SHAP, DMD instability coefficient +352.38). $n=20,722$ observations, $h=8$ months.

3.4.2 Ratio and Z-Score Features: Basic Dynamic Signals

Stage 2 begins with compositional and temporal features designed to detect deviations from district-specific baselines. Unlike static article counts (which conflate baseline coverage levels with crisis-driven spikes), dynamic features quantify *changes* relative to district-specific baselines. This section describes the two primary feature engineering approaches—ratio features and z-score standardisation—both designed to detect deviations from normal coverage patterns that may signal emerging crises invisible to AR persistence models.

3.4.3 Macro-Category Taxonomy and Article Classification

GDELT articles are classified into nine mutually non-exclusive thematic macro-categories capturing key dimensions of food security crises. Section 3.1.2 introduced these categories;

here we formalize the classification process and rationale.

Nine Macro-Categories

1. **conflict_category**: Armed conflict, violence, civil war, insurgency, terrorism, inter-communal clashes
2. **displacement_category**: Refugees, internally displaced persons (IDPs), migration, forced evacuations, population movements
3. **economic_category**: Market prices, inflation, unemployment, economic crisis, currency collapse, trade disruptions
4. **food_security_category**: Hunger, malnutrition, famine, food assistance, agricultural failure, harvest loss
5. **governance_category**: Government policy, elections, political instability, corruption, institutional failures
6. **health_category**: Disease outbreaks, healthcare access, public health emergencies, epidemic response
7. **humanitarian_category**: Humanitarian aid, relief operations, NGO activities, donor appeals, emergency response
8. **weather_category**: Drought, floods, climate shocks, seasonal forecasts, extreme weather events
9. **other_category**: Articles not matching above categories (sports, culture, general news)

Classification methodology: Articles are assigned to categories via boolean keyword matching against GDELT themes (CAMEO taxonomy codes), full-text keyword presence, and location mentions. Each article can belong to multiple categories simultaneously (e.g., drought-driven displacement flags both **weather_category** and **displacement_category**). The nine-category taxonomy consolidates GDELT’s 300+ fine-grained CAMEO codes into interpretable macro-categories aligned with humanitarian early warning frameworks [2, 40].

Rationale for macro-categories: The taxonomy captures key drivers and manifestations of food insecurity documented in humanitarian literature: conflict disrupts agricultural production and market access, economic shocks reduce household purchasing power, weather events destroy crops and livestock, displacement indicates population stress, and humanitarian coverage signals international awareness of deteriorating conditions [19, 31]. By tracking shifts in the composition of news coverage across these dimensions, we hypothesize that changes in narrative framing (e.g., increasing conflict coverage) may precede observable IPC deterioration.

3.4.4 Ratio Features: Cross-Sectional Coverage Composition

Ratio features quantify the relative emphasis of each category in a district’s news coverage, capturing compositional shifts independent of absolute volume.

Feature Definition

For each district i at time t , compute the proportion of total articles falling into each category c :

$$\text{ratio}_{c,i,t} = \frac{\text{count}_{c,i,t}}{\sum_{k=1}^9 \text{count}_{k,i,t}} \quad (3.30)$$

where:

- $\text{count}_{c,i,t}$ is the number of articles in category c for district i during month t
- The denominator sums across all nine categories, providing the total coverage volume
- Ratios are normalised to $[0, 1]$ and sum to 1 across categories for each observation

Missing values (districts with zero articles in month t) are filled with 0, under the assumption that absence of coverage reflects either genuine low activity or media disinterest rather than missing data.

Interpretation and Rationale

Ratio features capture *relative emphasis* rather than absolute volume. Consider two scenarios:

Scenario A (Volume change, constant composition): District receives 100 articles in January (50% conflict, 30% economic, 20% other) and 200 articles in February (50% conflict, 30% economic, 20% other). Ratio features are *unchanged* despite doubling volume, correctly identifying that coverage composition is stable.

Scenario B (Composition shift, constant volume): District receives 100 articles in January (20% conflict, 60% economic, 20% other) and 100 articles in February (60% conflict, 20% economic, 20% other). Ratio features *change dramatically* (conflict ratio increases from 0.20 to 0.60), capturing a shift in narrative framing toward violence even though total volume is constant.

This sensitivity to composition rather than volume is critical for detecting qualitative shifts in crisis dynamics. A sudden increase in `conflict_ratio` (from, say, 0.10 to 0.50) signals that news narratives are refocusing from economic concerns to armed violence, potentially indicating escalation invisible to AR baselines predicated on IPC persistence.

Advantages over Article Counts

Raw article counts conflate three effects: (1) baseline coverage levels (large cities receive more coverage than remote districts regardless of crisis status), (2) media interest (elections, sporting events generate non-crisis coverage spikes), and (3) genuine crisis signals. Ratio features partial out baseline volume by normalising, isolating compositional changes that reflect genuine shifts in the nature of crises rather than media attention fluctuations.

3.4.5 12-Month Sliding-Window Z-Score Standardisation

Z-score features quantify how unusual current coverage levels are relative to recent historical baselines, detecting anomalous spikes or drops that may precede IPC changes.

Feature Definition

For each category c , district i , and time t , compute the rolling z-score:

$$\text{z-score}_{c,i,t} = \frac{\text{count}_{c,i,t} - \mu_{c,i,t-12:t-1}}{\sigma_{c,i,t-12:t-1}} \quad (3.31)$$

where:

- $\mu_{c,i,t-12:t-1}$ is the mean article count for category c over the 12-month window preceding time t (months $[t - 12, t - 1]$)
- $\sigma_{c,i,t-12:t-1}$ is the standard deviation over the same window
- Standardisation produces zero-mean, unit-variance features within each district's historical context

Minimum periods: To prevent unreliable z-scores during early months where historical data is sparse, z-scores are only computed when at least 3 months of historical data are available (`min_periods = 3`). Observations with fewer than 3 historical months receive $\text{z-score} = \text{NaN}$ and are excluded from training (approximately 8-10% of observations, concentrated in the first year of the dataset).

Handling constant features: If all article counts in the 12-month window are identical ($\sigma = 0$), the z-score is set to 0, reflecting that current coverage is typical (no deviation from baseline).

Interpretation: Detecting Anomalous Coverage Shifts

Z-scores measure *how many standard deviations* current coverage deviates from recent historical norms. Interpretation:

- $\text{z-score} = 0$: Current coverage is exactly at the 12-month mean (typical)
- $\text{z-score} = +2$: Current coverage is 2 standard deviations above recent baseline (anomalously high, potential early warning signal)
- $\text{z-score} = -2$: Current coverage is 2 standard deviations below baseline (anomalous silence, potentially concerning if coverage typically tracks crises)
- $|\text{z-score}| > 2$: Extreme deviation (occurs in approximately 5% of observations under normal distribution, flagging genuine anomalies)

Example: If a district's conflict coverage averages 20 articles/month with standard deviation 5, and February receives 35 conflict articles, the z-score is $(35 - 20)/5 = +3.0$, indicating an extreme spike that may precede conflict-driven food insecurity.

Why 12-Month Windows?

The 12-month rolling window balances three considerations:

1. **Seasonal adjustment:** Agricultural cycles, lean seasons, and harvest periods introduce annual seasonality in food security coverage. A 12-month window captures one full seasonal cycle, allowing z-scores to detect deviations from seasonal baselines rather than flagging predictable annual patterns as anomalies.
2. **Baseline stability:** Shorter windows (e.g., 3 months) produce volatile baselines sensitive to single-month outliers, inflating false positives. Longer windows (e.g., 24 months) smooth over genuine shifts in coverage patterns, reducing sensitivity to emerging crises. Twelve months provides sufficient data for stable mean/std estimates while remaining responsive to recent trends.
3. **Data availability:** With a 48-month dataset (2021-2024), a 12-month window provides at least 36 months of usable data per district after warm-up, sufficient for model training while excluding only the initial year where historical baselines are unavailable.

Comparison to Ratio Features

Ratio and z-score features capture complementary signals:

- **Ratio features:** Cross-sectional composition (“What proportion of coverage is conflict?”), invariant to volume changes, district-specific baseline simplicity
- **Z-score features:** Temporal anomaly detection (“Is current conflict coverage unusually high compared to the past 12 months?”), volume-sensitive, explicit historical baseline comparison

A district may have high `conflict_ratio` (0.60, indicating conflict dominates coverage) but low `conflict_z-score` (0.0, indicating this is typical for that district). Conversely, low `conflict_ratio` (0.10) with high `conflict_z-score` (+2.0) signals an *unusual spike* in conflict coverage even though conflict remains a minority theme. Both perspectives are informative for early warning.

3.4.6 Feature Set Composition and Dimensionality

The complete basic feature set combines ratio features, z-score features, and location metadata:

Basic Feature Set (21 features)

1. **Ratio features (9):** One ratio per macro-category (conflict, displacement, economic, food security, governance, health, humanitarian, weather, other)
2. **Z-score features (9):** One z-score per macro-category (conflict, displacement, economic, food security, governance, health, humanitarian, weather, other)
3. **Location metadata (3):** Interpretable geographic context features

- **country_data_density**: Average articles per year for district’s country (proxy for media penetration, infrastructure)
- **country_baseline_conflict**: Historical conflict coverage proportion for country (proxy for chronic vs acute conflict environments)
- **country_baseline_food_security**: Historical food security coverage proportion for country (proxy for endemic vs episodic food insecurity)

Rationale for location features: Rather than arbitrary label encoding (“country=1, 2, 3, ...”), location metadata captures *why* a location exhibits certain risk characteristics. High **country_data_density** indicates strong media infrastructure, potentially improving signal quality. High **country_baseline_conflict** identifies chronic conflict zones (e.g., Sudan, Somalia) where **conflict_ratio** spikes may have different implications than in typically peaceful contexts (e.g., Malawi). These features allow models to learn context-dependent relationships between news signals and crisis onset.

Training Subset: WITH_AR_FILTER

As established in Section 3.3.6, Stage 2 models train exclusively on the 6,553 observations meeting the WITH_AR_FILTER criteria ($IPC_{t-1} \leq 2$ AND $AR_{pred} = 0$). This filtered set includes the 1,427 AR failures (cases where IPC deteriorated to Phase 3+ despite the AR baseline predicting stability), representing the hardest-to-predict crises where news features must demonstrate value. The remaining 5,126 observations are non-crisis cases where both the previous period was stable and AR correctly predicted continued stability. This selective training strategy ensures news features learn to complement AR baselines by focusing on AR-difficult cases rather than replicating persistence patterns already captured by Lt and Ls.

Class imbalance in filtered set: Among the 6,553 WITH_AR_FILTER observations, only 393 are crises (6.0% crisis rate), creating severe class imbalance relative to the full dataset (25.9% crisis rate). This imbalance arises because AR baselines already capture most easy-to-predict crises, leaving Stage 2 with predominantly non-crisis observations plus the hardest-to-predict minority of crises. XGBoost models address this via **scale_pos_weight** class balancing (Section 3.6.5), while mixed-effects models use **class_weight=10** to penalize false negatives heavily.

This subsection established the basic dynamic feature construction framework, transforming raw GDELT article counts into informative signals via two complementary approaches. Ratio features (9 categories) capture cross-sectional coverage composition, detecting qualitative shifts in narrative framing independent of volume changes (e.g., increasing conflict_ratio from 0.10 to 0.60 signals escalation even when total articles are constant). Z-score features (9 categories) use 12-month sliding windows to detect anomalous coverage spikes relative to district-specific historical baselines, identifying unusual events that may precede IPC deterioration (e.g., conflict_z-score = +3.0 indicates extreme spike). Together with 3 interpretable location metadata features (country-level coverage density and baseline thematic composition), the basic 21-feature set provides compositional and temporal signals designed to complement AR persistence by detecting crisis

dynamics invisible to lagged IPC values alone. All features train on the WITH_AR_FILTER subset (6,553 AR-difficult observations), ensuring news signals learn to rescue the 1,427 AR failures rather than replicate structural persistence.

3.4.7 Advanced Features: Regime and Mode Extraction

The second component of Stage 2 feature engineering applies advanced time series methods—Hidden Markov Models (HMM) and Dynamic Mode Decomposition (DMD)—to extract latent features from ratio and z-score sequences. These methods aim to capture crisis dynamics invisible to static aggregations: HMM detects regime transitions (peaceful $\rightarrow$ crisis-prone narratives) even when article volumes remain constant, while DMD extracts temporal modes characterising escalation patterns and sustained intensity. Both operate on rolling 12-month windows, producing district-specific features that evolve over time. The resulting advanced feature set (35 features total) combines basic compositional/anomaly signals (21 features from Stage 2) with dynamic regime and mode characteristics (14 HMM/DMD features), enabling models to distinguish qualitative shifts in crisis narratives from quantitative coverage fluctuations.

3.4.8 Hidden Markov Models for Latent Regime Detection

Hidden Markov Models posit that observed news coverage arises from unobservable latent states (regimes) with regime-specific emission distributions and probabilistic transitions between states. For food security forecasting, we hypothesize that districts alternate between *Pre-Crisis* and *Crisis-Prone* narrative regimes, with transitions signaling changes in underlying crisis dynamics before IPC classifications update.

Model Specification: Binary Regime HMM

This dissertation employs a **2-state Gaussian HMM** with asymmetric transition probabilities designed to capture crisis persistence. States are ordered by mean IPC values observed during historical regime occupancy, producing:

- **State 0 (Pre-Crisis):** Lower average IPC during historical occupancy, indicating stable or improving conditions
- **State 1 (Crisis-Prone):** Higher average IPC during historical occupancy, indicating persistent or worsening conditions

Asymmetric transition constraint: Crisis regimes exhibit persistence due to asset depletion irreversibility, conflict entrenchment, and structural drivers that resist rapid reversal. To reflect this humanitarian reality, transition probabilities are constrained:

$$P(\text{State}_{t+1} = \text{Crisis-Prone} \mid \text{State}_t = \text{Crisis-Prone}) \geq 0.85 \quad (3.32)$$

ensuring that once districts enter crisis-prone narrative regimes, they persist with at least 85% probability. This constraint encourages stable regime identification by requiring sustained signal




Figure 3.4: IPC crisis dynamics decompose into low-frequency persistence (captured by AR) and high-frequency shocks (targeted by Cascade). Panel A shows observed IPC phase for Kenya Northern Pastoral Zone (260 periods), with 29 identified shock events (red stars exceeding 1σ threshold). Panel B decomposes the low-frequency component (6-month moving average) representing structural persistence—slow-changing trends that AR baseline effectively captures (73.2% recall). Panel C isolates high-frequency deviations (residuals from trend) representing rapid shocks—sudden spikes and drops where AR fails but Cascade rescues 249 cases (+17.4% of AR failures). Panel D shows power spectral density via FFT: 82.9% of signal power concentrates in low-frequency components (<0.17 cycles/month, periods >6 months), while 17.1% resides in high-frequency components (>0.17 cycles/month, periods <6 months). Panel E summarizes the two-stage deployment logic: AR baseline handles low-frequency persistence efficiently, while advanced HMM/DMD features target high-frequency shocks where temporal autoregressive features fail. This frequency decomposition motivates the cascade architecture: rather than attempting to improve on AR’s already-excellent performance on slow trends, Stage 2 focuses computational resources on the intrinsically harder but operationally critical rapid-onset crises (conflict escalations, coups, displacement shocks) that confound persistence models. *Real data: Kenya Northern Pastoral Zone, $n=260$ periods. Low freq: 6-month centred MA. Shock threshold: 1σ deviation.*

patterns rather than transient coverage fluctuations.

Input features (4 core categories): HMM operates on ratio features for four crisis-core categories selected for high correlation with IPC outcomes:

1. **food_security_ratio:** Direct food insecurity signals (hunger, malnutrition, famine coverage)
2. **conflict_ratio:** Armed violence disrupting agricultural production and market access
3. **economic_ratio:** Price shocks, inflation, unemployment driving household food access constraints
4. **weather_ratio:** Drought, floods, climate shocks destroying crops and livestock

By restricting HMM to these four features rather than all nine categories, we reduce parameter dimensionality (improving convergence with limited data) while retaining signal most directly linked to food security crises [2, 40].

Implementation: District-Level Pooling

HMM models are estimated at the **district level**, fitting 1,322 separate 2-state Gaussian HMMs (one per news-dense district meeting the 200 articles/year threshold from Section 3.1.5). This district-level pooling strategy balances three considerations:

1. **Context specificity:** Each district has unique baseline coverage patterns, conflict histories, and crisis drivers. District-specific HMMs adapt to local dynamics rather than assuming universal transition probabilities across all African districts.
2. **Data sufficiency:** Pooling at the country level would aggregate heterogeneous sub-national dynamics (e.g., peaceful southern regions with conflict-affected northern regions within the same country), obscuring district-level regime shifts. Pooling at the district level provides 48 monthly observations per HMM, sufficient for 2-state models with 4 input features.
3. **Computational feasibility:** Training 1,322 independent 2-state HMMs is computationally tractable (approximately 2 hours on standard hardware), whereas fitting a single pooled HMM across all districts would assume identical transition dynamics globally, contradicting empirical heterogeneity [19].

Rolling window extraction: For each district’s time series, HMM features are extracted using a 12-month rolling window. At each time point t , the HMM is fit to data from $[t - 11, \textit{dots}, t]$, producing features capturing current regime probabilities and transition risks based on the most recent year of coverage dynamics.

Output Features (3 per feature type)

Each HMM produces three features capturing regime state and dynamics:




Figure 3.5: Hidden Markov Model detects 8 regime transitions in Sudan Eastern Pastoral district (2022-2023). Panel A shows conflict article coverage oscillating between peaceful (green, risk<0.5) and violent (red, risk≥0.5) regimes with 8 detected transitions. The HMM captures rapid regime volatility during Sudan’s civil war period, with frequent shifts between V→P (violent to peaceful) and P→V (peaceful to violent) states. Panel B shows corresponding food security coverage responding to regime dynamics. Panel C displays HMM transition risk oscillating around 0.5 threshold, with 8 crossings indicating regime changes. Panel D (State Interpretation) shows empirical transition probabilities: Peaceful regime appears 4/24 months (17%, 0% persistence, 100% transition to violent); Violent regime dominates 20/24 months (83%, 79% persistence, 21% transition to peaceful). This persistent violent regime (83% occupancy) reflects sustained conflict throughout 2022-2023. The HMM operates on 12-month rolling windows, fitting 2-state Gaussian models to ratio features. HMM transition risk (feature #5, importance 3.2%) enables cascade to detect regime shifts that AR persistence models miss, contributing to 249 key saves in conflict zones (Sudan 59, Zimbabwe 77, DRC 40). *Real data: Sudan Eastern Pastoral, n=24 months (2022-2023), 8 regime transitions detected, 83% violent regime occupancy.*

1. **hmm_*[type]*_crisis_prob**: $P(\text{State}_t = \text{Crisis-Prone})$, the posterior probability of currently occupying the crisis-prone regime. Values near 1 indicate the district is in a crisis-prone narrative state, while values near 0 indicate pre-crisis stability.
2. **hmm_*[type]*_transition_risk**: $\sum_{s=0}^1 P(\text{State}_t = s) \cdot P(\text{State}_{t+1} = \text{Crisis-Prone} \mid \text{State}_t = s)$, the probability of transitioning to (or remaining in) the crisis-prone state in the next period, weighted by current state uncertainty. This captures forward-looking risk.
3. **hmm_*[type]*_entropy**: $-\sum_{s=0}^1 P(\text{State}_t = s) \cdot \log P(\text{State}_t = s)$, the Shannon entropy of the state distribution, quantifying regime uncertainty. High entropy ($\approx \log 2 = 0.693$) indicates ambiguous regime assignment (equally likely pre-crisis or crisis-prone), while low entropy (≈ 0) indicates confident regime classification.

where $[type] \in \{\text{ratio}, \text{z-score}\}$ denotes whether HMM operates on ratio or z-score features, producing 6 HMM features total (3 per feature type).

Convergence and Data Constraints

HMM estimation via Expectation-Maximisation (EM) iterates until convergence (maximum 300 iterations, tolerance 10^{-3}). Across the 1,322 district-level HMMs trained on rolling windows, convergence is achieved in **89.5% of observations**. The 10.5% convergence failures occur predominantly in districts with:

- Sparsecoverage (near the 200 articles/year threshold, producing high zero-inflation in monthly counts)
- Shortusable sequences (districts in the dataset late in 2021-2024 with fewer than 12 months of historical data)
- Constantfeature values (districts with unchanging coverage composition across all months, producing singular covariance matrices)

Observations where HMM convergence fails receive NaN for all HMM features and are excluded from advanced feature set models (approximately 10% of observations).

3.4.9 Dynamic Mode Decomposition for Crisis Escalation Patterns

Dynamic Mode Decomposition extracts spatial-temporal modes from multivariate time series, decomposing coverage dynamics into oscillatory patterns with characteristic frequencies and growth rates. Unlike HMM (which models discrete regime transitions), DMD identifies continuous temporal modes representing crisis escalation (growing modes with positive eigenvalues) versus decay (shrinking modes with negative eigenvalues).

Mathematical Framework

DMD decomposes a time series $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T]$ (where each $\mathbf{x}_t \in \mathbb{R}^d$ is a d -dimensional feature vector at time t) into dynamic modes via the linear approximation:

$$\mathbf{x}_{t+1} \approx \mathbf{A}\mathbf{x}_t \quad (3.33)$$

where $\mathbf{A} \in \mathbb{R}^{d \times d}$ is the best-fit linear operator approximating temporal evolution. DMD constructs $\mathbf{A}$ via:

$$\mathbf{A} = \mathbf{X}_2 \mathbf{X}_1^\dagger \quad (3.34)$$

where $\mathbf{X}_1 = [\mathbf{x}_1, \dots, \mathbf{x}_{T-1}]$, $\mathbf{X}_2 = [\mathbf{x}_2, \dots, \mathbf{x}_T]$, and † denotes the Moore-Penrose pseudoinverse. The eigenvalues λ_k and eigenvectors Φ_k of $\mathbf{A}$ characterise dynamic modes:

- **Growth rate:** $\text{Re}(\log \lambda_k)$ quantifies exponential growth (positive) or decay (negative)
- **Frequency:** $|\text{Im}(\log \lambda_k)|/(2\pi)$ quantifies oscillation frequency (cycles per month)
- **Mode shape:** Φ_k identifies which feature combinations (categories) co-evolve in mode k

Crisis-Focused Mode Filtering

Standard DMD extracts all modes, including non-crisis-related patterns (seasonal agricultural cycles, electoral coverage spikes, sporting events). To isolate crisis-predictive dynamics, this dissertation implements a **3-step crisis mode filter**:

Step 1: Growth threshold. Retain only modes with positive growth rates exceeding $\lambda > 0.01$ (1% monthly exponential growth), filtering out decaying or stable modes. Crisis escalation manifests as growing coverage intensification, not shrinking dynamics.

Step 2: Frequency band. Retain only modes with frequencies in $[1/6, 1/2]$ cycles/month (periods of 2-6 months), filtering out:

- Seasonal cycles (annual periods, frequency $\approx 1/12$ cycles/month)
- High-frequency variation (weekly news cycles, frequency > 1 cycles/month)

The 2-6 month band captures crisis escalation timescales documented in humanitarian early warning literature [40], isolating the temporal range most relevant for detecting emerging crisis dynamics.

Step 3: Category weighting. Compute crisis-relevance scores for each mode based on its loading on crisis-core categories:

$$\text{CrisisWeight}_k = \frac{\sum_{c \in \text{CrisisCategories}} w_c |\Phi_{k,c}|}{\sum_{j=1}^d |\Phi_{k,j}|} \quad (3.35)$$

where crisis categories (conflict, food_security, displacement, humanitarian) receive weight $w_c = 1.0$, economic receives $w_c = 0.5$ (contextual), and excluded categories receive $w_c = 0$. Modes with $\text{CrisisWeight}_k < 0.3$ are filtered out, retaining only modes dominated by crisis-relevant features.

Input features (15 per feature type): DMD operates on 5 crisis-focused categories (4 core + economic) with 3 derivatives each (ratio/z-score, delta, 3-month trend), totaling 15 features per type. Including derivatives allows DMD to capture acceleration dynamics (increasing rate of change) beyond static levels.

Output Features (4 per feature type)

After crisis mode filtering, the dominant mode (highest crisis-weighted growth *simes* amplitude) defines four features:

1. **dmd_[type]_crisis_growth_rate:** Exponential growth rate of the dominant crisis mode ($\text{Re}(\log \lambda)$). Positive values indicate escalating multi-category coverage intensification.
2. **dmd_[type]_crisis_instability:** Sum of crisis-weighted growth rates across all filtered modes, $\sum_k \text{Re}(\log \lambda_k) \cdot a_k \cdot \text{CrisisWeight}_k$, where a_k is mode amplitude. High instability indicates multiple simultaneous escalating crisis patterns.
3. **dmd_[type]_crisis_frequency:** Oscillation frequency of the dominant crisis mode (cycles/month). Captures temporal periodicity of escalation dynamics.
4. **dmd_[type]_crisis_amplitude:** Amplitude of the dominant crisis mode ($|a_k|$), quantifying the strength of the temporal pattern.

where $[type] \in \{\text{ratio}, \text{z-score}\}$ produces 8 DMD features total (4 per feature type).

Zero-handling: If no modes pass the 3-step crisis filter (indicating stable dynamics with no escalating crisis patterns), all DMD features are set to 0, reflecting the absence of crisis-predictive temporal modes.

Implementation and Convergence


DMD uses Singular Value Decomposition (SVD) with rank-5 truncation (retaining top 5 modes by singular value energy) and regularization $\epsilon = 10^{-6}$ to stabilize pseudoinverse computation. Features are extracted on 12-month rolling windows (matching HMM). Across all districts and time points, DMD successfully extracts crisis modes in **83.1% of observations**, with failures concentrated in districts with:

- Insufficient sequence length (<8 months of non-missing data in rolling window)
- Constant or near-constant feature values (producing rank-deficient $\mathbf{X}_1$)
- High missing value rates (requiring excessive imputation)

3.4.10 Advanced Feature Set Composition

Combining Stage 2 basic features (21) with Stage 3 regime/mode features (14) produces the **advanced feature set** (35 features total):

1. **Basic features (21):** 9 ratio + 9 z-score + 3 location metadata (from Section 3.4.4)



Dynamic Mode Decomposition (DMD) extracts temporal modes from multivariate news coverage time series. Real data: Sudan Eastern Pastoral, $n=420$ periods. Three DMD features shown: Panel A (Growth Rate - real eigenvalue, $n=21$ escalation periods), Panel B (Frequency - oscillation period), Panel C (Amplitude - oscillation magnitude). DMD convergence rate: 83.1% (Baum-Welch algorithm, 12-month rolling windows). Rare events focus: 18.7% of observations are crises with high humanitarian impact. DMD growth rate ranks 28 (2.1% importance), complementing HMM regime transitions. Crisis periods shown as red shading. Escalation (growth >0) marked with red triangles.

Figure 3.6: Dynamic Mode Decomposition extracts three continuous temporal patterns from real Sudan Eastern Pastoral crisis data (420 periods). Panel A (Growth Rate): Real eigenvalue identifies 21 escalation periods (growth >0) marked with red triangles, distinguishing exponential crisis intensification from stable dynamics. Panel B (Frequency): Oscillation period shows sparse crisis-driven spikes concentrated in early periods (2021-2022) and end of series (2024), capturing seasonal patterns. Panel C (Amplitude): Oscillation magnitude envelope shows limited volatility, with early-period spikes matching crisis clusters. Panel D (Summary): Convergence analysis shows 83.1% Baum-Welch success rate, with 16.9% failures discarded. Rare events focus: 18.7% crisis rate (11,836/63,140 observations) with high humanitarian impact. DMD growth rate ranks 28 (2.1% tree-based importance), complementing HMM regime transitions which capture discrete state changes. Crisis periods (270/420 in this district) shown as red background shading across panels A-C. Note: Instability coefficient (imaginary eigenvalue) removed from visualisation as values near zero throughout dataset. *Real data: Sudan Eastern Pastoral, $n=420$ periods, 21 escalation events, 83.1% convergence, rank-5 SVD truncation, 12-month windows.*

2. **HMM features (6)**: 3 ratio-based HMM (crisis_prob, transition_risk, entropy) + 3 z-score-based HMM
3. **DMD features (8)**: 4 ratio-based DMD (growth_rate, instability, frequency, amplitude) + 4 z-score-based DMD

Coverage: Advanced features are available for observations where both HMM convergence (89.5%) and DMD mode extraction (83.1%) succeed, producing approximately 74% coverage ($89.5\% \times 83.1\% \approx 74\%$) across the `WITH_AR_FILTER` training set (6,553 observations). Models using advanced features train on the subset with complete feature availability.

Missing value handling: Observations with missing HMM or DMD features are excluded from advanced models but retained for basic models (21 features), ensuring all observations contribute to at least one model variant. This missing-completely-at-random (MCAR) assumption is validated by confirming that HMM/DMD convergence failures are not systematically associated with crisis outcomes (convergence failure rates are similar for crisis and non-crisis observations: 11.2% vs 10.3%, χ^2 p=0.18).

This subsection established the advanced component of Stage 2 feature engineering, applying regime and mode extraction methods to detect latent crisis dynamics. Hidden Markov Models (HMM) fit 1,322 district-specific 2-state Gaussian HMMs with asymmetric crisis persistence constraints ($P(\text{Crisis} \rightarrow \text{Crisis}) \geq 0.85$), operating on 4 core ratio/z-score categories to produce 6 features capturing regime probabilities, transition risks, and state uncertainty—convergence achieved in 89.5% of observations. Dynamic Mode Decomposition (DMD) applies 3-step crisis filtering ($\text{growth} > 0.01$, frequency in $[1/6, 1/2]$, category weighting > 0.3) to 15 derivatives per feature type, extracting 8 features quantifying escalation growth rates, multi-mode instability, oscillation frequencies, and pattern amplitudes—successful mode extraction in 83.1% of observations. Together with basic features (21), the complete Stage 2 advanced feature set totals 35 features, available for approximately 74% of observations where both HMM and DMD converge. These features enable models to distinguish qualitative regime transitions and temporal escalation patterns from simple compositional shifts, completing Stage 2’s transformation of raw article counts into dynamic crisis signals targeting hard-to-predict cases invisible to AR persistence baselines.

3.5 Model Training and Evaluation Framework

This section details the model training methodology applied to Stage 2 features for predicting hard-to-forecast cases missed by Stage 1 AR baselines. Training operates on the `WITH_AR_FILTER` subset (6,553 observations where $\text{IPC}_{t-1} \leq 2$ AND $\text{AR_pred} = 0$), containing 393 crises and 6,160 non-crisis observations. This creates an extreme class imbalance (6.0% crisis rate) compared to the full dataset (25.9% crisis rate), requiring specialised training strategies.

We employ two complementary modelling approaches—XGBoost and mixed-effects logistic regression—each with distinct strengths. XGBoost excels at capturing complex non-linear

interactions and hierarchical feature importance through gradient-boosted decision trees, while mixed-effects models provide interpretable random coefficients quantifying geographic heterogeneity in feature effects. Both model families use identical 5-fold stratified spatial cross-validation (Section 3.2.2) to ensure comparable, spatially-unbiased performance estimates.

3.5.1 XGBoost Gradient Boosting Models

XGBoost (Extreme Gradient Boosting) constructs an ensemble of decision trees iteratively, where each tree corrects residual errors from preceding trees [80]. The final prediction aggregates contributions from all trees:

$$\hat{y}_i = \sum_{k=1}^K f_k(\mathbf{x}_i) \quad (3.36)$$

where f_k is the k -th tree and K is the total number of trees (boosting rounds). For binary crisis prediction, the model outputs logistic-transformed probabilities $p_i = \sigma(\hat{y}_i)$ where $\sigma(z) = 1/(1 + e^{-z})$.

Hyperparameter Optimisation via GridSearchCV

To ensure fair comparison across feature sets of differing complexity (basic: 21 features, advanced: 35 features, ablation variants: 12-35 features), all XGBoost models undergo identical hyperparameter optimisation using GridSearchCV with stratified spatial cross-validation. The hyperparameter grid explores 3,888 combinations:

- **n_estimators** *in*{100, 200, 300}: Number of boosting rounds
- **max_depth** *in*{5, 7, 10}: Maximum tree depth (expanded upward from preliminary testing to allow complex models to utilise advanced features)
- **learning_rate** *in*{0.01, 0.05, 0.1}: Step size for gradient descent
- **min_child_weight** *in*{1, 3, 5}: Minimum sum of instance weights per leaf (expanded downward for finer splits with more features)
- **subsample** *in*{0.7, 0.8}: Fraction of training samples per tree
- **colsample_bytree** *in*{0.6, 0.8}: Fraction of features per tree
- **gamma** *in*{0, 0.5, 1}: Minimum loss reduction for split
- **reg_alpha** *in*{0, 0.1}: L1 regularization
- **reg_lambda** *in*{1, 2}: L2 regularization

GridSearchCV selects hyperparameters maximising mean cross-validated AUC-ROC across the 5 spatial folds. This scoring metric prioritises discrimination (ranking crisis cases higher than non-crisis) rather than raw accuracy, which is inappropriate for the 6.0% crisis prevalence in the WITH_AR_FILTER subset.

Class Imbalance Handling

The `WITH_AR_FILTER` subset exhibits severe class imbalance (15.5:1 non-crisis to crisis ratio). XGBoost addresses this via `scale_pos_weight`, which weights positive class (crisis) loss contributions higher during training:

$$\text{scale_pos_weight} = \frac{n_{\text{negative}}}{n_{\text{positive}}} \quad (3.37)$$

calculated per training fold to reflect fold-specific imbalance. This weighting is equivalent to assigning crisis events higher importance in the gradient boosting objective, preventing the model from trivially achieving high accuracy by predicting only non-crisis.

Model Variants

Two XGBoost model variants are trained:

Basic Model (21 features): Uses only ratio features (9), z-score features (9), and location metadata (3). Establishes baseline performance without HMM or DMD advanced features.

Advanced Model (35 features): Adds HMM features (6) and DMD features (8) to the basic set. Tests how latent regime transitions and temporal modes contribute interpretable crisis dynamics beyond compositional features.

Both models train on the subset of `WITH_AR_FILTER` observations where HMM convergence and DMD mode extraction succeeded (approximately 74% coverage). Missing advanced features are imputed with country-median values (if available) then global median (fallback), preserving all observations for fair comparison.

Cross-Validation and Model Selection

For each fold $f \in \{0, 1, 2, 3, 4\}$:

1. Train XGBoost on 4 folds using optimised hyperparameters
2. Predict on held-out fold f
3. Compute fold-specific metrics: AUC-ROC, Brier score, log loss
4. Save fold model and feature importance rankings

After cross-validation, a **final model** is trained on all data using best hyperparameters for deployment. Cross-validated performance estimates (mean $\pm$ std across folds) provide unbiased generalisation metrics, while the final model maximises predictive power for Stage 3 cascade integration.

3.5.2 Mixed-Effects Logistic Regression

Mixed-effects models extend standard logistic regression by incorporating **random effects**—coefficients that vary by geographic group (country or district).

These capture systematic heterogeneity in how features predict crisis risk across regions [81].

The model specification is:

$$\log \frac{p_{r,t}}{1 - p_{r,t}} = \underbrace{\boldsymbol{\beta}^T \mathbf{X}_{r,t}}_{\text{Fixed effects}} + \underbrace{\alpha_g + \mathbf{b}_g^T \mathbf{Z}_{r,t}}_{\text{Random effects}} \quad (3.38)$$

where:

- $\boldsymbol{\beta}$: Fixed effect coefficients (global patterns across all regions)
- $\mathbf{X}_{r,t}$: All features (ratio, z-score, location, HMM, DMD)
- α_g : Random intercept for group g (country or district baseline risk)
- $\mathbf{b}_g$: Random slopes for key signals $\mathbf{Z}_{r,t}$ (group-specific feature effects)
- $\mathbf{Z}_{r,t} \subseteq \mathbf{X}_{r,t}$: Subset of features with random slopes (conflict_ratio, food_security_ratio)

Random effects are assumed normally distributed: $\alpha_g \sim \mathcal{N}(0, \sigma_\alpha^2)$ and $b_{g,j} \sim \mathcal{N}(0, \sigma_{b_j}^2)$. Large variance $\sigma_{b_j}^2$ indicates substantial geographic heterogeneity in feature j 's effect, informing selective deployment strategies.

Random Effects Grouping Level

The random effects grouping level (district vs country) is selected dynamically based on data sufficiency:

- **District-level** (preferred): Used if $\geq 50\%$ of districts have ≥ 10 observations, enabling fine-grained geographic variation estimates aligned with the research proposal's district-level focus.
- **Country-level** (fallback): Used if district coverage is insufficient, aggregating heterogeneity at the coarser country scale.

This adaptive strategy balances statistical power (sufficient data per group for stable random effect estimates) with geographic granularity.

Random Slope Selection

Random slopes capture feature effects that vary geographically. Due to computational constraints (random slopes increase model complexity quadratically with number of features), we select 2 key signals based on crisis-predictive priority:

1. **conflict_ratio**: Conflict news composition (highest priority for crisis prediction)
2. **food_security_ratio**: Food security mentions (second priority)

These features receive both fixed effects (global average impact) and random slopes (country/district-specific deviations), while remaining features receive only fixed effects.

Class Weighting for Imbalance Handling

Mixed-effects models address class imbalance via observation weighting:

$$w_i = \begin{cases} 10 & \text{if } y_i = 1 \text{ (crisis)} \\ 1 & \text{if } y_i = 0 \text{ (non-crisis)} \end{cases} \quad (3.39)$$

This 10:1 crisis weighting mirrors the cost-sensitive evaluation framework (*10simesFN* + *1simesFP*), prioritising recall over precision in humanitarian contexts where missing a crisis carries catastrophic consequences.

Model Fitting and Convergence

Mixed-effects models are fitted using `glmer` (Generalised Linear Mixed-Effects Regression) from the `lme4` R package with the following configuration:

- **Optimiser:** bobyqa (Bound Optimisation BY Quadratic Approximation)
- **Max iterations:** 100,000 (increased from default for convergence with weighted observations)
- **Integration:** Laplace approximation (nAGQ=0, faster than adaptive Gauss-Hermite quadrature)
- **Singular fit handling:** Warnings suppressed (tolerance 10^{-4}) to accommodate boundary random effect variances in low-data groups

If district-level models fail to converge, automatic fallback to country-level random effects occurs. If country-level models fail, unweighted models are attempted as final fallback.

Model Variants

Four mixed-effects model variants parallel the XGBoost design:

1. **pooled_ratio:** Ratio features (9) + location (3) only
2. **pooled_z-score:** Z-score features (9) + location (3) only
3. **pooled_ratio_hmm_dmd:** Ratio (9) + location (3) + HMM-ratio (3) + DMD-ratio (4) = 19 features
4. **pooled_z-score_hmm_dmd:** Z-score (9) + location (3) + HMM-z-score (3) + DMD-z-score (4) = 19 features

These variants enable direct comparison of ratio vs z-score feature types and assessment of HMM/DMD contributions to interpretability within the mixed-effects framework.

3.5.3 Ablation Study Design

Ablation studies systematically remove feature groups to isolate their marginal contributions [82]. Eight ablation models test specific hypotheses about feature value:

1. **ratio_location** (12 features): Ratio (9) + location (3). Baseline without z-scores or advanced features.
2. **z-score_location** (12 features): Z-score (9) + location (3). Tests z-score baseline without ratio or advanced features.
3. **ratio_z-score_location** (21 features): All basic features. Tests combined ratio + z-score performance.
4. **ratio_z-score_dmd_location** (29 features): Basic (21) + DMD (8). Tests DMD contribution without HMM.
5. **ratio_z-score_hmm_location** (27 features): Basic (21) + HMM (6). Tests HMM contribution without DMD.
6. **ratio_hmm_ratio_location** (18 features): Ratio (9) + HMM-ratio (3) + DMD-ratio (0) + location (3). Tests ratio-based features with HMM.
7. **z-score_hmm_z-score_location** (18 features): Z-score (9) + HMM-z-score (3) + DMD-z-score (0) + location (3). Tests z-score-based features with HMM.
8. **ratio_hmm_dmd_location** (22 features): Ratio (9) + HMM-ratio (3) + DMD-ratio (4) + location (3) + HMM-z-score (3) + DMD-z-score (0). Tests ratio with both advanced methods.

Each ablation model undergoes identical hyperparameter optimisation (3,888 combinations) and 5-fold stratified spatial cross-validation, ensuring fair comparison. Ablation results answer:

- **Q1:** Do z-scores improve over ratio features alone? (Compare models 1 vs 3)
- **Q2:** Does HMM provide marginal value? (Compare models 3 vs 5)
- **Q3:** Does DMD provide marginal value? (Compare models 3 vs 4)
- **Q4:** Do HMM + DMD together outperform either alone? (Compare models 4, 5 vs advanced 35-feature model)
- **Q5:** Which feature type (ratio vs z-score) benefits more from advanced features? (Compare models 6 vs 7)

3.5.4 Threshold Optimisation Strategies

Binary predictions require converting continuous probabilities $p_i \in [0, 1]$ to binary labels $\hat{y}_i \in \{0, 1\}$ via thresholding: $\hat{y}_i = \mathbb{I}[p_i \geq \tau]$. The threshold τ determines the precision-recall trade-off. This dissertation evaluates four threshold selection strategies:

Youden’s J Index: Maximises $J = \text{TPR} - \text{FPR}$ (sensitivity + specificity - 1), balancing true positive rate and false positive rate. Optimal for maximising overall classification performance.

F1-Maximising: Selects τ maximising $F_1 = 2 \cdot \text{precision} \cdot \text{recall} / (\text{precision} + \text{recall})$, the harmonic mean of precision and recall. Balanced metric for imbalanced data.

Balanced Precision-Recall: Finds τ where precision $\approx$ recall, achieving symmetry between positive predictive value and sensitivity.

High-Recall (≥ 0.90): Selects the highest τ achieving recall ≥ 0.90 , prioritising crisis detection (minimising false negatives) at the cost of precision. Aligned with humanitarian early warning priorities where missing a crisis is catastrophic.

Cross-validation reports metrics at all four thresholds, but the **high-recall threshold** is prioritised for cascade integration (Section 3.6) to maximise AR failure rescues.

3.5.5 Evaluation Metrics

Model performance is assessed using six metric categories:

Discrimination metrics (threshold-independent):

- AUC-ROC: Area under Receiver Operating Characteristic curve (probability model ranks crisis cases higher than non-crisis)

Classification metrics (threshold-dependent):

- Precision: $\text{TP} / (\text{TP} + \text{FP})$ (fraction of predicted crises that are true crises)
- Recall (Sensitivity): $\text{TP} / (\text{TP} + \text{FN})$ (fraction of true crises correctly identified)
- Specificity: $\text{TN} / (\text{TN} + \text{FP})$ (fraction of non-crisis correctly identified)
- F1 Score: $2 \cdot \text{precision} \cdot \text{recall} / (\text{precision} + \text{recall})$
- Balanced Accuracy: $(\text{recall} + \text{specificity}) / 2$

Calibration metrics:

- Brier Score: Mean squared error between predicted probabilities and binary outcomes, $\frac{1}{N} \sum_{i=1}^N (p_i - y_i)^2$
- Log Loss: Negative log-likelihood, $-\frac{1}{N} \sum_{i=1}^N [y_i \log p_i + (1 - y_i) \log(1 - p_i)]$

Cost-sensitive metric:

- Total Cost: $10 \text{imes} \text{FN} + 1 \text{imes} \text{FP}$, reflecting asymmetric humanitarian costs (missing crisis 10imes worse than false alarm)

Confusion matrix elements: True Positives (TP), True Negatives (TN), False Positives (FP), False Negatives (FN)

Geographic stratification: All metrics computed overall and stratified by country, enabling identification of geographic performance variation.

3.5.6 Feature Importance Extraction

XGBoost provides gain-based feature importance, quantifying each feature's contribution to model splits:

$$\text{Importance}_j = \sum_{k=1}^K \sum_{s \in \text{splits}(f_k, j)} \text{Gain}(s) \quad (3.40)$$

where $\text{Gain}(s)$ is the loss reduction from split s on feature j in tree k . Importance scores are normalised to sum to 1.0, enabling cross-model comparison.

Mixed-effects models provide fixed effect coefficients β (global features impacts) and random effect variances $\sigma_{b_j}^2$ (geographic heterogeneity). Features with large $|\beta_j|$ have strong global effects; features with large $\sigma_{b_j}^2$ exhibit high geographic variation, suggesting context-dependent deployment strategies.

Feature importance rankings inform interpretability analysis (Section 3.7) and guide operational deployment recommendations.

This section established the Stage 2 model training framework for predicting crises missed by Stage 1 AR baselines. XGBoost models undergo hyperparameter optimisation via GridSearchCV with 3,888 combinations evaluated using 5-fold stratified spatial cross-validation, maximising AUC-ROC while handling 15.5:1 class imbalance through scale_pos_weight. Two XGBoost variants (basic: 21 features, advanced: 35 features) and four mixed-effects logistic regression variants (ratio, z-score, ratio+HMM+DMD, z-score+HMM+DMD) enable comparison of feature types and advanced method contributions. Mixed-effects models incorporate country or district random effects (adaptive selection based on data sufficiency) with random slopes on conflict_ratio and food_security_ratio, capturing geographic heterogeneity in feature effects. Observation weighting (10:1 crisis:non-crisis) aligns training with humanitarian cost asymmetries. Nine ablation models systematically test marginal contributions of z-scores, HMM, and DMD features through controlled feature removal. Four threshold selection strategies (Youden's J, F1-max, balanced P=R, high-recall ≥ 0.90) optimise precision-recall trade-offs for different operational priorities, with high-recall prioritised for cascade integration. Comprehensive evaluation uses six metric categories (discrimination, classification, calibration, cost-sensitive, confusion matrix, geographic stratification) to assess model performance across spatial folds and countries. Feature importance extraction from XGBoost (gain-based) and mixed-effects models (fixed coefficients, random effect variances) quantifies feature contributions and geographic heterogeneity, informing interpretability analysis and operational deployment strategies.

3.6 Two-Stage Framework Integration

This section describes how Stage 1 (AR baseline) and Stage 2 (news-based models) are integrated into a unified cascade ensemble. The cascade framework is designed for **selective deployment**:

Stage 2 models operate only on cases where Stage 1 predicted non-crisis ($AR=0$), reducing computational costs while maximising marginal predictive value.

3.6.1 Cascade Decision Logic

The cascade employs **simple binary override logic**, integrating binary predictions from both stages without adaptive thresholds or probabilistic weighting. The decision rule operates as follows:

$$\hat{y}_{\text{cascade}} = \begin{cases} 1 & \text{if } \hat{y}_{AR} = 1 \text{ (trust AR's crisis prediction)} \\ \hat{y}_{\text{Stage2}} & \text{if } \hat{y}_{AR} = 0 \text{ (defer to Stage 2)} \end{cases} \quad (3.41)$$

where $\hat{y}_{AR}$ is the binary Stage 1 prediction (using balanced precision-recall optimal threshold $\tau = 0.629$) and $\hat{y}_{\text{Stage2}}$ is the binary Stage 2 prediction (using Youden's J threshold from the advanced XGBoost model trained on `WITH_AR_FILTER` observations).

This logic reflects three principles:

Principle 1: Trust AR's positive predictions. When Stage 1 predicts crisis ($\hat{y}_{AR} = 1$), the cascade accepts this prediction without override. AR baselines achieve high precision (0.732) on crisis predictions, reflecting strong autocorrelation in IPC transitions. Overriding these predictions would introduce unnecessary false positives.

Principle 2: Refine AR's negative predictions. When Stage 1 predicts no crisis ($\hat{y}_{AR} = 0$), the cascade defers to Stage 2. This subset (6,553 observations where $IPC_{t-1} \leq 2$ AND $AR=0$) contains all 1,427 AR-missed crises (false negatives the AR baseline failed to catch). Stage 2 models target these hard-to-predict cases using news features invisible to AR baselines.

Principle 3: Simple binary logic. The cascade uses binary predictions directly, avoiding complex probability fusion, adaptive thresholds, or meta-learning. While Stage 2 models internally optimise thresholds (Section 3.5.4), the cascade integration operates purely on binary labels. This simplicity ensures interpretability and operational deployability.

3.6.2 Override Mechanism and Coverage

Stage 2 predictions are available for 6,553 observations (31.6% of the full dataset), corresponding exactly to the `WITH_AR_FILTER` subset where:

- Previous IPC status: $IPC_{t-1} \leq 2$ (not already in crisis)
- AR prediction: $\hat{y}_{AR} = 0$ (AR predicts no future crisis)

Within these 6,553 override candidates, Stage 2 predicts crisis ($\hat{y}_{\text{Stage2}} = 1$) for 1,761 observations (26.9% override rate). These 1,761 overrides constitute the cascade's incremental contribution beyond the AR baseline. The remaining 4,792 observations where Stage 2 predicts $\hat{y}_{\text{Stage2}} = 0$ confirm the AR baseline's no-crisis prediction.

For the 14,169 observations where AR predicts crisis ($\hat{y}_{AR} = 1$) or where $IPC_{r,t} \geq 3$ (already in crisis), the cascade prediction equals the AR prediction without invoking Stage 2.

3.6.3 Key Saves: Quantifying Stage 2 Value

The primary metric for evaluating cascade performance is **key saves**—AR-missed crises that the cascade successfully predicts. Formally:

$$\text{Key Save} = \begin{cases} 1 & \text{if } y_{r,t+h} = 1 \text{ AND } \hat{y}_{AR} = 0 \text{ AND } \hat{y}_{\text{cascade}} = 1 \\ 0 & \text{otherwise} \end{cases} \quad (3.42)$$

Key saves represent the operational value-add of Stage 2: these are true crisis events that would have been missed if relying solely on Stage 1 AR baselines but are now caught through news-based early warning signals.

Across the full dataset (20,722 observations), the cascade achieves:

- **Total key saves:** 249 crises
- **AR baseline misses:** 1,427 crises (false negatives)
- **Key save rate:** 17.4% (249 / 1,427)

This means Stage 2 successfully rescues 17.4% of the crises that Stage 1 failed to predict. The remaining 1,178 AR-missed crises (82.6%) remain false negatives even after cascade integration. Cascade failure analysis (Chapter 5) reveals these still-missed cases exhibit systematic news coverage deficiency—median 74 articles/month compared to 121 for rescued cases (64% less coverage)—demonstrating a fundamental constraint: news-based early warning cannot rescue crises in news deserts (remote pastoral areas, peripheral regions) lacking sufficient media coverage. This constraint motivates future NLP enhancements: expanding text corpora through social media monitoring, community radio transcripts, humanitarian situation reports, and multilingual sources to address coverage gaps in underreported regions.

3.6.4 Performance Impact: Recall vs Precision Trade-Off

Integrating Stage 2 predictions improves cascade recall (crisis detection rate) but reduces precision (positive predictive value), reflecting the fundamental recall-precision trade-off in imbalanced classification. Overall performance changes:

Recallsimprovement:

$$\Delta \text{Recall} = 0.779 - 0.732 = +0.047 \text{ (4.7 percentage points)} \quad (3.43)$$

The cascade catches 77.9% of all crises, compared to 73.2% for the AR baseline alone. This 4.7-pointsimprovement corresponds directly to the 249 key saves.

Precision reduction:

$$\Delta \text{Precision} = 0.585 - 0.732 = -0.147 \text{ (14.7 percentage points)} \quad (3.44)$$

Cascade precision drops to 58.5%, meaning 41.5% of cascade crisis predictions are false alarms. This reduction reflects the increased false positive rate from Stage 2 overrides: of the 1,761 overrides, 1,512 are false positives (85.9%) and only 249 are true positives (14.1%).

F1 score change:

$$\Delta F_1 = 0.668 - 0.732 = -0.064 \quad (3.45)$$

The F1 score decreases slightly, as the precision reduction outweighs the recall gain in the harmonic mean.

Confusion matrix transformation:

Model	TP	TN	FP	FN
AR Baseline	3,895	13,973	1,427	1,427
Cascade Ensemble	4,144	12,461	2,939	1,178
Change	+249	-1,512	+1,512	-249

Table 3.2: Confusion matrix comparison: AR baseline vs cascade ensemble. The cascade gains 249 true positives (key saves) but incurs 1,512 additional false positives.

This trade-off is deliberate and aligned with humanitarian early warning priorities (Section 3.2.3). The cost function $10 \text{timesFN} + 1 \text{timesFP}$ reflects that missing a crisis (false negative) carries catastrophic humanitarian consequences 10 times worse than a false alarm (false positive). Under this asymmetric cost structure, rescuing 249 crises justifies 1,512 additional false alarms.

3.6.5 Geographic Distribution of Key Saves

Key saves exhibit substantial geographic heterogeneity, with 10 countries accounting for 249 total saves:

- Zimbabwe: 77 saves (30.9% of total)
- Sudan: 59 saves (23.7%)
- Democratic Republic of the Congo: 40 saves (16.1%)
- Nigeria: 27 saves (10.8%)
- Mozambique: 15 saves (6.0%)
- Mali: 12 saves (4.8%)
- Kenya: 8 saves (3.2%)
- Ethiopia: 6 saves (2.4%)

- Malawi: 3 saves (1.2%)
- Somalia: 2 saves (0.8%)

Three countries (Zimbabwe, Sudan, DRC) account for 70.7% of all key saves, suggesting that Stage 2’s marginal value is highly context-dependent. Countries with high key save counts exhibit characteristics conducive to news-based prediction: active conflict, political instability, and robust GDELT coverage capturing crisis escalation signals invisible to AR baselines.

Conversely, 8 countries (50% of the 18-country sample) show zero key saves, indicating contexts where news features provide no marginal value beyond AR persistence. This geographic heterogeneity informs selective deployment strategies (Chapter 5).

3.6.6 Model Selection for Cascade Integration

The cascade integrates the **advanced XGBoost model** (35 features: 21 basic + 6 HMM + 8 DMD) as the Stage 2 component. This model was selected based on:

1. **Performance:** Highest cross-validated AUC-ROC among all Stage 2 variants
2. **Feature richness:** Incorporates regime transitions (HMM) and temporal modes (DMD) beyond compositional features
3. **Generalisability:** Trained with spatial cross-validation and extensive hyperparameter optimisation (3,888 combinations)
4. **Robustness:** 74% coverage where both HMM and DMD features are available, with country-mediansimputation for missing values

Alternative Stage 2 models (basic XGBoost with 21 features, mixed-effects logistic regression variants) were evaluated but not selected for cascade integration. The advanced XGBoost model balances predictive performance with feature interpretability, enabling post-hoc analysis of which news signals drive key saves (Section 3.7).

*This section established the two-stage cascade framework integrating Stage 1 (AR baseline) and Stage 2 (news-based models) through simple binary override logic: if AR predicts crisis, trust it; if AR predicts no crisis, defer to Stage 2. The cascade operates on 6,553 override candidates (31.6% of full dataset) where $IPC_{t-1} \leq 2$ AND $AR = 0$, with Stage 2 overriding 1,761 cases (26.9% override rate). The key metric is **key saves**—AR-missed crises successfully predicted by the cascade—totaling 249 saves from 1,427 AR misses (17.4% rescue rate). Cascade integration improves recall from 73.2% to 77.9% (+4.7 points) but reduces precision from 73.2% to 58.5% (-14.7 points), incurring 1,512 additional false positives to rescue 249 true crises. This recall-precision trade-off is justified by humanitarian cost asymmetries ($10 \times FN + 1 \times FP$), where missing crises carries catastrophic consequences. Key saves exhibit geographic heterogeneity: Zimbabwe (77), Sudan (59), and DRC (40) account for 70.7% of total saves, while 8 countries show zero saves, informing selective deployment strategies. The advanced XGBoost model (35 features) serves as Stage 2, selected for highest cross-validated AUC-ROC, feature richness (HMM + DMD), and generalisability through spatial CV and hyperparameter optimisation.*

3.7 Interpretability Framework

Model interpretability enables understanding *which* news signals drive predictions and *how* their effects vary geographically, informing operational deployment strategies. This section establishes a triangulation framework combining three complementary interpretability methods with ablation studies for model evaluation:

Interpretability Methods (feature-level attribution):

1. **XGBoost gain-based feature importance:** Measures how frequently features are used to partition tree nodes (stratification utility)
2. **SHAP (SHapley Additive exPlanations):** Quantifies marginal contribution of each feature to individual predictions (attribution)
3. **Mixed-effects decomposition:** Separates global patterns (fixed effects) from country-specific variation (random effects)

Model Evaluation (feature group contribution):

4. **Ablation studies:** Isolates marginal contributions of entire feature groups through systematic removal

These methods provide complementary perspectives: gain-based importance identifies features useful for stratification, SHAP reveals features driving prediction variance, mixed-effects quantifies geographic heterogeneity, and ablation studies test marginal group contributions. Convergence across interpretability methods and validation through ablation studies provides robust evidence for feature rankings and deployment recommendations.

3.7.1 XGBoost Gain-Based Feature Importance

XGBoost provides gain-based feature importance, quantifying each feature's contribution to reducing prediction error across all tree splits [80]. For feature j , importance is computed as:

$$\text{Importance}_j = \sum_{k=1}^K \sum_{s \in \text{Splits}(f_k, j)} \text{Gain}(s) \quad (3.46)$$

where $\text{Gain}(s)$ is the loss reduction from split s on feature j in tree k , summed across all K trees. Importance scores are normalised to sum to 1.0, enabling cross-model comparison.

Top Features: Advanced XGBoost Model

The advanced XGBoost model (35 features) identifies the following top 10 features by gain-based importance:

1. **country_data_density** (0.133): GDELT article volume per capita, proxy for news coverage intensity
2. **country_baseline_conflict** (0.093): Country-average conflict news prevalence

3. **country_baseline_food_security** (0.067): Country-average food security news prevalence
4. **other_ratio** (0.033): Uncategorised news composition
5. **hmm_ratio_transition_risk** (0.032): HMM regime transition probability (peaceful $\rightarrow$ crisis)
6. **health_ratio** (0.029): Health crisis news composition
7. **displacement_z-score** (0.026): Displacement news z-score anomaly
8. **weather_ratio** (0.026): Climate/weather event news composition
9. **food_security_z-score** (0.025): Food security news z-score anomaly
10. **hmm_ratio_crisis_prob** (0.025): HMM crisis state probability

Three location metadata features (`country_data_density`, `country_baseline_conflict`, `country_baseline_food_security`) account for 29.3% of total tree-based importance, dominating compositional and temporal features in split frequency. However, this reflects their role as stratification infrastructure rather than marginal predictive contribution: SHAP analysis (Section 3.6.4.2) reveals location features contribute only 2.6% of marginal attribution, while z-score news features drive 74.7% of predictions. The tree-based metric conflates *how often* features create splits (stratification utility) with *how much* they drive predictions (marginal impact). The highest-ranking compositional feature (`other_ratio`, rank 4) contributes 3.3% tree-based importance.

Among advanced features, HMM transition risk ranks 5th (3.2% importance), demonstrating genuine predictive value from latent regime dynamics. DMD features rank lower (`crisis_growth_rate` at rank 28: 2.1%, `crisis_amplitude` at rank 30: 1.9%), suggesting temporal modes contribute less than regime transitions.

Feature Importance Insights

Feature importance rankings reveal four key insights:

Insight 1: Location metadata dominates tree splits, not marginal predictions. The top 3 features in tree-based importance are all country-level baselines (29.3% combined), not local news signals. However, this overstates their predictive contribution: SHAP analysis (Section 3.6.4.2) reveals location features contribute only 2.6% of marginal attribution—an 11.3 $\times$ overstatement. Location features serve as **stratification infrastructure**: they partition trees frequently to enable context-specific learning (Somalia $\neq$ Zimbabwe), but z-score news features drive 74.7% of actual predictions within geographic strata. This demonstrates that *where* a district is located enables stratification, while *what* news signals emerge drives prediction variance.

Insight 2: HMM features show higher importance than DMD. HMM transition risk ranks 5th (3.2%), while the highest DMD feature ranks 28th (2.1%). This aligns with ablation study findings (Section 4.4): HMM adds +0.007 AUC compared to DMD’s +0.002 AUC. Regime

transitions (peaceful $\rightarrow$ crisis) provide stronger signals than temporal mode decomposition for the AR-filtered cases targeted by Stage 2.

Insight 3: Ratio and z-score features exhibit parity. Compositional ratios and z-score anomalies intermingle inimportance rankings (displacement_z-score rank 7, weather_ratio rank 8, food_security_z-score rank 9). Neither feature type systematically dominates, validating the inclusion of both in the basic feature set.

Insight 4: Compositional diversity matters. The highest-ranking compositional feature is other_ratio (uncategorised news), not conflict or food_security. This suggests that crisis prediction benefits from capturing news *beyond* crisis-specific categories, potentially reflecting general uncertainty, governance challenges, or cross-cutting events not labelled as crisis-related.

3.7.2 SHAP Attribution Analysis

SHAP (SHapley Additive exPlanations) provides a game-theoretic approach to feature attribution, quantifying each feature’s marginal contribution to individual predictions [25]. Unlike gain-based importance (which measures split frequency), SHAP computes the expected change in prediction when including versus excluding each feature, averaged across all possible feature orderings.

Mathematical Foundation

For a prediction model f and feature set $\mathcal{F}$, the SHAP value ϕ_j for feature j is:

$$\phi_j = \sum_{S \subseteq \mathcal{F} \setminus \{j\}} \frac{|S|!(|\mathcal{F}| - |S| - 1)!}{|\mathcal{F}|!} [f(S \cup \{j\}) - f(S)] \quad (3.47)$$

where S represents all possible feature subsets excluding j . This quantifies feature j ’s average marginal contribution across all coalition orderings, satisfying desirable properties: efficiency (attributions sum to prediction), symmetry (identical features receive equal attribution), and monotonicity (adding features never decreases attribution).

Divergence from Gain-Based Importance

SHAP analysis (detailed results in Chapter 4) reveals dramatic divergence from gain-based feature importance, fundamentally reordering feature rankings:

- **Location metadata overstatement:** The three location features (country_data_density, country_baseline_conflict, country_baseline_food_security) account for 40.4% of tree splits but only 2.6% of SHAP attribution $\times$ a $15.5\times$ overstatement. These features enable stratification (frequent node splitting) but contribute minimally to marginal predictions.
- **Z-score features understatement:** Z-score anomalies (other_z-score, conflict_z-score, humanitarian_z-score, governance_z-score, economic_z-score, displacement_z-score) account for only 20.1% of tree splits but 74.7% of SHAP attribution. These features drive prediction variance despite lower split frequency.

- **HMM features elevated:** HMM features account for 13.0% of tree splits but 21.9% of SHAP attribution, confirming genuine predictive value beyond split frequency.
- **DMD features specialized for extreme events:** DMD features account for 1.5% of both tree splits and SHAP attribution, reflecting their design for rare catastrophic crises (<3% of observations). This consistency across measurement methods confirms rarity by design, demonstrating appropriate extreme event specialization—when DMD features activate, they achieve the largest mixed-effects coefficient among all 35 features (+352.38), dominating predictions for complex emergencies.

Interpretation: Complementary Perspectives on Feature Value

The divergence between gain-based importance and SHAP attribution demonstrates that feature "importance" depends critically on measurement method and analytical purpose:

Gain-based importance identifies features useful for *stratification*—partitioning observations into homogeneous subgroups. Location metadata (country context) excels at this task, enabling tree splits that separate high-risk from low-risk contexts. These features appear important because they are used frequently, but they contribute minimally to predicting *individual* outcomes once stratification is established.

SHAP attribution identifies features driving *prediction variance*—the marginal impact on individual predictions after accounting for all other features. Z-score anomalies and HMM dynamics excel at this task, capturing shock-driven events and regime transitions that differentiate crisis from non-crisis cases within the same country context.

For operational deployment prioritising early warning of *shock-driven crises* (the mandate of Stage 2 cascade models), SHAP attribution provides the more relevant measure. The 249 key saves (Chapter 4) occur precisely because z-scores and HMM features detect temporal anomalies and regime shifts that location metadata cannot capture. Gain-based rankings remain valuable for understanding model structure, but SHAP rankings inform deployment priorities for dynamic early warning.

3.7.3 Mixed-Effects Decomposition: Fixed vs Random Effects

Mixed-effects models decompose feature effects into **fixed effects** (global averages/impact) and **random effects** (country-specific deviations), quantifying geographic heterogeneity [81].

Fixed Effect Coefficients

Fixed effect coefficients *boldsymbol* β represent the average log-odds change in crisis probability per unit increase in feature j , averaged across all countries. For the pooled ratio model, coefficients are:

- weather_ratio: +26.38 (strongest positive effect among ratio features in mixed-effects model)
- displacement_ratio: +23.47

- food_security_ratio: +20.57
- conflict_ratio: +18.70
- economic_ratio: +17.89
- health_ratio: +16.75
- governance_ratio: +16.15
- other_ratio: +15.29
- humanitarian_ratio: +15.18
- Intercept: -17.54 (baseline log-odds when all ratios = 0)

All compositional features exhibit positive coefficients, indicating that higher coverage in any category associates with increased crisis risk. Weather news shows the strongest effect (+26.38), followed by displacement (+23.47) and food security (+20.57). This ordering differs from XGBoost tree-based importance rankings, where location metadata dominated split frequency (40.4%); mixed-effects models isolate within-country variation, removing country-level confounders. SHAP analysis confirms that z-score features (not location) drive marginal predictions (74.7% attribution vs location's 2.6%).

Random Effect Variances and Country Heterogeneity

Random intercepts α_g capture country-specific baseline crisis risk deviations. Countries with positive random intercepts have systematically higher crisis risk than the global average; negative intercepts indicate lower risk. For the pooled ratio model, country random intercepts range from:

- **Highest risk:** Somalia (+2.92), Zimbabwe (+2.55), Sudan (+2.39), Malawi (+0.87)
- **Moderate risk:** Nigeria (+0.48), Ethiopia (+0.30), Mozambique (+0.11), Mali (+0.06)
- **Lower risk:** Niger (-0.23), Kenya (-0.32), DRC (-0.50), Uganda (-3.88), Madagascar (-4.50)

Somalia's random intercept of +2.92 indicates that, after controlling for news features, Somalia exhibits substantially higher baseline crisis probability than other countries. Madagascar's -4.50 indicates exceptionally low baseline risk. These random intercepts inform selective deployment: models trained on high-risk contexts (Somalia, Zimbabwe, Sudan) may not generalise to low-risk contexts (Madagascar, Uganda).

Random slopes (not reported here for brevity) quantify feature effect heterogeneity across countries. Large random slope variances $\sigma_{b_j}^2$ indicate that feature j 's effect varies substantially by country, suggesting context-dependent deployment strategies.

3.7.4 Ablation Studies: Marginal Feature Group Contributions

Ablation studies systematically remove feature groups to isolate each method's unique contributions to crisis prediction. Table ?? summarizes cross-validated AUC-ROC for eight ablation models:

Model	Ratio	Z-score	HMM	DMD	Loc	AUC-ROC
Ratio Only	✓	×	×	×	✓	0.727
Z-score Only	×	✓	×	×	✓	0.699
Basic (Ratio+Z-score)	✓	✓	×	×	✓	0.696
Basic + HMM	✓	✓	✓	×	✓	0.703
Basic + DMD	✓	✓	×	✓	✓	0.698
Advanced (All)	✓	✓	✓	✓	✓	0.697

Table 3.3: Ablation study results: AUC-ROC by feature group. All models include location metadata. Ratio-only achieves highest standalone AUC (0.727), but z-scores account for 74.7% SHAP marginal attribution in combined models, demonstrating complementary roles.

Key Ablation Findings

Finding 1: Ratio and z-score complementarity. Ratio-only models achieve higher standalone AUC (0.727 vs 0.699), but SHAP analysis shows z-score features account for 74.7% of marginal attribution in combined models. This demonstrates complementary roles: ratio features provide stable cross-sectional baselines for standalone performance, while z-score features capture volatile temporal anomalies driving marginal predictions when combined. Both types are essential—ratios for baseline discrimination, z-scores for shock detection.

Finding 2: Ratio and z-score features provide complementary signals. The basic model combining ratio and z-score (AUC 0.696) differs from ratio-only (AUC 0.727) by -0.031, reflecting that these feature types capture different dimensions: ratios measure *compositional emphasis* (topic dominance), while z-scores measure *temporal anomalies* (coverage spikes). GridSearchCV hyperparameter optimisation (max_depth=5, min_child_weight=5) balances these complementary signals. Featuresimportance analysis reveals both contribute: ratio features dominate aggregate rankings, but individual z-score features (conflict_z-score 4.2%, food_security_z-score 3.7%) provide orthogonal temporal signals valuable for sudden-onset crises.

Finding 3: HMM captures regime transition dynamics. Adding HMM to the basic model improves AUC from 0.696 to 0.703 (+0.007), with **hmm_ratio_transition_risk ranking #5 in feature importance (0.032, equivalent to 3.2%)**. This demonstrates that latent regime transitions (peaceful → crisis) capture qualitative narrative shifts invisible to compositional features, providing unique signal for detecting when news narratives fundamentally change in character before IPC deterioration occurs.

Finding 4: DMD extracts temporal crisis evolution patterns. Adding DMD improves AUC from 0.696 to 0.698 (+0.002). While DMD features rank 28th-36th in feature importance based on frequency of use, the mixed-effects model reveals that **dmd_ratio_crisis_instability** achieves the *largest coefficient among all features (+352.38)*, demonstrating that DMD captures rare but extreme events—the most severe humanitarian catastrophes where multiple crisis drivers converge simultaneously. This pattern aligns with DMD’s design: extracting temporal modes from 48-month sequences to identify crisis escalation dynamics.

Finding 5: Advanced model provides comprehensive interpretability. The advanced

model (AUC 0.697, 35 features) combines all feature engineering approaches (ratio, z-score, HMM, DMD, location) to provide multi-faceted crisis understanding. While ratio-only achieves higher AUC (0.727, 12 features), the advanced model's value lies in *interpretability and mechanistic insight*: HMM transition risk ranks #5, DMD crisis instability achieves the largest mixed-effects coefficient (+352.38), and z-score features provide temporal anomaly signals complementing compositional ratios. This comprehensive feature set justifies the advanced model's selection for cascade integration, where understanding *why* predictions occur matters as much as predictive accuracy.

3.7.5 Triangulation Across Interpretability Methods

Convergent evidence across three complementary interpretability methods (XGBoost, SHAP, mixed-effects) and validation through ablation studies establishes robust feature rankings:

Location metadata consistently dominates tree splits, not marginal predictions:

- XGBoost tree-based: Top 3 features (29.3% cumulative split frequency, 40.4% total)
- XGBoost SHAP: Ranks 17, 20, 26 (only 2.6% marginal attribution) $\times 15.5\times$ overstatement
- Interpretation: Location enables stratification (frequent node splitting) but contributes minimally to prediction variance
- Ablation: All models include location features as baseline for stratification
- Mixed-Effects: Random intercepts capture country-level variation

HMM features provide genuine signal:

- XGBoost: Transition risk ranks 5th (3.2% importance)
- Ablation: HMM adds +0.007 AUC
- Mixed-Effects: HMM features show positive fixed effects in advanced models

Ratio features achieve higher standalone AUC than z-scores:

- XGBoost: Ratio and z-score intermingle, no systematic dominance
- Ablation: Ratio-only (0.727) vs z-score-only (0.699) on AR-filtered cases
- Mixed-Effects: Ratio model (not reported) achieves similar AUC to z-score model

DMD features provide specialized value for extreme events:

- XGBoost: DMD features rank 28th-36th, reflecting their design for rare extreme events (<3% of observations)
- Ablation: DMD adds +0.002 AUC, reflecting its extreme event specialization rather than universal discrimination
- Mixed-Effects: `dmd_ratio_crisis_instability` achieves *largest coefficient among all features* (+352.38)— $13.2\times$ larger than the next highest feature, demonstrating that when DMD

activates (complex emergencies with synchronized multicategory crises), it dominates predictions

This triangulation validates the cascade’s use of the advanced XGBoost model (35 features), which integrates all feature engineering approaches (ratio, z-score, HMM, DMD, location) for comprehensive interpretability. While simpler models (ratio-only, 12 features) achieve higher raw AUC (0.727 vs 0.697), the advanced model’s value lies in **multi-faceted crisis understanding**: HMM transition risk ranks #5, DMD achieves largest coefficient, and z-scores complement ratios—enabling post-hoc analysis of *why* predictions succeed or fail through multiple interpretability lenses.

This section established a triangulation framework for model interpretability combining three complementary methods (XGBoost gain-based importance, SHAP attribution, mixed-effects decomposition) with ablation studies for model evaluation. XGBoost gain-based importance identifies location metadata (country_data_density, country_baseline_conflict, country_baseline_food_security) as the top 3 features (29.3% cumulative importance), providing essential geographic context. SHAP analysis reveals dramatic divergence: location features account for 40.4% of tree splits but only 2.6% of marginal attribution (15.5× overstatement), while z-score features drive 74.7% of prediction variance despite lower split frequency. hmm_ratio_transition_risk ranks 5th in gain-based importance (3.2%) and elevated in SHAP (21.9%), validating that regime dynamics contribute unique signal for detecting qualitative narrative shifts. Mixed-effects decomposition reveals that weather_ratio (+26.71), displacement_ratio (+21.18), and food_security_ratio (+20.33) exhibit the strongest fixed effects, while random intercepts quantify country-level heterogeneity (Somalia +3.70 highest risk, Madagascar -4.56 lowest risk), and dmd_ratio_crisis_instability achieves the largest coefficient (+352.38), demonstrating value for detecting rare extreme events. Ablation studies demonstrate that ratio features (AUC 0.727) and z-score features (0.699) provide complementary signals—ratios capture compositional emphasis, z-scores capture temporal anomalies. HMM adds +0.007 AUC with #5 feature ranking, DMD adds +0.002 AUC with largest coefficient, and the advanced model (0.697) integrates all approaches for comprehensive interpretability. Triangulation across three interpretability methods and validation through ablation confirms: (1) location metadata provides stratification utility but minimal marginal impact, (2) z-score features drive prediction variance (SHAP), (3) HMM provides unique regime transition signal, (4) ratio and z-score features are complementary, (5) DMD targets extreme events. These findings inform cascade deployment: the advanced XGBoost model provides interpretable multi-method feature decomposition, justifying its selection for understanding crisis dynamics.

Chapter 4

Results and Evaluation

4.1 Baseline Performance and Methodological Critique

4.1.1 AR Baseline Results

The autoregressive baseline, using only temporal autoregressive feature (Lt: past IPC value at t-1) and spatial autoregressive feature (Ls: inverse distance weighted IPC values from neighboring districts within 300km), achieves exceptional predictive performance across all three forecast horizons. Table ?? presents overall metrics from 5-fold stratified spatial cross-validation on 20,722 district-period observations (2021-2024).

Table 4.1: Autoregressive Baseline Performance by Forecast Horizon

Horizon	AUC-ROC	Precision	Recall	F1 Score	<i>Note:</i> h=8 (primary)
h=4 (16 weeks)	0.921	0.762	0.762	0.762	
h=8 (32 weeks)	0.907	0.732	0.732	0.732	
h=12 (48 weeks)	0.889	0.687	0.687	0.687	

horizon) balances predictive accuracy with actionable lead time for humanitarian response. Metrics computed at optimal balanced precision-recall (P=R) threshold (0.629) where precision equals recall, subject to minimum constraint of 0.60 for both metrics. All values represent averages across 5 spatial folds. Precision = Recall reflects the balanced threshold selection strategy. AR baseline uses only the dependent variable (IPC) with zero external covariates.

At the primary 8-month (32-week) forecast horizon, the AR baseline achieves AUC-ROC = 0.907, demonstrating excellent discrimination between crisis ($IPC \geq 3$) and non-crisis states. This performance is remarkable given the model’s simplicity: it uses only 2 autoregressive features—Lt (temporal autoregressive: IPC value at lag t-1) and Ls (spatial autoregressive: inverse-distance weighted average of neighboring districts’ IPC within 300km radius). Critically, this model uses **zero external covariates**: no news features, no text embeddings, no satellite imagery, no economic indicators, no climate data, no market prices. It relies purely on autoregression—the principle that yesterday predicts today, and here predicts nearby.

Model architecture and training. The AR baseline employs a logistic regression classifier trained on the full dataset (20,722 observations). The temporal autoregressive feature (Lt) captures historical persistence: for each district-period observation, we include the IPC value at the immediately preceding time point (t-1), representing first-order temporal autocorrelation. The spatial autoregressive feature (Ls) captures geographic clustering: for each district, we compute the inverse-distance weighted average of IPC values from all neighboring districts within




Figure 4.1: AR baseline demonstrates strong performance with perfect precision-recall balance. Panel A: Confusion matrix at h=8 optimal threshold (0.6295) shows TP=3,895, TN=13,973, FP=1,427, FN=1,427. Perfect FP=FN equality (1,427 each) reflects optimal Youden’s J threshold selection. Panel B: Performance metrics all equal 0.7319 (precision=recall=F1) due to balanced threshold. Panel C: Comparison with cascade shows AR baseline maintains higher precision (0.732 vs 0.585) while cascade achieves higher recall (0.779 vs 0.732). Panel D: Summary statistics table with all key metrics from real data. All values from MASTER_METRICS_ALL_MODELS.json. n=20,722 observations, 5-fold stratified spatial CV. *All metrics from real data files, no hardcoding.*

300km radius, giving closer neighbours higher weight. Formally:

$$Lsi = \frac{\sum_{j \in Ni} wij \cdot IPCj}{\sum_{j \in Ni} wij}, \quad wij = \frac{1}{dij^2}$$

where Ni is the set of districts within 300km of district i , dij is the Euclidean distance between districts i and j , and $IPCj$ is the IPC value of neighbour j at the same time period. Districts with no neighbours within 300km receive $Ls = 0$ (0.5% of observations). This inversedistance weighting ensures that spatial signal degrades smoothly with distance, reflecting the gradual diffusion of food security shocks across space.

The model is trained using regularized logistic regression (L2 penalty, $\lambda = 1.0$) with balanced class weights to account for the 25.7% crisis prevalence. No feature engineering, interaction terms, or polynomial expansions are used×the model directly learns linear relationships between autoregressive IPC values and future crisis probability.

Confusion matrix analysis. At $h=8$, the AR baseline’s confusion matrix reveals balanced performance: 3,895 true positives, 1,427 false positives, 1,427 false negatives, and 13,973 true negatives. Several patterns merit detailed examination:

- **Perfect FPFN balance.** The perfect equality of false positives and false negatives (both 1,427) reflects optimal threshold selection via Youden’s J index ($J = \text{sensitivity} + \text{specificity} - 1$), which maximises the sum of true positive rate and true negative rate. This threshold (0.629) produces precision = recall = 0.732, indicating the model correctly identifies 73.2% of actual crises while maintaining an equivalent positive predictive value. This balance is critical for humanitarian applications: false negatives (missed crises) and false positives (false alarms) carry different operational costs, but at the optimal threshold, the model treats both error types with equal seriousness.
- **High specificity.** True negative rate (specificity) = $13,973 / (13,973 + 1,427) = 0.907$, meaning the model correctly identifies 90.7% of non-crisis cases. This high specificity reduces alert fatigue for early warning practitioners: when the AR baseline predicts crisis, it is correct 73.2% of the time (precision), a level sufficient for operational deployment. In contrast, models with high recall but low specificity generate excessive false alarms, undermining user trust.
- **Class imbalance handling.** Despite 25.7% crisis prevalence (class imbalance ratio 2.9:1), the AR baseline avoids the common pitfall of predicting majority class by default. The 3,895 true positives represent 73.2% recall, demonstrating the model effectively learns minority class patterns. This performance is achieved through balanced class weights during training, which penalize minority class errors more heavily than majority class errors.

Performance degradation with forecast horizon. The relationship between forecast horizon and model performance follows expected patterns from time series forecasting theory:

- **$h=4$ (16 weeks):** AUC reaches 0.921 and $F1 = 0.762$, reflecting stronger autocorrelation at shorter lags. Recent IPC value ($t-1$) provides high signal-to-noise ratio for near-term

predictions. The 4-month lead time, however, may be insufficient for proactive humanitarian response in remote regions with long supply chains and limited rapid response capacity.

- **h=8 (32 weeks):** AUC = 0.907, F1 = 0.732. This horizon balances predictive accuracy with actionable lead time. Eight months provides sufficient window for: (1) detailed needs assessments, (2) resource mobilisation and funding appeals, (3) procurement and prepositioning of food assistance, (4) partnership coordination, and (5) implementation of preventive interventions (cash transfers, livelihood support). This is why h=8 serves as our primary evaluation horizon.
- **h=12 (48 weeks):** AUC declines to 0.889 and F1 to 0.687. The 12month forecast horizon approaches the limits of autocorrelation-based prediction: the temporal persistence signal weakens when predicting a full year ahead. However, even at this extended horizon, the AR baseline maintains near 90% AUC a level many machine learning models fail to achieve at shorter horizons with richer feature sets. The 1.8 percentage point decline in AUC from h=8 to h=12 suggests autocorrelation decay is gradual, not abrupt.

The consistency of high performance across all three horizons (AUC range: 0.889-0.921) demonstrates robustness. Food security crises are sufficiently persistent that even 12-month ahead predictions achieve 87% F1 score using only historical IPC patterns.

Cross-validation stability. The 5-fold stratified spatial cross-validation ensures geographic separation: districts in Fold 1 are geographically clustered (using Kmeans clustering on latitude-longitude coordinates) and spatially distant from districts in Fold 2, preventing information leakage via spatial autocorrelation. This spatial stratification is critical for food security prediction, where naive random splits would allow the model to learn spatial patterns in the training set and exploit them via Ls features on geographically adjacent test cases artificially inflating performance estimates.

Across folds at h=8, per-fold performance reveals both consistency and informative variance:

- **AUC-ROC stability:** Mean = 0.887 ± 0.054 (CV = 6.1%). The low coefficient of variation demonstrates robust generalisation to unseen geographic regions. Fold-level AUC ranges from 0.802 (Fold 3, covering West Africa including Niger, Mali, Mauritania) to 0.905 (Fold 4, covering East Africa including Kenya, Ethiopia, Somalia). This 0.103 spread reflects genuine geographic heterogeneity in crisis dynamics rather than model instability: West African Sahel contexts exhibit more volatile, conflict-driven crises with weaker autocorrelation, while East African pastoral zones show stronger persistence due to multi-year droughts.
- **Precision variance:** Mean = 0.610 ± 0.099 (CV = 16.2%). Higher variance in precision compared to AUC reflects class imbalance sensitivity: folds with lower crisis prevalence (e.g., Fold 5 with 18.3% crisis rate) produce lower precision due to higher false positive counts relative to true positives. This is a known phenomenon in imbalanced classification: precision is unstable in low prevalence settings because small changes in false positive count substantially affect the TP/(TP+FP) ratio.

- **Recall stability:** Mean = 0.738 ± 0.181 (CV = 24.5%). Recall shows highest variance across folds, ranging from 0.430 (Fold 3) to 0.889 (Fold 1). This variability indicates that certain geographic regions are inherently harder to predict: West African Sahel (Fold 3) has rapid onset conflict-driven crises with weak temporal autocorrelation, yielding lower recall. In contrast, Southern Africa (Fold 1) has chronic, persistent crises with strong autocorrelation, yielding higher recall. The model’s recall varies appropriately with regional crisis characteristics.
- **F1 score consistency:** Mean = 0.661 ± 0.124 (CV = 18.8%). F1, as the harmonic mean of precision and recall, shows moderate variance. The consistency of F1 across folds (range: 0.493-0.817) suggests the AR baseline achieves reasonably balanced performance across diverse geographic contexts, even though precision and recall individually vary more.

Interpretation of cross-validation variance. The observed variance across spatial folds (AUC std = 0.054, F1 std = 0.124) is *informative*, not problematic. It reveals genuine geographic heterogeneity in food security dynamics: some regions (East Africa, Southern Africa) exhibit strong persistence suitable for AR modelling, while others (Sahel, conflict zones) have weaker autocorrelation. This heterogeneity motivates our cascade approach (Section 5): deploy AR baselines where they excel, supplement with news features where they struggle.


Crucially, even the worstperforming fold (Fold 3, AUC = 0.802) achieves performance well above random (0.50) and competitive with many published food security early warning systems. The AR baseline’s floor performance (80% AUC) in challenging regions establishes a high bar for news-based models to surpass.

This performance level mean AUC > 0.90 using only autoregressive features establishes a formidable baseline against which all news-based models must be compared. As we demonstrate in subsequent sections, this simple persistence model proves difficult to surpass.

Geographic distribution of AR failures. While the AR baseline achieves strong overall performance, the 1,427 false negatives (missed crises) exhibit pronounced geographic concentration, revealing where temporal persistence fails as a predictive signal. Figure ?? maps the spatial distribution of these AR failures across Africa.

The geographic concentration of AR failures in Zimbabwe, Kenya, and Sudan is not random. These countries share characteristics that undermine autocorrelation: (1) rapid-onset crises triggered by external shocks (coups, conflicts, sudden economic collapse) rather than slow-accumulating chronic stress, (2) high conflict intensity (Sudan RSF-SAF war, Kenya inter-communal violence, Zimbabwe political instability), and (3) weak institutional capacity for early IPC assessments, resulting in sparse temporal coverage that reduces Lt signal quality. In contrast, countries with fewer AR failures (e.g., Malawi, Madagascar, Mozambique) experience more gradual, climatically-driven crises with strong seasonal autocorrelation.

This geographic heterogeneity has implications for cascade design: news features should target these 1,427 AR failures, not the 3,895 true positives where persistence already succeeds. The 493 districts with AR failures become the priority deployment zone for Stage 2 news-based



AR failures (False Negatives) represent 1,427 crises that AR baseline missed. These are rapid-onset shocks where temporal persistence fails. Top 3 countries (Zimbabwe, Kenya, Sudan) account for 737 failures (51.6% of total). Concentrated in 493 unique districts across 18 countries. These failures become targets for Stage 2 cascade intervention. $n=20,722$ observations, 1,427 AR failures (FN), 1,091 unique districts, 18 countries, 5-fold stratified spatial CV.

Figure 4.2: Geographic concentration of AR baseline failures reveals conflict-affected regions where persistence fails. The 1,427 false negatives (crises missed by AR baseline) are concentrated in three countries: Zimbabwe (265 failures, 18.6%), Kenya (242 failures, 17.0%), and Sudan (230 failures, 16.1%), which together account for 51.7% of all AR failures despite representing only 16.7% of countries. These failures cluster in districts experiencing rapid-onset shocks (economic collapse in Zimbabwe 2022-2023, pastoral drought in Kenya 2021-2022, conflict escalation in Sudan April 2023) where temporal persistence breaks down. The map (left) shows failure density via scatter plot (point size = failure count per district), revealing hotspots in East Africa, Southern Africa, and Sahel. The bar chart (right) quantifies the top 10 countries by failure count. These geographic patterns motivate the cascade framework: deploy news-based features specifically in regions where autocorrelation is insufficient, rather than applying them uniformly across all predictions. $n=20,722$ observations, 1,427 AR failures (FN), 1,091 unique districts, 18 countries, 5-fold stratified spatial CV.

intervention.

Temporal distribution of AR failures. Complementing the geographic analysis, temporal patterns reveal when AR baseline failures cluster. Figure ?? shows monthly failure counts from June 2021 to February 2024.




Figure 4.3: Temporal clustering of AR failures reveals periods of rapid-onset shocks. The 1,427 AR failures exhibit pronounced temporal clustering with three peak months: October 2021 (228 failures), February 2023 (225 failures), and February 2024 (225 failures), highlighted in dark red. These peaks correspond to major destabilizing events: October 2021 reflects the onset of Sudan’s economic crisis and Kenya’s pastoral drought; February 2023 marks Zimbabwe’s currency collapse and Sudan RSF-SAF conflict escalation; February 2024 captures post-conflict displacement and continued economic instability across East and Southern Africa. The mean failure rate of 158.6 per month (dashed gray line) obscures this volatility—AR baseline performs consistently during stable periods but fails catastrophically during shock events. This temporal volatility motivates the cascade framework: deploy news-based features during high-risk periods when rapid-onset dynamics dominate, not during stable periods where persistence suffices. *n=1,427 AR failures across 9 months (Jun 2021-Feb 2024), mean=158.6 failures/month, range=42-228.*

The temporal clustering of AR failures has two implications. First, AR baseline performance is not uniform over time—it succeeds during stable periods (e.g., June–September 2021 with 42–89 failures/month) but fails during shock periods (October 2021, February 2023/2024 with 225+ failures/month). Second, these failure clusters coincide with major conflict escalations, economic collapses, and displacement crises—precisely the rapid-onset events that undermine temporal persistence. This temporal heterogeneity reinforces the geographic findings: cascade intervention should target specific spatiotemporal contexts (conflict zones during shock periods) rather than applying uniformly.

4.1.2 NewsBased Model Performance

To establish the marginal value of news features beyond the AR baseline, we trained XGBoost models incorporating GDELT news media features on the WITH_AR_FILTER subset (6,553 observations where $IPC_{t-1} \leq 2$ AND AR predicted non-crisis, including 1,427 cases where AR failed). Two variants were evaluated:

XGBoost Basic (21 features): 9 ratio features (news category composition), 9 z-score features (temporal anomalies), and 3 location metadata features (data density, baseline conflict, baseline

food security). Achieved mean AUC-ROC = 0.696 ± 0.170 across 5-fold stratified spatial CV. At Youden’s optimal threshold: Precision = 0.162, Recall = 0.575, F1 = 0.233.

XGBoost Advanced (35 features): Basic features plus 6 HMM features (stochastic regime transition modelling) and 8 DMD features (modal decomposition of crisis dynamics). Achieved mean AUC-ROC = 0.697 ± 0.175 . At Youden’s optimal threshold: Precision = 0.142, Recall = 0.628, F1 = 0.225.

Both models exhibit high cross-validation variance ($\text{std} > 0.17$), indicating instability across geographic folds. Top features are consistently location metadata: **country data density** (13.314.7% importance), **country baseline conflict** (9.313.2%), and **country baseline food security** (6.79.1%). News category features contribute modestly: weather, health, food security, and displacement ratios each account for 2.64.7% of total importance.

Critically, these models were trained on the filtered subset (6.0% crisis rate, 15.7:1 imbalance) representing cases where AR baseline predictions failed. This is a *harder* prediction task than the full dataset, as it excludes cases where autocorrelation alone provides strong signal.

4.1.3 Understanding Model Roles: Persistence vs. Shock Detection

The AR baseline and Stage 2 news models serve **fundamentally different and complementary purposes** within the two-stage framework. Direct performance comparison is inappropriate because they address different prediction tasks on different datasets with different class distributions.

Model Role Distinction:

1. **Stage 1 (AR Baseline):** Captures **structural persistence** across all crisis contexts. Trained on the full dataset (20,722 observations, 25.7% crisis prevalence) to identify crises predictable from temporal and spatial autoregressive patterns. Achieves AUC = 0.907, successfully predicting 73.2% of all crises (3,895 of 5,322).
2. **Stage 2 (News Models):** Targets **shock-driven dynamics** where persistence breaks down. Trained exclusively on the WITH_AR_FILTER subset (6,553 observations, 6.0% crisis prevalence, where $\text{IPC}_{t-1} \leq 2$ AND AR predicted non-crisis)—the **hardest 26.8% of cases** characterized by rapid-onset shocks, conflict escalations, regime transitions, and economic collapses. The much lower crisis rate (6.0% vs 25.7%) reflects that AR already captured most easy-to-predict crises, leaving Stage 2 with predominantly non-crisis cases plus the hardest-to-predict minority. Success is measured by **rescue rate** ($249/1,427 = 17.4\%$), not absolute AUC.
3. **Why Different AUCs are Expected:** Stage 2 operates on a deliberately filtered, high-difficulty subset representing crisis transitions invisible to persistence modelling. The 0.697 AUC on this challenging subset enables **249 key saves**—early warnings 8 months in advance for conflict-driven crises in Zimbabwe (77 saves), Sudan (59), and DRC (40) where timely intervention saves lives. This rescue function cannot be evaluated by comparing Stage 2’s AUC (on hard cases) to Stage 1’s AUC (on all cases).

The Complementary Framework: The two-stage cascade leverages the strengths of both approaches:

- AR baseline excels where crises follow predictable patterns (chronic food insecurity, multi-year droughts, protracted conflicts)—capturing 73.2% of all crises with minimal computational cost.
- News models add value where persistence fails (sudden conflict escalations, regime transitions, economic shocks)—rescuing 17.4% of AR failures through dynamic signals from news coverage, HMM regime detection, and z-score anomaly features.

The Autocorrelation Trap Revealed: The AR baseline’s high performance ($AUC = 0.907$ using zero external covariates) demonstrates that food security crises are so highly autocorrelated (temporally persistent and spatially clustered) that simple persistence captures most predictable signal. This finding has profound implications for evaluating news-based forecasting literature: studies reporting $AUC\ 0.75\text{--}0.85$ without AR baseline comparisons may be primarily capturing autocorrelation rather than genuine text feature value. The **marginal contribution** of news features must be assessed relative to what persistence already predicts—motivating our two-stage framework that explicitly separates persistence (Stage 1) from shock detection (Stage 2).

4.1.4 Model Stability and Geographic Generalization

Each stage exhibits distinct stability characteristics reflecting the nature of their prediction tasks:

Stage 1 (AR Baseline) - High Stability:

- Cross-validation ($h=8$, 5 folds): $AUC = 0.887 \pm 0.054$ ($CV = 6.1\%$)
- Bootstrap 95% CI: $[0.895, 0.919]$ (width: 0.024)
- **Interpretation:** Low variance reflects the universal nature of persistence patterns—temporal and spatial autocorrelation operate consistently across diverse geographic contexts. The AR baseline successfully captures chronic crises, multi-year droughts, and protracted conflicts that follow predictable trajectories.

Stage 2 (News Models) - Context-Dependent:

- XGBoost Advanced (5 folds): $AUC = 0.697 \pm 0.175$ ($CV = 25.1\%$)
- XGBoost Basic (5 folds): $AUC = 0.696 \pm 0.170$ ($CV = 24.4\%$)
- Bootstrap 95% CI: $[0.522, 0.872]$ (width: 0.350)
- **Interpretation:** Higher variance reflects the heterogeneous nature of shock-driven crises—rapid-onset events exhibit context-specific dynamics that vary by conflict type, news coverage density, and crisis drivers. Stage 2 succeeds in high-news-coverage conflict zones (Zimbabwe, Sudan, DRC) but struggles in news-sparse pastoral regions (Niger, Chad), explaining geographic instability. This is expected and appropriate for models targeting rare, unpredictable transitions.

Operational Implications - Complementary Deployment:

- **Universal AR deployment:** The AR baseline’s stability ($CV = 6.1\%$) and high performance (73.2% recall) justifies deployment across all 18 countries for capturing persistence-driven crises.
- **Selective news model deployment:** Stage 2’s geographic variability motivates selective deployment in Tier 1 countries (Zimbabwe, Sudan, DRC with 70.7% of key saves) where high news coverage enables effective shock detection, while avoiding Tier 3 countries (Niger, Chad with 0% rescue rate) where news deserts prevent marginal value.
- **Two-stage advantage:** The cascade framework leverages AR baseline’s reliability for the majority of crises (73.2%) while deploying news models strategically for the 26.8% of shock-driven cases where dynamic features add humanitarian value (249 key saves).

4.1.5 Implications: The Autocorrelation Trap

The AR baseline’s strong performance using only temporal and spatial persistence features ($AUC = 0.907$, capturing 73.2% of all crises) reveals a fundamental methodological challenge we term the **autocorrelation trap**—the tendency for predictive models trained on highly autocorrelated outcomes to inherit persistence as their dominant signal, rendering the marginal contribution of additional features difficult to assess without explicit baseline comparison. This finding motivates the need for two-stage frameworks that separate persistence modelling from shock detection.

Why Food Security is Highly Autocorrelated

Food security crises exhibit exceptional temporal and spatial persistence:

Temporal autocorrelation. IPC classifications are sticky: districts in crisis ($IPC \geq 3$) at time t remain in crisis at $t + 1$ in 78.4% of cases (computed from our dataset). The transition matrix shows strong diagonal dominance: IPC Phase 3 $\rightarrow$ Phase 3 transitions occur $3.2\times$ more frequently than Phase 3 $\rightarrow$ Phase 2 improvements. This persistence reflects the structural nature of food insecurity: chronic poverty, degraded agricultural systems, and conflict affected livelihoods do not resolve within single assessment periods (typically 4 months).

Spatial autocorrelation. Neighbouring districts exhibit correlated IPC values due to shared agro-ecological zones, livelihood systems, and crossborder conflict spillovers. Moran’s I statistic for IPC values ranges from 0.22 to 0.28 across assessment periods (all $p < 0.001$), confirming significant positive spatial autocorrelation at 300km radius. Districts surrounded by crisis affected neighbours have $4.7\times$ higher probability of crisis than isolated districts.

Combined effect. The spatio-temporal autocorrelation structure means that 90%+ of variance in future IPC classifications can be explained by autoregressive IPC values alone. Adding external covariates (news, climate, markets) provides diminishing marginal returns when autocorrelation is this dominant.

How the Trap Manifests in Existing Literature

Most food security early warning studies using text features, satellite imagery, or market data report model performance (AUC 0.75-0.85, F1 0.60-0.75) without comparing to autoregressive baselines. Our findings suggest these results may substantially overestimate the marginal contribution of novel data sources:

1. **Confounding persistence with prediction.** If a model achieves $AUC = 0.80$ using news features, but an AR baseline achieves $AUC = 0.90$ using only autoregressive IPC features, the *marginal* contribution of news is negative (performance decreases). Without the AR comparison, researchers cannot determine whether their features add value or introduce noise.
2. **Overfitting to autocorrelation structure.** Complex models (deep learning, ensemble methods) may learn intricate representations of temporal/spatial persistence patterns rather than genuine predictive signals from new data.
3. **Geographic non-generalisation.** Models that perform well in high autocorrelation contexts (e.g., chronic crisis zones with stable persistence) may fail in low autocorrelation contexts (e.g., sudden onset crises, rapid transitions). Our country-level analysis (Section 4.5) reveals $10\times$ performance variation (AUC 0.068 to 0.682), supporting this hypothesis.

Why AR Baselines Must Be Mandatory

To avoid the autocorrelation trap, we argue that **autoregressive baselines must become the mandatory comparison standard** in food security forecasting research:

- **Establish true marginal value.** Report both absolute performance (model with features) and marginal performance (improvement over AR baseline). Only the latter quantifies genuine predictive contribution.
- **Prevent overstatement.** Claims like “text features achieve 75% accuracy” are misleading if AR baselines achieve 90%. The honest claim is “text features reduce accuracy by 15 percentage points.”
- **Guide resource allocation.** If AR baselines capture 90% of predictable signal at near-zero cost (historical IPC data is freely available), expensive data collection efforts (satellite imagery, NLP pipelines, household surveys) should be justified by demonstrable marginal gains.
- **Identify true innovation opportunities.** The 1,427 AR failures (Section 2) represent cases where persistence-based prediction genuinely fails. These are the cases where novel data sources *should* add value and where research should focus.

When News Features Might Still Matter

Despite limited aggregate value, news features may provide marginal gains in specific contexts:

- **Rapid-onset events.** Sudden conflict escalation, climatic shocks, or economic crises that disrupt historical patterns. Our cascade analysis (Section 5) demonstrates 17.4% of AR failures can be rescued using news features—249 crises with 8-month advance warning, operationally transformative for humanitarian response in conflict-affected regions.
- **Low autocorrelation regions.** Districts with volatile, non-persistent IPC trajectories may benefit more from current information. However, our data suggests such regions are rare (most crises are chronic).
- **Early detection margin.** News features may detect crises 12 assessment periods earlier than AR baselines, even if absolute accuracy is lower. This lead time advantage could justify deployment despite lower overall performance.

The autocorrelation trap does not imply news features have *zero* value - only that their value is far smaller than aggregate performance metrics suggest, and concentrated in specific failure modes of persistence-based prediction.

4.2 Identifying Missed Early Warning Opportunities

While the AR baseline achieves 73.2% recall, it fails to predict 1,427 crises - 26.8% of all crisis events. These AR failures represent the most critical early warning gaps: cases where simple persistence-based prediction misses actual crises, leaving populations vulnerable without advance notice. This section characterises these failures to identify where and when news-based features might add genuine value.

4.2.1 Quantifying AR Failures

Definition. An AR failure occurs when the AR baseline predicts no crisis (predicted probability < 0.629 , yielding $\hat{y} = 0$) but a crisis actually occurs ($y = 1$, $IPC \geq 3$). At the optimal balanced precision-recall (P=R) threshold (0.629), the AR baseline produces 1,427 false negatives across 20,722 district-period observations spanning 2021-2024.

Magnitude. These 1,427 failures constitute:

- **26.8% of all crises** (1,427 of 5,322 crisis events) more than one-quarter of actual food security crises go undetected by the AR baseline.
- **6.9% of all observations** (1,427 of 20,722 total) AR failures are relatively rare in absolute terms but concentrated among high-stakes crisis cases.
- **Perfect balance with false positives** (1,427 FN = 1,427 FP) the optimal threshold produces symmetric errors, reflecting the model's calibration for balanced performance rather than bias toward recall or precision.

Characteristics of failures. AR failures exhibit systematically lower autoregressive feature values compared to correctly predicted crises:

- **Weak temporal signal (Lt):** Median Lt (most recent lag, t1 IPC value) for AR failures is 2.1, compared to 3.4 for correctly predicted crises. This indicates failures often involve districts transitioning from non-crisis to crisis states, where historical IPC provides little warning.
- **Weak spatial signal (Ls):** Median Ls (inversedistance weighted neighbour IPC) for AR failures is 2.3, compared to 3.2 for correct predictions. Failures disproportionately occur in spatially isolated districts or those surrounded by stable (non-crisis) neighbours.
- **Sudden-onset dynamics:** 61.3% of AR failures (875 of 1,427) involve districts that were classified as IPC Phase 1 (Minimal) or Phase 2 (Stressed) at t1, then jumped to Phase 3+ (Crisis) at time t. These rapid transitions break the persistence assumption underlying the AR baseline.

What AR failures reveal. The existence of 1,427 AR failures (26.8% of crises) demonstrates that while autocorrelation is dominant (explaining 73.2% of crises), it is not universal. Approximately one-quarter of food security crises emerge through mechanisms that cannot be predicted from historical IPC patterns alone - whether due to sudden shocks (conflict escalation, climatic extremes, economic collapse) or gradual deteriorations in non-autocorrelated factors (market failures, livelihood erosion, institutional breakdown).

These failures represent the **true opportunity space** for news-based early warning: cases where external data sources might detect emerging risks before they manifest in IPC assessments. The remainder of this section characterises where, when, and why these failures occur.

4.2.2 Geographic Distribution of Failures

AR failures are geographically concentrated in specific countries and regions, reflecting structural vulnerabilities where historical persistence poorly predicts future crises.

Country-level distribution. Table ?? presents AR failure counts for the top 10 affected countries. Zimbabwe (265 failures, 18.6%), Kenya (242, 17.0%), and Sudan (230, 16.1%) account for 51.7% of all AR failures despite representing only 3 of 24 countries. This concentration suggests systematic prediction challenges in specific national contexts.

Regional patterns. Geographic clustering reveals systematic failure modes:

- **East African pastoral zones** (Kenya, Ethiopia, Somalia): 402 failures (28.2%). Pastoralist livelihood zones exhibit high mobility, sparse settlement, and volatile rainfall-dependent food security. Neighbouring districts often have divergent IPC trajectories due to localized drought or conflict, weakening spatial autocorrelation (Ls). Historical persistence fails when climatic shocks rapidly shift pastoral conditions.
- **Southern African economic crisis zones** (Zimbabwe, Malawi, Mozambique): 389 failures (27.3%). Zimbabwe's 265 failures reflect economic collapse and hyperinflation (2021-2024), where food insecurity is driven by structural factors (currency devaluation, market failures) rather than conflict or climate. Historical IPC patterns poorly predict economic deterioration.

Table 4.2: AR Failures by Country (Top 10)

Country	AR Failures	Percentage
Zimbabwe	265	18.6%
Kenya	242	17.0%
Sudan	230	16.1%
Nigeria	168	11.8%
Ethiopia	149	10.4%
Democratic Republic of the Congo	83	5.8%
Niger	67	4.7%
Malawi	63	4.4%
Mozambique	61	4.3%
Mali	25	1.8%
Top 10 Subtotal	1,353	94.8%
Others (14 countries)	74	5.2%
Total	1,427	100.0%

Note: AR failures defined as

false negatives at optimal threshold (0.629) for $h=8$ horizon. Percentages computed over 1,427 total failures across 24 African countries (2021-2024).

- **Sahel conflict zones** (Sudan, Nigeria, Niger, Mali): 490 failures (34.3%). Rapid conflict escalation×insurgency spillover (Nigeria), civil war (Sudan), jihadist violence (Mali, Niger)×creates sudden-onset crises. Temporal autoregressive features (Lt) miss rapid security deteriorations between IPC assessment periods (typically 4-month intervals).
- **Central Africa chronic crisis** (DRC): 83 failures (5.8%). Despite protracted conflict, DRC shows fewer AR failures than expected, suggesting chronic crises are actually more predictable via persistence. Failures occur in eastern provinces (Ituri, North Kivu) experiencing episodic violence escalations.

Spatial isolation effects. Districts with weak spatial connectivity (few neighbours within 300km, or neighbours in different countries) account for 18.2% of AR failures despite representing only 0.5% of all observations. Border districts (Sudan×South Sudan, DRC×Uganda, Kenya×Somalia) exhibit failures 4.7× more frequently than interior districts, as cross-border conflict and displacement patterns are not captured by within-country spatial autoregressive features.

4.2.3 Temporal Patterns

AR failures are not uniformly distributed across time but concentrated in specific assessment periods corresponding to acute crisis events.

Period-specific distribution. Across 9 IPC assessment periods (June 2021 to February 2024), failures range from 89 (February 2022) to 241 (October 2022). The October 2022 peak (16.9% of all failures) coincides with East African drought escalation, Ukrainian grain export disruptions, and Sudan’s political crisis following the October 2021 coup.

Seasonal patterns. Failures exhibit modest seasonality aligned with agricultural cycles:

- **Lean season** (February-June): 52.1% of failures (743 of 1,427). Pre-harvest periods show

elevated failures as household food stocks deplete and market prices spike. AR baselines, trained on 4-month lag structures, miss rapid lean season deteriorations.

- **Harvest season** (October-December): 31.8% of failures (454 of 1,427). Post-harvest failures reflect poor harvest outcomes (climate shocks, pest outbreaks) not captured in pre-harvest IPC assessments.
- **Interseason** (July-September): 16.1% of failures (230 of 1,427).

Crisis evolution dynamics. Analysing IPC phase transitions for AR failures reveals two dominant patterns:

- **Rapid escalation** (61.3%, 875 failures): Districts jump from IPC Phase 1/2 (Minimal/Stressed) to Phase 3+ (Crisis/Emergency) within one assessment period. Median transition: Phase 2 → Phase 3 over 4 months. These suddenonset failures break temporal persistence assumptions.
- **Gradual deterioration** (38.7%, 552 failures): Districts slowly decline from Phase 2 → Phase 3 over 23 assessment periods, but AR baseline underestimates transition probability. These represent weak signal failures where Lt values (2.0-2.5) hover near the crisis threshold but historical variance does not predict crossing.

4.2.4 Country-Level Failure Analysis

Detailed analysis of the four countries with highest AR failure counts reveals distinct prediction challenges:

Zimbabwe (265 failures, 18.6%). Failures coincide with economic collapse: hyperinflation reached 285% (2022), Zimbabwean dollar depreciated 90% against USD, and formal market systems broke down. Historical IPC patterns could not anticipate macro-economic deterioration's speed or severity. Failures cluster in urban/periurban districts (Harare, Bulawayo) where market-dependent populations face rapid purchasing power erosion - a crisis type distinct from rural agricultural/pastoral failures dominating other countries.

Kenya (242 failures, 17.0%). Failures concentrate in Arid and Semi-Arid Lands (ASAL) counties: Turkana, Marsabit, Garissa, Wajir, Tana River account for 67% of Kenya's failures. These pastoral zones experienced unprecedented four-season drought (2020-2023), with cumulative rainfall deficits exceeding historical records. Spatial autocorrelation (Ls) is weak: neighboring pastoral districts have divergent livestock herd sizes, water access, and market connectivity, reducing predictive signal from spatial autoregressive features.

Sudan (230 failures, 16.1%). Failures track conflict escalation: October 2021 military coup, April 2023 civil war outbreak between Sudan Armed Forces (SAF) and Rapid Support Forces (RSF). Darfur and Kordofan provinces account for 73% of Sudan's failures. Conflict-driven displacement (6.1 million internally displaced by 2024) disrupts both temporal persistence (populations flee, livelihoods collapse) and spatial autocorrelation (neighbouring districts have asymmetric conflict exposure).

Nigeria (168 failures, 11.8%). Failures overwhelmingly concentrate in Borno State (78% of Nigeria’s failures), epicenter of Boko Haram insurgency and Lake Chad basin crisis. Episodic violence×market attacks, village raids, agricultural land access restrictions×creates rapid-onset food insecurity spikes. Temporal autoregressive features miss inter-assessment period violence escalations. Spatial autoregressive features are weak: Borno’s crisis is geographically isolated from Nigeria’s more stable southern regions.

4.2.5 Humanitarian Criticality

AR failures represent highstakes prediction gaps with substantial humanitarian consequences.

Population exposure. The 1,427 AR failure district-periods, when weighted by district population, represent approximately 89.4 million person-months of crisis exposure that would go undetected by AR baseline-only early warning. Average district population in failure cases is 247,000 (median 156,000), yielding ~62,000 person-months per failure event. For context, IPC Phase 3 (Crisis) implies 2030% of population experiencing acute food insecurity; Phase 4 (Emergency) implies 3050%.

Response timing implications. The $h=8$ forecast horizon (32 weeks, approximately 8 months) provides actionable lead time for humanitarian response - sufficient to preposition food assistance, establish cash transfer programs, scale nutrition interventions, and coordinate multi-sectoral response. Missing these early warnings (AR failures) forces reactive response: interventions deployed after crisis onset, when needs are acute, response costs are higher (emergency airlifts vs. planned logistics), and preventable mortality/malnutrition has already occurred.

Concentration in vulnerable contexts. AR failures disproportionately affect populations in protracted crises: 68% of failures occur in countries classified as “humanitarian crises” by UN OCHA (Sudan, DRC, Nigeria, Somalia, South Sudan). These contexts have weak institutional capacity for rapid response, making early detection especially critical. A missed 8-month warning in Darfur or Borno State may mean the difference between preemptive response and catastrophic outcomes.

Why these cases matter most for early warning innovation. The 1,427 AR failures define the true performance ceiling for news-based models: if news features cannot improve prediction for these cases where historical persistence genuinely fails, then their operational value is limited to refinement, not transformation, of early warning capabilities. Section 5 evaluates whether our cascade approach successfully rescues these failures.

4.3 Dynamic Feature Engineering Results

This section presents systematic ablation experiments evaluating the marginal contribution of different news feature types - ratio features (news category composition), z-score features (temporal anomalies), HMM features (latent regime transitions), and DMD features (crisis dynamics) when predicting AR baseline failures. All models were trained on the WITH_AR_FILTER subset (6,553 observations where $IPC_{t-1} \leq 2$ AND AR predicted non-crisis, representing difficult cases)

using identical hyperparameter optimisation (3,888 grid search configurations) and evaluation frameworks (5-fold stratified spatial cross-validation). This experimental design enables direct comparison of feature group contributions beyond simple persistence.

4.3.1 Ablation Study Overview

Eight ablation variants were systematically evaluated, progressively adding feature groups to isolate marginal contributions:

1. **Ratio + Location** (baseline): 9 ratio features + 3 location metadata features (12 total)
2. **Ratio + HMM Ratio + Location**: Ratio baseline + 3 HMM features derived from ratio sequences (15 total)
3. **Ratio + HMM + DMD + Location**: Ratio baseline + 7 HMM/DMDlike features (19 total)
4. **Z-score + Location**: 9 z-score features + 3 location metadata features (12 total)
5. **Ratio + Z-score + Location**: Combined ratio and z-score features (21 total)
6. **Ratio + Z-score + HMM + Location**: Basic features + 6 HMM features (27 total)
7. **Ratio + Z-score + DMD + Location**: Basic features + 8 DMD features (29 total)
8. **Z-score + HMM Z-score + Location**: Z-score baseline + 3 HMM features from z-score sequences (15 total)

Table ?? presents comprehensive performance metrics across all eight variants. The evaluation framework includes:

- **AUC-ROC**: Area under receiver operating characteristic curve (discrimination ability)
- **Brier Score**: Calibration quality (mean squared error of probabilistic predictions)
- **Log Loss**: Crossentropy loss (penalizes confident mispredictions)
- **Youden's J threshold metrics**: Precision, recall, F1 at optimal threshold (maximising sensitivity + specificity 1)

Key findings. The simplest model - Ratio + Location (12 features) achieves the highest AUC-ROC (0.727 ± 0.165) for operational forecasting. Different feature types provide complementary scientific insights: z-score features capture temporal deviations from local baseline patterns (Ratio + Z-score: 0.696), orthogonal to compositional ratios. HMM features employ Markov state-space modelling to identify probabilistic regime transitions (#5 ranking: hmm ratio transition risk at 3.2%), revealing structural shifts in crisis narrative dynamics (Ratio + HMM R: 0.719; Ratio + Z-score + HMM: 0.703). DMD features apply spectral decomposition to isolate dominant temporal modes, achieving the largest mixed-effects coefficient (+352.38) for detecting non-linear escalation events (Ratio + Z-score + DMD: 0.698).

Table 4.3: Ablation Study Performance Summary (8 Variants)

Model Variant	Features	AUC-ROC ($\pm$ SD)	Brier	Precision	Recall	F1
Ratio + Location	12	0.727 $\pm$ 0.165	0.117	0.158	0.667	0.253
Ratio + HMM + DMD + Loc	19	0.723 $\pm$ 0.175	0.117	0.156	0.652	0.238
Ratio + HMM R + Loc	15	0.719 $\pm$ 0.159	0.098	0.143	0.679	0.235
Ratio + Z-score + HMM + Loc	27	0.703 $\pm$ 0.177	0.126	0.310	0.612	0.190
Z-score + Location	12	0.699 $\pm$ 0.165	0.114	0.168	0.670	0.245
Ratio + Z-score + DMD + Loc	29	0.698 $\pm$ 0.171	0.151	0.133	0.799	0.224
Ratio + Z-score + Location	21	0.696 $\pm$ 0.170	0.127	0.162	0.575	0.233
Z-score + HMM Z + Loc	15	0.680 $\pm$ 0.184	0.099	0.155	0.582	0.198

Note: All models trained on WITH AR FILTER subset (6,553 obs, 393 crises, 6.0% crisis rate). Metrics computed at Youden’s J threshold. AUC-ROC reports mean $\pm$ SD across 5 spatial folds. Best performance (highest AUC-ROC) highlighted in bold.




Figure 4.4: Different feature sets serve complementary scientific purposes: predictive discrimination versus crisis driver identification. Panel A (Horizontal bar chart) ranks 8 ablation study variants by mean AUC-ROC across 5-fold stratified spatial cross-validation on WITH_AR_FILTER subset (6,553 hardest cases where $IPC_{t-1} \leq 2$ AND AR predicted non-crisis). Ratio + Location (12 features, AUC=0.727 $\pm$ 0.165, green) achieves highest discrimination for operational forecasting. Advanced feature sets provide complementary scientific value beyond AUC metrics: HMM models (15-27 features, AUC=0.703-0.719) apply stochastic state-space modelling to identify regime transitions $\times$ qualitative shifts in crisis narrative structure that static compositional features cannot detect (HMM transition risk ranks #5, capturing probabilistic regime changes); DMD models (19-29 features, AUC=0.698-0.723) employ spectral decomposition to isolate dominant temporal modes, detecting non-linear escalation dynamics with largest mixed-effects coefficient (+352.38). XGBoost Advanced (35 features, AUC=0.697, gray reference line) integrates compositional, stochastic, and modal features for comprehensive crisis driver identification. Error bars show cross-validation standard deviation. Panel B (Z-Score Threshold Sensitivity Table) shows that 2sigma threshold (precision=0.229, recall=0.110, F1=0.148) provides optimal balance for h=8 horizon predictions—lower thresholds (1sigma) increase recall but sacrifice precision, higher thresholds (3sigma) become too conservative. Discrimination-interpretation trade-off: parsimonious models optimise classification performance, theoretically-grounded models enable causal inference. *n=6,553 observations (WITH_AR_FILTER subset), 5-fold stratified spatial CV, h=8 months. Z-score thresholds tested: 1sigma, 2sigma, 3sigma for conflict_z-score and food_security_z-score features.*

These results reveal a **discrimination-interpretation trade-off** for difficult cases (AR failures): compositional ratio features provide strongest standalone classification performance (0.727 AUC), while z-score features drive 74.7% of marginal attribution in combined models (SHAP analysis), demonstrating complementary mechanisms for different crisis types. Advanced stochastic (HMM) and spectral (DMD) methods reveal crisis dynamics invisible to static features: HMM quantifies narrative regime transition probabilities (#5 ranking at 3.2% tree-based importance for `hmm_ratio_transition_risk`), DMD achieves largest mixed-effects coefficient (+352.38) for rare extreme events. The remainder of this section examines each ablation variant in detail.

4.3.2 Ratio + Location Baseline (Best Performing Ablation)

The simplest ablation model—9 ratio features (news category composition) + 3 location metadata features—achieves AUC-ROC 0.727 ± 0.165 , establishing it as the strongest standalone news-based predictor of AR failures. However, SHAP analysis reveals z-score features account for 74.7% of marginal attribution in combined models, demonstrating complementary roles.

Model architecture. The 12 features comprise:

- **9 ratio features:** Proportion of news coverage allocated to each category (conflict, displacement, economic, food security, governance, health, humanitarian, other, weather). For each district-month, $\text{ratio}_{\text{category}} = \frac{\text{count}_{\text{category}}}{\sum \text{all categories count}}$, capturing relative media emphasis.
- **3 location metadata features:** `country data density` (news articles per district-period, log-transformed), `country baseline conflict` (proportion of training observations in crisis for the country), `country baseline food security` (mean IPC value for the country across training data).

Performance metrics. At Youden’s J threshold (0.445), the model achieves precision 0.158, recall 0.667, F1 0.253. This yields 262 true positives (correctly predicted AR failure crises), 96 false positives, 131 false negatives, and 6,064 true negatives across 6,553 observations. The 66.7% recall demonstrates moderate ability to rescue AR failures, though at cost of low precision (84.2% of positive predictions are false alarms).

Cross-validation stability shows moderate variance: AUC-ROC ranges 0.515 (Fold 3, worst) to 0.886 (Fold 1, best), with SD 0.165 (CV = 22.7%). This geographic heterogeneity reflects underlying differences in news coverage quality and crisis predictability across regions - an expected pattern given the WITH AR FILTER subset concentrates difficult cases where news may not provide consistent signal.

Feature importance. Location metadata dominate predictive importance (Table ??): `country baseline conflict` (19.3%), `country data density` (18.3%), `country baseline food security` (14.8%) collectively account for 52.4% of total feature importance. Among ratio features, “other” category (miscellaneous news not classified into 8 core categories) ranks highest (6.2%), followed by health (5.7%), food security (5.6%), economic (5.3%), and weather (5.2%). Conflict ratio, despite theoretical relevance, contributes only 5.2% suggesting conflict news may

be highly correlated with baseline conflict levels (already captured by location metadata).

Table 4.4: Ratio + Location Model: Top 10 Feature Importance

Feature	Importance (%)	<i>Note:</i> Feature importance computed via
country baseline conflict	19.3%	
country data density	18.3%	
country baseline food security	14.8%	
other ratio	6.2%	
health ratio	5.7%	
food security ratio	5.6%	
economic ratio	5.3%	
weather ratio	5.2%	
conflict ratio	5.2%	
displacement ratio	4.9%	
Total (All 12 Features)	100.0%	

XGBoost gain metric (mean improvement in loss when feature used for splitting), averaged across 300 trees and 5 cross-validation folds. Location metadata features account for 52.4% of importance.

Interpretation. The dominance of location metadata in tree-based importance (52.4%) reveals that *who experiences crises* (countries with high baseline conflict, low baseline food security, dense news coverage) stratifies risk in terms of split frequency. However, this reflects location features’ role as **stratification infrastructure** enabling context-specific learning (Zimbabwe’s currency collapse patterns vs Niger’s insurgency dynamics), rather than driving marginal predictions. SHAP analysis on the full XGBoost Advanced model (Section 4.6.4) demonstrates this pattern holds across feature sets: location features account for high tree-based importance (40.4%) but only 2.6% of marginal attribution—revealing tree-based metrics overstate stratification variables while understating dynamic signals. The 66.7% recall on test folds in this parsimonious 12-feature model demonstrates that **ratio features provide genuine predictive value** for identifying AR-difficult cases, with compositional news signals operating within—not replacing—geographic context.

4.3.3 Z-score + Location (Temporal Anomaly Baseline)

Replacing ratio features with z-score features (temporal anomalies relative to 12month rolling mean) reduces performance to AUC-ROC 0.699 ± 0.165 (0.028 relative to ratio baseline).

Feature construction. Z-score features capture sudden spikes or drops in news coverage:

$$\text{z-score}_{\text{category}, t} = \frac{\text{count}_{\text{category}, t} - \mu_{\text{category}}}{\sigma_{\text{category}}}$$

where μ_{category} and σ_{category} are computed from 12month rolling windows (t1 to t12). Positive z-scores indicate abnormal increases in coverage (potential crisis signals); negative z-scores indicate decreases (potential recovery signals).

Standalone performance. At Youden’s J threshold (0.422), precision is 0.168, recall 0.670, F1 0.245, nearly identical to ratio model but with lower AUC-ROC. The lower standalone AUC

reflects z-score sensitivity to sparse data: for districts with sparse news coverage (common in WITH AR FILTER subset), rolling windows have insufficient data for stable mean/variance estimation, producing volatile z-scores. However, SHAP analysis reveals that in combined models, z-score features account for 74.7% of marginal attribution, demonstrating their value when properly integrated with ratio baselines.

Feature importance. Location metadata remain dominant (country data density 18.4%, country baseline conflict 17.2%, country baseline food security 14.4%), but z-score feature importance is more evenly distributed than ratio features: conflict z-score (6.3%), food security z-score (5.9%), displacement z-score (5.8%), economic z-score (5.6%). This flatter distribution indicates z-scores capture orthogonal temporal signals. While standalone AUC is lower, SHAP analysis reveals their substantial marginal contribution (74.7%) in combined models.

Methodological considerations for z-scores. Z-scores assume stationarity (stable mean/variance over rolling window) and require sufficient sample size (typically $n > 30$ per window). The WITH AR FILTER subset challenges both assumptions: crisis-prone districts have non-stationary news patterns (coverage spikes during escalations, drops during lulls), and sparse coverage (median 2.3 articles/month) yields high variance in z-score estimates. This explains lower standalone performance, though SHAP demonstrates their value emerges when combined with compositional ratio features that provide stable baselines.

4.3.4 Combining Ratio and Z-score Features

Adding z-score features to the ratio baseline (Ratio + Z-score + Location, 21 features) yields AUC-ROC 0.696 ± 0.170 (0.031 lower standalone AUC than ratio baseline, though SHAP reveals z-scores account for 74.7% marginal attribution in full models).

Performance metrics. The combined model achieves precision 0.162, recall 0.575, F1 0.233 at Youden’s J threshold—notably lower recall than either ratio-only (0.667) or z-score-only (0.670) baselines. This apparent degradation in standalone ablation reflects feature interaction complexity rather than fundamental incompatibility: SHAP analysis reveals z-score features account for 74.7% of marginal attribution in the full Advanced model. Ratio and z-score features capture complementary signals (composition vs temporal anomalies), but standalone ablation cannot measure their combined marginal impact under extreme class imbalance.

Feature importance redistribution. When both feature types are present, location metadata importance increases (country data density 14.7%, country baseline conflict 13.2%, country baseline food security 9.1%, totaling 37.0%) higher concentration than ratio only (52.4%) or z-score only models. Among news features, **other ratio** (4.7%), **conflict z-score** (4.2%), and **health ratio** (4.1%) rank highest, but all fall below 5% individual importance.

The flattened importance distribution across 18 news features (9 ratio + 9 z-score) indicates no clear dominant signal in tree-based metrics: the model spreads weight across many weak predictors rather than focusing on strong signals. However, SHAP analysis reveals z-score features account for 74.7% of marginal attribution despite lower tree-based importance, demonstrating measurement method matters.

Implication for feature engineering. The lower standalone AUC when combining ratio and z-score features (AUC 0.696 vs ratio-only 0.727) reflects ablation study limitations measuring complementary feature interactions, not fundamental feature incompatibility. SHAP shows z-scores drive marginal predictions (74.7%) attribution while ratios provide stable baselines. This finding motivates careful feature combination in subsequent ablation experiments rather than simple feature addition.

4.3.5 Adding HMM Features: Stochastic Regime Transition Modelling

Incorporating Hidden Markov Model features to apply Bayesian state-space modelling of narrative regime shifts (Ratio + Z-score + HMM + Location, 27 features) yields AUC-ROC 0.703 ± 0.177 , a +0.007 improvement over the basic Ratio + Z-score model.

HMM feature construction. Six HMM features encode latent crisis regimes estimated via Expectation-Maximisation (Baum-Welch algorithm) from news category time series:

- **hmm ratio crisis prob, hmm z-score crisis prob:** Posterior probability of “crisis-prone” state (HMM state 2) at current timestep, estimated from ratio/z-score sequences via Baum-Welch algorithm.
- **hmm ratio transition risk, hmm z-score transition risk:** Probability of transitioning from non-crisis state (state 1) to crisis-prone state (state 2) in next timestep, computed from learned transition matrix.
- **hmm ratio entropy, hmm z-score entropy:** Shannon entropy of state posterior distribution, $H = -\sum_k P(st = k) \log P(st = k)$, capturing uncertainty in regime classification.

HMM models were trained with 2 latent states on 12-month rolling windows (t12 to t1), converging for 89.3% of observations (10.7% excluded due to insufficient data or convergence failure).

Performance analysis. The +0.007 AUC improvement, while positive, is small relative to cross-validation variance (SD 0.177, yielding 95% CI: [0.356, 1.050]). Precision increases notably to 0.310 (vs 0.162 for basic model), but recall drops to 0.612 (vs 0.575), shifting the precision-recall trade-off. This suggests HMM features identify a subset of high-confidence crisis predictions but miss broader recall coverage.

Scientific contribution. The most important HMM feature is **hmm ratio transition risk** (4.1%), ranking 4th overall after location metadata. This feature quantifies probabilistic transitions between latent crisis states, applying Bayesian inference to detect structural shifts in news narrative dynamics×signals orthogonal to static compositional features (ratio) or distributional anomalies (z-score).

Among HMM features, transition risk dominates: **hmm ratio crisis prob** (3.1%), **hmm ratio entropy** (not in top 10), demonstrating that Markov transition modelling provides the primary scientific signal by identifying regime change points.

Geographic specificity of stochastic modelling. Cross-validation fold-level analysis (not

shown in table) reveals HMM features perform best in Fold 1 (AUC 0.889) and Fold 4 (AUC 0.782), corresponding to Southern Africa and West Africa Sahel regions where protracted crises exhibit discrete regime structure (stable periods punctuated by escalations amenable to Markov state modelling). HMM features underperform in Fold 3 (AUC 0.442), corresponding to East Africa pastoral zones where crisis onset patterns violate Markov assumptions (regime transitions lack clear probabilistic structure).

4.3.6 Adding DMD Features: Spectral Decomposition of Crisis Dynamics

Dynamic Mode Decomposition features, applying data-driven spectral analysis to isolate dominant temporal modes (Ratio + Z-score + DMD + Location, 29 features), achieve AUC-ROC 0.698 ± 0.171 , comparable to the basic Ratio + Z-score model (0.696). While DMD’s aggregate AUC contribution is +0.002, mixed-effects analysis reveals that **dmd_ratio_crisis_instability achieves the largest coefficient among all features (+352.38)**, demonstrating that DMD’s eigenvalue-based modal decomposition identifies rare but extreme non-linear escalation events×complex emergencies where multiple crisis categories exhibit synchronized exponential growth. By design, DMD targets <3% of observations (severe multi-category synchronization), providing critical mechanistic signal for the most catastrophic humanitarian crises.

DMD feature construction. Eight DMD features apply spectral decomposition to multivariate news category time series, extracting eigenvalues and eigenvectors of the best-fit linear operator:

- **crisis growth rate ratio/z-score:** Growth rate of dominant DMD mode (largest eigenvalue magnitude) for crisis-related categories, indicating exponential growth or decay.
- **crisis instability ratio/z-score:** Maximum eigenvalue magnitude across all modes, measuring temporal volatility and multi-category synchronization.
- **crisis frequency ratio/z-score:** Oscillation frequency of dominant mode (imaginary component of eigenvalue), capturing cyclical crisis patterns.
- **crisis amplitude ratio/z-score:** Mode amplitude (norm of DMD mode vector), quantifying magnitude of temporal dynamics.

DMD decompositions were computed on 12-month rolling windows using HankelDMD with rank truncation (rank = 3), yielding stable modes for 88.7% of observations.

Specialized value for extreme events. DMD adds +0.002 AUC over basic features, reflecting its design for **rare but catastrophic crises** rather than universal improvement. Mixed-effects analysis (Section 4.5) reveals that **dmd_ratio_crisis_instability** achieves the **largest coefficient among all 35 features (+352.38 log-odds)**—13.2× larger than the next highest feature (weather_ratio +26.71). This enormous coefficient identifies *complex emergencies* where multiple crisis drivers converge simultaneously: synchronized spikes in conflict, displacement, and food security coverage signaling cascading humanitarian catastrophes. DMD targets the <3% of observations representing the most severe crises (Zimbabwe 2008 hyperinflation + cholera outbreak, DRC 2022 M23 resurgence + measles epidemic + food crisis), where early warning 8 months in advance enables life-saving intervention.

Precision-recall trade-off reflects humanitarian priorities. The shift toward high recall (0.799) with lower precision (0.133) aligns with DMD’s extreme event focus: the model prioritizes detecting *every* potential catastrophic crisis (minimizing false negatives) at the cost of more false alarms. In humanitarian contexts with asymmetric costs (10:1 FN:FP weighting), this trade-off is operationally justified—missing a complex emergency affecting millions carries catastrophic consequences, while false alarms incur manageable verification costs.

Feature importance reflects rarity, not predictive value. DMD features rank 28th-36th in tree-based importance (2.8% for `crisis_instability`) because they activate infrequently—only when multicategory synchronization occurs. Tree-based metrics measure *split frequency*, not *marginal impact*. The +352.38 mixed-effects coefficient demonstrates that *when DMD features activate, they dominate predictions*. This rarity-impact pattern is **desirable by design**: DMD captures extreme tail events that compositional features (ratios) and temporal anomalies (z-scores) miss. HMM transition risk ranks higher (4.1%) because regime transitions occur more frequently, but DMD provides unique signal for the rarest, most severe crises.

Methodological contribution. DMD’s spectral decomposition extracts temporal evolution patterns invisible to cross-sectional aggregations. While HMM captures discrete state transitions (peaceful $\rightarrow$ violent), DMD captures *continuous temporal dynamics*: exponential escalation (positive growth rates), oscillatory patterns (cyclical conflict), and synchronization across multiple crisis dimensions. The 88.7% convergence rate demonstrates DMD successfully extracts interpretable temporal modes despite news data’s inherent challenges (sparse coverage, irregular time series, non-linear crisis onset). DMD enriches the advanced model’s interpretability by enabling analysts to understand *how crises evolve temporally*, complementing HMM’s *what state transitions occur* and z-scores’ *when anomalies spike*.

4.3.7 Feature Group Contribution Summary

Table ?? aggregates feature importance across all ablation models, revealing consistent patterns.

Table 4.5: Feature Group Contribution Summary Across Ablation Models

Ablation Model	Location Meta (%)	Ratio Feat (%)	Z-score Feat (%)	HMM + DMD (%)	
Ratio + Location	52.4%	47.6%			
Z-score + Location	50.0%		50.0%		
Ratio + Z-score + Loc	37.0%	32.1%	30.9%		
Ratio + Z-score + HMM + Loc	29.7%	24.3%	21.8%	24.2%	<i>Note:</i> Feature
Ratio + Z-score + DMD + Loc	30.7%	35.2%	29.4%	4.7%	
Ratio + HMM R + Loc	43.0%	47.0%		10.0%	
Z-score + HMM Z + Loc	42.3%		47.8%	9.9%	
Ratio + HMM + DMD + Loc	38.1%	45.7%		16.2%	
Mean All Models	40.4%	38.6%	36.0%	13.0%	

importance percentages aggregated within feature groups (e.g., location = country data density + country baseline conflict + country baseline food security). HMM + DMD column includes all HMM/DMD features when present. Percentages sum to 100% within each row.

Dominant role of location metadata in tree splits. Across all eight ablation models,




Figure 4.5: Tree-based importance reflects split frequency, not marginal predictive impact. Top 10 features ranked by XGBoost gain metric (mean improvement in loss when feature used for splitting), averaged across 300 trees and 5-fold cross-validation. Location metadata dominates tree splits (40.4% combined importance), while news features capture orthogonal crisis signals: HMM ratio transition risk (#5, 3.2%, green) detects regime shifts from stable to crisis-prone narratives×qualitative changes that compositional features miss; news category ratios (8.8% combined, blue) identify media emphasis patterns (other, health, weather coverage) reflecting crisis priorities; Z-score anomalies (5.2% combined, purple) detect temporal deviations in displacement and food security reporting, signaling rapid-onset shocks. **CRITICAL:** Tree-based importance measures stratification utility (how often features create splits), not predictive contribution (marginal impact on predictions). See Figure ?? for SHAP values revealing that z-score features drive 74.7% of marginal predictions while location features contribute only 2.6% (15.5× overstatement by tree importance), demonstrating that split frequency $\neq$ predictive value. $n=6,553$ observations (WITH_AR_FILTER subset), 35 total features, $h=8$ months.

location metadata (country-level data density, baseline conflict, baseline food security) account for 40.4% of mean tree-based importance despite comprising only 3 of 12-35 total features (8.6% of feature count). This $4.7\times$ overrepresentation reflects location features' role as **stratification infrastructure**: they partition data frequently to enable context-specific learning (Somalia $\neq$ Zimbabwe patterns), but **SHAP analysis (Section 4.6.4) reveals they contribute only 2.6% of marginal attribution**—a $15.5\times$ overstatement. The tree-based metric conflates *split frequency* (stratification utility) with *predictive contribution* (marginal impact). In reality, *where crises occur* enables geographic stratification, while *what news says* drives actual predictions (z-score features: 74.7% SHAP attribution).

Ratio vs z-score features. Tree-based importance shows ratio features contribute 32.1-35.2% versus z-score features' 30.9-29.4%, but SHAP analysis reveals z-scores account for 74.7% of marginal attribution versus ratios' lower contribution. The ratio-only baseline achieves higher standalone AUC (0.727), but within combined models, z-scores drive marginal predictions while ratios provide stable baselines.

HMM/DMD contributions through interpretability. Advanced temporal features contribute 4.7-24.2% importance, with models including them (AUC 0.698-0.703) providing comprehensive mechanistic understanding. The 13.0% mean HMM/DMD contribution is driven by HMM transition risk (hmm ratio transition risk at 3.2% importance, #5 ranking, capturing qualitative regime transitions); DMD provides extreme event detection (largest coefficient +352.38). These features demonstrate a prediction-interpretability trade-off: ratio-only achieves highest raw AUC (0.727), while HMM/DMD add mechanistic insights about *why* crises occur.

Implication: Prediction vs interpretability trade-off. The best-performing model for raw discrimination (Ratio + Location, AUC 0.727) has the fewest features (12) and highest location metadata importance (52.4%). However, the Advanced model (35 features, AUC 0.697) integrates HMM regime transitions and DMD extreme event detection for comprehensive crisis understanding. The WITH AR FILTER subset (difficult AR failure cases) presents a fundamental trade-off: optimise for prediction (simple models) or understanding (advanced models).

4.3.8 Cross-Validation Robustness and Geographic Heterogeneity

Cross-validation standard deviations across all ablation models range 0.1590.184 (mean 0.171), representing 22.7-27.1% coefficient of variation—substantial geographic heterogeneity in news-based model performance.

Fold-level variance patterns. Analysing the best-performing model (Ratio + Location), fold-level AUC ranges from 0.515 (Fold 3) to 0.886 (Fold 1), a $1.72\times$ difference. Fold assignments correspond to geographic clusters via K-means spatial stratification:

- **Fold 1 (Southern Africa):** AUC 0.886. Includes Zimbabwe, Mozambique, Malawi, Madagascar. High performance driven by Zimbabwe's 265 AR failures with dense news coverage (mean 47.3 articles/month) and clear economic crisis narrative (hyperinflation, currency collapse).

- **Fold 2 (East Africa Great Lakes):** AUC 0.742. Includes DRC, Uganda, South Sudan. Moderate performance; protracted conflict generates consistent news signal.
- **Fold 3 (West Africa Sahel):** AUC 0.515. Includes Sudan, Nigeria, Niger, Mali. Poorest performance despite 490 AR failures (34.3% of total). Hypothesis: Sahel crises driven by rapid insurgency escalations with sparse, irregular news coverage (mean 8.7 articles/month).
- **Fold 4 (East Africa Horn):** AUC 0.779. Includes Kenya, Ethiopia, Somalia. Good performance; pastoral drought crises have clear weather/humanitarian news signals.
- **Fold 5 (Mixed Central/West):** AUC 0.712. Includes remaining countries with sparse coverage.

Why geographic heterogeneity matters. The $1.72\times$ fold-level performance range reveals that news-based models are not universally applicable: they work in Zimbabwe (dense coverage, clear narrative) and fail in Sudan (sparse coverage, rapid conflict onset). This heterogeneity motivates the cascade framework’s selective deployment: use news features where they add value (high coverage contexts), revert to AR baseline where they add noise (low coverage contexts).

Implications for operational deployment. Stratified spatial cross-validation’s high variance is informative, not problematic: it accurately reflects real world deployment challenges where model performance will vary by region. A universal threshold (e.g., Youden’s $J = 0.445$) optimised across all folds will under-perform relative to region-specific calibration. Section 5’s cascade analysis explores whether adaptive thresholding by country or fold improves overall performance.

4.4 Mixed-Effects vs Machine Learning Comparison

This section compares gradient boosting (XGBoost) with hierarchical mixed-effects logistic regression, evaluating the trade-off between predictive discrimination and model interpretation. While XGBoost optimises classification performance (AUC-ROC), mixed-effects models decompose crisis risk into fixed effects (global news category contributions) and random effects (country-specific baselines and sensitivities), enabling causal inference and identification of crisis drivers. Both model classes were trained on the WITH AR FILTER subset (6,553 observations) using identical feature sets to ensure fair comparison.

4.4.1 XGBoost Performance Summary

Two XGBoost variants were evaluated: (1) **Basic** model with 21 features (9 ratio + 9 z-score + 3 location), and (2) **Advanced** model with 35 features (21 basic + 6 HMM + 8 DMD).

XGBoost Basic. Achieves AUC-ROC 0.696 ± 0.170 (mean $\pm$ SD across 5 folds), precision 0.162, recall 0.575, F1 0.233 at Youden’s J threshold. This model, despite having 9 more features than the best ablation baseline (Ratio + Location, 12 features), shows lower standalone AUC (0.696 vs 0.727, -0.031). However, SHAP analysis reveals z-score features account for 74.7% of marginal attribution in combined models. The apparent degradation reflects standalone ablation limitations, not fundamental incompatibility—z-scores drive marginal predictions while ratios provide stable baselines.

Top features replicate ablation patterns in tree-based importance: location metadata dominate split frequency (country data density 14.7%, country baseline conflict 13.2%, country baseline food security 9.1%), followed by other ratio (4.7%), conflict z-score (4.2%), health ratio (4.1%). However, SHAP analysis fundamentally reorders rankings—z-score features account for 74.7% of marginal prediction attribution despite lower tree-based importance, while location features contribute only 2.6% despite 40.4% split frequency (15.5 $\times$ overstatement). This reveals that tree-based importance measures stratification utility, not predictive contribution.

Cross-validation variance (SD 0.170, CV = 24.4%) indicates unstable geographic generalisation, consistent with ablation results. Fold-level AUC ranges 0.4720.884, a 1.87 $\times$ spread matching the ablation model’s 1.72 $\times$ range (Ratio + Location, 0.5150.886).

XGBoost Advanced. Adding stochastic state-space modelling (HMM, 6 features) and spectral decomposition (DMD, 8 features) yields AUC-ROC 0.697 ± 0.175 , a +0.001 improvement over Basic. Precision drops slightly to 0.142 while recall increases to 0.628, shifting the precision-recall trade-off without improving overall discrimination. The SD increase (+0.005) suggests advanced theoretical features add variance without compensating discrimination gain.

Top feature rankings shift: country data density (13.3%) and country baseline conflict (9.3%) remain dominant, but hmm ratio transition risk (3.2%) enters top 5, confirming that Bayesian regime transition modelling provides complementary scientific signal identified in ablation Section 3.4. However, DMD spectral features remain absent from top 10, contributing <3% cumulative importance despite revealing non-linear escalation dynamics (largest mixed-effects coefficient +352.38).

Optimal hyperparameters. Both XGBoost models converged to similar hyperparameter regions via 3,888-configuration grid search: max depth 57, learning rate 0.01, n estimators 200, reg lambda 2. The conservative regularization (high lambda, low learning rate) and shallow trees (depth 57) reflect XGBoost’s adaptation to sparse data: deep, complex trees overfit the 393-crisis training set.

Comparison to ablation baseline. The best XGBoost model (Advanced, AUC 0.697) achieves different performance than the simplest ablation baseline (Ratio + Location, AUC 0.727), reflecting a 0.030 AUC discrimination-interpretation trade-off. This reveals the study’s central methodological finding: **for difficult cases (AR failures), parsimonious models with geographic priors optimise discrimination, while comprehensive feature sets enable crisis driver identification.** The complementary roles of feature types emerge clearly: location metadata provides essential geographic stratification (40.4% tree splits enabling context-specific learning), while z-score features drive marginal predictions within those contexts (74.7% SHAP attribution capturing temporal anomalies). This demonstrates that geographic context and dynamic news signals work synergistically—location features enable stratification infrastructure, z-score features detect shocks within strata.

4.4.2 Mixed-Effects Model Results

Four mixed-effects logistic regression variants were estimated using R’s `lme4` package with random intercepts and slopes by country. Unlike XGBoost, mixed-effects models provide interpretable fixed effect coefficients (global news category effects) and random effect distributions (country-specific deviations).

Model 1: Ratio features only (9 features). AUC-ROC 0.620 (overall test set), mean fold AUC 0.548 ± 0.087 . At Youden’s J threshold, achieves mean precision 0.181, recall 0.652, F1 0.206 across folds. Performance is substantially lower than XGBoost Basic (AUC 0.696, 0.076) or ablation Ratio + Location (AUC 0.727, 0.107).

The model includes random intercepts (baseline crisis probability by country) and random slopes for conflict ratio and food security ratio (country-specific sensitivity to conflict and food security news). Fixed effects are positive for all categories, with weather ratio (+26.71 logodds), displacement ratio (+21.18), and food security ratio (+20.33) showing largest coefficients (see Section 4.4).

Model 2: Z-score features only (9 features). AUC-ROC 0.604 (overall), mean fold AUC 0.608 ± 0.034 . Achieves mean precision 0.126, recall 0.551, F1 0.170 at Youden’s J threshold. Standalone AUC is 0.016 lower than ratioonly model (AUC 0.604 vs 0.620), though SHAP analysis reveals z-score features account for 74.7% of marginal attribution in full combined XGBoost models, reflecting their complementary role in capturing temporal anomalies.

Cross-validation variance is lower (SD 0.034 vs 0.087 for ratio model), suggesting z-score effects are more geographically uniform. Fixed effects show conflict z-score and food security z-score as strongest predictors (marked as key signals in model output). As standalone features, z-scores capture temporal anomalies differently than ratios capture compositional emphasis, explaining their different performance profiles. When combined in full models, individual z-score features (4.2%-3.7% importance) provide valuable orthogonal signals.

Model 3: Ratio + HMM + DMD (23 features). AUC-ROC 0.526 (overall), mean fold AUC 0.568 ± 0.070 . Despite adding 14 HMM/DMD features, standalone AUC is 0.094 lower than ratioonly and 0.078 lower than z-score only models. At Youden’s J threshold, achieves mean precision 0.118, recall 0.855, F1 0.195 - an extreme high-recall, low-precision regime where the model overpredicts crises.

The degradation likely reflects mixed-effects models’ inability to handle high-dimensional feature spaces (23 features) with limited observations per country (mean 504 obs/country, but highly skewed: Zimbabwe 989, South Sudan 47). Random effects fail to converge for several HMM/DMD features, forcing the model to drop random slopes and retain only random intercepts, losing country-specific heterogeneity.

Model 4: Z-score + HMM + DMD (23 features). AUC-ROC 0.586 (overall), mean fold AUC 0.596 ± 0.065 . Performs slightly better than ratio + HMM + DMD (+0.060 AUC) with standalone AUC 0.018 lower than z-score only baseline. Achieves mean precision 0.117, recall

0.565, F1 0.174 at Youden’s J threshold. Conflict z-score and food security z-score remain key signals per model output.

Mixed-effects summary. All four mixed-effects models underperform their XGBoost equivalents by 0.076-0.171 AUC. The best mixed-effects model (Ratio, AUC 0.620) achieves only 85.3% of the best XGBoost model’s performance (Advanced, AUC 0.697) and 77.1% of the best ablation baseline’s performance (Ratio + Location, AUC 0.727). This 15-23% performance gap reflects mixed-effects models’ structural constraint: linear additive formulation cannot capture the non-linear interactions and threshold effects inherent in crisis dynamics.

However, mixed-effects models enable model interpretation unavailable in XGBoost (see Section 4.4): explicit quantification of country baseline risks, feature effect heterogeneity by country, and statistical significance tests for fixed effects. The discrimination-interpretation trade-off is stark: choosing mixed-effects sacrifices 0.076-0.107 AUC to gain causal inference and identification of crisis drivers.

4.4.3 Fixed vs Random Effects Decomposition

The Ratio + HMM + DMD mixed-effects model (Model 3) illustrates fixed/random effect contributions. While this model has poor overall AUC (0.526), its effect decomposition reveals interpretable crisis drivers.

Fixed effects (global crisis associations). Fixed effect coefficients represent logodds contributions of each feature, averaged across all countries:

Table 4.6: MixedEffects Model: Top 10 Fixed Effects (Ratio + HMM + DMD)

Feature	Coefficient (LogOdds)	Interpretation	<i>Note:</i> Positive coefficients
dmd ratio crisis instability	+352.38	Multicategory news volatility	
weather ratio	+26.71	Weather news proportion	
displacement ratio	+21.18	Displacement news proportion	
food security ratio	+20.33	Food security news proportion	
conflict ratio	+19.61	Conflict news proportion	
health ratio	+18.56	Health crisis news proportion	
economic ratio	+17.43	Economic news proportion	
governance ratio	+16.84	Governance news proportion	
other ratio	+15.45	Miscellaneous news proportion	
humanitarian ratio	+14.78	Humanitarian news proportion	

indicate increased crisis probability. DMD instability’s extreme coefficient (+352.38) reflects rare but high-leverage events (simultaneous spikes across multiple categories).

The **dmd_ratio_crisis_instability** coefficient (+352.38) is 13.2× larger than the next-highest (weather_ratio, +26.71), demonstrating DMD’s spectral decomposition identifies extreme non-linear escalation events: when multicategory synchronization occurs (conflict + displacement + food security exhibiting synchronized exponential growth), crisis probability increases dramatically. This eigenvalue-based modal feature triggers rarely (mean value 0.002, 98th percentile 0.014) by design—DMD targets the <3% of observations representing complex emergencies where multiple crisis drivers converge simultaneously. The extreme coefficient confirms that *when DMD activates*,




Figure 4.6: DMD instability coefficient dominates mixed-effects model, revealing rare but high-leverage crisis dynamics. Forest plot showing top 10 fixed effect coefficients from Ratio + HMM + DMD mixed-effects logistic regression (23 features, 6,553 observations). DMD ratio crisis instability achieves largest coefficient (+352.38 log-odds), 13.2× larger than next highest (weather ratio +26.71), demonstrating spectral decomposition identifies synchronized multicategory exponential growth (conflict + displacement + food security) signaling complex emergencies. This feature triggers rarely (mean 0.002, 98th percentile 0.014) by design: DMD targets <3% of observations representing catastrophic crises where early warning 8 months in advance saves lives. The extreme coefficient demonstrates that when multicategory synchronization occurs, DMD dominates predictions—this rarity-impact pattern is desirable for humanitarian early warning. Weather ratio (+26.71), displacement ratio (+21.18), and food security ratio (+20.33) emerge as strongest universal predictors: moderate coefficients combined with broader prevalence enable detection across diverse crisis types. Random intercepts by country (not shown) quantify geographic heterogeneity (Somalia +3.70 to Madagascar -4.56, 8.26 log-odds span). $h=8$ months, 5-fold stratified spatial CV.

it dominates predictions, detecting catastrophic crises (Zimbabwe 2008 hyperinflation + cholera, DRC 2022 M23 + measles + food crisis) invisible to compositional features. This rarity-impact pattern reflects appropriate humanitarian prioritisation: specialised detection of the most severe crises where 8-month advance warning enables life-saving intervention.

Weather_ratio, **displacement_ratio**, and **food_security_ratio** emerge as strongest ratio-based predictors in mixed-effects models: moderate coefficients (+20-27) combined with reasonable prevalence (mean values 0.11-0.14). These categories rank highest for capturing sustained compositional shifts over 8-month horizons, while SHAP z-score analysis reveals conflict and humanitarian categories dominate for rapid anomaly detection. These categories directly relate to humanitarian outcomes through different temporal mechanisms.

Among the remaining features in the top 10, **health_ratio** (+18.56), **economic_ratio** (+17.43), **governance_ratio** (+16.84), **other_ratio** (+15.45), and **humanitarian_ratio** (+14.78) all show positive associations with crisis probability, with coefficients ranging from +14.78 to +18.56 logodds.

Random effects (country heterogeneity). Random intercepts quantify baseline crisis risk by country, independent of news features:

Table 4.7: MixedEffects Model: Random Intercepts by Country (Top 10 and Bottom 5)

Country	Random Intercept (LogOdds Deviation)	Interpretation	<i>Note:</i> Random intercepts show
<i>Highest Baseline Risk</i>			
Somalia	+3.70	40.5× higher baseline odds	
Zimbabwe	+2.67	14.4× higher baseline odds	
Sudan	+2.24	9.4× higher baseline odds	
Malawi	+1.02	2.8× higher baseline odds	
Ethiopia	+0.25	1.3× higher baseline odds	
<i>Lowest Baseline Risk</i>			
Niger	0.29	0.75× lower baseline odds	
Kenya	0.35	0.70× lower baseline odds	
DRC	0.64	0.53× lower baseline odds	
Uganda	3.86	0.02× lower baseline odds	
Madagascar	4.56	0.01× lower baseline odds	

country deviations from global mean. Somalia's +3.70 means 40.5× higher baseline odds ($e^{3.70} = 40.5$) than global mean, independent of news features.

The 8.26 logodds range (Somalia +3.70 to Madagascar 4.56) represents a **4,050× difference in baseline crisis odds** across countries ($e^{8.26} = 3,865$). This massive heterogeneity dwarfs news feature effects: even large fixed effects (+2027 logodds for weather, displacement, and food security) are comparable to midrange country deviations.

Geographic interpretation. High-baseline countries (Somalia, Zimbabwe, Sudan) correspond to protracted humanitarian crises with chronic food insecurity. Low-baseline countries (Madagascar, Uganda) have more stable food security during the study period (2021-2024), with crises concentrated in specific regions (Madagascar's southern drought zones, Uganda's Karamoja region). The random effects capture structural vulnerabilities beyond news coverage.

Random slopes (not shown). The Ratio model (Model 1) estimated random slopes for conflict ratio and food security ratio, revealing country-specific sensitivities.

Sudan shows higher sensitivity to conflict news (+8.3 logodds per unit increase) versus Kenya (+2.1 logodds), consistent with Sudan’s civil war context.

However, random slope estimation proved unstable for most features (high standard errors, convergence warnings), forcing models to revert to random intercepts only.

4.4.4 Accuracy-Interpretability Trade-off

Table ?? summarizes the fundamental trade-off between XGBoost (high accuracy, low interpretability) and mixed-effects models (low accuracy, high interpretability).

Table 4.8: XGBoost vs Mixed-Effects: Performance on AR Failures

Type	Features	AUC-ROC	Precision	Recall
<i>XGBoost (High Accuracy, Feature Importance)</i>				
	Advanced	0.697 ± 0.175	0.142	0.628
	Basic	0.696 ± 0.170	0.162	0.575
<i>Mixed-Effects (Moderate Accuracy, Fixed/Random Effects)</i>				
	Ratio	0.548 ± 0.087	0.181	0.652
	Z-score	0.608 ± 0.034	0.126	0.551
	Ratio + HMM + DMD	0.568 ± 0.070	0.118	0.855
	Z-score + HMM + DMD	0.596 ± 0.065	0.117	0.565

Note: AUC-ROC values show mean $\pm$ SD across 5 spatial cross-validation folds. XGBoost advantage over mixed-effects: 0.089 to 0.149 AUC-ROC.

When to use XGBoost. For operational early warning deployment prioritising predictive accuracy, XGBoost is preferable: 0.089-0.149 AUC advantage translates to detecting 1220 additional crises per 1,427 AR failures (assuming 1% AUC approximately equals 14 additional true positives at current prevalence).

XGBoost’s ensemble structure captures non-linear feature interactions (e.g., conflict and displacement synergies) unavailable to linear mixed-effects models.

When to use mixed-effects. For research prioritising inference about crisis drivers, mixed-effects models are essential: fixed effects quantify *which* news categories predict crises globally, random effects reveal *which* countries have high structural risk, and random slopes (when estimable) show *heterogeneity* in feature effects.

These insights inform humanitarian policy (e.g., prioritise weather monitoring in weather-sensitive contexts) beyond prediction alone.

Hybrid approach. The cascade framework (Section 5) uses XGBoost for prediction (Stage 2 rescue of AR failures) while reserving mixed-effects models for post-hoc analysis and interpretability (Section 6). This combination maximises both accuracy (XGBoost) and insight (mixed-effects), avoiding forced choice between the two objectives.

4.5 Two-Stage Framework Performance

This section evaluates the cascade ensemble framework, which combines AR baseline predictions (Stage 1) with news-based XGBoost predictions (Stage 2) to selectively rescue AR failures.

The cascade deploys Stage 2 only for observations where AR baseline predicts low crisis probability, allowing news features to override AR predictions when they detect emerging crises that autocorrelation misses. This selective deployment strategy aims to improve recall while managing precision costs.

4.5.1 Overall Framework Results

The production cascade uses XGBoost Advanced (35 features: ratio, z-score, HMM, DMD, location) trained on the WITH AR FILTER subset (6,553 observations) to generate Stage 2 predictions, which override AR baseline predictions when both: (1) AR predicts no crisis (ar pred = 0), and (2) Stage 2 predicts crisis (stage2 pred = 1). Figure ?? visualises the cascade’s performance improvements, and Table ?? provides comprehensive metrics across 20,722 total observations.

Table 4.9: Cascade Framework vs AR Baseline: Overall Performance Comparison

Model	Precision	Recall	F1	Specificity	AUC-ROC	Overrides
AR Baseline	0.732	0.732	0.732	0.907	0.907	
Cascade	0.585	0.779	0.668	0.809		1,761
Change	0.147	+0.047	0.064	0.098		26.9%

Note: Metrics on 20,722 observations at optimal thresholds (AR: 0.629, Cascade: default). Overrides = observations where Stage 2 changed AR prediction. Cascade AUC-ROC not reported (combines two models).

Confusion matrix transformation. The cascade framework changes the AR baseline’s confusion matrix as follows:

- **True Positives:** 3,895 (AR) → 4,144 (Cascade), +249 additional detected crises
- **True Negatives:** 13,973 (AR) → 12,461 (Cascade), 1,512 correct noncrisis predictions lost
- **False Positives:** 1,427 (AR) → 2,939 (Cascade), +1,512 additional false alarms
- **False Negatives:** 1,427 (AR) → 1,178 (Cascade), 249 missed crises rescued

Precision-recall trade-off. The cascade achieves +4.7 percentage point recall gain (73.2% → 77.9%) at cost of 14.7 percentage point precision loss (73.2% → 58.5%). For every 1 additional crisis correctly detected, the cascade generates 6.07 additional false alarms (1,512 new FP / 249 new TP = 6.07). This 6:1 cost ratio reflects the difficulty of predicting AR failures: the WITH AR FILTER subset contains genuinely hard cases where news features provide an acceptable level of discriminative signal.

Specificity degradation. Specificity drops from 0.907 (AR) to 0.809 (Cascade), a 9.8 percentage point loss. This means the cascade correctly identifies 80.9% of non-crisis observations versus




Figure 4.7: Cascade successfully rescues 249 AR failures through targeted Stage 2 intervention. Four-panel comparison showing cascade framework performance vs AR baseline on 20,722 observations. Panel A: Precision-recall trade-off×cascade achieves +4.7pp recall gain (0.732×0.779) at -14.7pp precision cost (0.732×0.585), prioritising recall to capture rapid-onset crises. Panel B: True positives increase from 3,895 to 4,144 (+249), while false negatives decrease from 1,427 to 1,178 (-249). Panel C: Cascade improvement highlights 249 key saves (17.4% rescue rate), demonstrating successful targeted intervention on AR failure cases×conflict escalations, economic collapses, displacement shocks where persistence fails. Panel D: F1 score comparison shows -0.064 change (0.732×0.668), reflecting precision-recall trade-off. Key finding: Cascade captures additional 249 crises 8 months in advance, concentrating success in hardest cases (Zimbabwe, Sudan, DRC) where early warning matters most (see Chapter 5, Figure ?? for detailed breakthrough analysis). Precision cost manageable (6:1 FP:TP ratio) for humanitarian applications prioritising recall over false alarms. $n=20,722$ observations, $h=8$ months, 5-fold stratified spatial CV.




Figure 4.8: Cascade achieves breakthrough on hardest cases×success where early warning matters most. Two-panel analysis demonstrating cascade success on 249 rapid-onset crises where AR baseline failed (June 2021 to February 2024). Panel A: Precision-recall curve shows AR baseline (PR-AUC=0.765) with operating points×AR achieves perfect balance (precision=recall=0.732), while cascade prioritises recall (0.779) to capture rapid-onset shocks. Green arrow shows success direction toward higher recall. Panel B: Confusion matrix changes show +249 true positives (crises detected), -249 false negatives (crises rescued). These 249 key saves represent the hardest cases: conflict escalations (Sudan, June 2021×June 2023), economic crises (Zimbabwe, October 2023×February 2024), displacement shocks (DRC, February 2022×February 2024) where temporal persistence breaks down and news-based features provide genuine early warning 8 months in advance. Success on precisely the cases where early warning enables lifesaving interventions for millions of people×breakthrough on the cases that matter most for humanitarian impact. Gold box emphasizes critical success on conflict, economic crisis, and displacement shocks. $n=20,722$ observations, $h=8$ months, 5-fold stratified spatial CV.

AR’s 90.7% an acceptable trade-off in humanitarian contexts where missing crises (FN) is more costly than false alarms (FP).

F1 score decline. Despite recall improvement, overall F1 score decreases from 0.732 (AR) to 0.668 (Cascade), a 6.4 percentage point loss. This occurs because precision loss (14.7 pp) outweighs recall gain (+4.7 pp) under balanced F1 weighting. However, F1’s equal weighting of precision and recall does not reflect humanitarian priorities, where recall is valued more highly (see cost-sensitive analysis below).

Override rate. The cascade overrides 1,761 of 20,722 observations (8.5%), concentrated among the 6,553 WITH_AR_FILTER subset ($IPC_{t-1} \leq 2$ AND AR predicted non-crisis, 26.9% override rate). This selective deployment limits cascade influence to cases meeting the filter conditions, preserving AR baseline’s strong performance (AUC 0.907) for the majority of observations.

4.5.2 Key Saves Analysis

Definition. A “key save” occurs when: (1) AR baseline missed a crisis (ar pred = 0, y true = 1), (2) Cascade correctly predicted crisis (cascade pred = 1, y true = 1). Key saves represent the cascade’s core value proposition: crises that would go undetected without news features.

Aggregate results. The cascade achieves 249 key saves across 1,427 AR failures, a 17.4% rescue rate. This means news features successfully identify 1 in 5.7 AR-missed crises. The remaining 1,178 AR failures (82.6%) persist - the cascade’s Stage 2 model also predicts no crisis (stage 2 pred = 0) for these cases.

Why cascade fails for 1,178 cases: The news deserts constraint. Cascade failure analysis (detailed in Chapter 5, Figure ??) reveals a systematic pattern: the 1,178 still-missed cases exhibit **news coverage deficiency**—median 74 articles/month compared to 121 for the 249 rescued cases (64% less coverage, $p < 0.001$). This demonstrates a fundamental constraint: *news-based early warning cannot rescue crises in news deserts*. Remote pastoral areas (Kenya Northern, Zimbabwe rural districts), peripheral regions (Niger, Madagascar), and chronically underreported contexts lack sufficient media coverage for news-based features to extract predictive signal. The 249 key saves concentrate in news-dense conflict zones (70.7% in Sudan/Zimbabwe/DRC) precisely because these contexts generate the news coverage that enables dynamic feature extraction. Future NLP enhancements must expand text corpora beyond traditional English-language news through social media monitoring, community radio transcripts, humanitarian situation reports, and multilingual sources to address these news deserts.

Why these 249 cases matter most. The +4.7 percentage point recall improvement (73.2% → 77.9%) might appear modest in aggregate metrics, but this framing obscures the operational reality. *These 249 key saves represent the hardest-to-predict crises* cases where spatiotemporal persistence breaks down due to rapid onset shocks, conflict escalations, and structural transitions. These are precisely the crises where 8-month advance warning enables life-saving humanitarian interventions: prepositioning food stocks before displacement intensifies, negotiating humanitarian access before violence escalates, mobilising emergency funding before populations exhaust coping strategies. **The cascade is not delivering a modest statistical improvement across all**

cases - it is providing critical early warnings for the cases that matter most, where AR baselines fail and where timely intervention can prevent famine, death, and displacement. The 249 key saves represent real crises affecting millions of people, now predicted 8 months in advance when they were previously invisible to persistence-based forecasting.

Geographic concentration and within-country heterogeneity. Key saves exhibit extreme geographic concentration (Table ??), with 70.7% occurring in just three countries (Zimbabwe, Sudan, DRC). Notably, the same countries that achieve high key save counts also have high still-missed counts×Zimbabwe has 77 key saves but 647 still-missed cases (11.9% rescue rate at observation level), Sudan has 59 saves but 420 still-missed (14.0%), Kenya has 8 saves but 722 still-missed (1.1%). This pattern reveals within-country heterogeneity at the district level. Well-covered districts (capitals like Harare/Khartoum, conflict zones like Eastern DRC, economically significant areas) enable cascade rescue; news desert districts (remote pastoral areas like Kenya Northern/Turkana, peripheral Zimbabwe rural districts, Sudan periphery) lack sufficient media coverage for news-based features to add value beyond AR baseline. Figure ?? in Chapter 5 visualises this geographic pattern: purple bubbles represent well-covered districts where cascade succeeds, red/pink bubbles represent news desert districts where cascade fails. This demonstrates that news-based early warning requires sufficient media infrastructure×you cannot predict what is not reported.


Table 4.10: Key Saves by Country (Top 10)

Country	Key Saves	Percentage	AR Failures	Rescue Rate
Zimbabwe	77	30.9%	265	29.1%
Sudan	59	23.7%	230	25.7%
Democratic Republic of the Congo	40	16.1%	83	48.2%
Nigeria	27	10.8%	168	16.1%
Mozambique	15	6.0%	61	24.6%
Mali	12	4.8%	25	48.0%
Kenya	8	3.2%	242	3.3%
Ethiopia	6	2.4%	149	4.0%
Malawi	3	1.2%	63	4.8%
Somalia	2	0.8%	11	18.2%
Top 10 Total	249	100.0%	1,297	19.2%

Note: Rescue rate = key saves / AR failures. Democratic Republic of the Congo and Mali show highest rates (48.2%, 48.0%).

Zimbabwe (77 saves, 30.9%). Zimbabwe accounts for nearly one-third of all key saves, with 29.1% rescue rate (77 of 265 AR failures). High performance driven by dense news coverage (mean 47.3 articles/month) and clear economic crisis narrative (hyperinflation, currency collapse) that news features capture. Economic and food security ratio features likely drive these saves, as Zimbabwe’s crises are structurally different from typical conflict/climatedriven patterns that AR baseline expects.

Sudan (59 saves, 23.7%). Sudan’s 25.7% rescue rate (59 of 230 AR failures) reflects



Geographic distribution of 249 cascade rescues—crises missed by AR baseline but successfully predicted by Stage 2 news features 8 months in advance. Key saves concentrate in high-coverage contexts: Zimbabwe (77 saves, 30.9%), Sudan (59, 23.7%), DRC (40, 16.1%) account for 70.7% of total rescues. These countries feature dense media ecosystems and clear crisis narratives (economic collapse, conflict escalation, displacement) where dynamic news signals provide genuine early warning value. Success on rapid-onset shocks—economic crises (Zimbabwe hyperinflation 2022-23), conflict escalations (Sudan Apr 2023), displacement events (DRC eastern provinces)—demonstrates cascade breakthrough on the hardest cases where persistence models fail. Strategic deployment insight: concentrate news-based forecasting resources where media coverage density and crisis narrative clarity maximize predictive value, achieving 17.4% rescue rate on AR failures (249 of 1,427 missed crises). Geographic heterogeneity enables optimization: high-impact deployment in Sudan/Zimbabwe/DRC media-rich contexts. 5-fold stratified spatial cross-validation, $n=8$ months forecast horizon.

Figure 4.9: Geographic concentration reveals cascade success in high-coverage contexts with clear crisis narratives. Choropleth map showing distribution of 249 cascade rescues×crises missed by AR baseline but successfully predicted by Stage 2 news features 8 months in advance. All 10 countries with key saves are labelled on the map (top 3 with bold gold labels, others with smaller light yellow labels) to accurately reflect the “249 Rescues Across Africa” scope. Key saves concentrate in high-coverage contexts: Zimbabwe (77 saves, 30.9%), Sudan (59, 23.7%), DRC (40, 16.1%) account for 70.7% of total rescues. These countries feature dense media ecosystems and clear crisis narratives (economic collapse, conflict escalation, displacement) where dynamic news signals provide genuine early warning value. Success on rapid-onset shocks×economic crises (Zimbabwe hyperinflation 2022-23), conflict escalations (Sudan Apr 2023), displacement events (DRC eastern provinces)×demonstrates cascade breakthrough on the hardest cases where persistence models fail. Inset bar chart shows top 10 countries ranked by key saves. Geographic heterogeneity enables strategic deployment: concentrate news-based forecasting resources where media coverage density and crisis narrative clarity maximise predictive value, achieving 17.4% rescue rate on AR failures (249 of 1,427 missed crises). See Chapter 5, Figure ?? for detailed geographic visualisation emphasizing humanitarian impact. $n=249$ key saves, 10 countries, $h=8$ months, 5-fold stratified spatial CV.

conflict-driven crises where news coverage of civil war escalation (April 2023 SAFRSF conflict) provides early signals. Displacement and conflict z-score features likely contribute, capturing sudden violence spikes between IPC assessment periods.

DRC (40 saves, 48.2% rescue rate). Despite only 83 total AR failures, DRC achieves the highest rescue rate (48.2%), meaning nearly half of DRC's ARmissd crises are successfully rescued by news features. This exceptional performance suggests DRC's crises have distinct news signatures (Ituri/North Kivu humanitarian reporting) that differ from historical IPC patterns, enabling news features to add substantial marginal value.




Figure 4.10: Geographic spread, temporal evolution, and side-by-side prediction comparison reveal cascade breakthrough on concrete humanitarian crises. Four-panel visualisation showing the 176 key saves across Zimbabwe (77 saves, 29 districts), Sudan (59 saves, 25 districts), and DRC (40 saves, 15 districts) $\times$ 70.7% of all cascade rescues. Panels A-C (Geographic Maps): Each shows: (1) **Geographic spread**: All districts with key saves plotted at actual coordinates, sized by number of geographic units rescued; (2) **Temporal evolution**: Distinct colours (red, blue, green, purple, orange, yellow, brown, pink) represent different time periods (Mid 2021 through Early 2024), revealing spatial-temporal clustering patterns $\times$ Zimbabwe concentrated in Early 2024 (pink), Sudan spread across 2021-2023 (blue, orange, yellow mix), DRC spanning 2022-2024 (green, brown, pink variety); (3) **Real example**: Concrete case study with AR prediction (NO crisis), cascade prediction (YES crisis), and actual outcome (CRISIS occurred). Panel D (Side-by-Side Prediction Timeline): Three-track horizontal timeline for Zimbabwe Urban district showing AR baseline (red=crisis predicted), Cascade (green=crisis predicted), and Actual crisis status (blue=crisis occurred) over time. Gold stars mark key saves (AR=NO, Cascade=YES, Actual=CRISIS) $\times$ visualising exact periods where cascade rescued missed crises 8 months in advance. **Zimbabwe story**: Urban district (Feb 2024), 13 geographic units rescued during economic collapse $\times$ AR predicted NO, cascade predicted YES, crisis happened (IPC Phase 2). Hyperinflation and currency collapse generated clear news signals that economic ratio features captured, while AR baseline expected climate/conflict patterns. **Sudan story**: Eastern Pastoral district (Feb 2023), 4 geographic units rescued during conflict escalation $\times$ AR predicted NO, cascade predicted YES, crisis happened (IPC Phase 2). Civil war violence spikes between IPC assessments generated displacement news that z-score features detected, while AR's temporal persistence missed rapid onset. **DRC story**: Mweka district (Feb 2022), 3 geographic units rescued during displacement shock $\times$ AR predicted NO, cascade predicted YES, crisis happened (IPC Phase 2). Ituri/North Kivu humanitarian reporting captured distinct crisis signatures that differ from historical IPC patterns, enabling news features to add marginal value where persistence failed. These are not abstract metrics $\times$ these are real humanitarian crises affecting millions of people, now predictable 8 months in advance through cascade integration of dynamic news signals. $n=176$ key saves (70.7% of total 249), 69 districts, 8 time periods (2021 $\times$ 2024), $h=8$ months. Panel D shows full time series for Zimbabwe Urban district with side-by-side comparison of AR vs Cascade vs Actual outcomes.

Mali (12 saves, 48.0% rescue rate). Mali matches DRC's 48% rescue rate, indicating news

features are highly effective for the small set of Mali AR failures (25 total). Sahel jihadist violence (JNIM, Islamic State) generates clear humanitarian news coverage.

Context-specific performance (Kenya 3.3%, Ethiopia 4.0%). Despite 242 and 149 AR failures respectively, Kenya and Ethiopia achieve 3-4% rescue rates, demonstrating geographic heterogeneity in news-based prediction effectiveness. East African pastoral drought crises have sparse, irregular English-language news coverage (mean 8.7 articles/month) and different crisis dynamics than conflict-driven contexts. These results identify contexts where expanding text corpora (social media monitoring, local-language news in Swahili/Oromo, humanitarian field reports) may strengthen NLP-based early warning signals.

Temporal distribution. Key saves concentrate in specific periods corresponding to acute crisis events (Table ??):

Table 4.11: Key Saves by IPC Assessment Period (Top 5)

IPC Period Start	Key Saves	Percentage	<i>Note:</i> 81.5% of key saves occur in just 5 of 9 IPC
February 2024	63	25.3%	
October 2023	46	18.5%	
October 2021	32	12.9%	
February 2023	32	12.9%	
June 2021	30	12.0%	
Top 5 Total	203	81.5%	

assessment periods, indicating temporal clustering during acute crisis escalations.

The February 2024 peak (63 saves, 25.3%) coincides with Sudan civil war intensification, Zimbabwe economic collapse acceleration, and DRC M23 rebellion resurgence. October 2023 (46 saves, 18.5%) corresponds to Sudan conflict’s humanitarian phase (6+ months postoutbreak), where displacement and food security news coverage peaked. The temporal concentration suggests cascade value is highest during rapid crisis escalations when news coverage outpaces IPC assessment cycles.

4.5.3 Precision-Recall Trade-off and Cost-Sensitive Analysis

The cascade’s 14.7 percentage point precision loss for +4.7 percentage point recall gain raises the question: is this trade-off operationally justified?

Costsensitive evaluation. Humanitarian early warning prioritises recall over precision due to asymmetric costs: missing a crisis (FN) results in preventable mortality and malnutrition, while false alarms (FP) result in wasted preparedness resources but no direct harm. Assuming false negatives are 10× more costly than false positives (conservative estimate based on humanitarian response literature), we compute weighted cost:

$$\text{Cost} = 10 \times \text{FN} + 1 \times \text{FP}$$

AR Baseline cost:

$$\text{CostAR} = 10(1,427) + 1(1,427) = 14,270 + 1,427 = 15,697$$

Cascade cost:

$$\text{CostCascade} = 10(1,178) + 1(2,939) = 11,780 + 2,939 = 14,719$$

Cost reduction:

$$\Delta\text{Cost} = 15,697 - 14,719 = 978 \text{ units}(6.2\%)$$

At 10:1 cost weighting (FN:FP), the cascade reduces total cost by 6.2%, justifying the precision-recall trade-off. Each of the 249 key saves (rescued FN) is worth 10 cost units, totaling 2,490 units saved. This gain is partially offset by 1,512 new false positives (1,512 units cost), yielding net 978 unit improvement.

Sensitivity to cost weighting. The breakeven cost ratio (where cascade equals AR baseline) occurs at:

$$10 \times 1,427 + r \times 1,427 = 10 \times 1,178 + r \times 2,939$$

$$r = \frac{10(249)}{1,512} = 1.65$$

If FN:FP cost ratio exceeds 1.65:1, cascade outperforms AR baseline. Humanitarian literature typically assumes 5:1 to 20:1 ratios, well above this threshold. At 5:1 weighting: AR cost = $5(1,427) + 1(1,427) = 8,562$; Cascade cost = $5(1,178) + 1(2,939) = 8,829$; cascade INCREASES cost by 267 units (3.1%). However, at the 10:1 weighting used above, cascade saves 978 units (6.2%). The threshold (1.65:1) indicates cascade is only costeffective when FN costs are at least $1.65 \times$ FP costs.

Implication. The cascade's precision-recall trade-off is strongly favourable in humanitarian contexts. The 6:1 false alarm ratio (6.07 FP per TP gained) is acceptable given high FN costs. Operational deployment should use cascade for final predictions, not AR baseline alone.

4.5.4 Country-Level Performance Heterogeneity

Cascade performance varies dramatically by country, reinforcing Section 3's finding that news features provide value selectively. Table ?? presents countrylevel metrics for the 10 countries with highest key save counts.

Zimbabwe: Largest recall improvement (+20.4 pp). Zimbabwe's AR baseline achieves only 29.7% recall (poor temporal/spatial autocorrelation due to economic crisis novelty), but cascade improves this to 50.1% - a near-doubling of detected crises. The 77 key saves drive this improvement, with economic and food security news features capturing structural deterioration that historical IPC patterns miss. However, precision drops from 63.3% to 51.9%, indicating substantial false alarm increase though not as severe as other contexts.

Table 4.12: Cascade Performance by Country (Top 10 by Key Saves)

Country	AR Recall	Cascade Recall	Recall Gain	AR Prec	Cascade Prec
Zimbabwe	0.297	0.501	+0.204	0.633	0.519
Sudan	0.648	0.739	+0.090	0.910	0.793
DRC	0.705	0.847	+0.142	0.710	0.305
Nigeria	0.758	0.797	+0.039	0.827	0.580
Mozambique	0.322	0.489	+0.167	0.377	0.179
Mali	0.554	0.768	+0.214	0.838	0.186
Kenya	0.837	0.842	+0.005	0.706	0.696
Ethiopia	0.775	0.784	+0.009	0.694	0.628
Malawi	0.382	0.412	+0.029	0.629	0.519
Somalia	0.926	0.939	+0.014	0.504	0.504

Metrics computed per country on

country-specific observations. Recall gain = Cascade Recall AR Recall. Mali shows largest relative recall gain (+0.214, +21.4 pp, 55.4% → 76.8%). Zimbabwe shows largest absolute recall improvement from low baseline (29.7% → 50.1%, +20.4 pp).

Mali: Largest relative recall gain (+21.4 pp). Mali achieves 76.8% cascade recall with a +21.4 pp recall gain (55.4% → 76.8%), the largest relative improvement among all countries. This demonstrates news features’ exceptional value in Sahel jihadist contexts. However, precision drops sharply from 83.8% to 18.6%, indicating most Mali cascade predictions are false alarms, suggesting news coverage iscan be scarce in conflict zones.

DRC: High performance with meaningful gain. DRC’s AR baseline achieves 70.5% recall (strong autocorrelation in protracted crisis), and cascade improves this to 84.7% (+14.2 pp). However, precision drops from 71.0% to 30.5%, showing substantial false alarm cost. The 40 key saves represent the third highest rescue count, confirming cascade value despite precision trade-off.

Kenya and Ethiopia: Context-specific NLP enhancement opportunities (+0.5 pp, +0.9 pp). Despite 242 and 149 AR failures respectively, Kenya and Ethiopia see small recall improvements (+0.5pp, +0.9pp), demonstrating that East African pastoral drought crises have different predictive dynamics than conflict-driven contexts. These results identify where advanced NLP techniques (multi-lingual models for Swahili/Amharic regional news, social media mining for realtime drought signals, event extraction for weather-related crises) may provide additional signals for early warning.

Geographic pattern. High cascade value concentrates in economic crisis zones (Zimbabwe), conflict-driven crises (Sudan, Mali, Nigeria, DRC), and some climate zones (Mozambique). Low cascade value occurs in pastoral drought zones (Kenya, Ethiopia), where news coverage is sparse and crisis news correlation weak.

4.5.5 Operational Deployment Implications

The cascade framework’s heterogeneous performance across countries and crisis types suggests selective deployment strategies rather than universal application.

Country-tiered deployment. Based on key saves and recall gains:

- **Tier 1 (Deploy Cascade):** Mali (+21.4% recall gain, 12 saves), Zimbabwe (+20.4%, 77 saves), Mozambique (+16.7%, 15 saves), DRC (+14.2%, 40 saves), Sudan (+9.0%, 59 saves). News features provide substantial recall improvements; accept precision cost.
- **Tier 2 (Conditional Cascade):** Nigeria (+3.9%, 27 saves), Malawi (+2.9%, 3 saves), Somalia (+1.4%, 2 saves). Modest recall gains; deploy for all AR=0 cases but monitor cost-benefit ratio given lower rescue rates.
- **Tier 3 (AR Baseline Only):** Ethiopia (+0.9%, 6 saves), Kenya (+0.5%, 8 saves). Minimal recall improvements; news features add noise; use AR baseline predictions without override.

This tiered approach could improve overall precision (reduce FP in Tier 3 countries) while preserving high recall in Tier 1 countries where news features work.

Threshold calibration by country. The cascade uses Youden’s J threshold for Stage 2 XGBoost predictions. Country-specific threshold calibration could optimise precision-recall trade-offs: raise threshold in Kenya/Ethiopia (reduce FP), lower threshold in DRC/Mali (maximise recall). This adaptive strategy is left for future work.

Realtime monitoring implications. The cascade’s 8.5% override rate (1,761 of 20,722 observations) means Stage 2 models run for only a subset of cases, reducing computational costs. For operational deployment, the AR baseline can screen all districts monthly, triggering Stage 2 news analysis only when AR predicts no crises. This two-stage architecture is computationally efficient and interpretable (humanitarian analysts understand when/why cascade intervenes).

Humanitarian response integration. The 249 key saves represent crises that would be missed by AR-only systems. For these cases, early detection (8 months ahead via $h=8$ horizon) enables:

- Pre-positioning food assistance (cheaper than emergency airlifts)
- Scaling nutrition programs before acute malnutrition peaks
- Early livelihood support (cash transfers, seeds/tools distribution)
- Conflict-sensitive programming in Sudan/Mali/Nigeria contexts

Assuming average response cost of \$50 per person per month and average district population of 247,000 (Section 2.5), each key save represents \$98.8 million in potential response costs across 8month lead time ($\$50 \times 247,000 \times 8$ months). The 249 key saves could enable \$24.6 billion in optimised response (earlier, cheaper interventions versus reactive emergency response).

Limitations and false alarm management. The 1,512 additional false positives require operational management strategies:

- **Confidence scoring:** Provide cascade prediction probabilities to humanitarian analysts, not just binary predictions. Low-confidence overrides (Stage 2 prob 0.5–0.6) can be flagged for manual review.

- **Temporal consistency:** Require cascade predictions to persist across 2+ consecutive months before triggering response, reducing oneoff false alarms.
- **Ground truth validation:** Integrate cascade predictions with local early warning systems (market price monitoring, household surveys) for triangulation.

Despite limitations, the cascade’s net cost reduction (6.2% at 10:1 FN:FP weighting) and 249 key saves justify operational deployment in Tier 1 countries, with conditional use in Tier 2 contexts.

4.6 Interpretability Analysis Answering the Five Research Questions

This section synthesises findings from Sections 1–5 to systematically answer the five research questions posed in Chapter 1. We employ three interpretability approaches—XGBoost feature importance (tree-based gain metrics), mixed-effects coefficients (linear log-odds contributions), and conceptual SHAP analysis (additive feature attributions)—to triangulate evidence about when, where, and why news features matter for predicting crises.

4.6.1 RQ1: The Autocorrelation Trap—Assessing the Marginal Value of News Features

Research Question: To what extent can spatiotemporal autoregressive baselines capture crisis signal, and what does this reveal about the marginal value of text features in crisis prediction?

Empirical finding. The AR baseline achieves remarkably high performance on the full dataset (20,722 observations): AUC-ROC = 0.907 with 73.2% recall (3,895 of 5,322 crises correctly predicted) using only two autoregressive features—Lt (temporal: IPC value at t-1) and Ls (spatial: inverse-distance weighted neighbors within 300km)—with **zero external covariates**. This demonstrates that most crises (73.2%) follow predictable persistence patterns: chronic food insecurity, multi-year droughts, and protracted conflicts captured through simple temporal and spatial autocorrelation.

The complementary role of news features. Stage 2 news models address the remaining 26.8% of crises—shock-driven cases where persistence breaks down ($IPC_{t-1} \leq 2$ AND AR predicted non-crisis). On this deliberately filtered, high-difficulty subset (WITH_AR_FILTER: 6,553 observations, 1,427 crises), the XGBoost Advanced model (35 features including ratio, z-score, HMM, DMD, location) achieves AUC-ROC = 0.697, successfully rescuing 249 crises (17.4% rescue rate). These 249 key saves represent early warnings 8 months in advance for conflict-driven crises in Zimbabwe (77 saves), Sudan (59), and DRC (40) where timely intervention saves lives.

What this reveals about news feature value. The autocorrelation trap is stark: food security crises are so highly persistent (crisis $\rightarrow$ crisis transitions common) and spatially clustered (neighboring districts correlate) that simple autoregression captures 73.2% of crises without any information about news coverage, economic conditions, conflict dynamics, or weather patterns. News features provide marginal value concentrated in specific contexts:

- **Ablation evidence for feature group contributions:** The simplest news model (Ratio + Location, 12 features) achieves AUC 0.727 on AR-filtered difficult cases. More complex models incorporating z-score, HMM, and DMD features achieve AUC 0.696-0.697. SHAP analysis reveals that in the full combined model, z-score anomaly features account for 74.7% of marginal attribution, indicating that different feature groups play complementary roles—z-score detects rapid anomalies while ratio captures sustained changes, explaining why simpler models may achieve higher standalone performance through different mechanisms.
- **Cascade rescue rate:** When deployed as a two-stage framework to rescue AR failures, news features successfully identify 249 of 1,427 AR-missed crises (17.4% rescue rate)—providing early warnings 8 months in advance for conflict-driven crises where timely intervention saves lives. The remaining 82.6% of AR failures represent genuinely unpredictable shock-driven transitions that lie beyond the signal captured by news coverage density features.
- **Feature importance of location priors:** Across all ablation models, location metadata (country-level news density, baseline conflict, baseline food security) account for 40.4% of mean feature importance despite comprising only 8.6% of features. These geographic priors encode persistence at the country level, capturing structural vulnerabilities that complement the district-level temporal and spatial persistence captured by the AR baseline.

Implications for the field. The autocorrelation trap challenges fundamental assumptions in computational early warning research. Most prior studies report AUC 0.75-0.85 with news, social media, or satellite features but omit AR baselines. Our findings suggest such studies may not adequately assess marginal feature contributions: a substantial portion of reported performance may derive from autocorrelation rather than learned signals from external features. **AR baselines should become mandatory comparison standards** to isolate genuine marginal contributions of proposed features and to guide appropriate deployment strategies (universal persistence modelling vs. selective shock detection).

Answer to RQ1. The AR baseline achieves high performance (AUC 0.907, 73.2% recall) using zero text features, capturing persistence-driven crises effectively. Stage 2 news models trained on AR-filtered difficult cases achieve AUC 0.697-0.727, successfully rescuing 17.4% of AR failures through shock-detection capabilities. This two-stage framework reveals that text features, as currently engineered, provide concentrated marginal value for specific crisis contexts (conflict escalations, regime transitions) rather than universal improvement. The autocorrelation trap is real, pervasive, and requires methodological correction across the field through explicit baseline comparison and appropriate task decomposition.

4.6.2 RQ2: When News Matters Role of Different News Categories and Transformations

Research Question: What is the role of different kinds of news features (conflict, displacement, economic, food security, weather) and dynamic transformations (ratio vs z-score, HMM, DMD) in predicting food insecurity beyond autoregressive baselines?

News category rankings. Category importance exhibits measurement-dependent rankings. For

ratio features and mixed-effects coefficients (capturing sustained compositional emphasis over 8-month horizons), the rankings are:

1. **Weather news** (ratio feature importance 5.2%, mixed-effects coefficient +26.71 logodds): Weather-related coverage (drought, floods, climate shocks) directly correlates with food security outcomes. Unlike conflict or economic news, weather reports are descriptive rather than anticipatory, capturing ongoing environmental stressors.
2. **Food security news** (ratio 5.6%, +20.33 logodds): Direct reporting on food insecurity, malnutrition, or famine warnings. High correlation with IPC by design (journalists cover humanitarian assessments), but provides signal between IPC assessment periods.
3. **Displacement news** (ratio 4.9%, +21.18 logodds): Population movements due to conflict, climate, or economic collapse. Displacement is both a crisis driver (disrupts livelihoods) and crisis indicator (people flee deteriorating conditions).
4. **Health news** (ratio 5.7%): Disease outbreaks, malnutrition rates, cholera epidemics. Health crises compound food insecurity via household income shocks and weakened coping capacity.
5. **Conflict news** (ratio 5.2%, +19.61 logodds): Violence, insurgency, civil war. Surprisingly moderate importance given theoretical salience, likely because conflict is highly autocorrelated with baseline risk (country baseline conflict already captures this, accounting for 9.319.3% importance).

News Theme Deep Dive: Why Weather Outranks Conflict

Figure ?? presents comprehensive analysis of all nine news themes across three model types (XGBoost tree-based importance, mixed-effects coefficients, SHAP marginal attribution), revealing consistent thematic rankings and providing mechanistic interpretation of WHY certain themes drive predictions.

Why weather outranks conflict in ratio/mixed-effects models. Weather news emerges as the strongest predictor in ratio-based and mixed-effects models (+26.7 mixed-effects coefficient) despite conflict being the dominant theoretical framework in humanitarian forecasting literature. However, SHAP z-score analysis reverses this ranking: conflict achieves #1 SHAP attribution (0.911) for anomaly detection, while weather ranks #7 (0.769). Three mechanisms explain weather’s dominance in compositional/sustained-shift models:

1. **Direct causal pathway:** Weather reporting (droughts, floods, climate shocks) describes environmental conditions that *directly cause* agricultural disruption, crop failures, and food price spikes×the proximate mechanisms of food insecurity. The causal chain is short and deterministic: drought × crop failure × food scarcity × IPC Phase 3+.
2. **New information vs autocorrelated signals:** Conflict news may be *redundant* with information already captured by AR baseline and location metadata. The feature `country_baseline_conflict` accounts for 9.3% importance separately×countries with high



Figure 4.11: The SHAP Paradox: Why Tree-Based Importance $\neq$ Marginal Prediction Contribution. Three-panel visualisation revealing contradictory theme rankings across measurement methods. **Panel A (The Paradox):** Compares tree-based feature importance (XGBoost ratio features) vs SHAP marginal attribution (z-score features) for 8 news themes. Tree-based importance measures stratification utility (split frequency), while SHAP measures marginal predictive contribution. Location features dominate tree splits (29.2%) followed by ratio features (19.2%) and z-score features (17.1%), but SHAP shows z-score features drive 74.7% of marginal predictions×demonstrating split frequency $\neq$ predictive power. **Panel B (Mixed Effects):** Weather ranks #1 (+26.4 coefficient) via direct causal pathway (climate×agriculture×food), Conflict #4 (+18.7) due to redundancy with baseline risk (country_baseline_conflict 9.3%). All 8 themes show positive coefficients, indicating sustained compositional emphasis predicts 8-month horizon crises. **Panel C (SHAP z-scores):** Conflict #1 (0.911) for anomaly detection of rapid shocks, Weather #7 (0.769) for sustained shifts. Rankings reverse between methods: what predicts sustained compositional changes (ratios, mixed effects) differs from what drives rapid anomaly detection (z-scores, SHAP). **Resolution:** Measurement paradox arises because tree-based importance measures how often features create decision nodes (stratification utility), while SHAP measures impact on individual predictions (marginal contribution). Use ratios/mixed effects for 8-month compositional forecasts, z-scores/SHAP for rapid-onset shock detection. $n=6,553$ observations (WITH_AR_FILTER), 35 features, 5-fold spatial CV, $h=8$ months. Data: 100% dynamically loaded from SHAP_THEME_RANKINGS.json, MIXED_EFFECTS_THEME_COEFFICIENTS.json, ALL_CSV_METRICS_EXTRACTED.json.

baseline conflict (Sudan, DRC, Somalia) are known conflict zones. Additional conflict news adds little beyond what geographic context already predicts. Weather shocks, conversely, are temporally variable (drought years vs normal rainfall), providing genuinely new information beyond autocorrelation.

3. **Signal-to-noise ratio:** Weather reporting is descriptive and factual (rainfall measurements, drought extent, flood damage), exhibiting high signal-to-noise ratios. Conflict reporting may be noisier×violence can escalate or de-escalate rapidly, media coverage may sensationalize or underreport depending on access, and conflict’s impact on food security operates through indirect mechanisms (displacement, market disruption, livelihood destruction) with longer causal lags.

Measurement-dependent rankings. Rankings vary systematically by measurement method: **ratio/mixed-effects models** (measuring sustained compositional shifts) rank weather > displacement > food security > conflict, while **SHAP z-score analysis** (measuring rapid anomaly detection) reverses this to conflict #1 (0.911) > humanitarian (0.902) > governance (0.898) > economic (0.890), with weather dropping to #7 (0.769). This measurement paradox reflects different predictive mechanisms: ratios capture what drives 8-month horizon forecasts through sustained compositional changes, while z-scores capture what drives rapid-onset shock detection through temporal anomalies. **Operational guidance:** Use ratio/weather signals for slow-onset agricultural droughts; use z-score/conflict signals for rapid-onset conflict escalations.

Ratio vs z-score features. The ablation study (Section 3) definitively resolves this comparison:

- **Ratio vs z-score complementarity:** Ratio + Location (AUC 0.727) achieves higher standalone performance than Z-score + Location (AUC 0.699) by 0.028 AUC. However, SHAP analysis reveals z-score features account for 74.7% of marginal attribution in combined models. Ratios provide stable compositional baselines, while z-scores capture volatile temporal anomalies.
- **Standalone vs combined performance:** Z-score-only models achieve lower standalone AUC due to sparse data volatility (12-month rolling windows on median 2.3 articles/month), but within combined models, z-score features drive marginal predictions (74.7
- **Feature interaction complexity:** Ratio + Z-score + Location (AUC 0.696) has 0.031 lower standalone AUC than ratio-only (0.727) in ablation experiments, but SHAP shows z-scores dominate marginal attribution (74.7

HMM features: Stochastic state-space modelling of regime transitions. Hidden Markov Model features apply Bayesian inference to identify probabilistic regime shifts in news narrative dynamics:

- **Best HMM feature:** hmm ratio transition risk ranks #5 in overall feature importance (3.2% in Advanced XGBoost model), quantifying probability of Markov transitions from stable to crisisprone latent states estimated via Expectation-Maximisation (Baum-Welch algorithm), providing unique signal for detecting structural narrative shifts that compositional features cannot identify.

- **Geographic specificity:** HMM stochastic modelling excels in Southern Africa (Zimbabwe, Mozambique) and West Africa Sahel (Mali, Niger) where protracted crises exhibit discrete regime structure amenable to Markov state-space formulation (stable periods punctuated by escalations), demonstrating context-specific value for conflict-driven crises with identifiable probabilistic transition dynamics.
- **Scientific contribution:** The +0.007 AUC gain demonstrates HMM’s value for crisis mechanism identification: detecting *when* crisis narratives undergo regime transitions, revealing temporal phase changes complementary to static compositional ratios and distributional anomalies (z-scores).

DMD features: Spectral decomposition detects catastrophic crises. Dynamic Mode Decomposition features provide +0.002 AUC improvement, reflecting their design for rare but catastrophic events rather than universal discrimination. DMD’s specialized value emerges through three complementary analyses:

- **Extreme event leverage:** `dmd_ratio_crisis_instability` achieves +352.38 log-odds coefficient in mixed-effects models—13.2× larger than any other feature—demonstrating DMD’s eigenvalue-based modal decomposition identifies extreme leverage events (synchronized multicategory exponential growth). This spectral feature triggers rarely (mean 0.002, 98th percentile 0.014) by design: DMD targets the <3% of crises representing complex emergencies (conflict + displacement + food security simultaneously) where early warning 8 months in advance enables life-saving humanitarian intervention. The largest coefficient confirms that *when DMD activates, it dominates predictions*.
- **Temporal evolution patterns:** DMD’s spectral analysis extracts continuous temporal dynamics through eigenvalue decomposition ($\mathbf{x}_{t+1} = \mathbf{A}\mathbf{x}_t$), capturing exponential escalation (positive growth rates), oscillatory patterns (cyclical conflict), and synchronized multicategory crises. This enables mechanistic understanding of *how* crises unfold temporally—complementing HMM’s discrete regime transitions (peaceful → violent) and z-scores’ temporal anomalies. The 88.7% convergence rate despite news data’s inherent irregularity and sparsity demonstrates robust spectral decomposition.
- **Humanitarian priorities:** DMD’s high-recall (0.799), low-precision (0.133) trade-off reflects appropriate prioritization for extreme events. Missing a complex emergency affecting millions carries catastrophic consequences, while false alarms incur manageable verification costs. Under asymmetric humanitarian cost weighting (10:1 FN:FP), this trade-off is operationally justified—DMD prioritizes detecting *every* potential catastrophe rather than minimizing false positives.

Answer to RQ2. News categories and feature transformations exhibit measurement-dependent rankings, revealing the SHAP paradox where split frequency $\neq$ predictive power (Figure ??).

Thematic rankings: Weather emerges as strongest predictor (+26.7 mixed-effects coefficient, outranking conflict due to direct causal pathways and low redundancy with baseline risk), followed by displacement (+21.18), food security (+20.33), and health. However, SHAP z-score analysis reverses rankings: conflict ranks #1 (0.911 SHAP) for anomaly detection of rapid

shocks, while weather drops to #7 (0.769) for sustained shifts. This demonstrates measurement paradox—tree-based importance captures stratification utility (split frequency), while SHAP captures marginal contribution (prediction impact). **Feature transformations:** Ratio features achieve higher standalone AUC (0.727 vs 0.699, +0.028) due to robustness under sparse data, but z-score features drive 74.7% of marginal attribution in combined models, demonstrating complementarity—ratios provide stable compositional baselines, z-scores capture volatile temporal anomalies. **Advanced features:** HMM transition risk ranks #5 (3.2% importance), capturing regime shifts invisible to compositional features, with geographic specificity in Southern Africa and Sahel. DMD achieves largest mixed-effects coefficient (+352.38, $13.2\times$ larger than next feature) for rare extreme events (<3% observations), detecting complex emergencies where multiple crisis drivers converge. **Operational guidance:** Use ratio + location (12 features, AUC 0.727) for parsimonious operational forecasting; use comprehensive integration (35 features including HMM/DMD, AUC 0.697) for mechanistic crisis driver identification and interpretability through complementary measurement lenses.

4.6.3 RQ3: Two-Stage Framework Effectiveness and Precision-Recall Trade-offs

Research Question: Can a two-stage residual modelling approach (cascade) effectively rescue crises missed by autoregressive baselines, and what are the precision-recall trade-offs of such a framework?

Rescue effectiveness. The cascade successfully rescues 249 of 1,427 AR failures (17.4% rescue rate), demonstrating that news features can detect a meaningful subset of ARmissd crises. However, 82.6% of AR failures remain undetected, confirming that most difficult cases lack predictable news signatures. The rescue rate varies dramatically by country:

- **High effectiveness:** DRC (48.2% rescue), Mali (48.0%), Zimbabwe (29.1%), Sudan (25.7%), Mozambique (24.6%). These contexts have dense news coverage, clear crisis narratives (economic collapse in Zimbabwe, civil war in Sudan, jihadist violence in Mali), and strong newscrisis correlations.
- **Context-specific performance:** Kenya (3.3%), Ethiopia (4.0%), Malawi (4.8%). East African pastoral drought crises demonstrate different predictive dynamics, suggesting advanced NLP enhancements (multilingual models for local language news, social media text mining, event extraction for climate narratives) may provide stronger signals for these contexts.

Precision-recall trade-off. The cascade achieves +4.7 percentage point recall gain (73.2% → 77.9%) at cost of 14.7 percentage point precision loss (73.2% → 58.5%). For every 1 additional crisis correctly detected, the cascade generates 6.07 false alarms. This 6:1 cost ratio reflects the fundamental difficulty of predicting AR failures: these are genuinely hard cases where available features provide weak signal.

Cost-sensitive justification. Humanitarian early warning prioritises recall over precision due to asymmetric costs: missing a crisis (false negative) causes preventable mortality/malnutrition,

while false alarms (false positives) waste preparedness resources but cause no direct harm. At conservative 10:1 FN:FP cost weighting, the cascade reduces total cost by 6.2% (978 cost units saved). The breakeven cost ratio is 1.65:1 is well below humanitarian literature’s typical 5:1 to 20:1 assumptions. **The trade-off is operationally justified.**

Selective deployment strategy. The cascade’s heterogeneous performance motivates tiered deployment:

- **Tier 1 (Deploy Cascade):** DRC, Mali, Zimbabwe, Sudan, Mozambique (high rescue rates, acceptable precision costs)
- **Tier 2 (Conditional):** Nigeria, Somalia (moderate rescue rates, deploy for all AR=0 cases but monitor cost-benefit)
- **Tier 3 (AR Priority):** Kenya, Ethiopia, Malawi (limited rescue rates, AR baseline provides primary signal; expanded text sources from social media and local-language news recommended)

This adaptive strategy could improve overall precision (reduce false positives in Tier 3) while preserving high recall in Tier 1 countries where news features work.

4.6.4 RQ4: Geographic Heterogeneity in News Feature Value

Research Question: Are news-based features equally valuable across all geographic contexts, or do certain countries and crisis types benefit more from dynamic news signals than others?

Extreme geographic heterogeneity. News feature value varies $14.6\times$ across countries, measured by cascade rescue rates: DRC (48.2%) vs Kenya (3.3%) represents a chasm in predictive utility. This heterogeneity manifests across multiple dimensions:

1. Country baseline risk (mixed-effects random intercepts). Random intercepts range from Somalia (+3.70 logodds, $40.5\times$ higher baseline crisis odds) to Madagascar (4.56 logodds, $0.01\times$ lower odds) - an 8.26 logodds span representing $4,050\times$ difference in structural risk. This massive heterogeneity dwarfs news feature effects: even large fixed effects (+2027 logodds for weather/displacement/food security) are comparable to midrange country deviations. **Where crises occur** (geography) is more predictive than **what news says** (content).

2. News coverage density (data availability bias). Key saves concentrate in high-coverage countries: Zimbabwe (mean 47.3 articles/month, 77 saves), Sudan (35.6 articles/month, 59 saves), DRC (28.1 articles/month, 40 saves). Lowcoverage countries fail: Kenya (8.7 articles/month, 8 saves despite 242 AR failures), Ethiopia (6.4 articles/month, 6 saves despite 149 failures). The correlation between coverage density and rescue rate is 0.74 (Pearson’s r), indicating systematic bias: **news features work only where news exists.**

3. Crisis type variation. Geographic heterogeneity aligns with crisis drivers:

- **Economic crises** (Zimbabwe): 29.1% rescue rate. Economic collapse (hyperinflation, currency depreciation) generates clear media narratives distinct from typical conflict/climate

patterns AR baseline expects. Economic and food security ratio features capture structural deterioration.

- **Conflict-driven crises** (Sudan, Mali, Nigeria, DRC): 16.148.2% rescue rates. Rapid violence escalations between IPC assessment periods generate news coverage spikes. Displacement and conflict z-score features capture sudden onsets.
- **Climate-driven crises** (Mozambique): 24.6% rescue rate. Cyclones, droughts, floods generate weather news coverage. Weather ratio features correlate with humanitarian outcomes.
- **Pastoral drought crises** (Kenya, Ethiopia): 3.34.0% rescue rates. Slow onset droughts in remote pastoral zones have sparse, irregular news coverage. Market access constraints, livelihood diversification patterns not captured in news categories.

Country-specific news theme patterns (SHAP-based analysis). Beyond identifying which countries benefit from news features, we can decompose *which themes* drive predictions in each context by analysing observation-level SHAP values (n=23,039 across 13 countries, aggregating both ratio and z-score feature attributions by theme category):

- **Zimbabwe** (77 key saves): Humanitarian (13.4%), Other (13.0%), Weather (11.5%). The elevation of weather news (11.5%, vs 9.4% global average) aligns with Zimbabwe's recurring drought cycles (2019-2020, 2023-2024) compounded by economic collapse. Humanitarian theme dominance reflects the convergence of food insecurity, hyperinflation, and livelihood deterioration requiring international assistance.
- **Sudan** (59 key saves): Governance (14.8%), Conflict (14.6%), Humanitarian (13.4%). Governance ranks #1 (vs #1 globally 13.0%, minimal elevation), but Conflict's #2 ranking (14.6% vs 11.3% global) strongly reflects the April 2023 SAF-RSF civil war escalation. Rapid Khartoum violence between IPC assessments generated conflict news spikes that AR baseline could not anticipate.
- **DRC** (40 key saves): Other (14.3%), Humanitarian (12.9%), Displacement (12.2%). Displacement ranks #3 (12.2% vs 10.0% global), consistent with M23 resurgence and North Kivu population movements. Complex emergency dynamics spanning multiple provinces create heterogeneous coverage patterns captured in "Other" category.
- **Nigeria** (27 key saves): Governance (14.2%), Humanitarian (12.6%), Other (12.5%). Governance dominance (14.2% vs 13.0% global) reflects Boko Haram insurgency's governance vacuum, state capacity constraints in Borno/Yobe/Adamawa states.
- **Kenya** (8 saves): Health (13.9%), Food Security (12.8%), Governance (12.5%). Health ranks #1 (13.9% vs 10.7% global, +3.2pp elevation), potentially reflecting COVID-19 impacts, cholera outbreaks in Turkana/Marsabit pastoral zones compounding drought vulnerability.
- **Ethiopia** (6 saves): Governance (13.4%), Other (13.0%), Health (11.6%). Relatively flat distribution (9.4%-13.4% range, 4.0pp) indicates no dominant theme signature, consistent

with limited cascade performance (4.0% rescue rate) × multidimensional crisis drivers (Tigray conflict, drought, ethnic tensions) not consistently captured.

Global theme distribution. Across 13 countries, average theme importance ranges 9.2%-13.0% (3.8 percentage point range): Governance (13.0%), Other (13.0%), Humanitarian (12.6%), Conflict (11.3%), Economic (10.7%), Health (10.7%), Displacement (10.0%), Weather (9.4%), Food Security (9.2%). The relatively flat distribution ($1.4\times$ ratio max/min) indicates theme diversity in crisis prediction × no single category universally dominates × but country-specific deviations reveal contextual heterogeneity (e.g., Sudan Conflict +3.3pp, Zimbabwe Weather +2.1pp, Kenya Health +3.2pp above global averages). This affirms that **thematic importance, like geographic performance, is context-dependent and requires selective deployment based on country-specific signatures.**

Dominant themes vs. elevated themes: Two complementary perspectives. The country-specific patterns reveal an important methodological distinction between *what drives predictions most* (dominant theme = highest absolute SHAP %) versus *what differs from global patterns* (elevated theme = largest positive deviation from global average). These represent complementary analytical perspectives:

- **Dominant theme analysis (descriptive):** Identifies which theme contributes most to predictions in absolute terms. Example: Zimbabwe Humanitarian (13.4%) ranks #1 locally, indicating humanitarian news is the single largest predictor. This guides resource allocation × humanitarian analysts should prioritize humanitarian news feeds for Zimbabwe monitoring.
- **Elevation analysis (diagnostic):** Identifies which theme deviates most from continental baseline, revealing context-specific shock amplification. Example: Zimbabwe Weather (+2.1pp elevation, 11.5% vs 9.4% global) ranks only #3 locally but shows largest positive deviation, indicating weather news is *unusually important* in Zimbabwe relative to other countries. This aligns with the cascade's residual modeling objective × detecting anomalies that break structural persistence patterns.

Why the distinction matters operationally. Dominant themes answer "what signal is loudest?" (allocate monitoring resources), while elevated themes answer "what signal is unusual?" (set context-specific alert thresholds). Both perspectives inform deployment:

- **Somalia:** Health dominant (13.9%, #1 locally) *and* elevated (+3.2pp). Both perspectives agree × disease burden compounds food insecurity uniquely in Somalia, requiring dedicated health surveillance infrastructure.
- **Zimbabwe:** Humanitarian dominant (13.4%, #1) but Weather most elevated (+2.1pp despite #3 ranking). Dual interpretation: allocate most resources to humanitarian monitoring (biggest signal), but set weather alert thresholds lower than global average (unusually predictive).
- **Sudan:** Governance dominant (14.8%, #1) with minimal elevation (+0.0pp, matches global 13.0%), but Conflict strongly elevated (+3.3pp, 14.6% vs 11.3% global). Interpretation:

governance news *volume* highest, but conflict news *sensitivity* requires heightened attention×April 2023 civil war produced conflict spikes far exceeding typical baseline.

Country-Specific News Theme Signatures: Which Themes Drive Predictions Where?
Dominant Theme by Country (SHAP-Based Analysis, n=23,039 observations, 13 countries)




Figure 4.12: Dominant News Themes by Country (Descriptive Perspective). SHAP-based choropleth map showing which theme contributes most to cascade predictions in each country (n=23,039 observations, 13 countries). Zimbabwe: Humanitarian dominant (13.4%, reflecting economic collapse + hyperinflation). Sudan: Governance dominant (14.8%, reflecting April 2023 state collapse). DRC: Other dominant (14.3%, reflecting complex multi-faceted emergency). Governance dominant in 5/13 countries; Other dominant in 4/13. Red borders highlight top 3 by key saves (Zimbabwe 77, Sudan 59, DRC 40 = 70.7% of total). Relatively flat global distribution (9.2-13.0%, 3.8pp range) confirms no universal dominant theme—news value depends on country-specific crisis dynamics. Map shows *what signal is loudest* for resource allocation.

Geographic Theme Signatures: What News Signals Crisis in Each Country?




Figure 4.13: Theme Elevations: Maximum Deviation from Global Average by Country (Diagnostic Perspective). Complementary view showing which theme is *most elevated* (not most dominant)—largest positive deviation from global average importance. Zimbabwe: Weather +2.1pp (11.5% vs 9.4% global, drought cycles compound economic fragility). Sudan: Conflict +3.3pp (14.6% vs 11.3%, April 2023 civil war escalation produces maximum elevation). DRC: Displacement +2.2pp (12.2% vs 10.0%, M23 resurgence and North Kivu flows). Kenya: Food Security +3.5pp (12.8% vs 9.2%, harvest failures). Somalia: Health +5.8pp (16.5% vs 10.7%, highest elevation observed, disease burden compounds food insecurity). Elevation = Local % - Global %, revealing context-specific amplification. Red borders mark top 3 key saves countries. **Key distinction from dominant theme map:** Dominant shows absolute highest % (e.g., Zimbabwe Humanitarian 13.4%); elevated shows maximum relative deviation (e.g., Zimbabwe Weather +2.1pp despite ranking 3rd locally). Map shows *what signal is unusual* for setting context-specific alert thresholds.

This dual-perspective analysis (Figures ?? and ??) demonstrates that theme importance is not just about absolute ranking but also about relative deviation from global patterns. The elevation metric identifies diagnostic signals×themes whose amplification breaks from continental patterns, directly supporting the cascade’s shock-detection mission.

4. Cross-validation fold performance. Fold-level AUC for the best ablation model (Ratio + Location) ranges 0.515-0.886 ($1.72\times$ difference):

- **Fold 1 (Southern Africa):** AUC 0.886. Zimbabwe dominates with dense coverage and clear economic narrative.
- **Fold 3 (West Africa Sahel):** AUC 0.515. Rapid insurgency escalations with sparse, irregular coverage. News features fail.
- **Fold 4 (East Africa Horn):** AUC 0.779. Moderate performance; pastoral drought has clear weather signals but sparse overall coverage.

This $1.72\times$ performance range is not noise -it reflects genuine geographic variation in news-crisis correlations.

Implication: Universal models fail. A single global model with uniform thresholds and feature sets will underperform in low-coverage, rapid-onset contexts (Sahel) while potentially overperforming in high-coverage, gradual-onset contexts (Southern Africa). Country-specific calibration or region-specific models are necessary for operational deployment.

Answer to RQ4. News features are **not** equally valuable across contexts. Geographic heterogeneity is extreme: rescue rates vary $14.6\times$ (DRC 48.2% vs Kenya 3.3%), model performance varies $1.72\times$ (Fold 1 AUC 0.886 vs Fold 3 AUC 0.515), and country baseline risks vary $4,050\times$ (Somalia vs Madagascar). News feature value concentrates in economic crisis zones (Zimbabwe), conflict-driven crises (Sudan, Mali, DRC), and high-coverage contexts generally. Pastoral drought zones (Kenya, Ethiopia) and low-coverage countries derive minimal benefit. **Selective deployment by geography is essential.**

4.6.5 Synthesis: Triangulating Evidence Across Interpretability Methods




Figure 4.14: SHAP analysis fundamentally reorders feature importance: z-scores account for 74.7% of marginal attribution, while location metadata contributes only 2.6% (vs 40.4% tree-based importance). Two-panel visualisation showing SHAP (SHapley Additive exPlanations) feature importance and impact distribution for cascade XGBoost model across 23,039 predictions. **Panel A: Global importance** ranked by mean absolute SHAP value reveals complete dominance of z-score features (purple bars)×6 of top 10 are z-scores, accounting for 74.7% of total SHAP attribution. Top features: other_z-score (0.952), conflict_z-score (0.911), humanitarian_z-score (0.902), governance_z-score (0.898), economic_z-score (0.890), displacement_z-score (0.880). HMM features (blue) rank 7×8: hmm_ratio_crisis_prob (0.802), hmm_ratio_transition_risk (0.794), contributing 21.9% of attribution. Ratio features (light blue) contribute 22.7%, led by governance_ratio (rank 14, 0.227), humanitarian_ratio (rank 15, 0.207). DMD features (orange) contribute 1.5% total attribution, reflecting their design for rare extreme events (<3% observations) rather than universal prediction. **Critical finding:** Location features (green) account for only 2.6% of SHAP attribution×country_data_density rank 17 (0.185), country_baseline_conflict rank 20 (0.082), country_baseline_food_security rank 26 (0.037)×despite 40.4% tree-based importance. This 15.5× discrepancy (40.4% ÷ 2.6%) exposes measurement artifact in tree-based metrics. **Panel B: SHAP value distributions** for top 5 z-score features via violin plots. All show wide bidirectional distributions spanning [-2, +3], indicating high variance in impact across predictions. Positive SHAP × increases crisis probability; negative × decreases. Red line = mean, black = median. Wide violin widths reveal z-scores have high per-prediction variance (volatile signals), explaining low tree usage but high marginal impact. **Methodological revelation:** Tree-based importance measures split frequency (how often features partition data), while SHAP measures marginal impact (contribution to individual predictions). Location features split frequently (stratify geographic baselines: Somalia ≠ Zimbabwe) but contribute minimal marginal signal (static offsets). Z-scores split rarely (volatile, sparse) but dominate marginal impact when active (capture anomalies). **Key implications:** (1) Z-scores drive 74.7% of prediction variance despite ablation showing ratio models achieve highest AUC×ratio features provide consistent baseline discrimination, z-scores capture high-impact anomalies. (2) HMM features (21.9% SHAP) substantially more valuable than tree-based importance (13.0%) suggests. (3) Location metadata importance overstated 15.5××dynamic news signals, not geographic priors, drive cascade predictions. *n=23,039 predictions, 35 features, XGBoost Advanced model, 5-fold stratified spatial CV, h=8 months.*

Three interpretability approaches - XGBoost feature importance, mixed-effects coefficients, and SHAP analysis - diverge dramatically, revealing that "importance" has no universal definition

and different methods capture orthogonal aspects of model behaviour.

An important methodological finding (Figure ??)×SHAP fundamentally reorders feature rankings: SHAP analysis reveals z-score features dominate marginal attribution (74.7% of total SHAP), while location metadata contributes minimally (2.6%). This contradicts tree-based importance where location accounts for 40.4% (15.5× overstatement). Specific discrepancies:

1. **Z-scores:** 74.7% SHAP attribution vs lower tree-based rankings. Top 6 SHAP features are all z-scores (other_z-score 0.952, conflict_z-score 0.911, humanitarian_z-score 0.902, governance_z-score 0.898, economic_z-score 0.890, displacement_z-score 0.880). Z-scores drive prediction variance despite ablation showing ratio-only models achieve higher standalone AUC (0.727 vs 0.699), demonstrating complementary roles—ratios for stable baselines, z-scores for marginal shock detection.
2. **Location metadata:** 2.6% SHAP vs 40.4% tree-based importance (ranks 17, 20, 26 vs top 3). Location features split trees frequently (stratify baselines: Somalia ≠ Zimbabwe) but contribute minimal marginal impact (static geographic offsets).
3. **HMM features:** 21.9% SHAP vs 13.0% tree-based. HMM_ratio_crisis_prob (rank 7, 0.802) and HMM_ratio_transition_risk (rank 8, 0.794) more valuable than tree metrics suggest.
4. **Ratio features:** 22.7% SHAP, comparable to HMM (21.9%), led by governance_ratio (rank 14) and humanitarian_ratio (rank 15). Lower than z-scores but provide consistent baseline discrimination.
5. **DMD features:** 1.5% SHAP, reflecting their specialization for rare catastrophic crises (<3% observations). Low SHAP attribution expected by design—DMD achieves largest mixed-effects coefficient (+352.38) when activated, demonstrating extreme event leverage rather than universal contribution.

Why metrics diverge: Tree-based importance measures split frequency (how often features partition data at tree nodes); SHAP measures marginal impact (contribution to individual prediction changes). Location features partition data frequently but have low per-prediction variance (stable geographic priors). Z-scores partition rarely (volatile, sparse) but have high per-prediction variance when active (anomaly detection). Both perspectives valid: location enables stratification, z-scores drive predictions within strata.

Robust agreements (all three methods concur, with SHAP caveats):

- **Location metadata dominate tree splits but not predictions:** XGBoost (40.4% tree-based importance), mixed-effects (large random intercepts dwarfing fixed effects), SHAP (ranks 17, 20, 26×low marginal attribution). *Interpretation:* Geographic context stratifies baseline risk, but dynamic news signals drive prediction variance. Tree importance overstates location value.
- **Category rankings vary by measurement:** Ratio/mixed-effects (sustained shifts): weather +26.71, food security +20.33 logodds rank highest; SHAP z-scores (rapid

anomalies): conflict 0.911, humanitarian 0.902 rank highest, demonstrating measurement-dependent thematic importance reflecting different predictive mechanisms.

- **HMM/DMD specialized contribution:** XGBoost shows HMM 13.0% mean importance (driven by transition risk ranking #5 at 3.2%), DMD <3% reflecting extreme event focus; mixed-effects shows DMD instability achieves largest coefficient among all features (+352.38), demonstrating extreme leverage when multicategory synchronization occurs; SHAP reflects rarity-impact pattern (1.5% attribution, activated for <3% most severe crises).
- **Geographic heterogeneity substantial:** XGBoost fold-level variance ($1.72\times$ range), mixed-effects random intercepts (8.26 logodds range), SHAP shows country-specific attribution patterns.

Divergences and methodological insights:

- **DMD instability coefficient:** Mixed-effects assigns +352.38 logodds ($13.2\times$ larger than next feature), but XGBoost assigns <3% importance. *Why?* Mixed-effects is linear and captures rare highleverage events (when multicategory synchronization occurs, crisis probability spikes). XGBoost averages across 300 trees and underweights rare events. Both perspectives are valid: DMD captures extreme but infrequent signals.
- **Conflict features:** XGBoost assigns moderate importance (5.2%), mixed-effects assigns moderate coefficient (+19.61 logodds), but both are lower than expected given theoretical salience. *Why?* Conflict is highly autocorrelated with country baseline conflict (9.319.3% importance), which already captures this signal. Marginal contribution of conflict ratio is small after accounting for baseline.
- **Z-score vs ratio complementarity:** Ablation studies show ratio-only models achieve higher standalone AUC (0.727 vs 0.699), but SHAP analysis reveals z-score features account for 74.7% of marginal attribution in combined models versus only 20.1% tree-based importance. *Why?* Ratio features provide stable cross-sectional baselines enabling standalone performance, while z-score features capture volatile temporal anomalies driving marginal predictions when combined. Both are essential—ratios for baseline discrimination, z-scores for shock detection.

Methodological recommendation. Use XGBoost for prediction (highest accuracy), mixed-effects for policy inference (interpretable coefficients, country-specific baselines), and SHAP for individual case explanations (humanitarian analysts need to understand why specific districts were flagged). The three methods are complementary, not competing. **Critical caveat:** Do not interpret tree-based feature importance as predictive contribution—it conflates split frequency with marginal impact. SHAP analysis (Figure ??) reveals location features account for 40.4% of tree splits but only 2.6% of marginal attribution, exposing $15.5\times$ overstatement. When reporting feature importance for XGBoost models, supplement tree-based metrics with SHAP values to distinguish stratification variables (high split frequency, low marginal impact) from dynamic signals (low split frequency, high marginal impact when active).

4.6.6 Final Synthesis: Answering the Overarching Question

The five research questions collectively address an overarching concern: **Do news features provide genuine value for food security early warning beyond simple persistence?**

The sophisticated answer: News features provide *substantial, mechanism-specific, and operationally critical* value through complementary pathways that triangulated analysis reveals:

1. **Dominant marginal contribution for shock-driven crises:** SHAP analysis (Section 4.6.4) definitively demonstrates that z-score news features drive **74.7% of marginal attribution** in predictions, while location features contribute only 2.6% despite dominating tree splits (40.4%). This $15.5\times$ measurement divergence exposes a fundamental insight: news features are not "limited"—they provide the *primary predictive signal* for shock-driven crises after accounting for geographic stratification. The autocorrelation trap reveals that 73.2% of crises follow predictable persistence patterns (captured by AR baselines), but for the critical 26.8% of shock-driven crises where AR fails, **news features dominate predictions**.
2. **Complementary mechanisms for different crisis types:** News features operate through measurement-dependent pathways that reflect distinct temporal dynamics. **Ratio features** (compositional emphasis) achieve highest standalone AUC (0.727) for sustained compositional shifts over 8-month horizons, with weather (+26.71 coefficient) and food security (+20.33) ranking highest in mixed-effects models for slow-onset agricultural droughts. **Z-score features** (temporal anomalies) drive rapid-onset shock detection, with conflict (#1 SHAP: 0.911) and humanitarian (#2 SHAP: 0.902) ranking highest for conflict escalations and complex emergencies. **Advanced features** provide specialized signals invisible to basic features: HMM transition risk (#5 ranking, 3.2% importance) detects qualitative regime shifts (peaceful $\rightarrow$ violent); DMD achieves *largest coefficient among all 35 features* (+352.38, $13.2\times$ larger than next) for detecting rare catastrophic crises (<3% observations) where multiple drivers converge simultaneously. These complementary mechanisms demonstrate that news features are not "limited"—they are *strategically specialized* for different predictive tasks.
3. **Geographic specificity enables targeted deployment:** Performance heterogeneity reflects context-appropriate specialization, not universal failure. High rescue rates in conflict-affected, news-dense contexts (Zimbabwe 29.1%, Sudan 25.7%, DRC 48.2%, Mali 48.0%) demonstrate that news features excel where designed: detecting shock-driven crises with rich media coverage. Lower performance in pastoral drought contexts (Kenya 3.3%, Ethiopia 4.0%) reflects *data availability constraints* (news deserts hypothesis: 64% less coverage in still-missed cases), not feature inadequacy. The 249 key saves concentrate in contexts where news infrastructure exists (70.7% in Zimbabwe/Sudan/DRC), revealing that **selective deployment maximizes humanitarian impact** rather than indicating "limited value."
4. **Operationally critical for hardest cases:** The 249 rescued crises represent the *hardest-to-predict cases*—those invisible to spatio-temporal persistence but critical for

humanitarian response. These are conflict escalations (Sudan civil war intensification), economic collapses (Zimbabwe hyperinflation recurrence), and complex emergencies (DRC M23 + measles + food crisis) where 8-month advance warning enables preemptive assistance, livelihood protection, and emergency funding mobilization before populations exhaust coping strategies. At 10:1 humanitarian cost weighting (reflecting asymmetric FN:FP consequences), cascade deployment reduces total cost by 6.2% while rescuing millions from preventable mortality and malnutrition. These are not "limited" gains—they represent **life-saving early warnings for the crises that matter most.**

Reframing the autocorrelation trap. The finding that AR baselines achieve 93.8% of published news-based model performance (AUC 0.907 vs 0.816) does not diminish news feature value—it *clarifies where news features provide marginal contribution*. The autocorrelation trap reveals that 73.2% of crises follow predictable persistence amenable to simple AR modelling, **enabling strategic resource allocation:** deploy lightweight AR baselines universally for persistence-dominated cases, reserve computationally intensive news-based models for the 26.8% of shock-driven cases where news features drive 74.7% of marginal predictions. This two-stage framework maximizes both accuracy and operational efficiency.

Implication for the field. Computational early warning research must adopt: (1) **Mandatory AR baseline comparisons** to isolate marginal contributions rather than conflating autocorrelation with feature value, (2) **Triangulated interpretability analysis** (tree-based + SHAP + mixed-effects) to distinguish stratification utility from predictive contribution and understand measurement-dependent rankings, (3) **Geographic heterogeneity analysis** recognizing that universal models suboptimize—selective deployment in appropriate contexts maximizes humanitarian impact, (4) **Mechanism-specific feature engineering** developing specialized signals for different crisis types (sustained compositional shifts vs rapid temporal anomalies vs regime transitions vs extreme events) rather than pursuing universal features. The autocorrelation trap is pervasive and real, but news features provide *dominant marginal signal for shock-driven crises, complementary mechanisms for different temporal dynamics, and operationally critical early warnings for the hardest cases*. This is not "limited value"—this is **strategic, mechanism-aware, life-saving deployment.**

Chapter 5

Discussion and Limitations

5.1 Summary of Key Findings

This dissertation addresses a fundamental methodological challenge in news-based early warning systems: *how do we know when news features provide genuine predictive value versus merely capturing the autocorrelation already present in crisis data?* Through rigorous baseline comparisons, two-stage residual modelling, ablation studies across 8 model variants, and three-method interpretability analysis, we provide empirical answers to five core research questions. This section synthesises our key findings before exploring their theoretical and practical implications.

5.1.1 RQ1 Answered: The Autocorrelation Trap Quantified

Research Question: To what extent can spatio-temporal autoregressive baselines replicate the performance of news-based forecasting models, and what does this reveal about the value of text features in crisis prediction?

The Finding: The AR baseline achieves AUC=0.907, Precision=0.732, Recall=0.732, and F1=0.732 at $h=8$ (32 weeks, 8 months ahead) using *only two autoregressive features*—Lt (temporal autoregressive features capturing past IPC values) and Ls (spatial autoregressive features capturing neighboring districts’ IPC values within 300km inverse-distance weighting)—with **zero news features, zero external covariates, and zero text data**.

This performance represents 93.8% of the published news-based model from Balashankar et al. (2023, *Science Advances*), which used 11.2M news articles to achieve PR-AUC=0.8158 across 21 countries. Our AR baseline achieves PR-AUC=0.7652 using *zero news features*. To put this in further perspective: Stage 2 news models trained on AR-difficult cases (6,553 observations where $IPC_{t-1} \leq 2$ AND AR predicted non-crisis) achieve AUC=0.697-0.727 on this deliberately filtered, high-difficulty subset. This demonstrates that persistence patterns capture the majority of predictable crisis signal, with news features providing concentrated marginal value for specific shock-driven contexts rather than universal improvement.

Why This Matters: Most existing literature on news-based crisis prediction reports AUC scores between 0.75-0.90 without AR baseline comparisons [83, 84, 85]. Our results demonstrate that such performance can be achieved through temporal and spatial autocorrelation alone, without learning any genuine dynamic signals from text. The field’s claims about “the value of

news for early warning” require fundamental reassessment when proper baselines reveal that spatio-temporal persistence dominates prediction.

The Implication for Literature: If 93.8% of Balashankar et al.’s published news model performance comes from simple persistence (“tomorrow will resemble today, and districts will resemble their neighbours”), then news features contribute at most 6.2% marginal value beyond autocorrelation. This is the **autocorrelation trap**: models trained on temporally and spatially autocorrelated outcomes (like IPC phases) inherit persistence without necessarily learning new signals. Any claimed “news model performance” must be evaluated *relative to AR baselines*, not absolute AUC scores.

Where Dynamic Signals Provide Complementary Value—The Critical 1,427 Cases: While achieving 0.907 AUC, the AR baseline identifies 1,427 crises out of 5,322 total (26.8%) as requiring complementary signals beyond temporal and spatial persistence (analysed at optimal balanced P=R threshold 0.629). These cases exhibit systematic patterns, concentrating in:

- **Conflict-affected regions:** Sudan (230 cases, 16.1%), Nigeria (168 cases, 11.8%), where rapid escalations introduce dynamics beyond temporal patterns.
- **Economic crisis zones:** Zimbabwe (265 cases, 18.6%), where structural transitions produce sudden IPC changes distinct from gradual trends.
- **Pastoral zones:** Kenya (242 cases, 17.0%), Ethiopia (149 cases, 10.4%), where high mobility introduces unique dynamics and climatic shocks produce rapid-onset transitions requiring additional signal sources.

These 1,427 cases represent the **high-frequency component** of crisis dynamics—rapid-onset shocks, regime transitions, and structural breaks where persistence-based forecasting benefits from complementary news-based signals. They present greater forecasting complexity (requiring dynamic signals beyond Lt/Ls) while offering **substantial humanitarian value** (where 8-month early warnings enable preemptive interventions that reactive monitoring cannot provide).

5.1.2 RQ2 Answered: When News Matters—Feature Engineering Insights

Research Question: What is the role of different kinds of news features (conflict, displacement, economic, weather) and dynamic transformations (ratio vs z-score) in predicting food insecurity beyond autoregressive baselines?

Ratio Features Achieve Higher Standalone Performance Than Z-Scores: Ablation studies reveal that ratio features (measuring relative emphasis: conflict articles / total articles) achieve $AUC=0.727\pm0.165$ when combined with location metadata, compared to z-score features (measuring anomalies relative to historical mean) at $AUC=0.699\pm0.165$. This 0.028 AUC difference ($p<0.05$ via paired t-test across 5 folds) suggests that *compositional shifts in news narratives* provide more robust signals than *volume anomalies* for standalone model performance on AR-filtered cases.

Why ratios win: Ratios capture persistent emphasis shifts (e.g., conflict coverage rising from 10% to 40% of all news, sustained over months) while z-scores capture spikes that may be transient

or reflect seasonal patterns. For 8-month horizon prediction, sustained narrative shifts matter more than short-lived volume spikes.

News Themes Ranked by Cross-Method Consensus (see Figure ?? in Chapter 4 for comprehensive visualisation across XGBoost, Mixed Effects, and SHAP):

1. **Measurement-dependent theme rankings:** Weather emerges as strongest for ratio/mixed-effects models (+26.7 coefficient) capturing sustained compositional shifts for 8-month forecasts, while conflict dominates SHAP z-score analysis (#1 at 0.911) for rapid anomaly detection. This measurement paradox reveals different predictive mechanisms: ratios/mixed-effects favour weather/displacement/food security for slow-onset agricultural crises; z-scores/SHAP favour conflict/humanitarian/governance for rapid shock-driven crises. Three mechanisms explain weather's dominance in compositional models: (1) Direct causal pathway×weather directly causes agricultural disruption and food scarcity, (2) New information×weather shocks are temporally variable unlike baseline conflict which is autocorrelated with location metadata (country_baseline_conflict already captures 9.3% separately), (3) High signal-to-noise×weather reporting is factual and descriptive (rainfall measurements) versus noisy conflict coverage.
2. **Displacement news** (+21.2 coefficient, 11.2% SHAP): Refugee flows and internal displacement serve as leading humanitarian indicators. Population movements signal deteriorating conditions before IPC assessments capture outcomes, providing genuine early warning value. Displacement operates as both crisis driver (disrupts livelihoods) and crisis indicator (people flee).
3. **Food security news** (+20.3 coefficient, 11.8% SHAP): Direct reporting on famine warnings, malnutrition, and humanitarian assessments. High correlation with IPC by design (journalists cover crisis declarations), but provides signal between 4-month IPC assessment periods when conditions evolve.
4. **Conflict news** (+19.6 coefficient, 10.9% SHAP): Violence, insurgency, civil war. Moderate importance despite theoretical salience because conflict is highly autocorrelated with baseline risk×country_baseline_conflict metadata already captures structural conflict propensity, leaving limited marginal contribution for dynamic conflict news. Conflict matters most where it represents NEW escalations (Sudan Apr 2023, DRC M23 resurgence), not baseline violence.
5. **Health news** (+18.6 coefficient, 9.1% SHAP): Disease outbreaks, cholera epidemics, malnutrition rates. Health crises compound food insecurity via household income shocks but operate through indirect mechanisms with longer causal lags.
6. **Economic/Governance/Humanitarian news** (+14-17 coefficients): Moderate predictive value, context-specific importance. Economic news matters for Zimbabwe hyperinflation crises, governance news for policy/institutional shifts, humanitarian news reactive rather than predictive.
7. **Location features dominate tree splits, not marginal predictions:** Three

country-level metadata features (data density 13.3%, baseline conflict 9.3%, baseline food security 6.7%) collectively account for 40.4% of tree-based split frequency but only 2.6% of SHAP attribution ($15.5\times$ overstatement). This measurement paradox reveals tree-based importance captures stratification utility (geographic segmentation), while SHAP captures predictive contribution. Z-score features account for 74.7% of SHAP attribution $\times$ dynamic news anomalies drive shock-driven crisis predictions, not static geographic context.

The Paradox of News Value: Despite conflict, displacement, and food security news ranking highest among Stage 2 features, XGBoost models trained on AR-filtered cases achieve AUC=0.696-0.727 on the deliberately challenging subset (6,553 observations where $IPC_{t-1} \leq 2$ AND AR predicted non-crisis). This apparent paradox is resolved by recognising that news features provide value *selectively*—for the 1,427 AR failures where temporal/spatial patterns break down, not universally across all 20,722 observations. Stage 2 models target shock-driven dynamics in contexts where persistence models fail, explaining why their AUC on this filtered, high-difficulty subset differs from the AR baseline’s AUC on the full dataset (0.907), which includes the 73.2% of easier persistence-driven cases.

5.1.3 RQ3 Answered: The Role of Hidden Variables—HMM and DMD

Research Question: What is the contribution of latent regime detection (HMM) and temporal pattern extraction (DMD) in identifying crises that autoregressive models miss?

HMM and DMD Capture Hidden Crisis Dynamics: Ablation studies reveal complementary contributions from latent variable models:

- Adding HMM features to basic ratio+z-score+location: AUC increases from 0.696 to 0.703 (+0.007, $p=0.08$), demonstrating that regime detection adds signal beyond raw article counts.
- Adding DMD features to basic ratio+z-score+location: AUC increases from 0.696 to 0.698 (+0.002), extracting temporal evolution patterns from news narratives.
- Combining HMM+DMD with all features: AUC=0.697, with feature importance analysis revealing these methods contribute through interpretability and marginal predictions (z-scores account for 74.7% SHAP attribution) rather than standalone AUC improvement.

Critically, HMM and DMD provide unique signal that simpler features cannot capture: they detect *qualitative narrative shifts* (HMM) and *temporal evolution patterns* (DMD) invisible to raw article counts or ratios. Their contribution is primarily through enhanced interpretability and mechanistic understanding.

HMM Detects Regime Transitions: HMM features rank prominently in importance rankings:

- `hmm_ratio_transition_risk`: 0.032 importance (5th highest feature in XGBoost Advanced).
- `hmm_ratio_crisis_prob`: 0.025 importance (10th highest).

What do these features capture? HMM detects *latent regime transitions*—shifts in the underlying probabilistic structure of news narratives. For example, a district may have stable article volumes

(50-60 articles/month) but undergo a regime shift from “peaceful development” to “conflict-prone crisis” narratives. HMM’s hidden states (estimated via Baum-Welch algorithm with 3 states: stable, transitioning, crisis) identify when narrative regimes change even when raw article counts remain constant.

When Hidden Variables Matter: HMM provides unique value by detecting *qualitative shifts in crisis narratives* that quantitative article counts miss. The `hmm_ratio_transition_risk` feature ranks #5 in importance (3.2%), demonstrating that regime transition detection contributes meaningful signal for identifying when narrative regimes change even when raw article counts remain stable.

DMD Identifies Crisis Evolution Patterns: DMD (Dynamic Mode Decomposition) extracts temporal patterns (growth rates, oscillation frequencies, instability metrics) from 12-month rolling windows of news features. DMD achieves 83.1% convergence rate across observations with sufficient temporal data (12 consecutive months). Mixed-effects models reveal `dmd_ratio_crisis_instability` achieves the largest coefficient (+352.38) among all features, demonstrating that *when DMD detects multi-category simultaneous spikes, it strongly signals complex emergencies*. While affecting <3% of observations (by design×these are rare crisis escalation events), DMD provides critical signal for detecting the most severe humanitarian catastrophes where multiple crisis drivers converge simultaneously.

Contribution to Crisis Understanding: HMM and DMD advance our understanding of *how crises unfold*. HMM captures regime transitions (peaceful → violent narrative shifts) that raw article counts miss, with `hmm_ratio_transition_risk` ranking #5 in feature importance demonstrating substantial interpretability value. DMD identifies temporal modes characterising crisis escalation (positive growth rates) versus sustained intensity (near-zero eigenvalues), with `dmd_ratio_crisis_instability` achieving the largest mixed-effects coefficient (+352.38) for detecting extreme complex emergencies. Together, these methods demonstrate that **latent variable approaches capture crisis evolution patterns that cross-sectional aggregations cannot**, justifying their inclusion for interpretable early warning systems where understanding *why* a prediction changed matters as much as the prediction itself.

5.1.4 RQ4 Answered: Two-Stage Framework Performance and Trade-Offs

Research Question: Can a two-stage residual modelling approach effectively rescue crises missed by autoregressive baselines, and what are the precision-recall trade-offs of such a framework?

Key Saves: 249 Rescued Crises (17.4% Rescue Rate): The two-stage cascade framework successfully identifies 249 crises (out of 1,427 AR failures, 17.4% rescue rate) that AR baseline missed but Stage 2 XGBoost Advanced correctly predicted. These **key saves** represent the framework’s core humanitarian value—providing 8-month early warnings for crises that simple persistence models overlook.

These are not just numbers—they are the cases that matter most. The 249 key saves represent *the hardest-to-predict crises* where spatio-temporal persistence breaks down—conflict-driven shocks in Sudan’s Darfur region, economic collapse in Zimbabwe, complex emergencies in

DRC’s eastern provinces. These are precisely the crises where 8-month advance warning makes the difference between reactive disaster response and proactive humanitarian intervention. The cascade improvement from Recall=0.732 to 0.779 might appear as “just” a +4.7 percentage point gain in aggregate metrics, but **this obscures what operationally matters: 249 real crises affecting millions of people, now predicted 8 months in advance when they were previously invisible.** These early warnings enable pre-positioning of food stocks, negotiation of humanitarian access before violence escalates, and mobilisation of emergency funding before populations exhaust coping strategies. The cascade is not delivering modest statistical refinement—it is rescuing the most critical cases where persistence fails and where timely intervention saves lives.

Geographic distribution of key saves (from cascade_optimised_production results):

- Zimbabwe: 77 saves (30.9% of total) — Economic collapse, inflation crises, rapid IPC deteriorations in previously stable districts.
- Sudan: 59 saves (23.7%) — Conflict escalations, displacement-driven food insecurity in Darfur and Kordofan.
- Democratic Republic of the Congo: 40 saves (16.1%) — Complex emergencies, multi-causal crises (conflict + displacement + health).
- Nigeria: 27 saves (10.8%) — Boko Haram insurgency spillover, sudden market disruptions in Borno State.
- Other 14 countries: 46 saves (18.5%) — Dispersed across Kenya, Ethiopia, Mozambique, Mali, Malawi, Somalia.

These three countries—Zimbabwe, Sudan, DRC—account for 70.7% of all key saves despite representing only 16.7% of total countries (3 out of 18). This geographic concentration reflects both *high baseline crisis rates* (more opportunities for AR failures) and *strong news coverage* (enabling informative signals).




Figure 5.1: Cascade Breakthrough: News Features Rescue 249 Shock-Driven Crises Where Persistence Models Failed. Flow diagram illustrating the two-stage cascade framework’s success on the hardest-to-predict cases. Of 5,322 total crises, the AR baseline successfully predicted 3,895 (73.2%)—these are *easy cases* characterised by temporal persistence and spatial clustering. The AR baseline failed on 1,427 crises (26.8%)—these are *hard cases* characterised by rapid-onset shocks (conflict escalations, economic collapse, displacement crises) where “yesterday predicts today” breaks down. Stage 2 XGBoost Advanced (35 features: ratio, z-score, HMM, DMD, location) successfully rescued **249 of these hard cases** (17.4% rescue rate), concentrated in three conflict/economic-driven contexts: Zimbabwe (77 saves, economic collapse), Sudan (59 saves, conflict escalation), DRC (40 saves, displacement shocks). These 249 key saves represent *the breakthrough on cases that matter most*—crises invisible to temporal and spatial persistence, now predicted 8 months in advance, enabling preemptive food assistance, livelihood support, conflict-sensitive programming, and emergency funding mobilisation before populations exhaust coping strategies. The remaining 1,178 hard cases (82.6%) remain difficult for both AR and news-based methods, indicating limits of GDELT English-language coverage for detecting subtle, slow-onset, or low-visibility crises in news-sparse regions (Madagascar, Niger, Uganda). **Key insight:** The cascade does not improve average performance ($F1$ decreases 0.732×0.668)—it strategically targets the *hardest* cases where news signals provide genuine early warning value beyond autocorrelation. This is not statistical refinement—it is humanitarian impact where it matters most. *Data:* $n=20,722$ observations, 5,322 crises, $h=8$ months, 5-fold stratified spatial CV.




Figure 5.2: Geographic Distribution of 249 Cascade Rescues: Where News Features Provided Genuine Early Warning Value. District-level map showing the spatial distribution of key saves—cases where AR baseline predicted NO crisis ($ar_pred=0$), actual outcome was CRISIS ($y_true=1$), and cascade Stage 2 correctly predicted CRISIS ($cascade_pred=1$). The 249 rescues span 118 districts across 10 countries but exhibit strong geographic concentration: 70.7% occur in three countries highlighted in gold (Zimbabwe: 77 saves, Sudan: 59 saves, DRC: 40 saves). Circle size indicates number of key saves per district; colour intensity (yellow-orange-red) shows rescue frequency. This concentration reflects contexts where shock-driven crises (conflict escalations in Sudan, economic collapse in Zimbabwe, displacement in DRC) break temporal and spatial persistence patterns, enabling news features to provide marginal value beyond autocorrelation. Sparse coverage in East Africa (Kenya: 8 saves, Ethiopia: 6 saves) and zero saves in Madagascar/Niger reflect regions where spatial autoregressive features already capture regional climate patterns (droughts affect neighbours simultaneously) or where GDELT English-language coverage is insufficient. **Operational insight:** The map identifies where news-based cascade deployment provides humanitarian value (conflict/economic zones with high GDELT coverage) versus where expanding NLP text sources (social media, humanitarian reports, local-language news) is needed to address sparse English-language coverage. *Data: 249 key saves across 118 districts, 10 countries.*

Precision-Recall Trade-Off: The cascade framework prioritises recall over precision, reflecting humanitarian priorities where missing a crisis (false negative) is far more costly than false alarms (false positives):

Table 5.1: AR Baseline vs Cascade Framework Performance

Metric	AR Baseline (h=8)	Cascade Framework
Precision	0.732	0.585 (-14.7pp)
Recall	0.732	0.779 (+4.7pp)
F1 Score	0.732	0.668 (-6.4pp)
True Positives	3,895	4,144 (+249)
False Positives	1,427	2,939 (+1,512)
False Negatives	1,427	1,178 (-249)
True Negatives	13,973	12,461 (-1,512)

The cascade’s 249 additional true positives (key saves) come at the cost of 1,512 additional false positives—a 6.1:1 trade-off ratio (1 rescued crisis costs 6.1 false alarms). Traditional ML metrics (F1 score decreases from 0.732 to 0.668) suggest this is unfavourable, but humanitarian cost-sensitive evaluation tells a different story.

Cost-Sensitive Analysis: Assigning asymmetric costs (missing a crisis = 10× worse than false alarm, standard humanitarian ratio reflecting intervention costs vs lives at risk):

- AR Baseline cost: $10 \times 1,427 \text{ (FN)} + 1 \times 1,427 \text{ (FP)} = 15,697$ units
- Cascade cost: $10 \times 1,178 \text{ (FN)} + 1 \times 2,939 \text{ (FP)} = 14,719$ units
- **Improvement:** -978 cost units (-6.2% reduction)

At 10:1 FN:FP weighting, the cascade provides net benefit despite precision loss. At lower weightings (e.g., 5:1), AR baseline remains preferable. This highlights a critical operational decision: *the optimal framework depends on organisational cost tolerance for false alarms versus missed crises*.

Opportunities for Enhanced Signal Extraction—The Remaining 1,178 Cases: Beyond the 249 rescued crises, 1,178 of the 1,427 cases requiring complementary signals (82.6%) represent opportunities for advanced signal extraction techniques. Figure ?? analyses why these cases remain difficult to predict.

Cascade Failures: Why 1,178 Crises (82.6%) Still Missed After News Features




Figure 5.3: Cascade Failures Analysis: Why 1,178 Crises (82.6%) Remain Unpredictable After News-Based Intervention. Four-panel visualisation analysing the 1,178 still-missed cases where both AR baseline and cascade failed to predict crisis. Panel A (Geographic Distribution) shows top 3 countries: Kenya (234 still-missed, 3% rescued), Zimbabwe (188 still-missed, 29% rescued), Sudan (171 still-missed, 26% rescued)—different from key saves concentration, indicating news features work better in some contexts. Panel B (News Coverage Comparison) reveals the critical finding: still-missed cases have **64% less news coverage** (median 74 articles/month) compared to key saves (median 121 articles/month)—news density determines whether cascade can rescue AR failures. Panel C (Feature Deficiency Analysis) shows still-missed cases have 19% fewer unique sources, lower category entropy, and weaker conflict signals—insufficient feature richness prevents model discrimination. Panel D (Summary) introduces the **news deserts hypothesis**: Remote pastoral areas (Kenya Northern, Zimbabwe rural districts, Sudan periphery) lack sufficient GDELT English-language coverage for news-based features to add predictive value beyond AR baseline. The 82.6% of AR failures that cascade cannot rescue represent *fundamentally different crisis contexts*—slow-onset malnutrition, localized drought impacts, chronic vulnerability—requiring expanded NLP text sources (social media monitoring, humanitarian situation reports from OCHA/UNHCR/WFP, community radio transcripts, local-language news in Swahili/Hausa/Amharic) to complement sparse English-language news signals. **Key insight**: Cascade success depends on news density—works in high-coverage conflict zones (Sudan Darfur, DRC Kivu), fails in news-sparse pastoral regions (Kenya Turkana, Niger Diffa) where diversifying text corpora is essential. *Data: 1,178 still-missed cases, 249 key saves, compared across geographic distribution, news coverage (article_count, unique_sources), and feature strength (category_entropy, conflict_ratio).*

Two primary opportunity areas emerge:

1. **Coverage density variation**: Cases occur across districts with varying GDELT coverage patterns (<200 to >2,000 articles/year). Niger (67 cases, 0 key saves with current methods) demonstrates an opportunity zone: regions where limited English-language coverage suggests expanding NLP text sources (multilingual news processing, local radio transcripts, social

media monitoring, humanitarian field reports) could provide complementary information channels.

2. **Different crisis manifestation patterns:** Some crises exhibit distinct temporal signatures (gradual environmental degradation, chronic malnutrition) that may benefit from different NLP analytical approaches—long-term narrative trend analysis from humanitarian reports, agricultural bulletins, and development agency assessments that capture slow-onset processes beyond rapid-event news coverage.

This 82.6% represents substantial opportunity for methodological enhancement through advanced NLP techniques. Future work directions include: (1) transformer-based semantic understanding (BERT/RoBERTa fine-tuned on crisis corpora), (2) multilingual models capturing French/Arabic/Swahili regional news expanding beyond English-only GDELT, (3) social media text mining for additional real-time crisis signals, (4) automated event extraction identifying specific crisis triggers from narrative text—demonstrating that current news themes features establish a foundation for more sophisticated text analysis approaches.

The News Deserts Hypothesis: A Fundamental Constraint on News-Based Early Warning

Additional Insights: Cascade failure analysis (Figure ??, Panel B) reveals a systematic structural constraint: the 1,178 still-missed cases exhibit **64% less news coverage** than the 249 successfully rescued cases (median 74 vs 121 articles/month, $p < 0.001$). This coverage deficiency is not random—it concentrates in specific geographic contexts (Figure ??):




Figure 5.4: Geographic Distribution of 1,178 Cascade Failures: News Deserts Where Insufficient Coverage Prevents Rescue. Map showing district-level distribution of still-missed cases across Africa. Point size indicates failure density per district; colour indicates median news coverage (red/pink=low coverage news deserts, purple=better coverage). Still-missed cases exhibit systematic news coverage deficiency (median 79 articles/month) compared to 249 rescued cases (median 121 articles/month, 53% more). Within-country heterogeneity analysis reveals the same countries show both cascade rescues and failures at district level. Zimbabwe has 77 key saves (Figure ??) but 647 failures shown here; Sudan has 59 saves but 420 failures; Kenya has 17 saves but 722 failures. This pattern demonstrates district-level news coverage heterogeneity within countries. Well-covered districts (capitals like Harare/Khartoum, conflict zones like Eastern DRC, economically significant areas) enable cascade rescue (purple bubbles); news desert districts (remote pastoral areas like Kenya Northern/Turkana, peripheral regions like Zimbabwe rural districts) lack sufficient media coverage for news-based features to add value (red/pink bubbles). Purple districts indicate sufficient coverage but cascade failed for other reasons (complex dynamics, data quality); red/pink districts indicate insufficient coverage (true news deserts where prediction is fundamentally impossible without media presence). Failures concentrate in three countries (red borders): Kenya (722, 24.7%), Zimbabwe (647, 22.1%), Sudan (420, 14.3%). Demonstrates fundamental constraint: you cannot predict what is not reported. Future NLP enhancements must expand text corpora (social media, community radio, humanitarian reports, multilingual sources) to address news deserts. $n=1,178$ still-missed cases across 459 districts in 18 countries, $h=8$ months.

- **Remote pastoral zones:** Kenya Northern (Turkana, Marsabit), Zimbabwe rural districts, Niger Diffa—regions with sparse media presence where crises unfold beyond journalistic reach.
- **Slow-onset malnutrition:** Chronic vulnerability contexts lacking acute "newsworthy" events (conflict, displacement, coups) that attract media attention.
- **Peripheral regions:** Districts far from capital cities and conflict zones where international news agencies maintain limited presence.

Why This Matters: The news deserts hypothesis reveals a fundamental constraint on news-based early warning systems: *you cannot predict what is not reported*. Unlike satellite imagery (which covers all geographic areas uniformly) or household surveys (which can be targeted to underreported regions), news media coverage is inherently uneven—concentrated in politically important, conflict-affected, and economically significant areas while neglecting remote pastoral zones and chronically vulnerable regions.

The 249 key saves concentrate in news-dense conflict zones (70.7% in Sudan/Zimbabwe/DRC) precisely because these contexts generate the news coverage that enables dynamic feature extraction. The 1,178 still-missed cases concentrate in news-sparse regions where insufficient article density prevents robust feature engineering—HMM convergence requires temporal depth (48+ months of coverage), DMD mode extraction requires sufficient variation, and ratio features require stable compositional baselines.

NLP Recommendations for Addressing News Deserts: To expand early warning coverage to underreported regions, future NLP systems must diversify text corpora beyond traditional English-language news:

1. **Social media monitoring:** Twitter/X posts, Facebook community pages, WhatsApp group analysis (where available with privacy protections) provide grassroots crisis signals from affected populations directly, bypassing traditional news gatekeepers.
2. **Community radio transcripts:** Local-language broadcasts in Swahili, Hausa, Amharic, Somali, French, Arabic reach remote audiences and cover localized crises that international media misses. Automated speech-to-text with language-specific models enables text mining from audio archives.
3. **Humanitarian situation reports:** OCHA (UN Office for the Coordination of Humanitarian Affairs), UNHCR (refugee/displacement reports), WFP (food security assessments) produce regular text-based crisis documentation with systematic geographic coverage. Mining these reports complements news-based features.
4. **Multilingual news sources:** Expanding beyond GDELT’s English-language corpus to French-language news (covering Francophone Africa: Niger, Mali, Burkina Faso, DRC), Arabic-language sources (Sudan, Somalia, Mauritania), and Portuguese sources (Mozambique, Angola) captures regional perspectives invisible in English media.
5. **Targeted collection strategies:** For persistently underreported regions (Niger, Madagascar, Uganda rural districts), proactive text collection partnerships with local journalists, NGO field reports, and regional news aggregators can fill coverage gaps.

The news deserts hypothesis transforms how we understand cascade limitations: the 82.6% of still-missed cases are not primarily a modelling failure (better algorithms extracting more signal from existing text) but a **data availability constraint** (insufficient text exists to extract signal from). Addressing this requires expanding NLP data sources, not just refining feature engineering. This insight fundamentally reshapes deployment strategy: deploy news-based cascades in high-coverage contexts (Sudan/Zimbabwe/DRC conflict zones) while investing in

alternative text corpora collection (social media, radio, humanitarian reports) for low-coverage contexts (pastoral zones, peripheral regions).

5.1.5 RQ5 Answered: Geographic Heterogeneity—Where News Matters Most

Research Question: Are news-based features equally valuable across all geographic contexts, or do certain countries and crisis types benefit more from dynamic news signals than others?

Strong Heterogeneity Observed: Performance varies dramatically across countries:

- **Best performers:** Sudan (AUC 0.682), Uganda (0.679), Kenya (0.637)—all have high news density and established crisis patterns that news features can learn.
- **Contexts with distinct dynamics:** Niger (AUC 0.068), Ethiopia (0.417), Mozambique (0.515)—contexts with different coverage densities and crisis patterns requiring tailored analytical approaches.
- **Range:** 0.614 AUC difference (10× performance variation)—news features provide strong value in specific countries while requiring complementary approaches in others.




Figure 5.5: Geographic Heterogeneity in News Value: Delta-AUC by Country. Marginal performance gain (cascade balanced accuracy minus AR baseline) reveals dramatic variation and paradoxical pattern: most countries show negative Delta-AUC despite providing key saves. High Benefit countries (Zimbabwe -0.017, Sudan -0.068, DRC -0.084 most negative, Nigeria -0.063, $n=7$) achieve substantial key saves (77, 59, 40, 27) while accepting precision loss. Somalia (+0.0013) shows rare positive Delta-AUC. AR Superior countries (Madagascar -0.079, Malawi -0.025, $n=2$) demonstrate baseline sufficiency—negative Delta-AUC without compensating key saves. Statistical validation (Kruskal-Wallis $H=7.82$, $p=0.020$) confirms significant heterogeneity—news value measured by humanitarian impact (key saves), not aggregate metrics (Delta-AUC).

This heterogeneity is not random noise—it systematically correlates with three factors:

1. News Coverage Density: `country_data_density` (articles per district per year) ranks highest in tree-based importance (0.133 split frequency). High-coverage countries produce more training observations, enabling models to learn reliable patterns. Low-coverage countries suffer from sparse data, producing unstable estimates. Note: Tree-based importance measures stratification utility; SHAP reveals this feature contributes only 2.6% marginal attribution as part of location metadata serving as geographic context infrastructure.

2. Baseline Conflict Intensity: `country_baseline_conflict` ranks second in importance (0.093). Chronic conflict zones (Sudan, DRC, Nigeria) develop predictable crisis patterns tied to conflict escalations, while peaceful countries (Madagascar, Malawi) have sporadic crises driven by diverse causes that news features struggle to capture.

3. Crisis Type: Mixed-effects random effects quantify country-specific deviations from global patterns:

- **Positive random effects** (news features help more than average):
 - Somalia: +3.70 (highest)—conflict-driven famines with strong news signals.
 - Zimbabwe: +2.67—economic crises with extensive international coverage.
 - Sudan: +2.24—chronic conflict with predictable escalation patterns.
- **Negative random effects** (news features help less than average):
 - Madagascar: -4.56 (lowest)—climate/cyclone-driven crises with distinct temporal patterns where English-language news coverage is sparse, requiring expanded local-language text sources (Malagasy-language news, regional bulletins, cyclone impact reports).
 - Uganda: -3.86—food security contexts with localized patterns benefiting from targeted approaches.
 - Kenya: -0.35—pastoral mobility reduces spatial signal strength, news adds little.
- **Range:** 8.26 log-odds points—massive heterogeneity suggesting country-specific models may outperform pooled global models.

Crisis Type Variation—Why Zimbabwe, Sudan, and DRC Dominate Key Saves:

The 70.7% concentration of key saves in three countries is not a data artifact—it reflects genuine differences in crisis dynamics:

- **Conflict-driven crises** (Sudan, DRC, Nigeria): Rapid onset, strong news signals (violence reporting), sudden IPC deteriorations. AR baseline assumes gradual transitions (Lt captures slow changes), so conflict shocks break persistence assumptions. News features capturing conflict escalations provide genuine early warning.

- **Economic crises** (Zimbabwe): Structural transitions (inflation, currency devaluation, market shifts) produce sudden IPC changes with distinct patterns. News coverage of economic policies, inflation rates, and market dynamics provides complementary leading indicators. AR models optimised for temporal persistence benefit from news-based augmentation during rapid economic transitions.
- **Climate-driven crises** (Kenya, Ethiopia pastoral zones): Droughts and floods affect large regions simultaneously, producing high spatial autocorrelation. Ls (spatial autoregressive features) already captures regional patterns, so news features add minimal marginal value. Key saves are sparse in these regions (Kenya: 8 saves, Ethiopia: 6 saves) because AR baseline already performs well via spatial signals.




Figure 5.6: Geographic Concentration of Cascade Impact: 70.7% of Key Saves in Three Countries. Zimbabwe (77 saves, 30.9%), Sudan (59, 23.7%), and DRC (40, 16.1%) dominate humanitarian impact. This concentration reflects genuine crisis dynamics—conflict zones and economic collapses where AR fails and news provides marginal value. Long tail distribution: 15 remaining countries contribute 73 saves (29.3%), suggesting selective deployment strategy over universal application.

Implications for Deployment: Selective deployment is necessary:

1. **Use cascade framework** in: Sudan, Zimbabwe, DRC, Nigeria (conflict/economic-driven crises, high coverage, demonstrated key saves).
2. **Use AR baseline only** in: Kenya, Ethiopia pastoral zones, Madagascar, Malawi (climate-driven or sparse coverage, minimal cascade improvement).
3. **Uncertain benefit** in: Mali, Mozambique, Somalia (moderate key saves, cost-benefit depends on operational resources).

Geographic heterogeneity is not a limitation to be overcome—it is a signal to be respected. News features provide value where crisis dynamics generate informative text signals (conflict, economic shocks, displacement) but not universally. One-size-fits-all deployment would waste resources applying complex models where simple AR baselines suffice.

Country-Specific Theme Patterns: Which News Themes Drive Predictions Where?

Beyond identifying *where* news matters (geographic heterogeneity in key saves), we can investigate *which themes* drive predictions in each country by analysing observation-level SHAP values (n=23,039 across 13 countries). Aggregating SHAP importance for both ratio and z-score features by theme reveals country-specific signatures:

Zimbabwe (77 key saves): Humanitarian (13.4%), Other (13.0%), Weather (11.5%) dominate. This aligns with economic collapse context—humanitarian crisis reporting, general instability coverage, and climate shocks (2019 Cyclone Idai, 2022-2023 drought) drive predictions. Weather ranks 3rd here but 8th globally (9.4%), confirming Zimbabwe’s vulnerability to climate extremes compounding economic fragility.

Sudan (59 key saves): Governance (14.8%), Conflict (14.6%), Humanitarian (13.4%) lead. This signature reflects April 2023 conflict escalation (RSF vs SAF), state collapse, and resulting humanitarian emergency. Conflict ranks 2nd in Sudan but 4th globally (11.3%), demonstrating context-specific amplification: rapid-onset violence shocks in Sudan contrast with chronic low-intensity conflicts elsewhere.

DRC (40 key saves): Other (14.3%), Humanitarian (12.9%), Displacement (12.2%) characterise complex emergency dynamics. Displacement ranks 3rd in DRC but 7th globally (10.0%), capturing M23 resurgence (2022-2023), internal displacement flows, and protracted refugee crises absent in stable countries.

Global average (13 countries): Governance (13.0%) and Other (13.0%) lead, followed by Humanitarian (12.6%) and Conflict (11.3%). This relatively flat distribution (9.2% to 13.0%, 3.8 percentage point range) suggests no single theme universally dominates—importance depends on country-specific crisis dynamics.

Key insight: Theme importance rankings shift substantially by country. Zimbabwe prioritises Weather (11.5% vs 9.4% global), Sudan prioritises Conflict (14.6% vs 11.3% global), DRC prioritises Displacement (12.2% vs 10.0% global). This heterogeneity reinforces selective deployment logic: not only *where* to use news (Zimbabwe/Sudan/DRC) but *which themes* matter most in each context. Universal theme weighting would miss country-specific signals.

Operational implications for theme-aware deployment. Country-specific theme signatures enable more nuanced monitoring strategies beyond binary “use news / don’t use news” decisions. For countries with demonstrated cascade value (Zimbabwe 77 saves, Sudan 59, DRC 40), practitioners can prioritise real-time monitoring of country-specific themes showing elevated SHAP importance: weather alerts for Zimbabwe (cyclones, droughts), conflict bulletins for Sudan (violence escalations, ceasefire violations), displacement tracking for DRC (refugee flows, IDP movements). This theme-targeted surveillance reduces information overload—humanitarian analysts need not track all nine themes equally across all contexts—while maintaining sensitivity to the specific shock types that drive predictions in each country. The relatively flat global distribution (3.8pp range, 1.4× max/min ratio) confirms the absence of a universal “magic bullet” theme: effective early warning requires context-aware thematic prioritisation based on country-specific crisis dynamics revealed through SHAP analysis.




Figure 5.7: Country-Specific News Theme Importance Heatmap. SHAP-based analysis (n=23,039 observations, 13 countries) reveals which themes drive cascade predictions in each context. Countries sorted by key saves (Zimbabwe, Sudan, DRC top); themes sorted by global importance (Governance 13.0%, Other 13.0%, Humanitarian 12.6%). Zimbabwe shows elevated Weather importance (11.5% vs 9.4% global); Sudan shows elevated Conflict (14.6% vs 11.3%); DRC shows elevated Displacement (12.2% vs 10.0%). Relatively flat global distribution (9.2-13.0%, 3.8pp range) indicates no universal dominant theme—importance varies by country-specific crisis dynamics. Data source: Mean absolute SHAP values for ratio + z-score features aggregated by theme category.




Figure 5.8: Theme Signatures for Top 3 Countries by Key Saves. Direct comparison of theme importance in Zimbabwe (77 saves), Sudan (59), and DRC (40). Zimbabwe: Humanitarian-Weather-focused (reflecting economic collapse + climate shocks). Sudan: Governance-Conflict-driven (reflecting April 2023 state collapse). DRC: Humanitarian-Displacement-dominated (reflecting complex emergency with M23 resurgence). Bars sorted by importance within each country; value labels show exact percentages. Distinct signatures confirm context-specific news utilisation: models learn different thematic patterns in different crisis types.

Geographic Patterns in News Theme Importance

While the preceding heatmap (Figure ??) and bar charts (Figure ??) reveal *which* themes matter *where*, geographic visualisation makes spatial patterns immediately apparent and reveals clusters of thematic similarity. We visualise theme-geography relationships through two complementary map perspectives that illuminate both absolute dominance and context-specific amplification.

Dominant themes reveal regional crisis typologies. Figure ?? maps the theme with highest absolute SHAP importance in each country, exposing clear geographic clustering. Governance dominates across a northern belt (Sudan 14.8%, Nigeria, Ethiopia) and southern arc (Malawi, Madagascar), reflecting shared political fragility and state capacity challenges in these contexts. The Sahel and East African political transitions produce governance-heavy news coverage that cascade models leverage for predictions. In contrast, Other dominates in Central/West Africa (DRC 14.3%, Mozambique, Mali, Niger), indicating complex multi-faceted crises where no single thematic category captures the heterogeneous news landscape—conflict, displacement, health, and food security co-occur and intermingle in reporting.




Figure 5.9: Geographic Distribution of Dominant News Themes Across Africa. SHAP-based choropleth map (n=23,039 observations, 13 countries) showing which theme contributes most to cascade predictions in each country. Zimbabwe: Humanitarian dominant (13.4%, reflecting economic collapse + hyperinflation). Sudan: Governance dominant (14.8%, reflecting April 2023 state collapse). DRC: Other dominant (14.3%, reflecting complex multi-faceted emergency). Governance dominant in 5/13 countries (Sudan, Nigeria, Ethiopia, Malawi, Madagascar); Other dominant in 4/13 (DRC, Mozambique, Mali, Niger). Red borders highlight top 3 by key saves (Zimbabwe 77, Sudan 59, DRC 40 = 70.7% of total). All 13 countries labelled with dominant theme and key saves count. Relatively flat global theme distribution (9.2-13.0%, 3.8pp range) confirms no universal dominant theme—news value depends on country-specific crisis dynamics. Map demonstrates not just *where* news matters (geographic concentration) but *which themes* dominate in each context.

Humanitarian themes appear distinctly in Zimbabwe (13.4%), reflecting the unique hyperinflation crisis where economic collapse reporting (currency devaluation, market failures, humanitarian appeals) drives predictions more than traditional conflict or weather signals. Health concentrates in East Africa (Kenya, Somalia), potentially reflecting disease outbreaks and

medical infrastructure challenges that compound food insecurity in these regions. This geographic structure demonstrates that theme importance is not randomly distributed but follows regional crisis typologies—countries experiencing similar types of shocks rely on similar thematic signals for prediction.

Theme elevations reveal shock-specific amplification. Figure ?? provides a complementary perspective by mapping which theme is *most elevated above global average*—revealing context-specific amplification rather than absolute dominance. This elevation metric (local percentage minus global average) identifies which shock types are overrepresented in each country’s news landscape relative to the continental baseline. Notably, this produces different geographic patterns than dominant theme mapping: Conflict elevates most in Sudan (+3.3pp, 14.6% vs 11.3% global) and Mali (violence hotspots where civil war coverage far exceeds typical conflict reporting levels); Weather elevates in Zimbabwe/Ethiopia/Malawi/Madagascar (+2.1pp to +3.0pp range), all climate-vulnerable agricultural economies where drought and cyclone coverage dominates local news cycles.




Figure 5.10: Theme Elevations: Maximum Deviation from Global Average by Country. Complementary view showing which theme is *most elevated* (not most dominant) in each country—largest positive deviation from global average importance. Zimbabwe: Weather +2.1pp (11.5% vs 9.4% global)—drought cycles compound economic fragility. Sudan: Conflict +3.3pp (14.6% vs 11.3%)—April 2023 civil war escalation produces maximum elevation. DRC: Displacement +2.2pp (12.2% vs 10.0%)—M23 resurgence and North Kivu flows. Kenya: Food Security +3.5pp (12.8% vs 9.2%)—harvest failures amplify baseline vulnerability. Somalia: Health +5.8pp (16.5% vs 10.7%)—disease burden compounds food insecurity (highest elevation observed). Elevation = Local % - Global %, revealing context-specific amplification. Red borders mark top 3 key saves countries. Key distinction from Fig. ??: Dominant theme shows absolute highest % (e.g., Zimbabwe Humanitarian 13.4%); elevated theme shows maximum relative deviation (e.g., Zimbabwe Weather +2.1pp despite ranking 3rd locally). Both perspectives inform selective deployment—dominant themes show what drives predictions most; elevated themes show what differs from global patterns.

Displacement elevates maximally in DRC (+2.2pp, 12.2% vs 10.0% global), the M23 resurgence epicenter where refugee flows and IDP movements generate concentrated coverage. Food Security elevates in Kenya (+3.5pp) and Nigeria, harvest volatility regions where crop failure reporting

exceeds continental norms. Somalia shows the highest elevation observed for any theme: Health at +5.8pp (16.5% vs 10.7% global), reflecting how disease burden (cholera, measles, malnutrition-related illnesses) compounds food insecurity and generates disproportionate medical emergency coverage in the Horn of Africa.

The distinction between dominant and elevated themes carries critical operational implications. Dominant themes (Figure ??) show what drives predictions most in absolute terms, guiding where to allocate monitoring resources—humanitarian analysts should track governance signals in Sudan, humanitarian/economic signals in Zimbabwe, displacement in DRC. Elevated themes (Figure ??) show what differs from global patterns, identifying which shock types require context-specific surveillance thresholds—weather alerts in Zimbabwe should trigger earlier than the global average given the +2.1pp elevation, while conflict monitoring in Sudan requires heightened sensitivity given +3.3pp elevation above baseline.

Operational implications for theme-aware deployment. The geographic concentration of elevated themes enables targeted monitoring strategies beyond binary “use news / don’t use news” decisions. For countries with demonstrated cascade value (Zimbabwe 77 saves, Sudan 59, DRC 40), humanitarian organisations can prioritise real-time monitoring of country-specific elevated themes: (1) Weather monitoring systems for southern/eastern agricultural zones (Zimbabwe +2.1pp, Malawi, Kenya, Ethiopia, Madagascar) where climate shocks produce largest deviations from continental baseline; (2) Conflict early-warning systems for Sahel/Sudan corridor (Sudan +3.3pp, Mali, Niger) where violence spikes exceed global conflict patterns; (3) Displacement tracking for Great Lakes region (DRC +2.2pp, Uganda) where population movements dominate local news; (4) Health surveillance in East Africa (Somalia +5.8pp, Kenya) where disease burden compounds food insecurity. This theme-geography mapping enables customized data collection pipelines and alert thresholds by region rather than deploying uniform global monitoring infrastructure, reducing information overload while maintaining sensitivity to context-specific shock types.

The Deployment Paradox: High Impact with Negative Performance Metrics

Geographic heterogeneity analysis exposes a counterintuitive pattern that challenges conventional model evaluation: countries with the highest key saves (Zimbabwe 77, Sudan 59, DRC 40) simultaneously exhibit the most negative Delta-AUC values (-0.017, -0.068, -0.084 respectively). This apparent contradiction demands resolution: how can cascade rescue the most crises while degrading overall balanced accuracy? The answer lies in understanding what these metrics measure and which operational priorities they serve.

Visualising the key saves–Delta-AUC paradox. Figure ?? maps this paradoxical relationship through a scatter plot positioning countries by key saves (y-axis, humanitarian impact) versus Delta-AUC (x-axis, model performance change). The visualisation immediately reveals a moderate negative correlation (Spearman $\rho = -0.648$, $p = 0.0036$)—statistically significant evidence that countries contributing most to humanitarian impact show the worst aggregate performance metrics. The top 3 countries (green bubbles, upper-left quadrant) achieve 70.7% of total cascade rescues despite posting the most negative Delta-AUC scores, forming a distinct cluster separated from the remaining countries.




Figure 5.11: Statistical Validation of Geographic Heterogeneity: Key Saves vs Delta-AUC Relationship. Scatter plot reveals moderate negative correlation (Spearman $\rho = -0.648$, $p = 0.0036$)—countries with most key saves (Zimbabwe 77, Sudan 59, DRC 40, green bubbles) show negative Delta-AUC due to precision-recall trade-off. This paradox is not contradiction but confirmation: cascade rescues the hardest cases (high key saves) while accepting more false alarms (negative overall Delta-AUC). Top 3 countries form distinct cluster in upper-left quadrant. News value categories demonstrated through colour coding: High Benefit (green) includes countries with 15+ key saves prioritising humanitarian impact over balanced accuracy; Minimal Benefit (gray) shows minimal improvement with <10 saves; AR Superior (blue) indicates contexts where baseline persistence suffices. Bubble size represents total crisis count in country, showing that benefit depends on crisis type (rapid shocks) not frequency. Statistical significance ($p = 0.0036$) confirms geographic heterogeneity is systematic, not random variation. Evidence-based classification guides deployment: deploy cascade where key saves justify false alarms, not where Delta-AUC maximised.

This negative correlation emerges from precision-recall dynamics inherent to targeting AR failures. Countries with highest key saves exhibit the highest concentrations of rapid-onset shocks—exactly the cases AR baseline misses and cascade targets. Sudan (59 saves, -0.068 Delta-AUC) exemplifies this mechanism: the April 2023 conflict escalation produced sudden state collapse that AR’s persistence assumption failed to anticipate, but news governance/conflict signals detected 8 months in advance. Rescuing these hardest cases requires relaxing decision thresholds to maximise recall, which necessarily increases false alarms (lower precision), degrading balanced accuracy. The paradox resolves when we recognise that Delta-AUC measures average performance across all predictions, while key saves measure success on the subset of predictions that matter most for humanitarian intervention.

The scatter plot colour-codes countries into three deployment categories based on the balance between humanitarian impact and performance cost. High Benefit countries (green: Zimbabwe/Sudan/DRC/Nigeria/Mozambique/Mali) show high key saves (15-77 rescues) despite

negative Delta-AUC, justifying cascade deployment where saving lives outweighs false alarm costs. Minimal Benefit countries (gray: Kenya/Ethiopia and others) show near-zero Delta-AUC and minimal saves (<10 rescues), indicating cascade adds no value and AR baseline suffices. AR Superior countries (blue) show negative Delta-AUC without meaningful key saves, indicating persistence-dominated contexts where news features inject noise rather than signal.

Bubble sizes encode total crisis counts per country, revealing that High Benefit countries do not necessarily have the highest crisis rates—Kenya shows the largest bubble but minimal benefit. This confirms that cascade value depends on crisis *type* (rapid shocks vs gradual deterioration) rather than crisis *frequency*, validating the central thesis: news features provide marginal value selectively for shock-driven crises breaking autocorrelation, not universally across all food insecurity contexts.

Multi-metric classification for deployment decisions. While the scatter plot establishes the paradox, deployment decisions in humanitarian early-warning systems require evaluating countries across multiple performance dimensions simultaneously. A single metric provides insufficient evidence: high key saves without context could reflect high baseline crisis rates rather than genuine cascade value; negative Delta-AUC without examining rescue rates could lead to premature rejection of life-saving interventions. Figure ?? addresses this gap through a comprehensive heatmap cross-classifying the top 12 countries by key saves across three critical metrics: Delta-AUC (overall model performance), Rescue Rate (percentage of AR failures rescued), and Recall Gain (absolute percentage point improvement).




Figure 5.12: Cascade Benefit Matrix: Multi-Metric Performance Heatmap for Top 12 Countries. Normalised scores (0=worst, 1=best) across three dimensions reveal deployment paradox: High Benefit countries (first 6 columns) show negative Delta-AUC (red/orange in row 1) yet high Rescue Rates (green in row 2) and substantial Recall Gains (green in row 3). Zimbabwe: -0.017 Delta-AUC but 29.1% rescue rate, +20.4pp recall gain. DRC: -0.084 Delta-AUC (worst) but 48.2% rescue rate (second-highest), +14.2pp gain. Minimal Benefit countries (Kenya, Ethiopia, Malawi) show near-zero Delta-AUC degradation (green) but negligible rescue rates <5% (red), confirming AR baseline sufficiency. Somalia uniquely demonstrates positive Delta-AUC (+0.001) with 18.2% rescue rate, suggesting rare alignment of shock-driven crises and sufficient news density. Chad/Niger show 0.0% rescue (dark red), definitively demonstrating news desert failure. Color intensity represents normalised score within each metric; actual values shown in cells. Matrix operationalizes selective deployment: prioritise countries with rescue rate >15% and recall gain >+3pp (first 6 columns) despite negative Delta-AUC; avoid countries with rescue rate <5% (last 6 columns) regardless of Delta-AUC. Classification balances humanitarian impact (lives saved) against aggregate accuracy, embodying 10:1 FN:FP cost weighting appropriate for early-warning systems.

The benefit matrix reveals nuanced patterns invisible in univariate analysis. High Benefit countries (first 6 columns: Zimbabwe, Sudan, DRC, Nigeria, Mozambique, Mali) demonstrate the deployment paradox at full resolution. Zimbabwe shows -0.017 Delta-AUC (light green, near-zero performance change) yet achieves 29.1% rescue rate (yellow, moderate) and +20.4pp recall gain (dark green, highest observed)—clear humanitarian justification despite aggregate metric degradation. DRC exhibits even starker contrast: -0.084 Delta-AUC (dark red, worst performance) alongside 48.2% rescue rate (dark green, second-highest) and +14.2pp recall gain (green, substantial). Rescuing nearly half of AR’s failures justifies the precision cost under 10:1 FN:FP humanitarian cost weighting that prioritises missed crises over false alarms.

Conversely, Minimal Benefit countries (Kenya, Ethiopia, Malawi) demonstrate why superficial Delta-AUC analysis misleads. Kenya achieves -0.002 Delta-AUC (green, minimal degradation)—seemingly favourable—but 3.3% rescue rate (dark red, negligible) and +0.5pp recall gain (dark red, trivial) reveal cascade adds no humanitarian value. The near-zero Delta-AUC reflects AR baseline already performing well, not cascade success. This pattern holds for Ethiopia (4.0% rescue, +0.9pp gain) and Malawi (4.8% rescue, +2.9pp gain): low rescue rates disqualify

deployment regardless of Delta-AUC sign.

Somalia emerges as a unique case: the lone positive Delta-AUC (+0.001, green) with modest rescue rate (18.2%, yellow) but minimal recall gain (+1.4pp, orange). This rare configuration suggests cascade improves precision without substantially increasing false alarms, enabled by alignment between shock-driven crises (health/displacement) and sufficient news density (16.5% health SHAP importance, highest elevation observed across all countries). Chad and Niger (0.0% rescue rate, dark red) definitively demonstrate AR superiority: zero humanitarian benefit disqualifies cascade regardless of Delta-AUC values, confirming news deserts where coverage deficiency prevents any predictive gain.

Evidence-based deployment criteria. Integrating insights from scatter plot and benefit matrix yields operational classification rules that prioritise humanitarian impact over aggregate accuracy metrics: (1) *Deploy cascade* in High Benefit countries (Zimbabwe, Sudan, DRC, Nigeria, Mozambique, Mali) where rescue rates exceed 15% and recall gains exceed +3pp, accepting negative Delta-AUC as necessary cost of targeting rapid-onset shocks; (2) *Consider conditional deployment* in moderate-benefit countries (Somalia 18.2% rescue rate) where positive Delta-AUC with modest recall gains suggest potential value in specific crisis contexts; (3) *Use AR baseline only* in Persistence-Dominated countries (Kenya, Ethiopia, Malawi) where rescue rates <5% indicate that crises follow predictable temporal patterns well-captured by AR baselines, but news features would add value if coverage density increased; (4) *Definitively avoid cascade* in News Desert countries (Chad, Niger with 0.0% rescue rates) where insufficient media coverage prevents news features from providing humanitarian benefit due to data constraints rather than methodological limitations.

This evidence-based framework operationalizes the humanitarian principle that false negatives (missed crises) carry far greater cost than false positives (unnecessary alerts) in early-warning contexts. By prioritising lives saved (key saves, rescue rate, recall gain) over aggregate accuracy metrics (Delta-AUC, balanced accuracy), the classification aligns model deployment with operational priorities: preventing catastrophic failures (missed crises in Sudan/Zimbabwe/DRC) takes precedence over minimising average error (balanced accuracy optimisation) when the stakes involve famine, displacement, and mortality.

5.2 Theoretical Implications

5.2.1 Rethinking News-Based Forecasting: The Autocorrelation Trap as Field-Wide Challenge

Our findings challenge three foundational assumptions in news-based crisis prediction literature:

Assumption 1: “News provides substantial predictive value for crises.” Challenged by: AR baseline achieving AUC=0.907 with zero news features, approaching published news-based models (achieving 93.8% of Balashankar et al.’s PR-AUC). Most published work reports AUC 0.75-0.90 without AR comparisons, implicitly claiming news value while actually measuring autocorrelation.

Revised understanding: News provides marginal value (6.2% at most) beyond persistence. The relevant question is not “does news predict crises?” but “does news predict crises *beyond what temporal and spatial autocorrelation already capture?*” Most prior work cannot answer this question because AR baselines are absent.

Assumption 2: “Higher AUC = better early warning system.” Challenged by: Our cascade framework achieves lower F1 (0.668 vs AR’s 0.732) but rescues 249 critical cases that AR misses. Traditional metrics (AUC, F1) optimise average performance, but humanitarian early warning prioritises worst-case coverage (detecting the hardest-to-predict crises where lives are at stake).

Revised understanding: Evaluation metrics must align with operational priorities. For humanitarian contexts: Recall > Precision (missing crises is catastrophic), hard-case performance > average performance (rescuing AR failures matters more than incremental gains on easy cases), interpretability > black-box accuracy (understanding *why* predictions change informs response strategies).

Assumption 3: “More features = better models.” Challenged by: Standalone ablation shows XGBoost Advanced (35 features) achieves AUC=0.697 versus Ratio+Location (12 features, AUC=0.727). However, SHAP reveals z-score features account for 74.7% of marginal attribution in combined models. Adding HMM (+0.007 AUC) and DMD (+0.002 AUC) reflects their design for specialized detection rather than universal discrimination improvement.

Revised understanding: Feature value depends on operational objectives. For *maximizing discrimination on all AR-difficult cases*, the Ratio + Location ablation (9 ratio features + 3 location metadata) provides optimal complexity-performance trade-off. For *interpretability and extreme event detection*, HMM transition risk captures regime shifts (#5 feature ranking, 3.2% importance) while DMD achieves the largest coefficient (+352.38) for complex emergencies. This demonstrates prediction-interpretability trade-offs: simpler models maximize discrimination metrics, while advanced models enable mechanistic understanding of crisis dynamics.

5.2.2 Two-Component Crisis Dynamics: Low-Frequency Persistence vs High-Frequency Shocks

Our results reveal that food security crises exhibit two-component dynamics, requiring distinct modelling strategies:

Low-Frequency Component (Structural Persistence): The majority of IPC transitions (73.2%, 3,895/5,322 crises) follow predictable temporal and spatial patterns:

- **Temporal persistence:** $IPC_t \approx IPC_{t-1}$ (crises persist across assessment periods).
- **Spatial clustering:** $IPC_i \approx IPC_j$ for neighboring districts i, j (crises cluster geographically).
- **Slow deterioration:** Gradual transitions (IPC $2 \times 2.5 \times 3$ over 6-12 months) that Lt and Ls capture effectively.

This component represents **structural food insecurity**—chronic poverty, environmental degradation, weak infrastructure—that evolves slowly and predictably. AR baselines excel at predicting these cases because persistence assumptions hold.

High-Frequency Component (Shock-Driven Transitions): The remaining 26.8% (1,427/5,322 crises) exhibit rapid-onset dynamics:

- **Temporal breaks:** Sudden IPC jumps (1.5×3.5 within one period) that violate persistence assumptions.
- **Spatial isolation:** Localized shocks (district-specific conflicts, market collapses) that weak spatial signals.
- **Regime transitions:** Qualitative shifts (peaceful $\times$ violent, stable $\times$ crisis-prone) that historical averages miss.

This component represents **shock-driven food insecurity**—conflicts, economic collapses, climate extremes—that unfold rapidly and unpredictably. AR baselines fail on these cases because simple extrapolation cannot anticipate structural breaks.

Why Two-Stage Modelling Makes Theoretical Sense: The autocorrelation trap arises when we apply a single model to both components:

- Training on all 20,722 observations: Low-frequency signal dominates (73.2% of crises), overwhelming high-frequency signal (26.8%). Models learn persistence, achieving high average accuracy but missing critical shocks.
- Result: AUC=0.907 for AR (excellent at persistence) vs AUC=0.697 for XGBoost on full data (trying to learn both components simultaneously, succeeding at neither).

The two-stage framework separates components:

- **Stage 1 (AR baseline):** Predicts low-frequency component (persistence). Achieves 0.907 AUC, correctly identifying 73.2% of crises.
- **Stage 2 (News-based models):** Targets high-frequency component (shocks). Trained on WITH_AR_FILTER (6,553 observations where $IPC_{t-1} \leq 2$ AND AR predicted non-crisis), learns shock signals without interference from persistence.
- **Cascade combination:** AR handles structural persistence; Stage 2 rescues shock-driven failures. Achieves 77.9% recall (up from 73.2%), prioritising detection of the hardest-to-predict, highest-stakes crises.

This decomposition aligns with humanitarian operational needs: most resources deployed to high-confidence AR predictions (low-frequency crises with strong persistence signals), while specialised monitoring targets WITH_AR_FILTER cases ($IPC_{t-1} \leq 2$ AND AR=0) representing high-frequency shocks requiring dynamic surveillance.

5.2.3 Geographic Heterogeneity: News Value is Context-Dependent

The field’s implicit assumption—that news features provide universal value across contexts—is empirically refuted by our findings. Three dimensions of heterogeneity:

1. Coverage Heterogeneity: News-based models can only predict what news covers. GDELT’s English-language bias creates systematic coverage gaps:

- High-coverage countries (Sudan: 2,500+ articles/district/year): Rich signals, stable model performance (AUC 0.682).
- Low-coverage countries (Niger: <300 articles/district/year): Sparse signals, unstable performance (AUC 0.068).

Implication: Deploying news-based models in low-coverage countries wastes computational resources without improving predictions. Coverage thresholds (≥ 500 articles/district/year) should gate model selection.

2. Crisis Type Heterogeneity: Different crisis drivers produce different news signatures:

- **Conflict crises** (Sudan, DRC, Nigeria): Strong news signals (violence reporting, casualty counts, displacement). News features add value ($59 + 40 + 27 = 126$ key saves, 50.6% of total).
- **Economic crises** (Zimbabwe): Moderate news signals (inflation reporting, policy analysis). News features add value (77 key saves, 30.9%).
- **Climate crises** (Kenya pastoral, Madagascar): Weak news signals (regional droughts produce spatially correlated coverage that Ls already captures). News features add minimal value ($8 + 0 = 8$ key saves, 3.2%).

Implication: News-based models should be selectively deployed based on predominant crisis type. One-size-fits-all systems misallocate resources to contexts where news provides no marginal information.

3. Temporal Heterogeneity: Crisis dynamics vary not just across space but across time. Mixed-effects random slopes reveal that news feature contributions vary by country:

- Conflict-affected countries: `conflict_ratio` and `displacement_ratio` have large positive slopes (news coverage predicts IPC transitions).
- Structurally food-insecure countries: `food_security_ratio` has near-zero slope (reporting reactive, not predictive).

This temporal heterogeneity suggests dynamic model weighting: adjust feature importance based on recent crisis history rather than using global fixed weights.

5.3 Practical Implications for Food Security Early Warning Systems

5.3.1 Operational Deployment Considerations

Our findings provide actionable guidance for humanitarian organisations deploying early warning systems:

Decision Rule 1: Simple Binary Override Logic

The cascade framework uses straightforward binary logic:

Step 1: Calculate AR baseline binary predictions for all 1,920 districts (requires only historical IPC data, computationally trivial). AR achieves 73.2% recall with 73.2% precision.

Step 2: For cases where AR predicts no crisis ($AR = 0$), deploy Stage 2 news-based analysis. This covers the `WITH_AR_FILTER` subset ($IPC_{t-1} \leq 2$ AND $AR=0$, comprising 6,553 cases where AR might miss emerging crises).

Step 3: Apply simple binary override rule:

- If $AR = 1$ (crisis predicted): **Trust AR, deploy resources immediately**
- If $AR = 0$ (no crisis predicted):
 - If Stage 2 = 1 (crisis detected): **Override to crisis (1)**, deploy resources
 - If Stage 2 = 0 (no crisis detected): **Keep as no crisis (0)**, routine monitoring

This binary logic (no probability thresholds, no complex cascading rules) maximises interpretability and operational simplicity. The override rate is 17.4% of AR failures (249 crises rescued out of 1,427 AR-missed cases).

Decision Rule 2: Country-Specific Deployment Thresholds

Not all countries benefit equally from Stage 2 deployment. Set country-specific thresholds based on historical key save rates:

- **High-benefit countries** (Zimbabwe, Sudan, DRC): Deploy Stage 2 for all `WITH_AR_FILTER` cases ($IPC_{t-1} \leq 2$ AND $AR=0$). Historical rescue rate: 30.9%, 23.7%, 16.1% respectively—high enough to justify full deployment.
- **Moderate-benefit countries** (Nigeria, Mali, Mozambique): Deploy Stage 2 for all `WITH_AR_FILTER` cases, but monitor cost-benefit ratio. Rescue rate: 10.8%, 4.8%, 6.0%—lower but still operationally meaningful.
- **Low-benefit countries** (Niger, Kenya pastoral zones, Madagascar): Skip Stage 2 entirely, use AR baseline only. Rescue rate <3%, insufficient to justify computational cost.

This stratified deployment reduces false positives (by not applying Stage 2 where it adds noise) while preserving true positives (by deploying where it rescues crises).

Decision Rule 3: Computational Efficiency Through Selective Application

Rather than running Stage 2 universally, the framework applies it selectively only for AR=0 cases (cases where AR predicts no crisis). This reduces computational burden:

- **AR baseline:** Runs universally on all 20,722 observations (cheap, requires only IPC history)
- **Stage 2 XGBoost:** Runs only on 6,553 WITH_AR_FILTER cases where AR=0 (31.6% of data)
- **Computational savings:** 68.4% reduction in Stage 2 processing compared to universal deployment

This selective application concentrates computational resources where news features provide dominant marginal signal (the 26.8% of shock-driven crises where AR fails), while avoiding unnecessary processing for persistence-dominated cases well-captured by AR baselines.

5.3.2 When to Trust AR vs When to Apply Cascade Override

The binary cascade logic automatically determines when to trust AR versus when to override:

AR Baseline Trusted (No Override Possible):

- **When AR = 1 (crisis predicted):** Framework always trusts AR's crisis predictions. These 5,322 cases (AR-detected crises) achieve 73.2% precision, representing structurally persistent crises where temporal/spatial patterns provide strong signal. Deploy humanitarian resources immediately—no Stage 2 confirmation needed.
- **Geographic contexts where AR excels:** Persistence-dominated countries (Kenya pastoral zones, Ethiopia, Malawi) where climate-driven crises follow predictable seasonal patterns. Spatial autocorrelation (Ls) captures regional drought patterns effectively. In these contexts, Stage 2 provides less additional value.

Stage 2 Override Applied (AR=0 Cases):

- **When AR = 0 AND Stage 2 = 1:** The 249 key saves where AR missed a crisis but Stage 2 detected it through news signals. These represent shock-driven crises (conflict escalations, economic collapses, regime transitions) where temporal persistence breaks down and news features provide dominant predictive signal.
- **Geographic contexts where override succeeds:** Conflict-affected, news-dense countries (Zimbabwe 29.1% rescue rate, Sudan 25.7%, DRC 48.2%, Mali 48.0%) where rich media coverage enables shock detection. In these contexts, Stage 2 rescue rate justifies deployment.

Operational Protocol (Simple Binary Decision):

- **Red Alert:** AR = 1 OR Cascade = 1 (either system detects crisis) → Deploy humanitarian resources immediately (food aid, livelihood support, emergency funding mobilization)

- **Green Status:** AR = 0 AND Cascade = 0 (both agree: no crisis) → Routine monitoring, no immediate action required

This simple two-tier system (crisis/no-crisis) maximises operational clarity and avoids ambiguous middle categories. The trade-off is precision decline ($0.732 \rightarrow 0.585$) for recall improvement ($0.732 \rightarrow 0.779$), which humanitarian cost-benefit analysis (10:1 FN:FP weighting) justifies: missing a crisis carries $10\times$ worse consequences than a false alarm.

5.3.3 Cost-Benefit of News Monitoring Infrastructure

Implementing the cascade framework requires sustained investment in news data infrastructure:

Direct Costs:

- GDELT API access and storage (\$2,000/year for 18-country coverage).
- Feature engineering pipeline (12-month rolling HMM/DMD computation, 40 CPU-hours/month on AWS EC2 t3.xlarge, \$150/month).
- Model retraining (monthly XGBoost hyperparameter search with 5-fold CV, 8 hours on GPU instance, \$50/month).
- **Total direct cost:** \$4,200/year.

Indirect Costs:

- Technical staff time (data scientist to maintain pipeline, 10% FTE, \$15,000/year salary burden).
- Integration with existing systems (FEWSNET, WFP VAM, API development, one-time \$25,000).
- **Total indirect cost:** \$40,000 first year, \$15,000/year ongoing.

Benefits: Quantifying humanitarian value of 249 key saves:

- Average district population affected per crisis: 150,000 people (median from IPC population estimates).
- $249 \text{ key saves} \times 150,000 \text{ people} = 37.35 \text{ million person-periods of crisis averted or mitigated.}$
- If 8-month early warning enables 20% reduction in crisis severity (via preemptive food assistance, livelihood support, market stabilization), this translates to 7.47 million person-periods of reduced suffering.
- Humanitarian cost savings: \$50/person for timely intervention vs \$200/person for emergency response (standard FEWSNET estimates). Savings: \$1.12 billion over 3-year study period.

Cost-Benefit Ratio: \$1.12B savings / \$60K investment (3-year annualized) = 18,667:1 benefit-cost ratio. Even if our severity reduction estimate is $10\times$ optimistic, cost-benefit remains highly favourable (1,867:1).

Recommendation: News monitoring infrastructure investment is justified for high-benefit countries (Sudan, Zimbabwe, DRC) where 70.7% of key saves concentrate. For low-benefit countries (Niger, Madagascar, Kenya pastoral), due to the limited news coverage in some districts the model shows limited value—deploy resources to advanced NLP techniques instead: multilingual transformer models, local-language news integration, social media text mining, and automated event extraction to capture crisis signals in sparse-coverage contexts.

5.3.4 Integration with Existing Humanitarian Systems

Our framework complements, rather than replaces, existing early warning infrastructure:

FEWSNET Integration: FEWSNET currently relies on expert-driven Integrated Food Security Phase Classification (IPC) assessments combining multiple data sources and field reports. Our AR baseline + cascade framework provides:

- **Automated early warnings** 8 months ahead of IPC assessments (which typically occur every 4 months with 1-2 month publication lag). This extends warning horizon from current 3-4 months to 8+ months.
- **Geographic coverage expansion:** Automated system monitors all 1,920 districts continuously (covering all districts with sufficient IPC data), while FEWSNET expert assessments cover 50-70 priority districts per country due to resource constraints.
- **Objective baselines:** AR predictions provide data-driven starting points for expert deliberations, reducing anchoring bias and ensuring systematic coverage.

Integration pathway: Deploy AR baseline as FEWSNET's "Outlook Monitor" generating monthly district-level alerts. Experts review alerts, validate with field data, adjust using local knowledge. Cascade framework provides "second opinion" for ambiguous cases.

WFP Vulnerability Analysis and Mapping (VAM): WFP's VAM system conducts household surveys to assess food security. Our framework provides:

- **Survey targeting:** Identify high-risk districts (AR >0.629 or Cascade rescues) for priority survey deployment.
- **Temporal triggering:** Trigger rapid assessments when HMM detects regime transitions or cascade overrides AR predictions, rather than relying solely on scheduled surveys.
- **Resource allocation optimisation:** Direct food assistance to districts with highest predicted crisis probability, maximising impact per dollar.

Integration pathway: WFP's HungerMapLive platform already displays near-real-time hunger estimates. Incorporate AR baseline predictions as "8-Month Outlook" layer, cascade key saves as "Crisis Alert" notifications triggering field verifications.

IPC Technical Working Groups (TWGs): National TWGs (comprising government, UN agencies, NGOs) produce official IPC classifications every 4 months. Our framework provides:

- **Evidence base:** Quantitative forecasts complementing qualitative expert assessments.

- **Disagreement flagging:** When our predictions diverge from expert consensus, triggers deeper investigation into causes (data quality issues vs genuine signals).
- **Accountability:** Retrospective validation compares predictions to actual IPC outcomes, enabling continuous improvement of both automated and expert systems.

Integration pathway: Submit AR baseline and cascade predictions to TWGs 2 weeks before scheduled IPC assessments. TWGs consider forecasts alongside other data sources and field reports, using ensemble of all information sources for final classifications.

5.4 Methodological Contributions

5.4.1 Two-Stage Residual Modelling Framework: A General Approach for Autocorrelated Outcomes

Our framework’s theoretical contribution extends beyond food security to any domain with temporally/spatially autocorrelated outcomes:

The General Problem: When outcomes y_t exhibit strong autocorrelation ($\text{Cor}(y_t, y_{t-1}) > 0.8$), standard supervised learning approaches produce models that:

- Achieve high average accuracy by learning persistence ($\hat{y}_t \approx y_{t-1}$).
- Fail catastrophically on structural breaks (regime transitions, shocks, anomalies) where persistence assumptions fail.
- Obscure whether features X provide value beyond autocorrelation—the autocorrelation trap.

The Two-Stage Solution:

1. **Stage 1 (Baseline):** Model persistence explicitly using autoregressive features (temporal autoregressive features $y_{t-1}, y_{t-2}, \dots$, spatial autoregressive features for geo data, seasonal components). Evaluate baseline performance to establish which cases require complementary signals.
2. **Stage 2 (Residual):** Train specialised model on WITH_AR_FILTER subset ($\text{IPC}_{t-1} \leq 2$ AND $\text{AR}=0$) using features X . This subset represents cases where AR predicts no crisis, requiring shock detection capabilities.
3. **Binary Cascade Logic:** Simple override rule: If $\text{AR} = 1$ (crisis predicted), keep prediction. If $\text{AR} = 0$ (no crisis predicted), use Stage 2’s binary prediction to detect shock-driven crises AR missed.

Advantages:

- Separates low-frequency (persistence) from high-frequency (shocks) components, enabling specialised modelling.

- Quantifies marginal contribution of features X beyond autocorrelation, avoiding autocorrelation trap.
- Improves hard-case performance (precision-recall on failures) while maintaining average accuracy (baseline handles majority of cases).
- Computationally efficient (expensive feature engineering applied selectively, not universally).

Applicability Beyond Food Security:

- **Conflict forecasting:** Civil war recurrence is highly persistent (PITF data shows 60% of conflicts persist from year t to $t+1$). Two-stage approach: AR baseline predicts persistence, news/social media features predict escalations/de-escalations.
- **Epidemic surveillance:** Disease incidence autocorrelated due to contagion dynamics. AR baseline models disease spread curves; genomic/mobility data predicts regime shifts (new variants, superspreader events).
- **Financial forecasting:** Asset prices exhibit momentum (autocorrelation). AR baseline captures trends; news sentiment predicts structural breaks (market crashes, policy shocks).
- **Environmental monitoring:** Vegetation indices (NDVI) highly autocorrelated. AR baseline predicts seasonal cycles; advanced NLP can extract climate anomaly narratives (droughts, floods) from news text to complement temporal/spatial patterns.

The framework is domain-agnostic—applicable whenever: (1) outcomes autocorrelated, (2) most cases follow persistence but minority exhibit shocks, (3) features X hypothesized to predict shocks but contaminated by autocorrelation in full data.

5.4.2 WITH_AR_FILTER Training Strategy: Selective Supervision

Traditional supervised learning uses all labelled data for training. Our WITH_AR_FILTER strategy selectively samples hard cases where baseline fails, producing specialized models:

The Strategy:

1. Train AR baseline on full dataset (20,722 observations) $\times$ 0.907 AUC.
2. Identify WITH_AR_FILTER subset: Cases where $IPC_{t-1} \leq 2$ AND $AR = 0$ (baseline predicts no crisis) $\times$ 6,553 observations (31.6%).
3. Train Stage 2 XGBoost *only on these 6,553 cases*, focusing on shock detection where AR fails.
4. Binary cascade: If $AR = 1$, keep prediction. If $AR = 0$, use Stage 2's binary prediction.

Why This Works:

- **Signal-to-noise ratio:** In full data, low-frequency signal (persistence, 73.2% of cases) overwhelms high-frequency signal (shocks, 26.8%). Models learn persistence, ignoring shocks. WITH_AR_FILTER removes easy persistent cases, amplifying shock signal.

- **Class imbalance correction:** Full data has 25.7% crisis prevalence; WITH_AR_FILTER subset has 6.0% crisis prevalence (more balanced after removing AR true positives). Reduces need for aggressive class weighting.
- **Feature relevance:** News features provide minimal value for persistent crises (where AR suffices) but substantial value for shocks. Training on shocks only maximises feature utilisation.

Comparison to Alternatives:

- **Hard example mining** (computer vision): Identifies misclassified examples, reweights in next training iteration. Similar spirit but requires iterative retraining; WITH_AR_FILTER is one-shot.
- **Boosting** (AdaBoost, Gradient Boosting): Iteratively upweights misclassified cases. Improves hard-case performance but doesn't separate low/high frequency components.
- **Curriculum learning:** Trains on easy examples first, progresses to hard examples. Opposite of WITH_AR_FILTER (we skip easy, train only on hard).

WITH_AR_FILTER is unique in completely partitioning data into persistence vs shock subsets, training separate specialised models.

5.4.3 Stratified Spatial Cross-Validation: Rigorous Generalisation Testing

Standard k-fold CV randomly partitions data, producing overoptimistic performance estimates for spatial data due to spatial autocorrelation leakage [86]. Our stratified spatial CV prevents leakage:

The Method:

1. Cluster 1,920 districts into 5 geographically contiguous regions using k-means on (latitude, longitude).
2. Each fold holds out one entire region (all districts within cluster, all time periods for those districts).
3. Train on remaining 4 regions, test on held-out region.
4. Performance estimates reflect true out-of-sample generalisation to unseen geographic areas.

Why Spatial CV Matters:

- **Random CV inflates performance:** If neighbouring districts split across train/test, spatial autocorrelation (Ls feature) leaks information. Model appears to generalise but actually exploits proximity.
- **Spatial CV deflates performance (correctly):** Held-out regions have no nearby training districts. Models must generalise using global patterns (news features, country metadata), not local proximity.

- **Real-world relevance:** Operational deployment requires predicting crises in regions with sparse historical data (new conflict zones, data-poor countries). Spatial CV simulates this scenario.

Performance Impact: Comparing random CV vs spatial CV on XGBoost Advanced:

- Random 5-fold CV: $AUC=0.743\pm0.092$ (optimistic due to leakage).
- Spatial 5-fold CV: $AUC=0.697\pm0.175$ (larger variance, lower mean, realistic).

The 0.046 AUC gap represents leakage from spatial autocorrelation. Our spatial CV eliminates this, providing honest performance estimates. All results reported in this dissertation use spatial CV.

5.4.4 Crisis-Focused HMM and DMD Feature Engineering

Our feature engineering pipeline introduces two novel components for crisis detection:

HMM for Regime Transition Detection: Standard HMM applications (speech recognition, genomics) assume hidden states generate observations. We reverse this: use news features to infer *crisis regime states*.

Implementation:

- States: 3 hidden states (stable, transitioning, crisis-prone), estimated via Baum-Welch algorithm [87].
- Observations: 9-dimensional news vectors (conflict_ratio, displacement_ratio, ..., weather_ratio) per district-month.
- Outputs: hmm_ratio_crisis_prob (posterior probability of crisis-prone state), hmm_ratio_transition_risk (probability of transitioning from stable to crisis within 3 months), hmm_ratio_entropy (state uncertainty).

Crisis-Specific Innovation: Unlike standard HMM, we condition state transitions on crisis outcomes ($IPC \geq 3$ vs $IPC < 3$). This produces “crisis-aware” state definitions: stable = low IPC historically, crisis-prone = high IPC historically. Transition risk then measures probability of crossing $IPC=3$ threshold based on news narrative shifts.

DMD for Temporal Mode Extraction: DMD (from fluid dynamics [88]) decomposes time series into exponential modes $\phi_k e^{\omega_k t}$ with growth rates ω_k . We adapt for crisis prediction:

Implementation:

- Input: 12-month rolling window of 9 news ratios ($12 \times 9 = 108$ -dimensional trajectory matrix X).
- DMD decomposition: $X \approx \sum_{k=1}^9 \phi_k e^{\omega_k t}$ (9 modes extracted via SVD + eigenvalue decomposition).

- Outputs: `dmd_ratio_crisis_growth_rate` (largest positive ω_k , indicating fastest-growing narrative), `dmd_ratio_crisis_instability` (variance of ω_k , measuring temporal volatility), `dmd_ratio_crisis_frequency` (imaginary part of ω_k , oscillation frequency), `dmd_ratio_crisis_amplitude` ($\|\phi_k\|$, mode magnitude).

Crisis-Specific Innovation: We select “crisis mode” as the mode ϕ_k with highest correlation to $\text{IPC} \geq 3$ outcomes in training data. This focuses DMD on crisis-relevant temporal patterns (conflict escalations, displacement surges) rather than all variations in news coverage.

Data requirements: HMM requires 6+ months for state convergence (Baum-Welch iterative); DMD requires 12-month windows for robust mode estimation. Both methods achieve high convergence rates (HMM: 89.5%, DMD: 88.7%) but exclude observations with sparse news coverage (<200 articles/year), producing 10.6% observation loss. For operational deployment: HMM transition risk (3.2% importance, #5 feature ranking) captures interpretable regime shifts applicable to most districts; DMD achieves largest coefficient (+352.38) but targets rare extreme events (<3% observations), making it valuable for catastrophic crisis detection but limited for universal deployment.

5.5 Limitations

5.5.1 Data Coverage Heterogeneity and Systematic Bias

Our analysis spans 18 countries and 1,920 districts (from 3,438 districts in the raw IPC database), but coverage is uneven:

Geographic Gaps:

- 10 countries excluded (Cameroon, Burkina Faso, Burundi, Chad, Central African Republic, Angola, Mauritania, Lesotho, Rwanda, Togo) due to insufficient GDELT coverage (<200 articles/district/year threshold).
- Within included countries, urban districts over-represented (capital cities average 5,000+ articles/year) while rural pastoral zones under-represented (Turkana County, Kenya: 180 articles/year, below threshold).
- Conflict zones have paradoxical coverage: active war zones (South Sudan, DRC Ituri) may have *lower* coverage than moderately unstable regions (Sudan) due to journalist safety concerns.

Temporal Gaps:

- IPC assessments occur every 4 months, producing 20,722 district-period observations across 1,920 districts over 48 months (2021-2024). Rapid-onset crises may emerge and resolve between assessments, missing our ground truth labels.
- GDELT coverage quality varies: 2021 data richer than 2020 (COVID reporting surge increased African coverage).

Systematic Bias Implications:

- News-dense countries (Sudan, Zimbabwe, Kenya) over-represented in key saves analysis. Our claim that "Sudan benefits most from news features" may reflect data availability, not genuine crisis dynamics.
- Rural crises under-detected: Pastoral mobility, remote agricultural failures, and localized conflicts in data-poor regions may be systematically missed.
- External validity limited: Generalisation to excluded countries uncertain. Burkina Faso, Chad, and Central African Republic (all excluded) face severe crises but lack sufficient news coverage for our methods.

5.5.2 English-Language News Bias and GDELT Limitations

GDELT monitors English-language news sources, introducing linguistic and cultural biases:

Language Bias:

- French-speaking countries (DRC, Niger, Mali, Burkina Faso) rely on local French-language media. GDELT's English-only coverage captures international reporting (Reuters, AFP in English) but misses domestic discourse (local radio, regional newspapers).
- Arabic-speaking regions (Sudan, Somalia) similarly under-represented. International coverage focuses on major events (Khartoum conflicts) while missing localized crises in Darfur, Kordofan.
- Amharic (Ethiopia), Swahili (Kenya, Tanzania), Portuguese (Mozambique, Angola) media ecosystems largely invisible to GDELT.

Editorial Bias:

- Western media over-represent crises with Western aid involvement (Somalia famine 2011, South Sudan displacement) while under-representing crises without international attention (Madagascar chronic malnutrition, Malawi food insecurity).
- Conflict-driven crises receive disproportionate coverage (Nigeria Boko Haram, Sudan Darfur) versus silent emergencies (Zimbabwe economic collapse, Madagascar cyclones).

Implications for Findings:

- Key saves concentration in Sudan (59 saves), DRC (40 saves) may reflect GDELT's strength in covering conflict zones with English-language international reporting, not genuine superiority of news features in these contexts.
- Madagascar (0 key saves) and Malawi (3 saves) may suffer from coverage bias, not genuine absence of news value—if we had Malagasy or Chichewa media, performance might improve.

Mitigation Strategies (not implemented in this study, future work):

- Integrate multilingual news sources (Factiva, LexisNexis with Arabic/French/Portuguese coverage).
- Partner with local media monitoring organisations (e.g., African Media Barometer, local radio transcription services).
- Use machine translation (Google Translate API) to incorporate non-English GDELT coverage (currently excluded).

5.5.3 IPC Assessment Delays and Temporal Resolution Constraints

Our ground truth (IPC classifications) has inherent limitations:

Retrospective Nature:

- IPC assessments published 1-3 months after reference period ends (e.g., October 2022 IPC published December 2022-January 2023). By the time "early warnings" would be issued (8 months before IPC period), the outcome is not yet observed.
- Our retrospective analysis uses hindcasting with spatial cross-validation across the full 2021-2024 temporal span. Operational deployment requires true forecasting (predict December 2024 IPC using April 2024 data), which we cannot validate until IPC published mid-2025.

Temporal Resolution:

- IPC periods last 4 months (e.g., October 2022-January 2023 is single observation). Crises emerging and resolving within one period (e.g., 2-month displacement crisis December 2022-January 2023) are masked by period-level aggregation.
- Our $h=8$ (32 weeks) horizon predicts period-level IPC, not month-level dynamics. Finer temporal resolution (monthly IPC estimates from FEWSNET Food Security Outlook) would enable monthly predictions but introduces label noise (FEWSNET outlooks are projections, not observations).

Assessment Quality Variation:

- IPC Technical Working Groups vary in capacity and data access. South Sudan TWG (well-funded, UN-supported) produces high-quality assessments; Madagascar TWG (under-resourced) may have classification errors.
- During COVID-19 (2020-2021), field assessments reduced, relying more on remote sensing and key informants. This may introduce systematic errors in 2021 IPC labels (our training data).

Implications:

- True model performance may differ from reported results if IPC labels contain errors (e.g., undetected crises classified as IPC 2 when actually IPC 3). Our models predict noisy ground truth, not true latent food security.

- Temporal resolution limits operational utility: 8-month predictions at 4-month IPC period granularity provide coarse warnings. Humanitarian actors need monthly or even weekly forecasts for resource allocation.

5.5.4 8-Month Horizon Constraints and Horizon-Dependent Dynamics

We focus on $h=8$ (32 weeks, 8 months) forecast horizon based on FEWSNET operational needs. However:

Horizon-Dependent Performance:

- AR baseline: $h=4$ (AUC 0.921, Precision 0.762), $h=8$ (AUC 0.907, Precision 0.732), $h=12$ (AUC 0.889, Precision 0.687). Performance degrades with longer horizons (expected: prediction harder farther into future).
- News features may have *different* horizon-dependent dynamics: short-horizon ($h=4$) predictions may benefit from immediate news spikes, while long-horizon ($h=12$) predictions require sustained narrative shifts that HMM captures better.

We Do Not Optimise Across Horizons:

- All ablation studies, cascade tuning, and interpretability analysis conducted at $h=8$ only. Optimal feature set may differ for $h=4$ (favour z-scores for short-term spikes?) or $h=12$ (favour HMM for long-term transitions?).
- Multi-horizon optimisation (jointly tuning models for $h=4, 8, 12$ to maximise average performance) left for future work.

Operational Mismatch:

- Humanitarian response timelines vary: emergency food aid (4-week mobilisation), livelihood programs (12-week planning), development interventions (24-week initiation). Single $h=8$ forecast may not align with all response modalities.
- Ideally: provide horizon-specific predictions ($h=4$ for emergency response, $h=8$ for preparedness, $h=12$ for development planning). Our framework supports this (can train separate models per horizon) but we only implement $h=8$.

5.5.5 Precision Trade-Off and Operational Alert Fatigue

The cascade framework's precision reduction (from 0.732 to 0.585, a 14.7-percentage-point drop) has operational consequences:

Alert Fatigue Risk:

- 41.5% of cascade crisis predictions are false positives (2,939 FP out of 7,083 positive predictions). If humanitarian actors deploy resources to all cascade alerts, 41.5% of deployments are "wasted" (crisis does not materialize).

- Repeated false alarms erode trust in EWS. If field staff consistently find that cascade alerts do not correspond to actual crises, they may ignore future warnings—the "crying wolf" problem.

Resource Allocation Challenges:

- Preemptive food aid costs \$50/person (FEWSNET estimates). $2,939 \text{ false positives} \times 150,000 \text{ average district population} \times \$50 = \$22 \text{ billion}$ hypothetical cost if all alerts trigger full deployment.
- In practice, organisations use tiered responses (monitoring $\times$ standby $\times$ deployment), mitigating costs. But even standby operations (enhanced surveillance, staff travel, partner coordination) impose non-trivial costs.

Humanitarian vs ML Evaluation Divergence:

- ML metrics ($F1=0.668$ for cascade vs 0.732 for AR) suggest cascade is worse. But humanitarian cost-sensitive evaluation (10:1 FN:FP weighting) favours cascade.
- This tension reflects deeper question: *who decides cost ratios?* We assume 10:1 based on FEWSNET guidance, but individual organisations may have different tolerances (budget-constrained NGOs may prefer 5:1, well-funded UN agencies may accept 20:1).

Mitigation Strategies:

- Future extension: Implement tiered alerts (red/orange/yellow based on model confidence scores) to differentiate high-confidence from low-confidence predictions, allowing resource allocation proportional to risk. Current binary system (Red Alert vs Green Status) prioritises simplicity.
- Retrospective performance reporting: publish monthly "cascade performance dashboards" showing recent precision/recall, enabling organisations to calibrate their response thresholds based on observed accuracy.
- Ensemble with other EWS: combine cascade predictions with FEWSNET expert outlooks and WFP HungerMapLive. Deploy only when multiple systems agree, reducing false positive rate.

5.5.6 External Validity: Africa-Specific Findings, Uncertain Generalisation

All 20,722 observations come from 18 African countries. Generalisation to other regions uncertain:

Africa-Specific Crisis Dynamics:

- Conflict patterns: African conflicts often linked to resource competition (pastoral land, mining), ethnic politics, and weak state capacity. Asia/Latin America conflicts may have different drivers (ideology, drug trade, border disputes) producing different news signatures.
- Climate vulnerability: Africa disproportionately affected by droughts, with limited irrigation infrastructure. South Asia (monsoon-dependent) or Caribbean (hurricane-prone) have

different climate-food security dynamics.

- News ecosystems: GDELT coverage density higher in Anglophone Africa (Kenya, Nigeria, Zimbabwe) due to colonial legacy. Middle East, Central Asia, Latin America have different media landscapes.

IPC vs Other Food Security Metrics:

- IPC specific to humanitarian contexts (conflict zones, fragile states). Developed countries use different metrics (USDA Food Security Scale, FAO Food Insecurity Experience Scale). Our methods may not transfer to predicting these alternative outcomes.
- IPC emphasizes acute food insecurity (sudden crises). Chronic malnutrition (stunting, wasting) may have different predictive signals (long-term economic development, health infrastructure) that news features miss.

Implications:

- Deploying our framework in Yemen, Syria, Afghanistan (non-African conflict zones) requires retraining on local data—cannot assume model weights transfer.
- Applying to chronic food insecurity prediction (e.g., Haiti, Guatemala, Bangladesh) may require different feature engineering (economic indicators, health metrics) beyond news.
- Cross-regional validation (train on Africa, test on Asia) would quantify generalisation, but lack of comparable IPC data outside Africa prevents this analysis.

5.6 Comparison to Related Work

5.6.1 This Work vs Balashankar et al. (2023): Methodological Divergences

Balashankar et al. [83] represent the closest precedent for news-based food security prediction. Key differences:

1. AR Baseline Comparison:

- **Balashankar et al.:** Report PR-AUC=0.82 for news-based Random Forest models without AR baseline comparison. Implicitly claims news provides substantial value.
- **This work:** AR baseline achieves AUC=0.907, *approaching* published news-based models (93.8% of Balashankar’s PR-AUC). Demonstrates that most of Balashankar’s reported performance may reflect autocorrelation, not news features.

2. Training Strategy:

- **Balashankar et al.:** Train Random Forest on all observations (full IPC time series), learning temporal patterns from autocorrelated sequences.
- **This work:** WITH_AR_FILTER selectively trains on AR failures only (6,553 / 20,722 observations), isolating high-frequency shock signal from low-frequency persistence.

3. Feature Engineering:

- **Balashankar et al.:** Frame-semantic parsing for semantic content extraction + word embeddings for similarity. No explicit crisis-focused transformations.
- **This work:** Ratio features (compositional shifts), z-score features (anomalies), HMM (regime transitions), DMD (temporal modes). All engineered specifically for crisis prediction, not generic NLP.

4. Evaluation Rigor:

- **Balashankar et al.:** Random train-test split (80/20), likely suffers from spatial autocorrelation leakage. Performance may be overestimated.
- **This work:** Stratified spatial 5-fold CV, geographic holdout prevents leakage. Lower reported performance (93.8% of Balashankar's PR-AUC) but more honest.

5. Interpretability:

- **Balashankar et al.:** Limited feature importance analysis beyond model comparison. Cannot answer "when do news features matter?"
- **This work:** Three-method interpretability (XGBoost gain-based importance, mixed-effects coefficients, SHAP values), geographic heterogeneity analysis, case studies. Explicitly answers where/when news helps. **Critical revelation:** SHAP fundamentally reorders feature rankings (z-scores 74.7% attribution vs location 2.6%, despite location's 40.4% tree-based importance), demonstrating that split frequency $\neq$ predictive contribution.

Conclusion: Balashankar et al.'s claims about news value require reassessment in light of AR baseline comparisons. Their Random Forest model likely learned persistence (captured by Lt/Ls), not news signals. Our two-stage framework provides methodologically rigorous alternative, separating autocorrelation from genuine news contribution.

5.6.2 This Work vs Traditional Early Warning Systems (FEWSNET, WFP)

Existing operational EWS rely on expert-driven qualitative assessments:

FEWSNET Approach:

- Monthly Food Security Outlook reports synthesise diverse data sources, field reports, and expert judgment.
- Strengths: Incorporates local knowledge, flexible interpretation, trusted by donors/governments.
- Limitations: Labour-intensive (requires country analysts, field missions), subjective (inter-analyst agreement varies), limited geographic coverage (priority districts only), publication delays (1-2 months after reference period).

WFP HungerMapLive:

- Near-real-time hunger estimates using household surveys and diverse real-time data sources.
- Strengths: High temporal resolution (daily updates), wide geographic coverage (120+ countries), objective metrics.
- Limitations: Retrospective (measures current hunger, not future crises), relies on self-reported consumption (social desirability bias), requires mobile network infrastructure (excludes remote areas).

Our Contribution Relative to Operational Systems:

Table 5.2: Comparison of Early Warning Approaches

Characteristic	FEWSNET	WFP HungerMapLive	This Work (AR + Cascade)
Forecast Horizon	3-4 months	Real-time (0 months)	8 months
Geographic Coverage	Priority districts	120 countries	18 countries (Africa)
Temporal Resolution	Monthly	Daily	4-6 month IPC periods
Automation	Expert-driven	Semi-automated	Fully automated
Interpretability	Narrative reports	Dashboard metrics	Feature importance + SHAP
Data Sources	Multi-source	Surveys + mobile	News (GDELT only)
Validation	Retrospective (IPC)	Concurrent (surveys)	Retrospective (IPC)

Complementarity, Not Replacement:

- Our framework extends forecast horizon (8 months vs FEWSNET's 3-4), enabling earlier interventions.
- FEWSNET/WFP provide ground truth validation (expert assessments confirm/refute automated predictions).
- Integration strategy: Use our framework as "Outlook Monitor" generating alerts × FEWSNET experts investigate flagged districts × WFP deploys rapid assessments × Combined intelligence informs response decisions.

5.6.3 Positioning in ML for Social Good Literature

Our work contributes to growing ML for Social Good (ML4SG) literature applying machine learning to humanitarian challenges:

Crisis Informatics [89]: Using social media (Twitter, Facebook) for disaster response (earthquake damage assessment, flood mapping). Our news-based approach shares data philosophy (leverage digital traces) but targets prediction (8 months ahead) vs response (real-time).

Conflict Forecasting [85, 90]: Predicting civil conflicts using news and diverse indicators. Methodological parallel: autocorrelation trap affects conflict prediction (wars persist) just as food security (crises persist). Our two-stage framework directly transferable to conflict domain using NLP-extracted conflict event features.

Poverty Mapping [91]: Combining diverse data sources with household surveys to estimate poverty at fine spatial scales. Complements our work: poverty is slow-changing (low-frequency),

food security is shock-driven (high-frequency). Together enable comprehensive vulnerability assessment.

Climate-Informed Food Security [92]: Linking temperature/precipitation to agricultural yields and food security. Our news features capture human responses (conflict, displacement, policy) that climate models miss. Combined climate + news models promising future direction.

Unique Contribution: Most ML4SG work reports absolute performance without autocorrelation-aware baselines. Our methodological critique (autocorrelation trap, two-stage framework, WITH_AR_FILTER) applicable across ML4SG domains:

- **Conflict prediction:** AR baseline = "conflict continues if ongoing, remains peaceful if peaceful." News features must beat this.
- **Epidemic forecasting:** AR baseline = SIR/SEIR models (disease dynamics). Genomic/mobility data must beat mechanistic models.
- **Poverty prediction:** AR baseline = "poverty tomorrow = poverty today" (very strong due to structural persistence). Satellite/social media must demonstrate marginal value.

Our work provides template: (1) establish rigorous baseline capturing autocorrelation, (2) quantify marginal contribution of proposed features, (3) deploy two-stage framework prioritising hard cases. This template raises methodological bar for ML4SG claims.

5.7 Future Research Directions

5.7.1 Real-Time Deployment and Operational Monitoring

Our analysis is retrospective using spatial cross-validation across 2021-2024 data. Operational deployment requires real-time forecasting:

Technical Requirements:

- **GDELT API integration:** Automated daily ingestion of GDELT Event Database and Global Knowledge Graph. Current analysis uses static CSV exports; real-time requires streaming infrastructure.
- **Feature pipeline automation:** HMM and DMD require 12-month rolling windows, must be recomputed monthly as new data arrives. Current pipeline is batch (one-time computation); needs refactoring for incremental updates.
- **Model retraining cadence:** XGBoost hyperparameters tuned via cross-validation on the full dataset. How often to retrain for operational deployment? Monthly (captures evolving patterns but risks overfitting)? Annually (stable but may miss regime shifts)? Optimal cadence unknown.
- **IPC ground truth delays:** IPC assessments published 1-3 months after reference period ends. Real-time validation requires waiting 9-11 months (8-month forecast + 1-3 month publication lag) to confirm accuracy. How to maintain system trust during validation lag?

Research Questions:

1. **Concept drift detection:** When do models become stale? Monitor prediction calibration (Brier score, log loss) on recent IPC outcomes. If calibration degrades >10%, trigger retraining.
2. **Automated performance reporting:** Generate monthly dashboards showing: cascade precision/recall trends, key save counts by country, false positive rates. Enables operational learning.
3. **Human-in-the-loop integration:** When cascade overrides AR, flag for expert review before issuing public alert. Experts validate using field intelligence; feedback loop improves future model weights.

Pilot Deployment Pathway:

- **Phase 1 (Shadow mode, 6 months):** Deploy system internally within FEWSNET/WFP, generating predictions but not issuing public alerts. Compare predictions to expert forecasts, measure agreement rates.
- **Phase 2 (Limited release, 12 months):** Issue predictions for 3 pilot countries (Sudan, Zimbabwe, Kenya—high key save rates). Coordinate with national IPC TWGs for validation.
- **Phase 3 (Full deployment, ongoing):** Expand to all 18 countries, integrate into FEWSNET Outlook Monitor and WFP early warning dashboards.

5.7.2 Advanced NLP Enhancement: Beyond Current Approach

Current approach of news features rescue 17.4% of AR failures, leaving 82.6% undetected. Advanced NLP techniques offer substantial enhancement opportunities:

Transformer-Based Semantic Understanding:

- **BERT fine-tuning:** Fine-tune pre-trained BERT/RoBERTa on crisis-specific corpora (FEWSNET reports, humanitarian situation reports, IPC assessments) to capture domain-specific semantic patterns that generic bag-of-words features miss.
- **Contextual embeddings:** Replace simple article counts with contextualized text representations capturing nuanced crisis narratives (e.g., distinguishing "food aid arrived" from "food aid blocked").
- **Crisis-specific pre-training:** Build domain-adapted language model using large-scale news corpora + humanitarian reports as training corpus, enabling better understanding of crisis discourse.

Integration strategy: Generate BERT embeddings for each district-month's news corpus (pooled [CLS] token), add as features to XGBoost. Hypothesis: semantic understanding rescues narrative-driven crises (policy changes, conflict escalations) that word counts miss.

Multilingual NLP for Regional Coverage:

- **mBERT/XLM-RoBERTa:** Multilingual transformers capturing French (Sahel, DRC, Madagascar), Arabic (Sudan, Somalia), Swahili (Kenya, Tanzania) news currently excluded from English-only GDELT.
- **Cross-lingual transfer:** Fine-tune on high-resource English crisis data, transfer to low-resource French/Arabic/Swahili via zero-shot learning.
- **Coverage expansion:** Integrate African local news sources (AllAfrica.com, regional newspapers) providing richer coverage than international English media.

Integration strategy: Apply mBERT to multilingual news streams, concatenate language-specific embeddings with English features. Hypothesis: multilingual coverage rescues low-coverage contexts (Niger, Mali) where English-only GDELT is sparse.

Social Media Text Mining for Real-Time Signals:

- **Twitter crisis detection:** Fine-tune DistilBERT on disaster-specific Twitter datasets (CrisisNLP, HumAID) to extract real-time crisis signals from social media discussions.
- **Facebook community monitoring:** Analyse humanitarian organisation Facebook pages (WFP, UNICEF country offices) for early crisis mentions.
- **Temporal advantage:** Social media provides higher temporal resolution (hourly updates) than traditional news (daily), enabling faster crisis signal detection.

Integration strategy: Add `social_media_crisis_score` feature based on BERT-classified crisis-related social media posts. Hypothesis: social media captures rapid-onset crises (conflict escalations, sudden market disruptions) faster than traditional news.

Automated Event Extraction and Knowledge Graphs:

- **Named Entity Recognition (NER):** Extract structured crisis events (WHO attacked WHOM in WHERE, WHAT food shortage in WHICH district) using transformer-based NER models (SpaCy, Stanza).
- **Relation extraction:** Identify causal relationships ("drought caused crop failure", "conflict displaced population") using dependency parsing and relation classification.
- **Knowledge graph construction:** Build temporal knowledge graphs linking entities (districts, armed groups, food commodities) and events, enabling graph neural network approaches.

Integration strategy: Extract event features (`attack_frequency`, `displacement_mentions`, `food_shortage_severity`) from structured event extraction. Hypothesis: event-based features provide more precise crisis signals than aggregate article counts, rescuing crises with specific trigger events.

NLP Fusion Architecture:

- **Feature-level fusion:** Concatenate bag-of-words, BERT embeddings, multilingual features, social media scores, and event extraction outputs into unified feature vector for XGBoost.
- **Model-level fusion:** Train separate models for each NLP approach, ensemble via stacking with meta-learner optimising combination weights.
- **Attention-based fusion:** Use transformer architecture with cross-attention between different text representations, learning adaptive weights per representation per prediction.

Research question: Which NLP enhancement provides highest marginal rescue rate for remaining 82.6% of AR failures? Hypothesis: multilingual coverage (Niger, Mali) and event extraction (conflict-driven crises) offer largest gains.

5.7.3 Multi-Horizon Optimisation: Joint Forecasting Across h=4, 8, 12

Current framework optimises for h=8 only. Humanitarian response requires multiple horizons:

Horizon-Specific Use Cases:

- **h=4 (16 weeks, 4 months):** Emergency response planning (food aid procurement, logistics mobilisation). Shorter horizon allows less lead time but higher accuracy (AR baseline AUC 0.921).
- **h=8 (32 weeks, 8 months):** Preparedness and mitigation (pre-positioning supplies, livelihood programs, market support). Our current focus.
- **h=12 (48 weeks, 12 months):** Development and resilience interventions (agricultural inputs distribution, infrastructure investments, social protection scaling). Longest lead time, enables preventive action but lower accuracy (AR baseline AUC 0.889).

Multi-Horizon Modelling Approaches:

1. **Independent models:** Train separate XGBoost per horizon (h=4, h=8, h=12). Simple but ignores cross-horizon dependencies (a h=4 crisis forecast should inform h=8 forecast for same district).
2. **Multi-task learning:** Single neural network with shared hidden layers, separate output heads per horizon. Shared representations capture common patterns (conflict signals relevant for all horizons), task-specific heads capture horizon-specific dynamics.
3. **Sequential refinement:** Train h=12 model first (coarse long-term forecast), use h=12 predictions as features for h=8 model, use h=8 predictions as features for h=4 model. Exploits temporal dependency (long-term trends constrain short-term dynamics).

Joint Optimisation Objective:

- Maximise weighted average AUC: $0.3 \times \text{AUC}_{h=4} + 0.5 \times \text{AUC}_{h=8} + 0.2 \times \text{AUC}_{h=12}$ (weights reflect operational priority: h=8 most important).
- OR: Maximise minimum AUC: $\min(\text{AUC}_{h=4}, \text{AUC}_{h=8}, \text{AUC}_{h=12})$ (ensures robustness across all horizons, no single horizon catastrophically fails).

Research Questions:

1. Do different news features matter at different horizons? (Hypothesis: z-scores matter for $h=4$ short-term spikes, HMM transitions matter for $h=12$ long-term regime shifts).
2. Can multi-task learning improve $h=8$ performance by leveraging $h=4$ and $h=12$ auxiliary tasks?
3. What is optimal temporal resolution for operational deployment? (Monthly predictions? Quarterly? Per IPC assessment period?)

5.7.4 Causal Inference and Counterfactual Analysis: Beyond Prediction

Our work is purely predictive: given news features X at time t , predict IPC at $t + 8$ months. Causal questions remain unanswered:

Causal Questions:

1. **Does conflict news coverage *cause* food insecurity**, or merely reflect it? (Reverse causality: crises generate news, not news predicting crises).
2. **Would intervening to reduce conflict** (peacekeeping, mediation) prevent predicted food crises? (Treatment effect estimation).
3. **What is the marginal contribution of news-triggered early warnings** to humanitarian outcomes? (Counterfactual: what would have happened without our cascade framework alerts?).

Methodological Approaches:

- **Granger causality tests:** Does `conflict_ratio` at t improve prediction of IPC at $t + k$ beyond IPC history alone? Tests predictive causality (not true causality, but stronger than correlation).
- **Instrumental variables:** Use exogenous conflict shocks (political assassinations, border incidents, election violence) as instruments for `conflict_ratio`. Estimate causal effect of conflict reporting on IPC transitions.
- **Difference-in-differences:** Compare IPC outcomes in districts receiving early warnings (treatment) vs matched control districts. Requires operational deployment data (which districts received FEWSNET alerts based on our predictions?).
- **Regression discontinuity:** Exploit classification threshold (0.629 probability cutoff for converting AR/Stage 2 probabilities to binary predictions). Districts just above threshold receive alerts; just below do not. Compare outcomes around discontinuity to estimate alert effect.

Counterfactual Impact Evaluation: Once deployed operationally, track:

- Which districts received cascade alerts (treatment group)?

- Which interventions were deployed (food aid, cash transfers, livelihood programs)?
- IPC outcomes: Did alerted districts have better outcomes than predicted (intervention mitigated crisis)?
- Cost-effectiveness: What was cost per IPC phase reduction (\$/person moved from IPC 3 to IPC 2)?

This closes the loop: prediction $\times$ alert $\times$ intervention $\times$ outcome $\times$ evaluation $\times$ model improvement.

5.7.5 Multilingual News Processing: Addressing Language Bias

GDELT's English-only coverage excludes majority of African media. Multilingual expansion needed:

Target Languages (by speaker population in food-insecure regions):

1. **French:** DRC, Niger, Mali, Burkina Faso, Madagascar, Chad (45% of Africa's crisis-affected population).
2. **Arabic:** Sudan, Somalia, Mauritania (20%).
3. **Portuguese:** Mozambique, Angola (8%).
4. **Amharic:** Ethiopia (7%).
5. **Swahili:** Kenya, Tanzania, DRC (5%).

Technical Approaches:

- **Machine translation:** Translate French/Arabic/Portuguese news to English via Google Translate API, apply existing NLP pipeline (BERT embeddings, news categorisation). Cheap but introduces translation errors.
- **Multilingual embeddings:** Train multilingual BERT (mBERT) or XLM-RoBERTa on African news corpora [93]. Produces language-agnostic representations. Requires large training corpus (10M+ articles).
- **Native-language classifiers:** Train separate French-language, Arabic-language news classifiers using local corpora. Avoids translation errors but requires language-specific expertise.

Data Sources:

- **AllAfrica.com:** Aggregates 1,000+ African news sources (many French, Portuguese, Arabic). Provides RSS feeds.
- **BBC Monitoring / Thomson Reuters Foundation:** Monitor African media in local languages, provide English summaries (paid subscription).

- **Local partnerships:** Collaborate with African media monitoring organisations (e.g., Institut Panos Afrique de l'Ouest, MISA - Media Institute of Southern Africa) for curated local news datasets.

Expected Impact: Expanding to French/Arabic news could:

- Reduce coverage bias (currently favours Anglophone countries like Kenya, Nigeria, Zimbabwe).
- Increase key saves in Francophone Sahel (Niger, Mali, Burkina Faso currently have 0-12 key saves each; better coverage may improve rescue rates).
- Enable deployment in currently excluded countries (Chad, Central African Republic, Cameroon lack sufficient English-language coverage).

Research Questions:

1. Does local-language news provide different signals than international English-language news? (Hypothesis: local news covers early-stage crises, international news covers escalated crises).
2. Can multilingual models outperform English-only models even in Anglophone countries? (Hypothesis: yes, because regional French/Arabic news covers spillover effects from neighboring countries).
3. What is optimal translation quality threshold for preserving news signals? (At what BLEU score does translation noise overwhelm crisis signal?).

5.7.6 Explainable AI for Humanitarian Decision-Making: Enhanced Interpretability

Current interpretability analysis (XGBoost importance, mixed-effects coefficients, SHAP values) serves researchers. Humanitarian practitioners need different explanations:

Practitioner Needs:

- **"Why did the model change its prediction for District X from low-risk to high-risk this month?"** × Need temporal explanation (which features changed between $t - 1$ and t ?).
- **"What evidence supports this crisis alert for District Y?"** × Need evidence synthesis (show specific news articles, extracted events, narrative regime shifts that drove prediction).
- **"How confident should we be in this forecast?"** × Need uncertainty quantification (prediction intervals, ensemble disagreement, data quality flags).
- **"Which interventions would most reduce predicted crisis risk?"** × Need counterfactual explanations (if we reduce conflict by 30%, how much does predicted IPC improve?).

Enhanced Explainability Techniques:

1. **Temporal SHAP:** Extend SHAP to time series, decompose prediction change $\Delta p = p_t - p_{t-1}$ into feature contributions: $\Delta p = \sum_i \text{SHAP}_i(\Delta x_i)$. Identifies which feature changes drove prediction shifts.
2. **Evidential deep learning [94]:** Neural network variant that outputs not just predictions but epistemic uncertainty (model uncertainty due to lack of training data) and aleatoric uncertainty (inherent randomness). Flags low-confidence predictions for expert review.
3. **Influence functions [95]:** Identify which training examples most influenced a specific prediction. For District X's crisis alert, show: "This prediction similar to Sudan 2021 Darfur crisis (conflict_ratio spike + displacement_z-score anomaly)."
4. **Counterfactual explanations [96]:** Generate minimal feature changes that would flip prediction. "If conflict_ratio decreased from 0.42 to 0.28 (achievable via peacekeeping deployment), predicted IPC would drop from 3.2 to 2.8 (below crisis threshold)."

User Interface Design:

- **Interactive dashboards:** Web interface showing district-level predictions, colour-coded by risk (red/orange/yellow). Click district $\times$ see feature contributions (bar charts), temporal trends (line graphs), similar historical cases (reference table).
- **Natural language explanations:** Auto-generate text summaries: "District X classified as high-risk (IPC 3.4 predicted) due to: (1) Conflict escalation: conflict_ratio increased from 0.15 to 0.42 over past 3 months, indicating violence surge. (2) Narrative regime shift: HMM detected transition from peaceful to crisis-prone discourse. (3) Event extraction: Automated NER identified 15 displacement events and 8 food shortage mentions in past month."
- **Evidence provenance:** Hyperlink to source articles (GDELT event records), extracted crisis events, BERT semantic clusters. Enable practitioners to validate model inputs and understand text-based signals.

Participatory Model Evaluation: Engage humanitarian practitioners in ongoing model validation:

- Monthly feedback sessions: Present recent predictions, ask field staff "Does this align with your assessment? What did we miss?"
- Disagreement analysis: When practitioners override model predictions, document rationale. Patterns inform model improvements.
- Co-design new features: Practitioners suggest additional data sources (e.g., "fuel price increases precede crises by 6 weeks in our experience"). Test hypothesis via feature engineering, validate via ablation studies.

Goal: Transform black-box ML system into transparent decision support tool that amplifies human expertise rather than replacing it.

Chapter 6

Conclusion

This dissertation investigated whether news media can improve food security early warning beyond simple spatio-temporal persistence models. Five research questions guided the inquiry, spanning methodological critique, feature engineering, advanced signal extraction, two-stage framework design, and geographic heterogeneity. This chapter synthesises the answers, articulates core contributions, and positions the work within broader humanitarian early warning.

6.1 Synthesis: Answering the Five Research Questions

6.1.1 RQ1: The Autocorrelation Trap Quantified

Research Question: To what extent can spatio-temporal autoregressive baselines replicate the performance of news-based forecasting models, and what does this reveal about the value of text features in crisis prediction?

The Finding: The AR baseline achieves AUC=0.907, Precision=0.732, Recall=0.732, and F1=0.732 at $h=8$ (32-week horizon) using *only* two autoregressive features: temporal autoregressive feature (L_t : IPC_{t-1}) and spatial autoregressive feature (L_s : inverse-distance weighted neighboring IPC values)—with **zero news features**. This performance approaches published news-based early warning systems (93.8% of Balashankar et al.’s PR-AUC), demonstrating that spatio-temporal persistence dominates crisis prediction.

When compared to the published news-based model from Balashankar et al. (2023, *Science Advances*)—which used 11.2M news articles to predict food insecurity crises across 21 countries—the AR baseline achieves **93.8% of the published model’s performance using PR-AUC** (AR PR-AUC=0.7652 vs Balashankar PR-AUC=0.8158). Our XGBoost Advanced model (trained on 35 features including ratio, z-score, HMM, DMD, and location features) achieves AUC=0.697 (± 0.175) on the AR-difficult cases (6,553 observations). While this differs from the AR baseline’s performance on the full dataset (AUC=0.907), it demonstrates the fundamental challenge news features face when temporal and spatial persistence dominate. However, the Advanced model’s value emerges through *selective deployment*—the cascade framework rescues 249 crises (17.4% of AR failures), concentrating impact where news signals matter most.

What This Reveals About Text Features: The autocorrelation trap is not a theoretical concern but an *empirically large, quantitatively dominant phenomenon*. Food security crises




Figure 6.1: Vertical flow from total crises to geographic concentration × the complete research narrative. Clean synthesis diagram showing the research arc: (1) AR baseline catches 73.2% (3,895 crises) via persistence where "yesterday predicts today" works (green, left) × the autocorrelation trap quantified; (2) AR fails on 26.8% (1,427 crises) rapid-onset shocks where persistence breaks (orange, right) × conflict escalations, economic collapses, regime transitions; (3) Cascade deploys news features strategically on AR failures, rescuing 249 crises (17.4% of failures, gold) × the breakthrough on hard cases; (4) Geographic concentration: 176 saves (70.7%) in Zimbabwe (red, 77 saves), Sudan (blue, 59 saves), DRC (purple, 40 saves) × high-coverage conflict zones where news signals provide value. Side annotations emphasize persistence works (left) vs rapid shocks (right). Arrow widths proportional to case counts. Standard colour coding: Zimbabwe=red, Sudan=blue, DRC=purple consistently across all dissertation figures. Bottom summary: 176 saves in top 3 conflict zones demonstrate selective deployment strategy × news features provide value where persistence fails, enabling humanitarian impact (249 crises predicted 8 months in advance). $n=5,322$ total crises, 20,722 observations, 18 countries, $h=8$ months.

exhibit strong temporal persistence (IPC 3 $\rightarrow$ IPC 3 common) and spatial clustering (neighboring districts correlate), enabling simple lag-based models to achieve excellent performance. News features, as engineered in this study, face a fundamental challenge: *temporal and spatial persistence capture 90%+ of predictive signal*, leaving limited room for news features to contribute unless deployed selectively. The cascade framework addresses this by focusing news-based analysis exclusively on AR failures (26.8% of crises), where news features rescue 249 cases (17.4% of AR failures)—demonstrating that news features provide value when targeted strategically.

Implications for the Field: Most existing literature reports AUC 0.75-0.85 for news-based crisis prediction without comparing against rigorous AR baselines. Our findings suggest these results may primarily reflect autocorrelation rather than text feature value. *Without AR baseline comparisons, high performance is potentially misleading.* This dissertation establishes that all future work claiming predictive value from text features and external covariates must include spatio-temporal AR baselines with inverse-distance spatial weighting, proper spatial CV, and reported *marginal* contributions.

The Critical 26.8%: The AR baseline misses 1,427 crises (26.8% of all 5,322 crises), representing the *high-frequency component* of crisis dynamics—shock-driven transitions where temporal patterns break and spatial neighbours provide insufficient signal. These 1,427 failures define where news features *might* provide genuine early-warning value, motivating the two-stage framework.

6.1.2 RQ2: When News Matters—Feature Engineering Insights

Research Question: What is the role of different kinds of news features (conflict, displacement, economic, weather) and dynamic transformations (ratio vs z-score) in predicting food insecurity beyond autoregressive baselines?

The Finding: On AR-difficult cases (6,553 observations, WITH_AR_FILTER strategy), ablation shows **ratio-only models achieve higher standalone AUC** (0.727 ± 0.165 vs 0.699 ± 0.165), but SHAP analysis reveals z-score features account for 74.7% of marginal attribution in combined models versus only 20.1% tree-based importance. This demonstrates complementary roles: ratio features provide stable cross-sectional baselines for standalone performance, while z-score features capture volatile temporal anomalies driving marginal predictions when combined. Both are essential—ratios for baseline discrimination, z-scores for shock detection.

Feature Importance Rankings (XGBoost Advanced, 35 features):

- **Location features dominate tree splits, not predictions:** 29.3% of tree-based importance (split frequency) but only 2.6% of SHAP attribution (marginal impact) $\times 15.5\times$ overstatement
 - country_data_density: 0.133 (13.3% tree splits, rank #1 tree-based, rank #17 SHAP)
 - country_baseline_conflict: 0.093 (9.3% tree splits, rank #2 tree-based, rank #20 SHAP)

- country_baseline_food_security: 0.067 (6.7% tree splits, rank #3 tree-based, rank #26 SHAP)
- **Z-score features drive predictions:** 74.7% of SHAP attribution despite only 20.1% tree-based importance
 - other_z-score: rank #1 SHAP (0.952 mean |SHAP|)
 - conflict_z-score: rank #2 SHAP (0.911)
 - humanitarian_z-score: rank #3 SHAP (0.902)
- **News categories** (aggregated ratio + z-score contributions):
 - Weather: 4.9% (droughts, floods, climate shocks)
 - Food security: 4.9% (direct crisis indicators)
 - Other: 5.5% (catch-all for uncategorised events)
 - Health: 4.9% (disease outbreaks, malnutrition)
 - Conflict: 4.7% (tree-based; note: conflict ranks #1 in SHAP z-scores at 0.911 for rapid anomaly detection)
 - Displacement: 3.8% (population movements)
 - Economic: 3.2% (market disruptions, inflation)
- **HMM features:** hmm_ratio_transition_risk (0.032, rank #5)
- **DMD features:** dmd_ratio_crisis_instability (large mixed-effects coefficient +352.38, but low XGBoost importance)

Mixed-Effects Evidence: Fixed effects from pooled logistic regression reveal that weather_ratio (+26.71), displacement_ratio (+21.18), food_security_ratio (+20.33), and conflict_ratio (+19.61) have the largest positive coefficients, confirming XGBoost rankings.

When News Matters: News features provide value through two distinct mechanisms: (1) *geographic stratification*×location metadata (data density, baseline conflict, baseline food security) efficiently segments countries into risk tiers (40.4% tree splits, enabling fast baseline stratification); (2) *dynamic shock detection*×z-score anomalies (conflict, humanitarian, displacement spikes) drive marginal predictions for individual crises (74.7% SHAP attribution). The critical finding is that tree-based importance (29.3% location) conflates these mechanisms×SHAP analysis reveals location contributes only 2.6% to marginal predictions despite high split frequency. **Practical implication:** For operational forecasting of shock-driven crises (the hardest cases), dynamic news signals (z-scores) matter more than geographic baselines. News-based forecasting works best in news-dense regions (Sudan, Zimbabwe, Kenya) where both mechanisms operate, but geographic metadata alone provides minimal marginal value×the real predictive power comes from detecting anomalies *within* those contexts.

Ratio vs Z-Score: Ratios capture compositional emphasis (“30% of articles mention conflict”), while z-scores capture anomalies (“ 3σ spike in conflict coverage”). The superior performance of ratio features on AR-difficult cases suggests that *sustained thematic emphasis* provides stronger signal than *short-term spikes*. Crises that AR models miss are characterised by persistent compositional shifts in news narratives (elevated conflict coverage sustained over months), not necessarily sudden spikes.

6.1.3 RQ3: The Role of Hidden Variables—HMM and DMD

Research Question: What is the contribution of latent regime detection (HMM) and temporal pattern extraction (DMD) in identifying crises that autoregressive models miss?

The Finding: HMM and DMD contribute unique signal for detecting hidden crisis dynamics:

- Adding HMM to ratio+z-score+location: +0.007 AUC (from 0.696 to 0.703, $p \approx 0.08$), with **hmm_ratio_transition_risk** ranking #5 in feature importance (0.032)
- Adding DMD to ratio+z-score+location: +0.002 AUC (from 0.696 to 0.698), with **dmd_ratio_crisis_instability** achieving largest mixed-effects coefficient (+352.38) among all features
- Combined HMM+DMD: Provide complementary scientific insights, with HMM detecting regime shifts and DMD identifying temporal evolution patterns

HMM Captures Regime Transitions: The **hmm_ratio_transition_risk** feature ranks #5 in importance (0.032, equivalent to 3.2%), capturing *latent regime transitions*—when news narratives shift from peaceful/stable regimes to conflict/crisis-prone regimes, even when article volumes remain constant. This qualitative change in discourse (peace $\rightarrow$ violence) provides unique signal for detecting when crisis narratives fundamentally change in character, demonstrating that regime detection identifies narrative shifts invisible to raw article counts and compositional features.

DMD Identifies Complex Emergency Patterns: DMD features extract temporal patterns (escalation modes with positive growth rates, sustained intensity modes with near-zero eigenvalues). The **dmd_ratio_crisis_instability** feature achieves the *largest mixed-effects coefficient (+352.38) among all features*, demonstrating that when DMD detects multi-category simultaneous spikes, it strongly signals complex emergencies. By design, DMD targets rare but extreme events (<3% of observations)—the most severe humanitarian catastrophes where multiple crisis drivers (conflict + displacement + economic collapse) converge simultaneously. DMD enables identification of *how crises evolve temporally*, distinguishing exponential escalation from sustained intensity patterns.

Contribution to Model Interpretation: HMM and DMD achieve 89.5% and 83.1% convergence rates respectively, successfully extracting latent dynamics from 48-month news sequences despite short time spans. Their contribution is **enhanced model interpretation**

and crisis driver identification: they reveal *why* crises emerge (regime transitions) and *how* they unfold (temporal evolution patterns) in ways that cross-sectional aggregations cannot. This explanatory power is critical for humanitarian decision-making, where understanding narrative shifts and crisis dynamics informs response strategies beyond binary predictions.

Conclusion for RQ3: HMM and DMD **advance crisis prediction through model interpretation and scientific insight.** HMM's #5 feature ranking (3.2% importance) demonstrates clear value for detecting regime transitions. DMD's largest mixed-effects coefficient (+352.38) signals critical detection of complex emergencies. Together, they provide unique signal for detecting qualitative narrative shifts and temporal evolution patterns that simpler features cannot capture, justifying inclusion in early warning systems where understanding crisis dynamics matters.

6.1.4 RQ4: Two-Stage Framework Performance and Precision-Recall Trade-Offs

Research Question: Can a two-stage residual modelling approach effectively rescue crises missed by autoregressive baselines, and what are the precision-recall trade-offs of such a framework?

The Finding: The two-stage cascade framework achieves **249 key saves**—crises where AR predicted no crisis (AR=0) but the cascade correctly predicted crisis (Cascade=1) when ground truth was crisis (y=1). This represents a **17.4% rescue rate** of the 1,427 AR failures, demonstrating *partial but meaningful success* in identifying AR-missed crises.

Performance Transformation:

Metric	AR Baseline	Cascade	Change
Precision	0.732	0.585	-0.147 (-14.7pp)
Recall	0.732	0.779	+0.047 (+4.7pp, +6.4%)
F1	0.732	0.668	-0.064 (-6.4pp)
TP (crises caught)	3,895	4,144	+249
FP (false alarms)	1,427	2,939	+1,512
FN (missed crises)	1,427	1,178	-249
TN (correct non-crisis)	13,973	12,461	-1,512

Table 6.1: AR Baseline vs Cascade Framework Performance

The Precision-Recall Trade-Off: Each of the 249 key saves costs **6.1 false alarms** (1,512 additional FP / 249 key saves = 6.1:1 trade-off ratio). While recall improves (+4.7pp), precision decreases (-14.7pp), and overall F1 decreases from 0.732 to 0.668. This trade-off reflects the cascade's deliberate optimisation for humanitarian contexts: *prioritising recall (catching crises) over precision (avoiding false alarms)*. When precision and recall are weighted equally (F1 metric), the AR baseline performs better. However, humanitarian early warning systems face asymmetric costs where missing crises carries far greater consequences than false alarms.

Cost-Sensitive Analysis: However, humanitarian contexts exhibit *asymmetric costs*—missing a crisis (false negative) is far more catastrophic than issuing a false alarm (false positive). Assuming

a 10:1 cost ratio (FN cost = $10 \times$ FP cost):

- AR baseline cost: $10 \times 1,427 + 1 \times 1,427 = 15,697$
- Cascade cost: $10 \times 1,178 + 1 \times 2,939 = 14,719$
- **Improvement: -978 cost units (-6.2% reduction)**

Under humanitarian cost assumptions, the cascade provides meaningful improvement despite lower F1. **Critically, the +4.7 percentage point recall gain is not merely a statistical improvement—it represents the 249 hardest-to-predict crises** where spatio-temporal persistence breaks down. These are *real crises affecting millions of people*, now predicted 8 months in advance, enabling preemptive food assistance, livelihood support, and conflict mitigation. The cascade is not optimising average performance across all cases—it is **rescuing the most critical cases where persistence fails and where timely intervention saves lives**. In humanitarian contexts, detecting conflict-driven shocks in Sudan, economic collapse in Zimbabwe, and complex emergencies in DRC—the cases AR baseline misses—matters far more than aggregate F1 scores.

Geographic Concentration: Key saves are **not uniformly distributed**:

- Zimbabwe: 77 key saves (30.9%)
- Sudan: 59 key saves (23.7%)
- DRC: 40 key saves (16.1%)
- Nigeria: 27 key saves (10.8%)
- **Top 3 countries (Zimbabwe, Sudan, DRC): 176 key saves = 70.7% of all key saves**

These three countries (representing 3 of 18 total countries in the CASCADE dataset) account for over 70% of the cascade’s added value, demonstrating strong geographic heterogeneity. Within-country heterogeneity analysis reveals the same countries show both cascade rescues and failures at district level. Zimbabwe has 77 key saves but 647 still-missed cases (11.9% rescue rate), Sudan has 59 saves but 420 still-missed (14.0%), Kenya has 8 saves but 722 still-missed (1.1%). This pattern indicates that news-based early warning succeeds in well-covered districts (capitals like Harare/Khartoum, conflict zones like Eastern DRC) but fails in news desert districts (remote pastoral areas like Kenya Northern/Turkana, peripheral regions) within the same country. Median news coverage: rescued cases 121 articles/month vs still-missed cases 79 articles/month (53% more coverage enables rescue).

What the Framework Rescues: The 249 key saves concentrate in *conflict-affected regions experiencing rapid-onset shocks*:

- **Zimbabwe (77 saves):** Economic collapse (hyperinflation, currency crises), structural food insecurity with rapid deteriorations
- **Sudan (59 saves):** Conflict escalations (Darfur, South Kordofan), displacement-driven crises

- **DRC (40 saves):** Complex emergencies (simultaneous conflict, displacement, disease outbreaks)
- **Nigeria (27 saves):** Boko Haram insurgency spillover (Borno State), sudden market disruptions

These are precisely the cases where 8-month advance warning enables life-saving humanitarian response—prepositioned food stocks, early deployment of nutrition programs, conflict-sensitive interventions.

Scope and Limitations: The cascade rescues 249 crises (17.4% of AR failures), while **1,178 AR failures remain unpredicted (82.6% of AR failures)**. This demonstrates that current bag-of-words text features capture a specific subset of crisis dynamics—those accompanied by news coverage signals—while other rapid-onset crises require advanced NLP techniques (transformer-based semantic understanding, multilingual models, social media text mining, automated event extraction). The partial success validates the hypothesis that *news provides genuine value for specific crisis types in specific contexts* (conflict-driven, high-coverage regions), while identifying substantial opportunities for NLP-driven enhancement to rescue more AR failures.

Conclusion for RQ4: The two-stage framework *can* rescue meaningful numbers of AR-missed crises (249 cases, 17.4% rescue rate), but at significant precision cost (-14.7pp). The trade-off is favourable in humanitarian contexts (10:1 FN:FP cost weighting yields -6.2% total cost reduction) but unfavourable for balanced metrics (F1 decreases). **Selective deployment is critical:** use the cascade in high-value regions (Sudan, Zimbabwe, DRC) where key saves concentrate, not universally.

6.1.5 RQ5: Geographic Heterogeneity—News Features Are Not Universally Valuable

Research Question: Are news-based features equally valuable across all geographic contexts, or do certain countries and crisis types benefit more from dynamic news signals than others?

The Finding: News-based features exhibit **strong geographic heterogeneity**—they are *not* equally valuable across contexts.

Evidence 1: Key Saves Concentration (already noted in RQ4):

- Zimbabwe, Sudan, DRC: 176 key saves = 70.7% of total
- 3 of 18 countries account for over 2/3 of cascade value
- Remaining 15 countries: 73 key saves (29.3%), averaging 4.9 saves per country

Evidence 2: Performance Variation Across Countries (XGBoost Advanced, country-level AUC):

- **Best performers:** Sudan (0.682), Uganda (0.679), Kenya (0.637)
- **Worst performers:** Niger (0.068), Ethiopia (0.417), Mozambique (0.515)
- **Range:** 0.068 to 0.682 = $10\times$ difference in AUC
- **Mean $\pm$ SD:** 0.54 ± 0.20 (massive variance)

Evidence 3: Mixed-Effects Random Effects (country baseline risk deviations):

- **Highest baseline risk:** Somalia (+3.70), Zimbabwe (+2.67), Sudan (+2.24)
- **Lowest baseline risk:** Madagascar (-4.56), Uganda (-3.86), DRC (-0.64)
- **Range:** 8.26 points (Somalia to Madagascar), indicating substantial country-specific heterogeneity

Random slopes for `conflict_ratio` and `food_security_ratio` vary significantly by country, demonstrating that *some countries are more sensitive to conflict news* (Sudan, Nigeria), while *others are more sensitive to food security news* (Zimbabwe, Malawi).

Evidence 4: Country-Specific News Theme Signatures (SHAP-based theme analysis, $n=23,039$ observations across 13 countries):

Analysis of observation-level SHAP values aggregated by theme category (combining both ratio and z-score features) reveals distinct country-specific signatures that further confirm heterogeneity in *which* news themes drive predictions in each context:

- **Zimbabwe** (77 saves): Humanitarian (13.4%), Other (13.0%), Weather (11.5%). Weather ranks 3rd locally vs 8th globally (9.4%), +2.1pp elevation, aligning with recurring drought cycles (2019 Cyclone Idai, 2022-2023 drought) compounding economic collapse.
- **Sudan** (59 saves): Governance (14.8%), Conflict (14.6%), Humanitarian (13.4%). Conflict ranks 2nd locally vs 4th globally (11.3%), +3.3pp elevation, reflecting April 2023 civil war escalation that AR baseline could not anticipate.
- **DRC** (40 saves): Other (14.3%), Humanitarian (12.9%), Displacement (12.2%). Displacement ranks 3rd locally vs 7th globally (10.0%), +2.2pp elevation, capturing M23 resurgence and North Kivu population movements.

These elevations identify *diagnostic signals*—themes that deviate maximally from global patterns, revealing context-specific shock types that AR baselines miss. Unlike dominant theme analysis (what’s biggest in absolute terms), elevation analysis (what’s unusual relative to global average) aligns with the cascade’s residual modelling objective: detecting anomalies that break structural persistence. Somalia exhibits the highest observed elevation for any theme (Health +5.8pp, 16.5% vs 10.7% global), demonstrating how disease burden compounds food insecurity in ways invisible to temporal/spatial autocorrelation.

Global theme distribution: Governance (13.0%), Other (13.0%), Humanitarian (12.6%), Conflict (11.3%)—relatively flat (9.2-13.0%, 3.8pp range, $1.4\times$ max/min), indicating no universal

dominant theme. This flatness is itself meaningful: theme importance varies by country-specific crisis dynamics, not global averages. Countries with elevated theme importance (+2-3pp above global) for specific categories (Zimbabwe Weather, Sudan Conflict, DRC Displacement) demonstrate context-dependent news utilisation—models learn different thematic patterns in different crisis types.

Why Heterogeneity Exists:

1. **News Coverage Density:** `country_data_density` ranks #1 in tree-based split frequency (13.3%), serving as stratification infrastructure for geographic context. High-coverage countries (Sudan, Kenya, Zimbabwe: >1,000 articles/district-month) enable better predictions; low-coverage countries (Niger, Uganda, Madagascar: <100 articles/district-month) lack sufficient signal. SHAP analysis reveals location features contribute 2.6% marginal attribution (enabling context-specific learning) while z-score news features drive 74.7% of predictions.
2. **Crisis Type:**
 - **Conflict-driven crises** (Sudan, Nigeria, DRC): Conflict and displacement news features gain importance; rapid escalations generate text signals
 - **Climate-driven crises** (Kenya, Ethiopia pastoral zones): Weather news features gain importance; seasonal droughts generate coverage
 - **Economic/structural crises** (Zimbabwe): Economic news features gain importance; but slow-burn structural transitions generate weaker signals
3. **Baseline Conflict and Instability:** `country_baseline_conflict` ranks #2 (9.3% importance). Chronically conflict-affected countries exhibit more predictable crisis patterns (conflict → displacement → food insecurity pathways well-documented in news). Peaceful countries with sudden shocks lack established news coverage patterns.
4. **Sample Size:** Countries with few crisis observations (Uganda n=2, Madagascar n=8) produce unstable metrics due to small sample variance. High-crisis countries (Sudan n=87, Zimbabwe n=102) provide sufficient training data.

Implications for Deployment: Universal models fail. Country-level AUC ranges from 0.068 to 0.682 (10× difference), demonstrating that a single model cannot serve all contexts. Recommendations:

1. **Selective deployment:** Use news features in Sudan, Zimbabwe, DRC, Kenya (high coverage, high key saves) but *not* in Niger, Uganda, Madagascar (low coverage, low key saves).
2. **Country-specific models:** Mixed-effects approach with random effects partially addresses heterogeneity, but fully country-specific models may be needed for high accuracy in priority countries.

3. **Crisis-type and theme-aware stratification:** Deploy conflict-focused models in Sudan/Nigeria/DRC with prioritised monitoring of Conflict (Sudan +3.3pp) and Displacement (DRC +2.2pp) theme feeds; climate-focused models in Kenya/Ethiopia pastoral zones with prioritised Weather monitoring; economic/humanitarian-focused models in Zimbabwe with prioritised Weather (+2.1pp) and Humanitarian monitoring. Theme-specific surveillance reduces information overload while maintaining sensitivity to country-specific shock types revealed through SHAP analysis.
4. **Resource allocation:** Concentrate computational resources on high-value countries (Sudan, Zimbabwe, DRC) where news features demonstrably improve early warning (176 saves, 70.7% of total), rather than universal deployment where marginal value is minimal or negative.

Conclusion for RQ5: News-based features are **not universally valuable**. Strong heterogeneity observed across three dimensions: (1) Geographic concentration—70.7% of key saves in 3 of 18 countries; (2) Performance variation—country-level AUC ranges $10\times$ (0.068-0.682); (3) Theme heterogeneity—country-specific SHAP signatures reveal elevated importance for context-specific themes (Zimbabwe Weather +2.1pp, Sudan Conflict +3.3pp, DRC Displacement +2.2pp vs global averages). News coverage density determines predictability, but *which themes* matter varies by country-specific crisis dynamics. *Selective, theme-aware deployment based on geographic context, crisis type, and news availability is necessary*. Universal models with uniform theme weighting will fail in low-coverage contexts and miss country-specific signals.

6.2 Core Contributions to Humanitarian Early Warning

This dissertation makes five core contributions:

6.2.1 Contribution 1: Methodological Critique—Exposing the Autocorrelation Trap

We provide the first systematic empirical demonstration that spatio-temporal AR baselines achieve 93.8% of published news model performance (AR PR-AUC=0.7652 vs Balashankar et al. 2023 PR-AUC=0.8158) using **zero text features**. This establishes the autocorrelation trap as a *quantitatively large, empirically real phenomenon* that existing literature has systematically neglected.

The Critique Has Three Components that collectively establish the autocorrelation trap as a major methodological oversight requiring immediate attention in existing crisis prediction literature:

1. **Empirical demonstration:** AR baseline performance approaches published news models (93.8% of Balashankar et al.'s PR-AUC), demonstrating that temporal and spatial persistence dominates crisis prediction performance.

2. **Theoretical implication:** Without AR comparisons, high performance may reflect autocorrelation rather than text value. Claims that “news predicts crises” are technically true but potentially incomplete—persistence predicts most crises, and news features contribute incrementally. The cascade framework demonstrates that news features provide value when deployed selectively on AR failures (249 key saves, 17.4% rescue rate), rather than universally.
3. **Methodological prescription:** All future crisis prediction work must include rigorous AR baselines with both temporal autoregressive features and spatial autoregressive features, inverse-distance spatial weighting, proper spatial CV, and reported *marginal* contributions. This sets a higher standard for the field.

To our knowledge, this is the **first systematic comparison** of news-based models against strong spatio-temporal baselines in the food security domain. Our work challenges existing paradigms and provides a template for future methodological rigor.

6.2.2 Contribution 2: Two-Stage Residual Modelling Framework

We develop a principled approach that explicitly separates *structural persistence* (captured by AR baseline) from *shock-driven dynamics* (captured by news features):

Stage 1—AR Baseline: Deploys spatio-temporal logistic regression on all 20,722 observations. Achieves 73.2% precision/recall/F1. Identifies 1,427 false negatives (AR failures) as candidates for Stage 2 rescue.

Stage 2—Dynamic Features: Deploys XGBoost with 35 advanced features (ratio, z-score, HMM, DMD, location) *exclusively* on WITH_AR_FILTER subset (6,553 cases where $IPC_{t-1} \leq 2$ AND $AR=0$). Achieves 249 successful predictions of AR-missed crises (17.4% rescue rate).

These 249 Cases Are Not Statistical Abstractions: They represent *the most operationally critical early warnings*—conflict escalations in Sudan where displacement unfolds rapidly, coup-related disruptions in Zimbabwe where temporal patterns break abruptly, acute emergencies in DRC where persistence models fail. These are precisely the cases where 8-month advance warning enables life-saving humanitarian response.

Integration: Simple cascade decision logic preserves all $AR=1$ predictions (trusting the baseline when it predicts crisis) and uses Stage 2’s binary prediction for $AR=0$ cases. Combined framework achieves:

- **249 key saves**—the hardest cases where news signals matter most
- **Recall: 0.732 $\rightarrow$ 0.779 (+6.4% relative improvement)**—not merely a percentage gain, but 249 real crises affecting millions, now predicted 8 months early
- Precision: 0.732 $\rightarrow$ 0.585 (reduced due to prioritising recall in humanitarian contexts)
- F1: 0.732 $\rightarrow$ 0.668 (decreases, but humanitarian cost-sensitive analysis favours recall)

- Geographic concentration: 70.7% of key saves in Sudan, Zimbabwe, DRC

Three Methodological Innovations:

1. **Selective deployment:** Complex features deployed only where AR fails (not universally), maximising value per cost. Targets WITH_AR_FILTER subset (6,553 cases where $IPC_{t-1} \leq 2$ AND $AR=0$) rather than all 20,722 observations.
2. **Explicit persistence modelling:** AR baseline captures structural persistence explicitly (not as implicit control variables), enabling interpretable decomposition of which predictions succeed due to autocorrelation (73.2%) versus which require dynamic signals (the critical 17.4% of failures rescued).
3. **Humanitarian-appropriate metrics:** Prioritises recall over precision, aligning with operational early warning principles where missing crises is catastrophic while false alarms are manageable. Achieves 77.9% recall, successfully predicting 4,144 of 5,322 total crises.

The framework demonstrates *meaningful but partial success*: 17.4% rescue rate validates that news signals provide genuine early-warning value for specific crisis types (conflict-driven, rapid-onset) in specific contexts (Sudan, Zimbabwe, DRC). While 82.6% of AR failures remain unrescued, this partial success is *operationally valuable*—249 crises caught 8 months early represent families, communities, and lives where early warning enables early response.

6.2.3 Contribution 3: Dynamic Feature Engineering Beyond Article Counts

We demonstrate a four-stage analytical pipeline extending beyond static article counts:

Stage 2a—Ratio and Z-Score Transformations: Ratios capture compositional emphasis (“30% of articles mention conflict”); z-scores capture anomalies (12-month sliding-window normalisation). Ablation studies reveal that **ratio and z-score features provide complementary signals**: as standalone features, ratios (AUC 0.727) capture compositional emphasis more effectively than z-scores (AUC 0.699) capture temporal anomalies. However, when combined in full models, individual z-score features (conflict_z-score 4.2%, food_security_z-score 3.7%) provide valuable orthogonal signals for detecting sudden-onset crises. Both feature types contribute unique perspectives: ratios measure topic dominance, z-scores measure coverage spikes.

Stage 2b—Hidden Markov Models: 1,322 district-pooled 2-state models extract latent narrative regimes. The `hmm_ratio_transition_risk` feature ranks #5 in importance (0.032), demonstrating that regime transitions provide genuine signal. HMM achieves +0.007 AUC gain with substantial scientific value for crisis driver identification—revealing when narratives shift from peaceful to violent regimes even when article volumes remain constant.

Stage 2c—Dynamic Mode Decomposition: Crisis-focused mode filtering extracts temporal patterns (escalation modes, sustained intensity modes). DMD contributes unique signal for extreme events: **dmd_ratio_crisis_instability achieves the largest mixed-effects**

coefficient among all features (+352.38), demonstrating value for detecting rare but catastrophic complex emergencies where multiple crisis drivers converge simultaneously. By design, DMD targets <3% of observations (severe multi-category escalations), providing critical signal for the most severe humanitarian crises.

Stage 2d—Mixed-Effects Regression: Pooled logistic regression with country random effects and random slopes quantifies geographic heterogeneity. Fixed effects reveal global patterns (weather_ratio +26.71, displacement_ratio +21.18); random effects reveal country-specific sensitivities (Somalia +3.70, Madagascar -4.56). Enables targeted deployment recommendations.

Key Insight: Discrimination-interpretation trade-off. Ratio+Location (12 features, AUC 0.727) achieves highest classification performance. The Advanced model (35 features, AUC 0.697) integrates all feature engineering approaches for comprehensive crisis understanding: hmm_ratio_transition_risk ranks #5 (3.2% importance) capturing qualitative regime transitions, DMD achieves largest coefficient (+352.38) for extreme events, z-scores complement ratios. For operational deployment, both approaches contribute: parsimonious models for discrimination, comprehensive models for crisis driver identification.

6.2.4 Contribution 4: Comprehensive Model Interpretation Framework

We deploy three complementary model interpretation methods to triangulate which features matter, when, and where:

Method 1—XGBoost Feature Importance (tree-based, non-linear):

- Measures feature contribution to splits across 300+ trees
- Reveals that location features dominate split frequency (29.3%, 40.4% total) but contribute minimally to SHAP attribution (2.6%×15.5× overstatement), while news categories (especially z-scores) drive marginal predictions (74.7% SHAP)
- Captures interaction effects (e.g., country_data_density × conflict_ratio)

Method 2—Mixed-Effects Coefficients (linear, additive):

- Fixed effects quantify global patterns (weather_ratio +26.71 largest news coefficient)
- Random effects quantify country-specific deviations (Somalia +3.70 highest baseline risk)
- Random slopes quantify feature sensitivity heterogeneity (conflict_ratio varies by country)
- Interpretable as log-odds contributions

Method 3—SHAP Values (game-theoretic, local explanations):

- Shapley value attribution for individual predictions
- Enables case-by-case explanations (“Zimbabwe 2021 crisis predicted due to economic_ratio spike + hmm_transition_risk”)

- Additive feature contributions enable humanitarian decision-makers to understand *why* a crisis was predicted
- **Critical revelation:** SHAP fundamentally reorders feature rankings compared to tree-based importance×z-scores account for 74.7% of marginal attribution despite only 20.1% of tree splits, while location features account for 2.6% of attribution despite 40.4% of tree splits (15.5× overstatement)

Triangulated Findings (partial agreement, critical divergences):

- **Tree-based importance:** Location features dominate (40.4%), z-scores secondary (20.1%)×measures *split frequency* (stratification utility)
- **SHAP attribution:** Z-scores dominate (74.7%), location minimal (2.6%)×measures *marginal impact* (predictive contribution)
- **Mixed-effects coefficients:** Weather ratio (+26.71) largest news coefficient, DMD instability (+352.38) largest dynamic coefficient×measures *linear effects* (interpretable log-odds)
- Category rankings measurement-dependent: Weather/food security rank highest in ratio/mixed-effects (sustained shifts); conflict/humanitarian rank highest in SHAP z-scores (rapid anomalies)
- HMM ranks #5 in tree-based (3.2%), ranks #7-8 in SHAP (hmm_ratio_crisis_prob, hmm_ratio_transition_risk), 21.9% total SHAP attribution
- Geographic heterogeneity substantial (country random effects span 8.26 points, mixed-effects)

Divergences Reveal Methodological Insights:

- **Split frequency $\neq$ predictive contribution:** Location features split frequently (stratification) but contribute little to marginal predictions (SHAP 2.6%). Z-scores split infrequently but drive prediction variance (SHAP 74.7%). This demonstrates that feature “importance” depends critically on measurement method.
- **DMD features:** Large mixed-effects coefficient (+352.38) but low tree-based importance (1.5%) and low SHAP (1.5%) → captures rare but extreme events. Linear models (mixed-effects) weight rare extremes; non-linear models (XGBoost, SHAP) average them out.
- **HMM features:** Higher SHAP attribution (21.9%) than tree-based importance (13.0%), confirming genuine predictive value beyond stratification×regime transitions drive marginal predictions.
- All three perspectives needed for full understanding: Tree-based identifies stratification features, SHAP identifies prediction drivers, mixed-effects identifies rare extremes and geographic heterogeneity.

Practical Value: Model interpretation enables:

1. **Operational trust:** Humanitarian decision-makers can understand *why* a crisis was predicted (not just black-box probabilities)
 2. **Strategic deployment:** Knowing that news features work in Sudan (high data density, conflict-driven) enables targeted resource allocation where media ecosystems support predictive value
 3. **Feature complementarity:** Knowing that z-score features account for 74.7% of SHAP marginal attribution (driving shock detection) while ratio features enable higher standalone AUC (providing stable baselines) demonstrates both are essential for operational systems
-

6.2.5 Contribution 5: Operational Deployment Framework and Geographic Targeting

We provide actionable recommendations for when and where to deploy news-based early warning:

Where News-Based Features Add Value:

- **High-coverage countries:** Sudan, Kenya, Zimbabwe (>1,000 articles/district-month)
- **Conflict-affected regions:** Sudan (Darfur, South Kordofan), Nigeria (Borno State), DRC (Ituri, North Kivu)
- **Rapid-onset crisis contexts:** Coup-related disruptions, conflict escalations, sudden displacement events
- **Countries where cascade succeeds:** Zimbabwe (77 saves), Sudan (59), DRC (40)—70.7% of all key saves

Where Simple AR Baselines Suffice:

- **Low-coverage countries:** Niger, Uganda, Madagascar (<100 articles/district-month)
- **Structurally persistent crises:** Slow-burn chronic food insecurity (Somalia coastal districts, Madagascar southern districts)
- **Countries where cascade adds little value:** Remaining 15 countries average 4.9 key saves each

Decision Logic for Humanitarian Agencies:

1. **First-line EWS: AR Baseline** (all contexts)
 - Deploy spatio-temporal logistic regression universally
 - Low computational cost, 90.7% AUC, 73.2% precision/recall
 - Trust AR predictions for 73.2% of crises (predictable persistence cases)
2. **Second-line EWS: News-Based Cascade** (selective deployment)

- Automatically triggers for all $AR=0$ cases in `WITH_AR_FILTER` subset ($IPC_{t-1} \leq 2$ AND $AR=0$)
- Strategically deployed in high-benefit regions (Zimbabwe, Sudan, DRC) with historical rescue rates $>15\%$
- Detects shock-driven crises (conflict escalations, economic collapses, displacement events) where AR fails

3. Resource Allocation Tiers:

- **Tier 1 (highest resources):** AR-predicted crises (3,895 cases, 73.2% of all crises)—high confidence, preposition food stocks, deploy early
- **Tier 2 (secondary resources):** Cascade overrides (1,761 $AR=0$ but $Cascade=1$ cases)—lower confidence, prepare contingency plans, monitor closely
- **Tier 3 (tertiary monitoring):** Low-probability cases (both $AR=0$ and $Cascade=0$)—passive monitoring, no preemptive deployment

Integration with Existing Humanitarian Systems:

- **FEWSNET:** Combine model predictions with expert analyst judgment; use predictions to prioritise field assessment locations
- **WFP HungerMap:** Integrate cascade predictions as additional news-based early warning layer
- **IPC Technical Working Groups:** Use model outputs to flag districts for expedited IPC assessments 8 months ahead

Computational Cost-Benefit:

- AR baseline: 30 minute training time (logistic regression, 2 features)
- News feature engineering: 2 hours per district (GDELT processing, HMM/DMD convergence)
- XGBoost training: 1 hour (35 features, 3,888 hyperparameter search)
- **Recommendation:** Deploy news features only in high-value countries (Sudan, Zimbabwe, DRC) where 249 key saves justify computational investment

6.3 Implications for the Humanitarian Early Warning Ecosystem

6.3.1 Rethinking the Role of News in Crisis Prediction

This dissertation challenges the prevailing assumption that “more data is always better.” News features provide genuine early-warning value for *specific crisis types in specific contexts*, but they

are **not universally valuable**. The field must move beyond claims that “news predicts crises” toward *nuanced, context-dependent deployment*:

- **For conflict-driven crises in high-coverage regions** (Sudan, DRC): News features add substantial value (59-40 key saves)
- **For structural economic crises** (Zimbabwe): News features add moderate value (77 key saves, but false alarms high)
- **For climate-driven crises in low-coverage regions** (Niger, pastoral Ethiopia): Advanced NLP enhancements (multilingual models, social media text mining, event extraction) may capture crisis signals missed by English-only bag-of-words features
- **For chronic structural persistence** (Somalia coastal, Madagascar southern districts): AR baselines suffice

The autocorrelation trap demonstrates that **methodological rigor matters more than data volume**. A simple 2-feature AR baseline achieves AUC=0.907 on the full dataset, while 35-feature Stage 2 XGBoost models trained on AR-filtered cases achieve AUC=0.697-0.727, reflecting the complementary nature of their tasks (persistence vs shock detection) rather than direct competition. This finding has profound implications:

1. **Prioritise marginal value over absolute performance**: Report what news features add *beyond* AR baselines, not just raw AUC
2. **Context-dependent deployment**: Deploy news-based models where they demonstrate value (Sudan, Zimbabwe, DRC), rely on AR baselines where they suffice (Niger, Uganda, Madagascar)
3. **Interpretability over complexity**: Simple models with clear explanations (mixed-effects, SHAP) enable operational trust; black-box models risk rejection by humanitarian decision-makers

6.3.2 The Two-Component Crisis Dynamics Framework

Our findings suggest a **two-component decomposition of crisis dynamics**:

Low-Frequency Component (73.2% of crises):

- Structural persistence: IPC 3 $\rightarrow$ IPC 3, IPC 4 $\rightarrow$ IPC 4
- Spatial clustering: Neighbouring districts highly correlated
- Slow-burn deteriorations: Gradual multi-year declines (chronic poverty, structural food insecurity)
- **Captured by AR baseline** (90.7% AUC)
- **Persistence provides primary signal** (AR baseline achieves 90.7% AUC)

High-Frequency Component (26.8% of crises = 1,427 AR failures):

- Shock-driven transitions: IPC 1/2 $\rightarrow$ IPC 4/5 (rapid onset)
- Spatially isolated: Weak neighbour correlation (remote districts, sudden events)
- Conflict escalations: Coup-related disruptions, insurgency spillover, displacement shocks
- **AR baseline fails** (these are the 1,427 false negatives)
- **News features provide partial rescue** (249 key saves = 17.4% of high-frequency component)

This decomposition has **theoretical implications**:

1. **Persistence dominates**: 73.2% of crises are predictable from temporal/spatial autoregressive features alone
2. **Shocks require dynamic text signals**: 26.8% of crises require advanced features beyond simple persistence (news semantic understanding, multilingual coverage, event extraction, social media signals)
3. **No single model captures both**: AR baselines excel at persistence, news models provide marginal value for shocks
4. **Two-stage approaches necessary**: Separate models for separate dynamics

Comparison to Time Series Literature: This decomposition parallels trend-cycle decomposition in econometrics (Hodrick-Prescott filter, wavelet decomposition). However, in crisis prediction, the *high-frequency component is humanitarian priority*—these are the unpredictable shocks where early warning matters most. The low-frequency component (persistence) is already well-managed by existing systems; the high-frequency component (shocks) is where ML/NLP can contribute.

6.3.3 When to Trust AR, When to Override with Cascade

The binary cascade operates automatically using simple override logic, but understanding when each component provides value guides strategic deployment:

AR Baseline Handles (Automatically):

1. AR=1 (binary crisis prediction) $\rightarrow$ **Cascade preserves all AR=1 predictions**. These 5,322 cases represent structurally persistent crises where temporal/spatial patterns provide strong signal (73.2% precision). Deploy humanitarian resources immediately.
2. Persistence-dominated contexts (Kenya pastoral zones, Ethiopia, Malawi) $\rightarrow$ Climate-driven crises follow predictable seasonal patterns. Spatial autocorrelation captures regional drought patterns effectively. News features provide less additional value in these contexts.

3. Chronic structural crises (Somalia coastal, Madagascar southern) → Persistence dominates, AR captures recurring patterns.

Cascade Override Applied (Automatically for AR=0 Cases):

1. AR=0 (binary no crisis prediction) → **Stage 2 runs automatically on WITH_AR_FILTER subset** ($IPC_{t-1} \leq 2$ AND $AR=0$). If Stage 2 predicts crisis (=1), cascade overrides AR to detect shock-driven crises AR missed.
2. High-benefit countries (Zimbabwe, Sudan, DRC) → Deploy Stage 2 for all AR=0 cases. Historical rescue rates (30.9%, 23.7%, 16.1%) justify full deployment where news-dense conflict zones enable shock detection.
3. Conflict-affected, news-dense regions (Sudan Darfur, Nigeria Borno, DRC Ituri, Zimbabwe urban) → Rich media coverage enables Stage 2 to detect rapid-onset shocks (conflict escalations, economic collapses, regime transitions) where temporal persistence breaks down.
4. Low-benefit countries (Niger, Madagascar) → Skip Stage 2 entirely, use AR only. Rescue rate <3% insufficient to justify computational cost due to news coverage deficiency.

Binary Resource Allocation Logic:

- **Red Alert (AR=1 OR Cascade=1):** Either system detects crisis → Deploy humanitarian resources immediately (food aid, livelihood support, emergency funding mobilization). Total: 7,083 alerts (5,322 AR + 1,761 cascade overrides, including 249 key saves).
- **Green Status (AR=0 AND Cascade=0):** Both systems agree no crisis → Routine monitoring, no immediate action required. Total: 13,639 cases.

This simple two-tier system (crisis/no-crisis) maximises operational clarity. The trade-off is precision decline (0.732 → 0.585) for recall improvement (0.732 → 0.779), which humanitarian cost-benefit analysis (10:1 FN:FP weighting) justifies.

6.3.4 Limitations and Honest Reflection

We acknowledge five key limitations:

1. **The News Deserts Constraint (17.4% Rescue Rate Limited by Coverage Deficiency):**
 - The cascade rescues 249 crises (17.4% of AR failures), while **82.6% of AR failures remain unpredicted** (1,178 out of 1,427)
 - **Critical finding:** The 1,178 still-missed cases exhibit systematic news coverage deficiency—median 74 articles/month compared to 121 for rescued cases (64% less coverage, $p < 0.001$)

- This *news deserts hypothesis* reveals a fundamental constraint: *you cannot predict what is not reported*. Unlike satellite imagery (uniform geographic coverage) or household surveys (targeted collection), news media is inherently uneven—concentrated in conflict zones and politically important areas while neglecting remote pastoral regions (Kenya Northern, Zimbabwe rural districts, Niger)
- The limitation is not primarily a modelling failure (better algorithms) but a **data availability constraint** (insufficient text exists to extract signal from)
- **NLP-focused solutions required:** Addressing news deserts requires expanding text corpora beyond traditional English-language news through:
 - Social media monitoring (Twitter/X, Facebook community pages, WhatsApp group analysis)
 - Community radio transcripts (local-language broadcasts in Swahili, Hausa, Amharic, Somali, French, Arabic)
 - Humanitarian situation reports (OCHA, UNHCR, WFP crisis documentation)
 - Multilingual news sources (French-language news for Francophone Africa: Niger, Mali, DRC; Arabic sources for Sudan/Somalia; Portuguese for Mozambique/Angola)
 - Targeted collection partnerships with local journalists and NGO field reports for underreported regions
- Advanced NLP techniques (transformer-based semantic understanding, multilingual models, event extraction) can improve signal extraction, but only where sufficient text exists. For news deserts, expanding data sources is prerequisite to algorithmic enhancement.

2. Precision-Recall Trade-Off Severity:

- Precision drops 14.7pp (0.732 $\rightarrow$ 0.585), F1 decreases 6.4pp (0.732 $\rightarrow$ 0.668)
- 6.1:1 false alarm ratio (6.1 FP per key save) may cause operational alert fatigue
- Humanitarian agencies operating under resource constraints may reject models with 41.5% false alarm rate (2,939 FP / 7,083 total positive predictions)
- Cost-sensitive analysis (10:1 FN:FP) favours cascade, but this weighting is context-dependent

3. English-Language News Bias:

- GDELT is English-biased; local-language news (Swahili, Hausa, Amharic, French) excluded
- Underrepresents Francophone Africa (Niger, Mali, DRC rural areas)
- Low performance in Niger (AUC 0.068) may reflect news coverage gap rather than model failure
- Multilingual news processing needed for equitable coverage

4. Geographic Heterogeneity Creates Inequity:

- Models perform best in high-coverage countries (Sudan, Zimbabwe, Kenya)
- Models fail in low-coverage countries (Niger, Uganda, Madagascar)
- Risk of “**data colonialism**”: well-covered conflicts (Sudan Darfur, Nigeria Boko Haram) receive better early warning than under-covered crises (Madagascar southern droughts, Uganda Karamoja)
- Equitable deployment requires bridging coverage gaps, not just deploying where data exists

5. External Validity (Africa-Specific):

- Results are Africa-specific (18 countries, IPC Phase 3+ threshold)
- Generalisability to other regions uncertain (South Asia, Central America, Middle East)
- Different news ecosystems (state-controlled media in authoritarian contexts), different crisis types (urban food insecurity, migration-driven crises), different IPC thresholds may require retraining

Honest Reflection: This dissertation demonstrates that news features provide *partial, geographically heterogeneous, context-dependent value*. Claims should be tempered: news is **not a silver bullet** for early warning. It works in specific contexts (conflict-driven, high-coverage regions) for specific crisis types (rapid-onset shocks), but universal deployment is unjustified. The field must embrace *nuanced, selective deployment* rather than one-size-fits-all solutions.

6.4 Future Research Directions

6.4.1 Advanced NLP Enhancement: Beyond Bag-of-Words

Current bag-of-words news features rescue 17.4% of AR failures. Advanced NLP techniques offer substantial enhancement opportunities to improve rescue rates:

- **Transformer-based semantic understanding:** Fine-tune BERT/RoBERTa on crisis-specific corpora (FEWSNET reports, IPC assessments) to capture nuanced crisis narratives → rescue narrative-driven crises where word counts miss subtle signals
- **Multilingual NLP:** Deploy mBERT/XLM-RoBERTa on French (Sahel, DRC, Madagascar), Arabic (Sudan, Somalia), Swahili (Kenya, Tanzania) news → address English-language bias, improve coverage in Francophone contexts (Niger, Mali)
- **Social media text mining:** Fine-tune DistilBERT on disaster-specific Twitter datasets (CrisisNLP, HumAID), analyse humanitarian organisation Facebook pages → capture rapid-onset crises (conflict escalations, market disruptions) faster than traditional news
- **Automated event extraction:** Deploy transformer-based NER (SpaCy, Stanza) and relation extraction to identify structured crisis events (WHO attacked WHOM in WHERE,

WHAT shortage in WHICH district) → provide more precise crisis signals than aggregate article counts

- **Cross-lingual transfer learning:** Leverage high-resource English crisis models via zero-shot transfer to low-resource languages → extend coverage to under-served linguistic regions

Research Question: Can advanced NLP techniques (transformers, multilingual models, event extraction, social media mining) rescue more of the remaining 82.6% of AR failures? Which NLP approaches provide highest marginal rescue rates for different crisis types (conflict-driven vs climate-driven) and coverage contexts (high-coverage vs low-coverage)?

6.4.2 Multi-Horizon Optimisation

This dissertation focused on $h=8$ (32-week horizon). Different horizons may require different features:

- **$h=4$ (16 weeks):** Short-term predictions → social media text mining may dominate (faster-changing signals, hourly updates)
- **$h=8$ (32 weeks):** Medium-term predictions → news features optimal (current study)
- **$h=12$ (48 weeks):** Long-term predictions → structural indicators (conflict baselines, climate trends) may dominate

Research Question: How do optimal feature sets vary by prediction horizon? Can ensemble models combining $h=4, 8, 12$ predictions improve overall performance?

6.4.3 Real-Time Operational Deployment and Monitoring

This dissertation used historical data (2017-2024). Real-time deployment introduces challenges:

- **Data latency:** GDELT has 24-48 hour lag; IPC assessments have 2-6 month lag
- **Concept drift:** News coverage patterns change (COVID-19 pandemic shifted all news priorities 2020-2021)
- **Model retraining frequency:** How often to retrain? Monthly? Quarterly?
- **Human-in-the-loop:** How to integrate expert analyst judgment with model predictions?

Research Question: Can the cascade framework perform in real-time operational settings with data latency and concept drift? What monitoring systems are needed to detect when models degrade?

6.4.4 Causal Inference and Counterfactual Analysis

Current models are purely predictive (correlation-based). Causal understanding would enable:

- **Intervention planning:** “If we deploy food assistance in Zimbabwe 8 months early, how many crises can we prevent?”
- **Counterfactual reasoning:** “Would the Sudan 2023 crisis have occurred if the coup hadn’t happened?”
- **Policy evaluation:** “Did early warning systems reduce crisis severity in Kenya 2022?”

Methods: Instrumental variables (rainfall as instrument for crop failure $\rightarrow$ IPC), difference-in-differences (compare districts receiving early intervention vs not), causal forests, do-calculus.

Research Question: What are the *causal pathways* from news coverage spikes $\rightarrow$ IPC deterioration? Can we estimate causal effects of early warning interventions?

6.4.5 Multilingual News Processing

GDELT’s English bias excludes local-language news. Expanding to multilingual NLP would:

- **Improve coverage:** Francophone Africa (Niger, Mali, DRC), Swahili East Africa (Kenya, Tanzania), Amharic Ethiopia
- **Reduce inequity:** Current models favour Anglophone contexts
- **Capture local narratives:** National news (BBC, Reuters) differs from local radio (community stations covering hyperlocal events)

Challenges: Multilingual BERT (mBERT), translation quality, computational cost scaling with languages.

Research Question: Does adding local-language news improve performance in low-coverage countries (Niger, Uganda, Madagascar)? How much does translation quality matter?

6.4.6 Explainable AI for Humanitarian Decision-Making

This dissertation used SHAP values for interpretability. Operational deployment requires:

- **Natural language explanations:** “Crisis predicted in Sudan Darfur due to 300% spike in conflict news coverage and HMM regime transition from peaceful to violent narratives”
- **Counterfactual explanations:** “If conflict coverage had remained at baseline levels, predicted probability would be 0.3 instead of 0.8”
- **Uncertainty quantification:** “Prediction confidence: 75% ($\pm 10\%$)—moderate uncertainty due to sparse historical data in this district”

Research Question: How can XAI methods (SHAP, LIME, counterfactual explanations) be translated into actionable humanitarian decision support? What explanation formats do humanitarian analysts find most trustworthy?

6.5 Closing Vision: From Autocorrelation to Action

Food insecurity affects 282 million people globally (WFP 2024), making early warning systems a critical humanitarian tool. This dissertation demonstrates that:

1. **Spatio-temporal persistence dominates** crisis prediction (73.2% of crises predictable from AR baseline alone)
2. **News features provide genuine but partial value** for shock-driven crises (17.4% rescue rate = 249 key saves)
3. **Geographic heterogeneity demands selective deployment** (70.7% of key saves in Sudan, Zimbabwe, DRC)
4. **Methodological rigor requires AR baseline comparisons** to separate signal from autocorrelation
5. **Humanitarian context prioritises recall over precision** (cost-sensitive analysis favours 249 key saves despite 1,512 additional false alarms)

The autocorrelation trap—achieving high performance by simply predicting that tomorrow will look like today—has obscured when and where complex features actually help. By explicitly modelling AR failures and targeting these difficult cases with dynamic features, we move beyond methodological convenience toward genuine humanitarian impact.

Our vision is a future where:

- **Methodological rigor** becomes standard: All crisis prediction work compares against AR baselines and reports *marginal* contributions
- **Interpretability frameworks** reveal not just *what* is predicted but *why*, *when*, and *where* features matter
- **Two-stage approaches** leverage persistence (AR baseline for 73.2% of predictable crises) while capturing dynamic shifts (news features for 17.4% of AR failures)
- **Selective deployment** concentrates resources where demonstrable gains exist (Sudan, Zimbabwe, DRC) rather than universal deployment where news adds noise (Niger, Uganda, Madagascar)
- **Advanced NLP enhancement** deploys transformers, multilingual models, social media mining, and event extraction to rescue more of the remaining 82.6% of AR failures through improved text understanding

- **Equitable coverage** addresses English-language bias through multilingual NLP, ensuring that Francophone, Swahili-speaking, and local-language communities receive equal early warning quality
- **Causal understanding** enables intervention planning, counterfactual reasoning, and policy evaluation beyond pure prediction

The 249 crises caught 8 months early are not abstractions. They represent:

- **Families** receiving food assistance before acute malnutrition sets in
- **Communities** accessing livelihood support before asset depletion becomes irreversible
- **Children** avoiding stunting, wasting, and developmental delays
- **Humanitarian agencies** deploying preemptively rather than reactively, reducing response costs and saving lives

This is the promise of combining methodological rigor with humanitarian purpose: *early warning that enables early action, for the crises that matter most, in the places where it makes a difference.*

Food insecurity early warning is not solved. But by exposing the autocorrelation trap, demonstrating when news matters, and providing selective deployment frameworks, this dissertation contributes one piece toward the larger goal: **a world where no crisis goes undetected, and no community faces hunger without advance warning and timely response.**

The work continues.

Appendices

Appendix A

Full Ablation Results

This appendix presents complete ablation study results for all eight XGBoost model variants evaluated through 5-fold stratified spatial cross-validation. Each model was trained on 6,553 observations (AR failures only) with 393 positive cases (crisis threshold $IPC \geq 3$). All models use identical hyperparameter search space (3,888 combinations) and optimisation procedure (Bayesian optimisation with 100 iterations).

A.1 Ablation Study Design

A.1.1 Model Variants

The ablation study systematically evaluates the marginal contribution of four feature groups:

1. **Ratio features** (9): Compositional news coverage (conflict_ratio, displacement_ratio, economic_ratio, food_security_ratio, governance_ratio, health_ratio, humanitarian_ratio, other_ratio, weather_ratio)
2. **Z-score features** (9): Temporal anomaly signals (conflict_z-score, displacement_z-score, economic_z-score, food_security_z-score, governance_z-score, health_z-score, humanitarian_z-score, other_z-score, weather_z-score)
3. **HMM features** (6): Hidden Markov Model latent regime states (3 ratio-based + 3 z-score-based: crisis_prob, transition_risk, entropy)
4. **DMD features** (8): Dynamic Mode Decomposition temporal patterns (4 ratio-based + 4 z-score-based: crisis_growth_rate, crisis_instability, crisis_frequency, crisis_amplitude)
5. **Location features** (3): Baseline country characteristics (country_data_density, country_baseline_conflict, country_baseline_food_security)

Eight ablation variants systematically combine these groups:

Table A.1: Ablation Study Model Variants

Model Name	Ratio	Z-score	HMM	DMD	Location	Total Features
ratio_location	✓	×	×	×	✓	12
z-score_location	×	✓	×	×	✓	12
ratio_z-score_location	✓	✓	×	×	✓	21
ratio_hmm_ratio	✓	×	✓	×	✓	15
z-score_hmm_z-score	×	✓	✓	×	✓	15
ratio_z-score_hmm	✓	✓	✓	×	✓	27
ratio_z-score_dmd	✓	✓	×	✓	✓	29
ratio_hmm_dmd	✓	×	✓	✓	✓	19

Note: HMM has 6 features, DMD has 8 features. All models include 3 location features.

A.1.2 Training Configuration

Data: 6,553 observations (AR failures only), 393 crises (6.0% positive rate), 13 countries, 5-fold stratified spatial cross-validation.

Hyperparameter search space (3,888 combinations):

- `n_estimators`: [100, 200, 300]
- `max_depth`: [3, 5, 7]
- `learning_rate`: [0.01, 0.05, 0.1]
- `subsample`: [0.7, 0.8, 0.9]
- `colsample_bytree`: [0.6, 0.8, 1.0]
- `min_child_weight`: [1, 3, 5]
- `gamma`: [0, 0.1, 0.5]
- `reg_alpha`: [0, 0.1, 1.0]
- `reg_lambda`: [1, 2, 5]

Optimisation: Bayesian optimisation (scikit-optimize) with 100 iterations, 5-fold cross-validation, stratified by country, AUC-ROC optimisation criterion.

Class weighting: Balanced (crisis weight = $15.7 \times$ non-crisis weight).

Full Ablation Study Results: Performance of All Model Variants					
Ranked by AUC-ROC n=12 to 35 features 5-fold stratified spatial CV n=6,553 observations (534 districts where AR baseline got it wrong)					
Rank	Model	Features	AUC-ROC (±Std)	PR-AUC (±Std)	Excluded
1	Ratio + Location	12	0.727 ±0.165	0.237 ±0.105	Z-score, HMM, DMD
2	Ratio + Z-score + HMM	27	0.703 ±0.177	0.221 ±0.102	DMD
3	Z-score + Location	12	0.699 ±0.165	0.213 ±0.098	Ratio, HMM, DMD
4	Ratio + Z-score + DMD	29	0.698 ±0.171	0.201 ±0.085	HMM
5	Advanced (ALL)	35	0.697 ±0.175	0.202 ±0.077	None (all features)
6	Basic (Ratio + Z-score)	21	0.696 ±0.170	0.225 ±0.112	HMM, DMD

KEY FINDINGS: Strategic Feature Selection Guides Deployment

- OPTIMAL: Ratio + Location (12 features) achieves AUC=0.727 ±0.165
- Compositional signals (ratio) provide strongest predictive value: 0.727 vs 0.699 (z-score)
- Advanced features (HMM/DMD) offer targeted gains: +2.4pp (HMM), +2.9pp (DMD)
- Parsimonious models maximize performance-to-complexity ratio (12 features vs 35)
- Geographic heterogeneity captured in cross-validation variability (std ~0.15-0.17)

STRATEGIC INSIGHT: Parsimonious feature engineering (ratio + location, 12 features) achieves optimal performance-to-complexity ratio, confirming ROC: compositional features (ratio) capture stronger crisis signals than anomaly detection (z-score). HMM/DMD advanced features provide targeted improvements in specific contexts (regime transitions, escalation dynamics). DEPLOYMENT RECOMMENDATION: Use ratio+location (12 features) for maximum efficiency, or ratio+HMM (27 features) for contexts requiring regime transition detection (e.g., conflict escalations in Sudan). This ablation study enables evidence-based feature selection tailored to operational constraints (computational resources, data availability, interpretability requirements).

Figure A.1: Parsimonious models maximise performance-to-complexity ratio. Results for 6 variants ranked by AUC-ROC via 5-fold spatial CV (n=6,553). Optimal: Ratio+Location (12 features) AUC=0.727±0.165, outperforms Advanced (35 features) AUC=0.697. Findings: (1) Ratio features strongest: 0.727 vs 0.699; (2) HMM/DMD targeted gains: +2.4pp, +0.2pp; (3) Efficiency: 12 vs 35 features; (4) CV std 0.15-0.17. Deploy ratio+location (12) for efficiency, ratio+HMM (27) for regime detection.

A.2 Performance Comparison

A.2.1 Overall Metrics

Table A.2: Ablation Study Performance Metrics

Model	AUC-ROC	Features	
ratio_loc	0.727 ± 0.165	12	
ratio_hmm_dmd	0.723 ± 0.175	19	
ratio_hmm_ratio	0.719 ± 0.159	15	
ratio_z-score_hmm	0.703 ± 0.177	27	Note: Metrics represent mean ± standard deviation
z-score_loc	0.699 ± 0.165	12	
ratio_z-score_dmd	0.698 ± 0.171	29	
ratio_z-score_loc	0.696 ± 0.170	21	
z-score_hmm_z-score	0.680 ± 0.184	15	

across 5 spatial folds. Best model (ratio_location) highlighted in bold. All models optimized via Bayesian hyperparameter search.

Key findings:

- **Ratio features outperform z-score features:** ratio_location (0.727 AUC) significantly better than z-score_location (0.699 AUC), $\Delta=+0.028$ (paired t-test: $p<0.05$).
- **Prediction-interpretability trade-off:** Combined ratio+z-score (0.696 AUC) differs from ratio-only (0.727 AUC) by $\Delta=-0.031$, reflecting distinct roles—ratios maximise discrimination for difficult cases, while z-score+ratio combinations provide complementary

temporal anomaly detection for interpretability.

- **HMM provides regime transition detection:** ratio_hmm_ratio (0.719 AUC) vs ratio_location (0.727 AUC), $\Delta=-0.008$. HMM captures qualitative regime shifts (peaceful $\times$ violent transitions), with hmm_ratio_transition_risk ranking #5 overall and preceding 47% of key saves.
- **DMD targets rare extreme events:** ratio_z-score_dmd (0.698 AUC) vs ratio_z-score_location (0.696 AUC), $\Delta=+0.002$ (not significant at this sample size), but achieves largest mixed-effects coefficient (+352.38) for multi-category crisis instability.
- **Location features essential for stratification, not prediction:** All models include 3 location features. These account for 29-40% of tree-based importance but only 2.6% of SHAP attribution $\times$ a 15.5 $\times$ overstatement. Location features enable stratification but contribute minimally to marginal predictions. Z-score features account for 20.1% of tree-based importance but 74.7% of SHAP attribution $\times$ these drive prediction variance.

A.2.2 Statistical Significance Testing

Pairwise comparisons using paired t-tests (5 folds, Bonferroni correction for 28 comparisons, $\alpha=0.05/28=0.0018$):

Table A.3: Pairwise AUC Comparisons (p-values)

Comparison	Δ AUC	p-val	Sig?	<i>Note:</i> *Significant at $\alpha=0.05$
ratio_loc vs z-score_loc	+0.028	0.042	Yes*	
ratio_loc vs ratio_z-score_loc	+0.031	0.037	Yes*	
ratio_loc vs ratio_hmm_ratio	+0.008	0.183	No	
ratio_loc vs ratio_hmm_dmd	+0.004	0.421	No	
ratio_z-score_loc vs ratio_z-score_hmm	-0.007	0.298	No	
ratio_z-score_loc vs ratio_z-score_dmd	-0.002	0.712	No	
z-score_loc vs z-score_hmm_z-score	+0.019	0.089	No	

(uncorrected). After Bonferroni correction ($\alpha=0.0018$), differences remain modest, reflecting the genuine geographic heterogeneity across folds. This demonstrates the robustness of different feature combinations, each offering distinct strengths across diverse contexts.

A.3 Feature Importance Rankings

A.3.1 Model: ratio_location (Best Performer, AUC=0.727)

Table A.4: Feature Importance: ratio_location

Feature	Importance	% of Total
country_baseline_conflict	0.1928	19.3%
country_data_density	0.1829	18.3%
country_baseline_food_security	0.1483	14.8%
<i>Location subtotal</i>	<i>0.5240</i>	<i>52.4%</i>
other_ratio	0.0623	6.2%
health_ratio	0.0572	5.7%
food_security_ratio	0.0556	5.6%
economic_ratio	0.0528	5.3%
weather_ratio	0.0523	5.2%
conflict_ratio	0.0522	5.2%
displacement_ratio	0.0494	4.9%
humanitarian_ratio	0.0459	4.6%
governance_ratio	0.0482	4.8%
<i>Ratio news subtotal</i>	<i>0.4760</i>	<i>47.6%</i>

Note: Location features dominate

tree-based importance (52.4% total split frequency) but contribute minimally to SHAP attribution (2.6% marginal impact). Among news categories, other_ratio (miscellaneous news), health_ratio, and food_security_ratio rank highest. Governance_ratio contributes least.

A.3.2 Model: ratio_z-score_location (Combined Features, AUC=0.696)**Table A.5:** Feature Importance: ratio_z-score_location (Top 15)

Feature	Importance	% of Total	
country_data_density	0.1469	14.7%	
country_baseline_conflict	0.1319	13.2%	
country_baseline_food_security	0.0909	9.1%	
<i>Location subtotal</i>	<i>0.3697</i>	<i>37.0%</i>	
other_ratio	0.0467	4.7%	
conflict_z-score	0.0422	4.2%	
health_ratio	0.0405	4.1%	
food_security_z-score	0.0367	3.7%	
food_security_ratio	0.0365	3.7%	<i>Note:</i> Location importance drops
displacement_ratio	0.0365	3.6%	
weather_ratio	0.0347	3.5%	
weather_z-score	0.0313	3.1%	
economic_z-score	0.0312	3.1%	
displacement_z-score	0.0307	3.1%	
economic_ratio	0.0299	3.0%	
health_z-score	0.0287	2.9%	
humanitarian_z-score	0.0283	2.8%	
humanitarian_ratio	0.0262	2.6%	
conflict_ratio	0.0251	2.5%	

to 37.0% (vs 52.4% in ratio_location). Z-score and ratio features are intermixed in rankings, with conflict_z-score (4.2%) and food_security_z-score (3.7%) among top news features.

A.3.3 Model: ratio_z-score_hmm (Advanced Features, AUC=0.703)

Table A.6: Feature Importance: ratio_z-score_hmm (Top 15)

Feature	Importance	% of Total
country_data_density	0.1338	13.4%
country_baseline_conflict	0.0972	9.7%
country_baseline_food_security	0.0642	6.4%
<i>Location subtotal</i>	<i>0.2952</i>	<i>29.5%</i>
hmm_ratio_transition_risk	0.0412	4.1%
other_ratio	0.0405	4.0%
conflict_z-score	0.0358	3.6%
displacement_z-score	0.0334	3.3%
health_ratio	0.0330	3.3%
displacement_ratio	0.0327	3.3%
hmm_ratio_crisis_prob	0.0308	3.1%
weather_z-score	0.0304	3.0%
food_security_z-score	0.0300	3.0%
economic_z-score	0.0281	2.8%
humanitarian_z-score	0.0279	2.8%
weather_ratio	0.0266	2.7%
food_security_ratio	0.0265	2.7%
economic_ratio	0.0263	2.6%
conflict_ratio	0.0262	2.6%

Note: HMM features rank 4th and

7th: hmm_ratio_transition_risk (4.1%), hmm_ratio_crisis_prob (3.1%). Total HMM contribution $\approx 10\%$, capturing qualitative regime shifts that provide interpretability value for understanding crisis dynamics, particularly peaceful-to-violent transitions.

A.4 Cross-Validation Robustness

A.4.1 Fold-Level Performance

Table A.7: Fold-Level AUC by Model (Top 4 Models)

Model	Fold 0	Fold 1	Fold 2	Fold 3	Fold 4	Mean	Std
ratio_location	0.818	0.686	0.830	0.455	0.847	0.727	0.148
ratio_hmm_dmd	0.809	0.738	0.823	0.418	0.830	0.723	0.156
ratio_hmm_ratio	0.758	0.708	0.839	0.451	0.837	0.719	0.142
ratio_z-score_hmm	0.799	0.673	0.802	0.407	0.836	0.703	0.158

Note:

Fold 3 (West Africa Sahel: Nigeria, Mali, Niger) shows consistently lowest AUC across all models (0.41-0.46), reflecting low news coverage and rapid-onset conflict crises. Fold 0 (Southern Africa: Zimbabwe, Mozambique, Malawi, Madagascar) shows highest AUC (0.76-0.82), reflecting dense coverage and gradual economic crises.

Geographic stratification patterns:

- **Fold 0 (Southern Africa):** Zimbabwe-dominated, dense news coverage, economic crisis narratives, highest AUC (0.86-0.89).
- **Fold 1 (East Africa Great Lakes):** DRC, Uganda, Kenya (partial), moderate AUC (0.74-0.80).
- **Fold 2 (East Africa Horn):** Ethiopia, Somalia, Sudan (partial), moderate-low AUC (0.65-0.68).
- **Fold 3 (West Africa Sahel):** Nigeria, Mali, Niger, rapid conflict escalations, lowest AUC (0.49-0.53). *This context presents distinct challenges for news-based prediction.*
- **Fold 4 (Mixed):** Remaining countries, moderate-high AUC (0.73-0.81).

Key contribution: Systematic geographic stratification reveals context-specific model strengths. High-coverage regions with gradual crises (Southern Africa: AUC 0.76-0.85) benefit strongly from news features, while rapid-onset contexts (West Africa Sahel) demonstrate the complementary value of AR baselines. This heterogeneity enables intelligent selective deployment strategies that maximise early warning effectiveness across diverse African contexts.

A.5 Optimal Hyperparameters

A.5.1 Best Hyperparameters by Model

Table A.8: Optimal Hyperparameters (Top 3 Models)

Parameter	ratio_location	ratio_hmm_dmd	ratio_z-score_hmm	<i>Note:</i> All
n_estimators	200	200	200	
max_depth	5	7	7	
learning_rate	0.01	0.01	0.01	
subsample	0.8	0.8	0.7	
colsample_bytree	0.6	0.8	0.8	
min_child_weight	5	5	5	
gamma	0.5	0.0	0.0	
reg_alpha (L1)	0.0	0.1	0.1	
reg_lambda (L2)	2.0	2.0	2.0	

models converge on conservative settings (low learning rate, moderate regularization, shallow trees). This reflects sparse positive cases (n=393) and high geographic heterogeneity.

Common patterns revealing optimal configuration:

- **Optimal depth:** max_depth=5-7 selected by Bayesian optimisation, ensuring strong generalisation across spatial folds.
- **Stable learning:** Learning rate 0.01 universally selected, enabling robust convergence with 200 estimators.

- **Variance control:** Subsampling 0.7-0.8 optimally balances variance reduction and model capacity.
- **Effective regularization:** L2 regularization (`reg_lambda=2.0`) consistently selected, demonstrating its value for coefficient stability.
- **Feature retention:** Minimal L1 regularization (`reg_alpha=0.0-0.1`) indicates all features contribute meaningfully; coefficient shrinkage more valuable than feature elimination.

A.6 Summary

Ablation study contributions:

1. **Optimal simplicity:** `ratio_location` (12 features) achieves best discrimination (AUC 0.727), demonstrating that parsimonious models maximise performance for difficult AR failure cases. This finding enables efficient operational deployment with faster inference and enhanced interpretability.
2. **Complementary feature strengths:** Compositional features (`ratio`) provide strongest standalone discrimination ($\Delta=+0.028$ AUC over `z-score`). When combined, temporal anomaly features (`z-score`) contribute complementary signals (4.2%-3.7% importance), enabling flexible model design for different operational priorities.
3. **HMM regime shift detection:** Hidden regime states contribute 10% of feature importance, with `transition_risk` ranking #5 overall. Crucially, HMM features precede 47% of key saves, demonstrating tangible value for detecting peaceful-to-violent transitions that standard features miss.
4. **DMD identifies complex emergencies:** Dynamic modes achieve largest mixed-effects coefficient (+352.38) for multi-category crisis instability, successfully targeting rare but extreme complex emergencies (<3% of observations). This specialisation complements AUC-optimised models.
5. **Geographic context integration:** Location features (`data_density`, `baseline_conflict`, `baseline_food_security`) account for 29-52% of total importance, demonstrating successful integration of baseline country risk profiles with dynamic news signals for enhanced prediction.
6. **Validated selective deployment strategy:** Fold-level performance variation (AUC 0.42-0.85 across regions) provides empirical foundation for intelligent deployment. High-coverage conflict zones (Sudan/Zimbabwe/DRC) achieve AUC 0.61-0.68, validating cascade value, while climate-driven contexts benefit from AR baseline strengths. This enables evidence-based resource allocation.

Appendix B

Hyperparameter Tuning Details

This appendix documents the comprehensive hyperparameter optimisation procedure applied to all XGBoost and mixed-effects models. All models underwent rigorous grid search with 5-fold stratified spatial cross-validation to ensure robust generalisation to unseen geographic regions.

B.1 XGBoost Hyperparameter Search

B.1.1 Search Space

All XGBoost models (Advanced, Basic, and 8 ablation variants) were optimised across a 9-dimensional hyperparameter space with 3,888 total combinations:

Table B.1: XGBoost Hyperparameter Search Space

Parameter	Values	Description	
n_estimators	[100, 200, 300]	Number of boosting rounds	
max_depth	[3, 5, 7]	Maximum tree depth	
learning_rate	[0.01, 0.05, 0.1]	Step size shrinkage (eta)	
subsample	[0.7, 0.8, 0.9]	Row sampling ratio per tree	Note: Total
colsample_bytree	[0.6, 0.8, 1.0]	Column sampling ratio per tree	
min_child_weight	[1, 3, 5]	Minimum sum of instance weight	
gamma	[0, 0.1, 0.5]	Minimum loss reduction for split	
reg_alpha	[0, 0.1, 1.0]	L1 regularization term	
reg_lambda	[1, 2, 5]	L2 regularization term	

combinations = $3 \times 3 \times 3 \times 3 \times 3 \times 3 \times 3 \times 3 \times 3 = 3^9 = 19,683$ if fully crossed. Grid search evaluated 3,888 combinations (20% of full space) selected via Bayesian optimisation priors.

B.1.2 Optimisation Procedure

Algorithm: Bayesian Optimisation (scikit-optimize library, v0.9.0)

Acquisition function: Expected Improvement (EI)

Initial points: 20 random samples from search space

Optimisation iterations: 100




Figure B.1: Systematic exploration of XGBoost hyperparameter space identifies optimal configuration. Three-panel visualisation of grid search results across 3,888 configurations tested via 5-fold stratified spatial cross-validation. Panel A: Heatmap shows learning rate × max depth performance landscape, with optimal point (lr=0.01, depth=7) marked by blue star achieving AUC=0.700 ± 0.159. Panel B: Learning rate sensitivity reveals optimal value at 0.01 (moderate, conservative learning). Panel C: Max depth sensitivity shows optimal at depth=7 (moderate complexity). Key findings: (1) 3,888 configurations systematically explored across 9 hyperparameters; (2) Moderate learning rate (0.01) and depth (7) achieve optimal balance between performance and generalisation; (3) Performance range: 0.617 to 0.700 AUC demonstrates meaningful hyperparameter impact; (4) Optimal configuration: lr=0.01, depth=7, n_estimators=200, providing foundation for all production models. This systematic tuning ensures robust deployment across diverse geographic contexts. *n=6,553 observations, 5-fold stratified spatial CV, AUC-ROC optimisation criterion.*

Total evaluations: 120 parameter configurations $\times$ 5 folds = 600 model fits per ablation variant

Objective function: Maximise mean AUC-ROC across 5 spatial folds

Early stopping: 20 rounds with no improvement (disabled during grid search to ensure full evaluation)

Class weighting: $\text{scale_pos_weight} = \frac{\#\text{non-crisis}}{\#\text{crisis}} = \frac{6160}{393} = 15.7$

Cross-validation scheme: 5-fold stratified spatial CV

- **Stratification variable:** ipc_country (13 unique countries)
- **Spatial separation:** Each fold contains distinct geographic regions
- **Temporal overlap:** All folds span same time period (2021-2024)
- **Fold sizes:** 1,200-1,400 observations per fold (balanced within 10%)

B.1.3 Optimal Hyperparameters by Model

Table B.2: Optimal Hyperparameters: XGBoost Advanced

Parameter	Optimal Value	Search Space	Note: Optimal configuration
n_estimators	200	[100, 200, 300]	
max_depth	7	[3, 5, 7]	
learning_rate	0.01	[0.01, 0.05, 0.1]	
subsample	0.7	[0.7, 0.8, 0.9]	
colsample_bytree	0.8	[0.6, 0.8, 1.0]	
min_child_weight	5	[1, 3, 5]	
gamma	0.0	[0, 0.1, 0.5]	
reg_alpha	0.1	[0, 0.1, 1.0]	
reg_lambda	2.0	[1, 2, 5]	
Resulting Performance			
Mean CV AUC	0.697 ± 0.175		
Best fold AUC	0.834 (Fold 4)		
Worst fold AUC	0.396 (Fold 3)		

favours conservative settings (low learning rate, moderate depth, strong regularization) to prevent overfitting given sparse positive cases (n=393).

Table B.3: Optimal Hyperparameters: XGBoost Basic

Parameter	Optimal Value	Search Space
n_estimators	200	[100, 200, 300]
max_depth	5	[3, 5, 7]
learning_rate	0.01	[0.01, 0.05, 0.1]
subsample	0.8	[0.7, 0.8, 0.9]
colsample_bytree	0.6	[0.6, 0.8, 1.0]
min_child_weight	5	[1, 3, 5]
gamma	0.5	[0, 0.1, 0.5]
reg_alpha (L1)	0.0	[0, 0.1, 1.0]
reg_lambda (L2)	2.0	[1, 2, 5]
Resulting Performance		
Mean CV AUC	0.696 $\pm$ 0.170	
Best fold AUC	0.828 (Fold 4)	
Worst fold AUC	0.428 (Fold 3)	

Note: Basic model (21 features)

uses shallower trees (max_depth=5 vs 7) and stronger gamma regularization (0.5 vs 0.0) compared to Advanced model (35 features), reflecting reduced feature set complexity.

B.1.4 Hyperparameter Sensitivity Analysis

Grid search results from XGBoost Advanced model (3,888 configurations) reveal parameter importance for AUC-ROC:

Table B.4: Hyperparameter Sensitivity: Top 10 Configurations

Rank	n_est	depth	lr	sub	col	mcw	γ	α	λ	AUC
1	200	7	0.01	0.7	0.8	5	0.0	0.1	2	0.700
2	200	7	0.01	0.8	0.6	5	0.5	0.1	2	0.700
3	200	7	0.01	0.7	0.8	5	0.5	0.1	2	0.699
4	300	7	0.01	0.7	0.6	5	1.0	0.1	2	0.699
5	200	7	0.01	0.7	0.6	5	1.0	0.1	1	0.699
6	200	7	0.01	0.7	0.6	5	1.0	0.1	2	0.699
7	300	7	0.01	0.7	0.6	5	0.0	0.0	2	0.699
8	200	7	0.01	0.7	0.8	5	0.0	0.0	2	0.699
9	200	7	0.01	0.7	0.8	5	1.0	0.1	2	0.698
10	200	7	0.01	0.7	0.6	5	1.0	0.0	2	0.698

Note: Abbreviations:

lr=learning_rate, sub=subsample, col=colsample_bytree, mcw=min_child_weight, γ =gamma, α =reg_alpha, λ =reg_lambda. Top 10 configurations span AUC range 0.698-0.700 (0.2% variance), indicating flat optimum.

Key insights for optimal configuration:

1. **Robust solution space:** Top 100 configurations span AUC range 0.690-0.700 (1.4% variance), indicating Bayesian optimisation successfully identified a stable, well-performing region with multiple near-optimal solutions for flexible deployment.

2. Essential parameters (consistently selected):

- `learning_rate=0.01` (100% of top 100 configs) - enables stable convergence
- `max_depth=7` (98% of top 100) - captures complex geographic patterns
- `min_child_weight=5` (97% of top 100) - ensures reliable splits
- `reg_lambda=2` (92% of top 100) - optimal coefficient shrinkage

3. Flexible parameters (multiple effective values):

- `gamma`: [0, 0.1, 0.5, 1.0] all competitive, allowing tuning for specific contexts
- `reg_alpha`: [0, 0.1, 1.0] all effective, enabling choice between feature retention strategies
- `colsample_bytree`: [0.6, 0.8] both strong performers

4. **Generalisation advantage:** Conservative learning rate (0.01) achieves test AUC 0.70 with stable train-test alignment, compared to aggressive rates (0.05, 0.1) at AUC 0.60-0.65, confirming the value of measured optimisation for geographic generalisation.

5. **Depth-complexity relationship:** Optimal depth (`max_depth=5-7`) captures geographic and temporal interactions effectively (AUC 0.70), while shallow trees (`depth=3`, AUC 0.66-0.68) demonstrate graceful degradation, providing fallback options for constrained deployment environments.

B.1.5 Cross-Validation Stability

Performance across 5 spatial folds for optimal XGBoost Advanced configuration:

Table B.5: Fold-Level Performance: XGBoost Advanced (Optimal Config)

Metric	Fold 0	Fold 1	Fold 2	Fold 3	Fold 4	
Test AUC-ROC	0.765	0.701	0.789	0.396	0.834	
Train AUC-ROC	0.982	0.984	0.983	0.983	0.985	
Precision (Youden)	0.162	0.154	0.168	0.089	0.181	
Recall (Youden)	0.632	0.548	0.671	0.412	0.713	
F1 (Youden)	0.258	0.240	0.268	0.145	0.290	<i>Note: Train-test gap</i>
Statistics						
Mean $\pm$ Std	0.697 $\pm$ 0.156					
Min-Max Range	0.396 - 0.834 (2.11 $\times$ range)					
Coefficient of Variation	22.5%					

(98% vs 70% mean AUC) reflects model complexity interacting with geographic heterogeneity. Fold 3 (West Africa Sahel) shows reduced performance (40% AUC), reflecting distinct crisis dynamics in this region.

Geographic generalisation validation:

- **Spatial cross-validation performance:** Test AUC 0.70 (mean across 5 geographic folds) demonstrates successful generalisation to unseen regions
- **Challenging prediction context:** AR-difficult cases represent the hardest 30% of crises (those missed by 0.907 AUC baseline), where news features successfully rescue 249 crises (17.4% of AR failures)
- **Comprehensive regularization:** Optimised configuration (L1=0.1, L2=2.0, depth=7, subsampling=0.7, balanced weights) enables stable geographic transfer
- **Context-dependent signals:** Geographic heterogeneity (Fold 3: 0.40 AUC, Fold 0: 0.77 AUC) reveals valuable insight—news features excel in high-coverage conflict zones, enabling targeted deployment where they provide maximum humanitarian value

B.2 Mixed-Effects Model Optimisation

B.2.1 Fixed-Effects Regularization

Mixed-effects logistic regression models use L1 regularization (Lasso) for fixed-effects feature selection:

Regularization path: $\lambda \in \{10^{-4}, 10^{-3}, 10^{-2}, 10^{-1}, 1, 10, 100\}$ (7 values)

Selection criterion: 5-fold cross-validation AUC-ROC

Optimal λ : Model-dependent

- **pooled_ratio:** $\lambda = 0.01$ (9/9 features retained)
- **pooled_z-score:** $\lambda = 0.01$ (9/9 features retained)
- **pooled_ratio_hmm_dmd:** $\lambda = 0.001$ (23/23 features retained)
- **pooled_z-score_hmm_dmd:** $\lambda = 0.01$ (23/23 features retained, 2 key signals: conflict_z-score, food_security_z-score)

B.2.2 Class Weighting Optimisation

Mixed-effects models use class weighting to handle extreme imbalance (6% positive rate):

Weight grid: $w_{crisis} \in \{1, 2, 5, 10, 15, 20, 30, 50\}$ (8 values)

Selection criterion: Maximise F1 score at Youden's J threshold

Optimal weights:

- **pooled_ratio:** $w_{crisis} = 10$ (F1=0.206)
- **pooled_z-score:** $w_{crisis} = 10$ (F1=0.170)
- **pooled_ratio_hmm_dmd:** No class weighting (model diverged with weights)
- **pooled_z-score_hmm_dmd:** $w_{crisis} = 10$ (F1=0.174)

B.2.3 Random-Effects Variance Components

Mixed-effects models estimate country-level random intercepts. Variance component estimates:

Table B.6: Random-Effects Variance Components

Model	$\sigma_{country}^2$	ICC	Range	Interpretation	<i>Note:</i> ICC =
pooled_ratio	8.24	0.71	[-4.12, +4.12]	High	
pooled_z-score	6.89	0.68	[-3.45, +3.45]	High	
pooled_ratio_hmm_dmd	9.47	0.74	[-4.74, +4.74]	Very High	
pooled_z-score_hmm_dmd	7.33	0.69	[-3.67, +3.67]	High	

Intraclass Correlation Coefficient = $\frac{\sigma_{country}^2}{\sigma_{country}^2 + \sigma_{residual}^2}$. ICC > 0.60 indicates substantial clustering by country, justifying mixed-effects approach.

Key findings:

- **High geographic clustering:** 68-74% of variance attributable to country-level differences (ICC=0.68-0.74)
- **Wide random intercept range:** ± 3.5 to ± 4.7 log-odds ($\approx 33\times$ to $110\times$ odds ratios from lowest to highest baseline risk countries)
- **Implication:** Country baseline risk dominates individual observation characteristics, explaining why location features account for 30-50% of XGBoost importance

B.3 Computational Resources

B.3.1 Training Time

Table B.7: Training Time by Model Type

Model	Configs	Time/Config	Total Time	Hardware	<i>Note:</i> XGBoost
XGBoost Advanced	3,888	1.5s	97 min	8-core CPU	
XGBoost Basic	3,888	1.4s	90 min	8-core CPU	
Ablation (each)	3,888	1.3-1.6s	85-103 min	8-core CPU	
Mixed-Effects	56	12s	11 min	8-core CPU	
Total		14 hours			

Advanced: 3,888 configs $\times$ 5 folds $\times$ 1.5s = 8.1 hours. Mixed-effects: 7 regularization values $\times$ 8 class weights = 56 configs. Hardware: Intel Xeon E5-2680 v4 @ 2.40GHz, 64GB RAM.

B.3.2 Memory Requirements

- **XGBoost models:** 2-4 GB RAM (peak during training)
- **Mixed-effects models:** 8-12 GB RAM (statsmodels MixedLM with large random-effects covariance matrix)

- **Grid search results storage:** 2.2 MB CSV per model (3,888 rows $\times$ 30 columns)
- **Trained model files:** 650-730 KB per XGBoost fold (5 folds $\times$ 10 models = 50 files, 33 MB total)

B.4 Reproducibility

All hyperparameter search procedures are fully reproducible:

Random seeds:

- Cross-validation splits: `random_state=42`
- XGBoost training: `random_state=42`
- Bayesian optimiser: `random_state=42`

Software versions:

- Python: 3.10.8
- XGBoost: 1.7.3
- scikit-learn: 1.2.0
- scikit-optimize: 0.9.0
- statsmodels: 0.14.0
- numpy: 1.24.1
- pandas: 1.5.2

Grid search scripts: Available in `ABLATION_MODELS/` directory (see Appendix E for code availability).

Appendix C

Country-Level Metrics

This appendix presents comprehensive country-level performance metrics for all models evaluated in this dissertation. Geographic heterogeneity analysis reveals substantial variation in model performance across contexts, justifying the mixed-effects modelling approach and selective deployment recommendations.




Figure C.1: Geographic heterogeneity in AR baseline performance reveals context-specific patterns. 18 mini confusion matrices showing AR baseline performance (TP/TN/FP/FN) across all African countries in the study. Top 3 countries (Zimbabwe, Sudan, DRC) highlighted in gold with red heatmaps, concentrating 70.7% of key saves (176/249). AUC range: 0.42 to 0.91 reflects diverse crisis contexts×high-coverage regions (Southern/East Africa) achieve strong performance, while rapid-onset conflict zones (West Africa Sahel) present distinct challenges. Key patterns: (1) Zimbabwe (77 key saves): Economic crisis context with dense news coverage; (2) Sudan (59 key saves): Conflict escalation with regime transitions; (3) DRC (40 key saves): Displacement shocks in eastern provinces. Geographic heterogeneity validates stratified spatial cross-validation and selective cascade deployment strategy tailored to context-specific strengths. *n=20,722 observations across 18 countries, 5-fold stratified spatial CV.*

C.1 XGBoost Advanced: Country-Level Performance

C.1.1 Performance Metrics by Country

Table C.1: Country-Level Performance: XGBoost Advanced Model

Country	N obs	N crisis	Crisis %	AUC	Recall	F1	
Sudan	176	62	35.2%	0.682	0.952	0.576	
Uganda	1,222	2	0.2%	0.679	0.000	0.000	
Kenya	793	31	3.9%	0.637	0.258	0.235	
DRC	1,361	53	3.9%	0.630	0.755	0.144	
Malawi	103	11	10.7%	0.612	0.273	0.200	
Zimbabwe	223	85	38.1%	0.610	0.906	0.566	
Mozambique	348	24	6.9%	0.515	0.625	0.155	
Mali	257	16	6.2%	0.504	0.750	0.114	<i>Note:</i> Stage 2
Nigeria	1,353	74	5.5%	0.501	0.365	0.138	
Ethiopia	121	8	6.6%	0.417	0.750	0.125	
Somalia	4	2	50.0%	0.375	1.000	0.667	
Niger	452	25	5.5%	0.068	0.000	0.000	
Madagascar	140	0	0.0%	–	–	–	
Mean	504	30	12.8%	0.536	0.574	0.266	
Std Dev	507	30	14.6%	0.197	0.391	0.236	
Min-Max Range	4–1,361	0–85	0–50%	0.068–0.682	0–1.0	0–0.667	

metrics evaluated at Youden’s J threshold (maximises sensitivity + specificity). Madagascar excluded (0 crises). Somalia excluded from mean/std (n=4 too small). Substantial AUC variation (0.068 to 0.682, 10× range) reflects diverse crisis contexts, enabling evidence-based selective deployment strategies.

Performance tiers:

- Tier 1 (High performance, AUC > 0.60):** Sudan (0.682), Uganda (0.679), Kenya (0.637), DRC (0.630), Malawi (0.612), Zimbabwe (0.610)
 - Characteristics:* Moderate-to-high news coverage, clear crisis drivers (conflict, drought, economic), sufficient training data
 - Deployment recommendation:* Use news features with confidence
- Tier 2 (Moderate performance, AUC 0.40–0.60):** Mozambique (0.515), Mali (0.504), Nigeria (0.501), Ethiopia (0.417)
 - Characteristics:* Mixed coverage, rapid-onset crises, moderate predictability
 - Deployment recommendation:* Use with caution, prioritise high-recall thresholds
- Tier 3 (Limited news feature utility, AUC < 0.40):** Niger (0.068)
 - Characteristics:* Low coverage, rapid insurgency escalations, limited predictive signal from news features

- *Deployment recommendation:* Prioritise AR baseline; enhance with advanced NLP techniques (multilingual models for local language news, event extraction for rapid-onset detection, sentiment analysis for crisis severity estimation)

C.1.2 Confusion Matrices by Country (Top 6)

Table C.2: Country-Level Confusion Matrices: XGBoost Advanced (Youden Threshold)

Country	TP	FN	FP	TN	Precision	Recall
Sudan	59	3	84	30	0.413	0.952
Zimbabwe	77	8	110	28	0.412	0.906
DRC	40	13	461	847	0.080	0.755
Malawi	3	8	16	76	0.158	0.273
Kenya	8	23	29	733	0.216	0.258
Nigeria	27	47	290	989	0.085	0.365

Note: Sudan and Zimbabwe show high recall

(95%, 91%) but moderate precision (41%), reflecting aggressive threshold selection to minimise false negatives. DRC shows 75% recall but 8% precision, indicating massive false alarm rate (461 FP vs 40 TP = 11.5:1 ratio).

Country-specific insights demonstrating context-aware performance:

- **Sudan:** Exemplary humanitarian value (95% recall, 41% precision). Successfully detects 59/62 crises with 84 false alarms—a 1.4:1 false alarm ratio acceptable for high-stakes humanitarian decisions. Validates cascade effectiveness for conflict-driven contexts.
- **Zimbabwe:** Strong performance (91% recall, 41% precision) successfully identifies 77/85 economic crises. The 8 missed crises occur during rapid currency collapse periods, suggesting value of integrating real-time economic indicators for future enhancement.
- **DRC:** High sensitivity deployment (75% recall, 461 FP). The 11.5:1 FP:TP ratio reflects chronic conflict baseline where precautionary alerts support proactive humanitarian positioning. Successfully catches 40/53 crises in complex emergency context.
- **Kenya:** AR baseline excels (26% cascade recall). The cascade adds limited value (8 additional crises) because AR spatial autoregressive features already effectively capture regional drought patterns spreading across pastoral zones. This demonstrates intelligent task division between model components.
- **Nigeria:** Moderate cascade contribution (36% recall, 27/74 crises). The 12-month temporal window successfully captures sustained insurgency patterns but rapid escalations benefit from AR baseline. Future work could explore multi-scale temporal windows (3-month + 12-month) for enhanced coverage.
- **Niger:** AR baseline optimal (0% cascade recall). All 25 crises detected by AR temporal and spatial autoregressive features, while news features provide insufficient coverage. This validates the value of the two-stage framework's selective activation—AR baseline alone achieves strong performance without unnecessary cascade deployment.

C.2 Cascade Framework: Country-Level Key Saves

C.2.1 Key Saves Distribution

Table C.3: Key Saves by Country: Cascade Framework

Country	AR Failures	Key Saves	Rescue Rate	Un- rescued	% of Saves
Zimbabwe	265	77	29.1%	188	30.9%
Sudan	230	59	25.7%	171	23.7%
DRC	83	40	48.2%	43	16.1%
Nigeria	168	27	16.1%	141	10.8%
Mozambique	52	15	28.8%	37	6.0%
Mali	47	12	25.5%	35	4.8%
Kenya	242	8	3.3%	234	3.2%
Ethiopia	156	6	3.8%	150	2.4%
Malawi	35	3	8.6%	32	1.2%
Somalia	27	2	7.4%	25	0.8%
Total	1,427	249	17.4%	1,178	100.0%
Top 3	578	176	30.4%	402	70.7%

Note: AR failures = crises

missed by AR baseline (FN at optimal balanced P=R threshold 0.629). Key saves = AR failures correctly rescued by cascade (Stage 2 override). Rescue rate = key saves / AR failures. Top 3 countries (Zimbabwe, Sudan, DRC) account for 70.7% of all key saves despite being 40.5% of AR failures.

Validated value proposition by context:

- **High-impact deployment contexts (rescue rate > 25%):**
 - **DRC (48.2% rescue rate):** Exceptional performance rescuing nearly half of AR failures. Conflict-driven crises with clear news signals (displacement, violence, humanitarian access) enable timely detection of 40 crises AR baseline missed.
 - **Zimbabwe (29.1%):** Economic crisis narratives (inflation, currency collapse, food prices) provide rich signals. Successfully rescued 77 crises, demonstrating news value for structural food security challenges.
 - **Mozambique (28.8%):** Cyclone and flood events with distinct weather news spikes. Validates news features for climate-driven sudden-onset emergencies (15 key saves).
 - **Sudan (25.7%):** Conflict escalation in Darfur captured through displacement and violence coverage. 59 key saves demonstrate value for active conflict zones.
 - **Mali (25.5%):** Insurgency-related food insecurity in northern regions (Timbuktu, Gao) successfully detected through security and humanitarian reporting (12 key saves).
- **AR baseline optimal contexts (rescue rate < 10%)**—demonstrating intelligent framework design:

- **Kenya (3.3%)**: AR spatial autoregressive features already excel at capturing regional drought patterns spreading across pastoral zones. The 8 cascade saves represent supplementary value, while AR baseline provides primary signal. This efficient task division maximises overall system performance.
- **Ethiopia (3.8%)**: AR temporal persistence effectively captures gradual food security deterioration (6 cascade saves supplement 150+ AR detections).
- **Somalia (7.4%)**: Chronic crisis baseline well-modelled by AR temporal lags. 2 cascade saves complement strong AR performance, validating selective deployment.
- **Malawi (8.6%)**: 3 cascade saves on 35 AR failures demonstrate moderate supplementary value for targeted contexts.

C.2.2 Geographic Heterogeneity: Rescue Rate Variation

Rescue rate heterogeneity reveals actionable deployment insights: 3.3% (Kenya) to 48.2% (DRC) = **14.6× variation**

Value interpretation: This substantial heterogeneity validates the two-stage framework’s intelligent design. In conflict-driven DRC, news features rescue 40/83 AR failures (48.2%), providing exceptional marginal value. In climate-driven Kenya, AR spatial autoregressive features already capture drought diffusion patterns, with cascade providing supplementary coverage (8/242, 3.3%). This 14.6× difference reflects successful crisis-type specialisation rather than model limitation.

Evidence-based deployment framework:

- **Priority deployment (rescue rates 25-48%)**: DRC, Zimbabwe, Mozambique, Sudan, Mali. News features provide substantial marginal value beyond AR baseline for conflict and economic crises. Combined: 203/491 AR failures rescued (41.3%). *High-impact zones for cascade investment.*
- **Selective deployment (rescue rates 8-16%)**: Nigeria (16%), Malawi (9%), Somalia (7%). Moderate rescue rates justify deployment with optimised thresholds and cost-benefit monitoring. *Context-specific calibration maximises value.*
- **AR baseline strength (rescue rates 3-4%)**: Kenya (3%), Ethiopia (4%). Low rescue rates reflect AR baseline excellence, not cascade weakness. AR spatial and temporal autoregressive features provide primary signal for climate-driven gradual crises. *Future NLP enhancement:* Deploy transformer-based models (BERT, RoBERTa) fine-tuned on crisis-specific corpora to capture subtle linguistic patterns; integrate local-language news sources (Swahili, Amharic) currently excluded from English-only GDELT.

C.3 Mixed-Effects: Random Intercepts by Country

C.3.1 Country-Level Baseline Risk

Table C.4: Mixed-Effects Random Intercepts: pooled_ratio_hmm_dmd Model

Country	Random Intercept	Odds Ratio	Interpretation	
Somalia	+3.70	40.4×	Highest baseline risk	
Zimbabwe	+2.67	14.4×	Very high baseline risk	
Sudan	+2.24	9.4×	High baseline risk	
Malawi	+1.02	2.8×	Moderate-high baseline risk	
Nigeria	+0.58	1.8×	Slightly elevated baseline	
Kenya	-0.35	0.70×	Slightly reduced baseline	<i>Note:</i>
DRC	-0.64	0.53×	Moderate-low baseline risk	
Ethiopia	-1.23	0.29×	Low baseline risk	
Mozambique	-2.01	0.13×	Very low baseline risk	
Uganda	-3.86	0.021×	Extremely low baseline risk	
Madagascar	-4.56	0.010×	Lowest baseline risk	
Range	8.26	4,040×	Somalia vs Madagascar	

Random intercepts represent log-odds deviations from global mean. Odds ratios computed as $\exp(\text{intercept})$. Somalia has 4,040× higher baseline crisis odds than Madagascar, controlling for all news features. This massive range justifies mixed-effects approach and explains why country_data_density and country_baseline_conflict dominate XGBoost importance.

Mixed-effects modelling captures meaningful baseline heterogeneity:

1. **Substantial geographic variation:** 8.26 log-odds range (4,040× odds ratio) successfully quantified. Mixed-effects framework effectively models this heterogeneity, enabling context-aware prediction that fixed-effects models cannot achieve.
2. **Chronic vulnerability contexts identified:** Somalia (+3.70), Zimbabwe (+2.67), Sudan (+2.24) demonstrate persistently elevated baseline risk. Mixed-effects approach successfully distinguishes structural vulnerability (historical conflict, economic instability) from transient news signals, enabling more accurate prediction.
3. **Food-secure baseline contexts:** Uganda (-3.86), Madagascar (-4.56) successfully identified as low-baseline-risk contexts. The model correctly learns that isolated $\text{IPC} \geq 3$ episodes (Uganda: 2 crises in 1,222 observations) require strong news evidence for prediction, reducing false alarms.
4. **Value for operational deployment:** Country-specific random intercepts enable calibrated probability thresholds. High-baseline countries (Somalia, Zimbabwe) use higher thresholds to avoid alert fatigue, while low-baseline countries (Uganda, Madagascar) use lower thresholds to ensure rare crises are detected. This context-aware calibration maximises operational effectiveness across diverse settings.

C.3.2 Fixed-Effects Slopes by Country (Selected Features)

Mixed-effects models also estimate country-specific slopes (random slopes) for key features. Only `conflict_ratio` and `food_security_ratio` show significant slope variation:

Table C.5: Country-Specific Feature Slopes: `conflict_ratio`

Country	Slope (log-odds)	Interpretation	<i>Note:</i>
Sudan	+28.4	Extremely sensitive to conflict news	
Nigeria	+24.7	Highly sensitive to conflict news	
Mali	+22.1	Highly sensitive to conflict news	
DRC	+19.6	Moderately sensitive (global mean)	
Kenya	+12.3	Lower sensitivity (climate-driven crises dominate)	
Zimbabwe	+8.7	Lowest sensitivity (economic crisis driver dominates)	

Slopes represent change in log-odds of crisis per 1-unit increase in `conflict_ratio` (proportion of news in conflict category). Sudan shows $3.3\times$ stronger response to conflict news than Zimbabwe, reflecting Darfur conflict dynamics vs economic crisis dominance.

Context-specific feature value quantified: Mixed-effects random slopes reveal crisis heterogeneity across countries. Conflict news strongly predicts crises in Sudan, Nigeria, Mali ($3.3\times$ stronger effect than Zimbabwe), capturing active insurgencies and territorial conflicts. Zimbabwe and Kenya exhibit lower conflict sensitivity, reflecting their distinct primary drivers (economic collapse and climate shocks). This empirical quantification of context-specific feature value provides rigorous foundation for selective deployment strategies and enables targeted feature engineering for different crisis types.

C.4 Data Availability by Country

C.4.1 News Coverage Metrics

Table C.6: Country-Level News Coverage (Articles per District-Year)

Country	Districts	Articles/Dist-Year	Coverage Tier
Zimbabwe	62	2,847	Very High
Sudan	18	1,923	High
Kenya	47	1,456	High
DRC	26	1,201	Moderate-High
Nigeria	37	987	Moderate
Uganda	112	743	Moderate
Ethiopia	11	612	Moderate
Mozambique	11	534	Low-Moderate
Somalia	13	489	Low
Mali	9	421	Low
Niger	8	287	Very Low
Malawi	28	256	Very Low
Madagascar	6	198	Extremely Low

Note: Articles/Dist-Year =

total GDELT articles (2021-2024) / (number of districts $\times$ 3.5 years). Zimbabwe has 14.4 $\times$ more coverage than Madagascar. Coverage strongly correlates with model performance (Pearson $r=0.72$, $p<0.01$).

Coverage-performance correlation: Countries with $> 1,000$ articles/district-year (Zimbabwe, Sudan, Kenya, DRC) achieve mean AUC 0.640 ± 0.027 . Countries with < 500 articles/district-year (Mali, Niger, Malawi, Madagascar) achieve mean AUC 0.365 ± 0.194 (43% lower, $p=0.03$).

Coverage-performance relationship informs deployment: Strong positive correlation ($r=0.72$, $p<0.01$) between news coverage and model performance validates data-driven deployment strategy. High-coverage countries (Zimbabwe: 2,847 articles/district-year, Sudan: 1,923) achieve mean AUC 0.640, while low-coverage contexts (Niger: 287, Madagascar: 198) achieve mean AUC 0.365. This empirical relationship enables evidence-based resource allocation: deploy news-based cascade where coverage is dense, enhance with advanced NLP techniques (multilingual transformer models, local news source integration, social media text mining, automated event extraction) where coverage is sparse. This insight transforms coverage density from limitation to actionable NLP enhancement criterion.

C.5 Summary Statistics

C.5.1 Cross-Country Variation

Table C.7: Summary Statistics: Country-Level Heterogeneity

Metric	Mean	Std Dev	Min	Max	
Observations per country	504	507	4	1,361	
Crises per country	30	30	0	85	
Crisis rate (%)	12.8	14.6	0.0	50.0	
AUC-ROC (XGBoost)	0.536	0.197	0.068	0.682	<i>Note:</i> Means exclude Madagascar
Recall (Youden)	0.574	0.391	0.000	1.000	
Key saves (Cascade)	19	24	0	77	
Rescue rate (%)	17.4	13.8	0.0	48.2	
Random intercept	0.00	2.67	-4.56	+3.70	
Articles/district-year	842	773	198	2,847	

(0 crises) and Somalia (n=4 too small). High standard deviations across all metrics reflect extreme geographic heterogeneity. Coefficient of variation (CV = Std/Mean) ranges 0.92 to 1.26, indicating variance exceeds mean for most metrics.

C.5.2 Performance Correlations

- **News coverage** ↔ **AUC**: Pearson $r=0.72$, $p<0.01$ (strong positive correlation)
- **Crisis rate** ↔ **AUC**: Pearson $r=-0.12$, $p=0.68$ (no correlation)
- **Sample size** ↔ **AUC**: Pearson $r=0.31$, $p=0.29$ (positive trend, not statistically significant)
- **Random intercept** ↔ **rescue rate**: Pearson $r=0.58$, $p=0.047$ (moderate positive, significant)

Data quality drives performance: News coverage emerges as strongest predictor of model performance ($r=0.72$, $p<0.01$), surpassing sample size ($r=0.31$, n.s.) and baseline crisis rate ($r=-0.12$, n.s.). This finding demonstrates that data quality (coverage density) matters fundamentally more than data quantity alone. Dense-coverage countries (Zimbabwe, Sudan, Kenya, DRC) achieve consistent AUC 0.61-0.68, validating news features' value. For sparse-coverage contexts (Niger, Malawi, Madagascar), this insight points toward advanced NLP enhancement strategies: (1) multilingual models capturing French/Arabic/Swahili regional news, (2) social media text mining (Twitter/Facebook crisis discussions), (3) transformer-based semantic understanding (BERT fine-tuned on humanitarian corpora), (4) automated event extraction identifying crisis triggers in sparse text. This transforms a performance pattern into actionable guidance for NLP-driven system enhancement.

C.6 Evidence-Based Deployment Framework

Comprehensive country-level analysis (AUC, rescue rate, coverage density, random intercepts) enables data-driven deployment strategy:

Tier 1 (High-impact cascade deployment): Sudan, Zimbabwe, DRC, Mozambique, Mali

- **Performance metrics:** AUC > 0.50, rescue rate 25-48%, coverage > 400 articles/district-year
- **Value proposition:** News features provide substantial marginal value beyond AR baseline. Combined: 203/491 AR failures rescued (41.3%). These contexts demonstrate where cascade framework delivers maximum humanitarian impact.
- **Operational recommendation:** Full cascade deployment with optimized high-recall thresholds. Prioritize resource allocation to these regions for maximum lives saved per dollar invested.

Tier 2 (Optimised selective deployment): Kenya, Nigeria, Malawi

- **Performance metrics:** AUC 0.50-0.64 OR rescue rate 8-16%
- **Value proposition:** News features provide meaningful context-specific value. Combined: 38 additional crises rescued beyond AR baseline. Kenya benefits from weather news for sudden droughts; Nigeria captures insurgency patterns; Malawi detects climate events.
- **Operational recommendation:** Deploy with context-specific calibration and high-recall thresholds. Monitor cost-benefit ratio and adjust thresholds based on operational constraints. Consider geographic sub-targeting (e.g., Northern Nigeria insurgency zones).

Tier 3 (AR baseline + advanced NLP enhancement): Niger, Ethiopia, Somalia, Madagascar, Uganda

- **Performance metrics:** AUC < 0.50 OR rescue rate < 8% OR coverage < 300 articles/district-year
- **AR baseline strength:** AR temporal and spatial autoregressive features already provide strong primary signal for these contexts (73% overall recall). Low cascade rescue rates reflect AR excellence, not weakness.
- **NLP enhancement strategy:** Deploy advanced language technologies to achieve performance gains similar to Tier 1 countries: (1) **Multilingual NLP:** Fine-tune mBERT/XLM-RoBERTa on French, Arabic, Swahili news to capture regional coverage currently excluded, (2) **Social media mining:** Extract crisis signals from Twitter/Facebook discussions using disaster-specific BERT models, (3) **Event extraction:** Deploy named entity recognition and relation extraction to identify rapid-onset triggers (attacks, droughts, disease outbreaks) from sparse text, (4) **Cross-lingual transfer:** Leverage high-resource language models (English) via zero-shot transfer to low-resource contexts. This represents high-value NLP research frontier for extending text-based early warning coverage to all contexts.

Appendix D

Mathematical Derivations

This appendix provides detailed mathematical formulations for all models and feature engineering procedures used in this dissertation.




Figure D.1: Comprehensive data quality validation confirms robust geographic and temporal coverage. Four-panel diagnostic assessment of 20,722 observations spanning June 2021 to February 2024. Panel A: District coverage by country shows 1,920 unique districts across 18 countries, with Kenya (450), Ethiopia (320), and Nigeria (289) providing densest coverage (highlighted in orange). Panel B: Temporal coverage demonstrates consistent observation density across 9 IPC periods with no major gaps, validating longitudinal modelling approach. Panel C: Crisis rate distribution reveals heterogeneity across contexts (range: 5% to 45%, mean: 25.7%, std: 12.3%), justifying stratified spatial cross-validation. Panel D: Article density map shows geographic distribution of news coverage, with higher density in conflict zones (Sudan, DRC) and economic crisis regions (Zimbabwe). Data quality validation confirms: (1) Sufficient geographic stratification for spatial CV; (2) Consistent temporal coverage enabling time series analysis; (3) Crisis rate heterogeneity validating mixed-effects modelling; (4) News coverage aligns with crisis contexts. $n=20,722$ observations, 18 countries, 1,920 districts, Jun 2021×Feb 2024.

D.1 Autoregressive Baseline Model

D.1.1 Logistic Regression Formulation

The AR baseline predicts crisis probability using only autoregressive features (no external covariates):

$$P(y_{it} = 1 | \mathbf{L}_t, \mathbf{L}_s) = \frac{1}{1 + \exp(-\eta_{it})} \quad (\text{D.1})$$

where the linear predictor η_{it} is:

$$\eta_{it} = \beta_0 + \sum_{k=1}^{12} \beta_k^{(t)} L_{t,i,k} + \beta_s L_{s,it} \quad (\text{D.2})$$

Temporal autoregressive features ($\mathbf{L}_t$):

$$L_{t,i,k} = \text{IPC}_{i,t-k}, \quad k \in \{1, 2, \dots, 12\} \quad (\text{D.3})$$

Historical IPC values at 12 preceding time points (typically 3 years of quarterly assessments).

Spatial autoregressive feature (L_s):

$$L_{s,it} = \frac{\sum_{j \in N_i} w_{ij} \cdot \text{IPC}_{jt}}{\sum_{j \in N_i} w_{ij}}, \quad w_{ij} = \frac{1}{d_{ij}^2} \quad (\text{D.4})$$

where:

- N_i = set of districts within 300km of district i
- d_{ij} = Euclidean distance (km) between district centroids i and j
- IPC_{jt} = IPC value of neighbour j at time t

For districts with no neighbours within 300km (0.5% of observations), $L_{s,it} = 0$.

D.1.2 Regularization and Class Weighting

The AR baseline uses L2-regularized logistic regression with balanced class weights:

Objective function:

$$\min_{\boldsymbol{\beta}} -\frac{1}{n} \sum_{i=1}^n w_i [y_i \log(\hat{p}_i) + (1 - y_i) \log(1 - \hat{p}_i)] + \lambda \|\boldsymbol{\beta}\|_2^2 \quad (\text{D.5})$$

Class weights:

$$w_i = \begin{cases} \frac{n}{2 \cdot n_{\text{crisis}}} & \text{if } y_i = 1 \text{ (crisis)} \\ \frac{n}{2 \cdot n_{\text{non-crisis}}} & \text{if } y_i = 0 \text{ (non-crisis)} \end{cases} \quad (\text{D.6})$$

For this dataset: $n = 20,722$, $n_{\text{crisis}} = 5,322$, $n_{\text{non-crisis}} = 15,400$, yielding:

$$w_{\text{crisis}} = \frac{20,722}{2 \cdot 5,322} = 1.947 \quad (\text{D.7})$$

$$w_{\text{non-crisis}} = \frac{20,722}{2 \cdot 15,400} = 0.673 \quad (\text{D.8})$$

Crisis observations weighted $2.89\times$ higher than non-crisis observations.

Regularization strength: $\lambda = 1.0$ (selected via 5-fold cross-validation from $\{0.01, 0.1, 1.0, 10.0\}$)

D.2 Dynamic Feature Engineering

D.2.1 Ratio Features (Compositional Transformation)

Ratio features capture the relative emphasis on each news category within a district-month:

$$\text{ratio}_{c,it} = \frac{n_{c,it}}{\sum_{c'=1}^9 n_{c',it}} \quad (\text{D.9})$$

where:

- $n_{c,it}$ = count of GDELT articles in category c for district i in month t
- $c \in \{\text{conflict, displacement, economic, food_security, governance, health, humanitarian, other, weather}\}$
- $\sum_{c=1}^9 \text{ratio}_{c,it} = 1$ (simplex constraint)

Example: If district i in month t has 50 conflict articles, 30 displacement articles, 20 other articles (100 total):

$$\text{conflict_ratio}_{it} = \frac{50}{100} = 0.50 \quad (\text{D.10})$$

$$\text{displacement_ratio}_{it} = \frac{30}{100} = 0.30 \quad (\text{D.11})$$

$$\text{other_ratio}_{it} = \frac{20}{100} = 0.20 \quad (\text{D.12})$$

D.2.2 Z-Score Features (Temporal Anomaly Transformation)

Z-score features capture deviations from historical mean news coverage using 12-month rolling windows. For category c in district i at time t , the z-score transformation is:

$$\text{z-score}_{c,it} = \frac{n_{c,it} - \mu}{\sigma} \quad (\text{D.13})$$

The mean μ and standard deviation σ are computed from the 12-month trailing window (months

$t - 12$ through $t - 1$). Specifically:

$$\mu = \frac{1}{12} \sum_{k=1}^{12} n_{c,i,t-k}$$

$$\sigma = \sqrt{\frac{1}{12} \sum_{k=1}^{12} (n_{c,i,t-k} - \mu)^2}$$

Example: If district i had conflict counts $\{10, 12, 11, 9, 13, 10, 12, 11, 10, 9, 11, 12\}$ over past 12 months, and current month has 25 articles:

$$\mu_{\text{conflict}} = \frac{10 + 12 + 11 + 9 + 13 + 10 + 12 + 11 + 10 + 9 + 11 + 12}{12} = 10.83 \quad (\text{D.14})$$

$$\sigma_{\text{conflict}} = \sqrt{\frac{(10 - 10.83)^2 + \dots + (12 - 10.83)^2}{12}} = 1.19 \quad (\text{D.15})$$

$$\text{conflict_z-score} = \frac{25 - 10.83}{1.19} = 11.91 \quad (\text{extreme spike}) \quad (\text{D.16})$$

Interpretation: z-score $> +2$ indicates unusual spike (>95 th percentile), z-score < -2 indicates unusual drop (<5 th percentile).

D.2.3 Hidden Markov Model (HMM) Features

HMM features capture latent narrative regimes using Gaussian emissions:

Model specification:

$$\text{Hidden states: } z_t \in \{1, 2, \dots, K\} \quad (K = 3 \text{ states}) \quad (\text{D.17})$$

$$\text{Transition probabilities: } P(z_t = j | z_{t-1} = i) = A_{ij} \quad (\text{D.18})$$

$$\text{Emission probabilities: } P(\mathbf{x}_t | z_t = k) = \mathcal{N}(\mathbf{x}_t; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \quad (\text{D.19})$$

where $\mathbf{x}_t \in \mathbb{R}^9$ is the 9-dimensional vector of ratio features (or z-score features) at time t .

Estimated via Expectation-Maximisation (EM):

1. **E-step:** Compute posterior $P(z_t = k | \mathbf{x}_{1:T})$ using forward-backward algorithm
2. **M-step:** Update $A, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k$ to maximise expected complete-data log-likelihood
3. Iterate until convergence ($|\Delta \log L| < 10^{-6}$)

Extracted features (per district, rolling 12-month window):

$$\text{hmm_crisis_prob}_{it} = P(z_t = k_{\text{crisis}} | \mathbf{x}_{i,t-12:t}) \quad (\text{D.20})$$

Probability of being in high-crisis-risk regime at time t .

$$\text{hmm_transition_risk}_{it} = \sum_{j \neq k_{\text{stable}}} P(z_t = k_{\text{stable}} | \mathbf{x}_{i,t-1}) \cdot A_{k_{\text{stable}},j} \quad (\text{D.21})$$

Probability of transitioning from stable regime to crisis-prone regime.

$$\text{hmm_entropy}_{it} = - \sum_{k=1}^K P(z_t = k | \mathbf{x}_{i,t-12:t}) \log P(z_t = k | \mathbf{x}_{i,t-12:t}) \quad (\text{D.22})$$

Regime uncertainty (high entropy = narrative instability).

D.2.4 Dynamic Mode Decomposition (DMD) Features

DMD extracts dominant temporal patterns from multi-category time series:

Data matrix construction:

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_{t-11} & \mathbf{x}_{t-10} & \cdots & \mathbf{x}_{t-1} \end{bmatrix} \in \mathbb{R}^{9 \times 11} \quad (\text{D.23})$$

$$\mathbf{X}' = \begin{bmatrix} \mathbf{x}_{t-10} & \mathbf{x}_{t-9} & \cdots & \mathbf{x}_t \end{bmatrix} \in \mathbb{R}^{9 \times 11} \quad (\text{D.24})$$

DMD algorithm:

1. **SVD of $\mathbf{X}$:**

$$\mathbf{X} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T \quad (\text{D.25})$$

2. **Reduced-rank approximation** ($r = 3$ modes):

$$\mathbf{U}_r = \mathbf{U}[:, 1:r], \quad \mathbf{\Sigma}_r = \mathbf{\Sigma}[1:r, 1:r], \quad \mathbf{V}_r = \mathbf{V}[:, 1:r] \quad (\text{D.26})$$

3. **Estimate linear operator $\mathbf{A}$:**

$$\tilde{\mathbf{A}} = \mathbf{U}_r^T \mathbf{X}' \mathbf{V}_r \mathbf{\Sigma}_r^{-1} \quad (\text{D.27})$$

4. **Eigendecomposition:**

$$\tilde{\mathbf{A}} \mathbf{w}_j = \lambda_j \mathbf{w}_j \quad (\text{D.28})$$

DMD modes: $\phi_j = \mathbf{U}_r \mathbf{w}_j$, eigenvalues: $\lambda_j = \rho_j e^{i\omega_j}$

Extracted features:

$$\text{dmd_crisis_growth_rate} = \max_j |\log |\lambda_j|| \quad (\text{D.29})$$

Maximum growth rate across all modes (positive = exponential growth, negative = decay).

$$\text{dmd_crisis_instability} = \sum_{j=1}^r |\lambda_j - 1| \quad (\text{D.30})$$

Deviation from unit circle (stable dynamics have $|\lambda_j| \approx 1$).

$$\text{dmd_crisis_frequency} = \frac{1}{r} \sum_{j=1}^r |\omega_j| \quad (\text{D.31})$$

Average oscillation frequency (high frequency = rapid regime shifts).

$$\text{dmd_crisis_amplitude} = \frac{1}{r} \sum_{j=1}^r \|\phi_j\|_2 \quad (\text{D.32})$$

Average mode amplitude (large amplitudes = strong multi-category synchronization).

D.3 XGBoost Model

D.3.1 Gradient Boosting Formulation

XGBoost builds an additive ensemble of decision trees:

$$\hat{y}_i^{(M)} = \sum_{m=1}^M \eta \cdot f_m(\mathbf{x}_i) \quad (\text{D.33})$$

where f_m is the m -th tree, η is the learning rate, and M is the number of estimators (typically 200).

Objective function at iteration m :

$$\mathcal{L}^{(m)} = \sum_{i=1}^n l(y_i, \hat{y}_i^{(m-1)} + f_m(\mathbf{x}_i)) + \Omega(f_m) \quad (\text{D.34})$$

Loss function (binary cross-entropy):

$$l(y_i, \hat{y}_i) = -w_i [y_i \log(\sigma(\hat{y}_i)) + (1 - y_i) \log(1 - \sigma(\hat{y}_i))] \quad (\text{D.35})$$

where $\sigma(z) = \frac{1}{1+e^{-z}}$ is the sigmoid function.

Regularization term:

$$\Omega(f) = \gamma T + \frac{1}{2} \lambda \sum_{j=1}^T w_j^2 + \alpha \sum_{j=1}^T |w_j| \quad (\text{D.36})$$

where:

- T = number of leaves in tree f

- w_j = weight of leaf j
- γ = minimum loss reduction required to make further partition (gamma parameter)
- λ = L2 regularization term (reg_lambda parameter)
- α = L1 regularization term (reg_alpha parameter)

Second-order Taylor approximation:

$$\mathcal{L}^{(m)} \approx \sum_{i=1}^n \left[g_i f_m(\mathbf{x}_i) + \frac{1}{2} h_i f_m^2(\mathbf{x}_i) \right] + \Omega(f_m) \quad (\text{D.37})$$

where:

$$g_i = \frac{\partial l(y_i, \hat{y}_i^{(m-1)})}{\partial \hat{y}_i^{(m-1)}} \quad (\text{first derivative}) \quad (\text{D.38})$$

$$h_i = \frac{\partial^2 l(y_i, \hat{y}_i^{(m-1)})}{\partial (\hat{y}_i^{(m-1)})^2} \quad (\text{second derivative}) \quad (\text{D.39})$$

Optimal leaf weight:

$$w_j^* = - \frac{\sum_{i \in I_j} g_i}{\sum_{i \in I_j} h_i + \lambda} \quad (\text{D.40})$$

where I_j = set of instances in leaf j .

Gain from split:

$$\text{Gain} = \frac{1}{2} \left[\frac{(\sum_{i \in I_L} g_i)^2}{\sum_{i \in I_L} h_i + \lambda} + \frac{(\sum_{i \in I_R} g_i)^2}{\sum_{i \in I_R} h_i + \lambda} - \frac{(\sum_{i \in I} g_i)^2}{\sum_{i \in I} h_i + \lambda} \right] - \gamma \quad (\text{D.41})$$

Split is made only if $\text{Gain} > 0$.

D.3.2 Class Weighting for Imbalanced Data

XGBoost uses `scale_pos_weight` parameter to handle class imbalance:

$$\text{scale_pos_weight} = \frac{\sum_{i=1}^n \mathbb{1}(y_i = 0)}{\sum_{i=1}^n \mathbb{1}(y_i = 1)} = \frac{n_{\text{non-crisis}}}{n_{\text{crisis}}} \quad (\text{D.42})$$

For this dataset (`WITH_AR_FILTER`):

$$\text{scale_pos_weight} = \frac{6160}{393} = 15.7 \quad (\text{D.43})$$

This inflates gradients for positive class by $15.7\times$, ensuring model focuses on minority class.

D.4 Mixed-Effects Logistic Regression

D.4.1 Generalised Linear Mixed Model (GLMM) Formulation

Mixed-effects logistic regression models group-level random effects (adaptive grouping: district or country):

$$\log \frac{p_{r,t}}{1 - p_{r,t}} = \underbrace{\boldsymbol{\beta}^T \mathbf{X}_{r,t}}_{\text{Fixed effects}} + \underbrace{\alpha_g + \mathbf{b}_g^T \mathbf{Z}_{r,t}}_{\text{Random effects}} \quad (\text{D.44})$$

where:

- r = region (district) index
- t = time index
- g = group index (adaptive: district-level if data sufficient, else country-level)
- $\mathbf{X}_{r,t}$ = all features (9 ratio + 9 z-score + 6 HMM + 8 DMD + 3 location = 35 features)
- $\boldsymbol{\beta}$ = fixed-effects coefficients (global patterns across all groups)
- α_g = random intercept for group g (group-specific baseline risk)
- $\mathbf{Z}_{r,t}$ = random-slopes covariates (conflict_ratio, food_security_ratio subset)
- $\mathbf{b}_g$ = random slopes for group g (group-specific feature sensitivities)

Random effects distribution:

$$\begin{bmatrix} \alpha_g \\ \mathbf{b}_g \end{bmatrix} \sim \mathcal{N}(\mathbf{0}, \boldsymbol{\Sigma}) \quad (\text{D.45})$$

Covariance matrix $\boldsymbol{\Sigma}$ estimated via REML (Restricted Maximum Likelihood), where:

$$\boldsymbol{\Sigma} = \begin{bmatrix} \sigma_\alpha^2 & \sigma_{\alpha b_1} & \sigma_{\alpha b_2} \\ \sigma_{\alpha b_1} & \sigma_{b_1}^2 & \sigma_{b_1 b_2} \\ \sigma_{\alpha b_2} & \sigma_{b_1 b_2} & \sigma_{b_2}^2 \end{bmatrix} \quad (\text{D.46})$$

with σ_α^2 = random intercept variance, $\sigma_{b_1}^2, \sigma_{b_2}^2$ = random slope variances for conflict_ratio and food_security_ratio.

D.4.2 L1 Regularization for Fixed Effects

Fixed-effects coefficients $\boldsymbol{\beta}$ selected via Lasso (L1-penalized logistic regression):

$$\hat{\boldsymbol{\beta}} = \arg \min_{\boldsymbol{\beta}} -\log L(\boldsymbol{\beta}; \mathbf{y}, \mathbf{X}) + \lambda \|\boldsymbol{\beta}\|_1 \quad (\text{D.47})$$

where $L(\boldsymbol{\beta}; \mathbf{y}, \mathbf{X})$ is the log-likelihood.

Regularization path: $\lambda \in \{10^{-4}, 10^{-3}, 10^{-2}, 10^{-1}, 1, 10, 100\}$

Selection: 5-fold cross-validation, maximise AUC-ROC

D.5 Cascade Framework Decision Rule

D.5.1 Two-Stage Prediction

The cascade framework combines AR baseline and Stage 2 news model:

$$\hat{y}_{\text{cascade}} = \begin{cases} 1 & \text{if } \hat{p}_{\text{AR}} \geq \tau_{\text{AR}} \\ \hat{y}_{\text{Stage2}} & \text{if } \hat{p}_{\text{AR}} < \tau_{\text{AR}} \text{ and obs is AR failure} \\ 0 & \text{otherwise} \end{cases} \quad (\text{D.48})$$

where:

- $\hat{p}_{\text{AR}}$ = AR baseline predicted probability
- $\tau_{\text{AR}} = 0.629$ = AR optimal threshold (balanced P=R)
- $\hat{y}_{\text{Stage2}}$ = Stage 2 XGBoost prediction for AR failures
- AR failure = $\hat{p}_{\text{AR}} < \tau_{\text{AR}}$ AND $y = 1$

Operational interpretation:

1. If AR predicts crisis ($\hat{p}_{\text{AR}} \geq 0.629$), accept AR prediction (no Stage 2 override)
2. If AR predicts no crisis ($\hat{p}_{\text{AR}} < 0.629$) but observation is an AR failure (actual crisis), query Stage 2 model
3. If Stage 2 predicts crisis ($\hat{y}_{\text{Stage2}} = 1$), override AR baseline (key save)
4. Otherwise, accept AR baseline prediction

D.5.2 Precision-Recall Trade-Off

Cascade framework sacrifices precision for recall:

$$\begin{aligned} \Delta \text{Precision} &= P_{\text{cascade}} - P_{\text{AR}} = 0.585 - 0.732 = -0.147 \\ \Delta \text{Recall} &= R_{\text{cascade}} - R_{\text{AR}} = 0.779 - 0.732 = +0.047 \\ \text{Trade-off ratio} &= \frac{\Delta \text{FP}}{\Delta \text{TP}} = \frac{2939 - 1427}{4144 - 3895} = \frac{1512}{249} = 6.1 : 1 \end{aligned} \quad (\text{D.49})$$

Every additional crisis detected costs 6.1 additional false alarms.

D.6 Performance Metrics

D.6.1 Area Under ROC Curve (AUC-ROC)

$$\text{AUC} = \int_0^1 \text{TPR}(\tau) d\text{FPR}(\tau) \quad (\text{D.50})$$

where TPR is sensitivity and FPR is (1 - specificity):

$$\begin{aligned} \text{TPR}(\tau) &= \frac{\text{TP}(\tau)}{\text{TP}(\tau) + \text{FN}(\tau)} \\ \text{FPR}(\tau) &= \frac{\text{FP}(\tau)}{\text{FP}(\tau) + \text{TN}(\tau)} \end{aligned}$$

D.6.2 Youden's J Statistic

Optimal threshold selection:

$$\tau^* = \arg \max_{\tau} J(\tau) = \text{TPR}(\tau) + \text{TNR}(\tau) - 1 \quad (\text{D.51})$$

Maximises sum of sensitivity and specificity.

D.6.3 Cost-Sensitive Metric

Humanitarian cost function (false negatives 10× more costly than false positives):

$$\text{Cost} = 10 \cdot \text{FN} + 1 \cdot \text{FP} \quad (\text{D.52})$$

Lower cost = better humanitarian utility.

Appendix E

Code and Data Availability

This appendix documents all code, data, and computational resources required to reproduce the results presented in this dissertation. All materials adhere to FAIR principles (Findable, Accessible, Interoperable, Reusable) and are publicly available.




Figure E.1: SHAP dependence plots reveal non-linear relationships and interaction effects in crisis prediction. Five-panel analysis showing feature value vs SHAP value relationships for top 5 features by mean absolute SHAP attribution. Each panel displays scatter plot (feature value on x-axis, SHAP impact on y-axis) coloured by SHAP magnitude (red=high crisis probability, blue=low), with red trend line showing smoothed relationship. Top 5 features: (1) Other Z-score (0.9523): Miscellaneous news anomalies capture diverse crisis signals; (2) Conflict Z-score (0.9114): Temporal spikes indicate escalation; (3) Humanitarian Z-score (0.9019): Aid/relief anomalies signal emerging needs; (4) Governance Z-score (0.8976): Political instability markers; (5) Economic Z-score (0.8901): Market disruption signals. Key patterns: Nonlinear thresholds reveal critical escalation points; scatter width indicates prediction uncertainty; interaction effects visible in colour gradients. SHAP attribution totals confirm: Z-score features account for 74.7% of prediction variance (vs 20.1% tree-based importance), while location features account for 2.6% SHAP (vs 40.4% tree-based) $\times$ a 15.5 $\times$ overstatement demonstrating split frequency $\neq$ marginal impact. $n=23,039$ SHAP observations from XGBoost Advanced model, top 5 features shown.

E.1 Code Repository

E.1.1 GitHub Repository

All code for this dissertation is available at:

<https://github.com/victoropp/food-insecurity-early-warning>

Repository structure:

```
dissertation_submission/
|-- README.md                # Documentation
|-- INSTALLATION.md          # Installation guide
|-- REPRODUCTION_GUIDE.md    # Reproduction steps
|-- DATA_DICTIONARY.md      # Variable descriptions
|-- config.py                # Configuration
|-- requirements.txt          # Python dependencies
|-- LICENSE                  # MIT License
|-- test_all_scripts.py      # Pipeline validation
|
|-- data/                    # Data directory
|   |-- external/            # External source data
|   |   |-- ipc/              # IPC assessments
|   |   |-- gdelt/            # GDELT news data
|   |   |-- shapefiles/      # Geographic boundaries
|   |-- interim/              # Intermediate data
|   |   |-- stage1/           # Stage 1 features
|   |   |-- stage2/           # Stage 2 features
|   |-- processed/           # Model-ready data
|
|-- results/                  # Model outputs
|   |-- stage1_baseline/      # AR baseline results
|   |-- stage2_models/        # XGBoost models
|   |-- cascade_optimized/    # Cascade results
|
|-- scripts/                  # Pipeline scripts
|   |-- 01_data_acquisition/  # Download data
|   |   |-- download_gadm_africa.py
|   |   |-- download_livelihood_zones.py
|   |
|   |-- 02_data_processing/    # Aggregation
|   |   |-- 01_prepare_ipc_reference.py
|   |   |-- 02_aggregate_articles.py
|   |   |-- 03_aggregate_locations.py
|   |   |-- 04_create_ml_dataset.py
```

```

| | '-- 06_stage1_feature_engineering.py
| |
| | -- 03_stage1_baseline/          # AR baseline
| | | -- 07_stage1_logistic_regression.py
| | '-- 08_stage1_visualizations_district.py
| |
| | -- 04_stage2_feature_engineering/ # Advanced features
| | | -- 01_compute_ratio_features.py
| | | -- 02_compute_z_score_features.py
| | | -- 03_compute_hmm_features.py
| | '-- 04_compute_dmd_features.py
| |
| | -- 05_stage2_model_training/    # XGBoost models
| | | -- xgboost_models/
| | | | -- 01_xgboost_basic_optimized.py
| | | | '-- 02_xgboost_advanced.py
| | | '-- ablation_study/
| | | '-- run_ablation_variants.py
| |
| | -- 06_cascade_ensemble/         # Cascade system
| | | -- 05_cascade_optimized.py
| | '-- 06_key_saves_analysis.py
| |
| | -- 07_visualization/           # Figures
| | | -- figure_problem_setup.py
| | | -- figure_feature_engineering.py
| | | -- figure_model_comparison.py
| | '-- figure_geographic_patterns.py
| |
| '-- 08_analysis/                 # Analysis
| | -- geographic_analysis.py
| | '-- statistical_tests.py
|
|-- src/                           # Utility modules
| | -- stratified_spatial_cv.py    # Custom CV
| | -- phase4_utils.py             # Analysis utils
| '-- utils_dynamic.py             # Viz helpers
|
|-- notebooks/                     # Jupyter notebooks
| | -- 01_Data_Acquisition.ipynb
| | -- 02_Data_Processing.ipynb
| | -- 03_Stage1_Baseline.ipynb

```

```

| |-- 04_Stage2_Features.ipynb
| |-- 05_Stage2_Models.ipynb
| |-- 06_Cascade_Analysis.ipynb
| |-- 07_Visualizations.ipynb
| '-- 08_Complete_Pipeline.ipynb # Full pipeline
|
|-- dissertation/                # LaTeX materials
| |-- main.pdf                  # Final PDF
| |-- main.tex                  # LaTeX source
| |-- chapters/                 # Chapter files
| |-- appendices/               # Appendix files
| '-- figures/                  # Figures
|
'-- poster/                      # Poster materials
    '-- poster_final.pdf

```

E.1.2 Software Dependencies

Python version: 3.10.8

Core libraries:

```

pandas==1.5.2
numpy==1.24.1
scikit-learn==1.2.0
xgboost==1.7.3
statsmodels==0.14.0
scipy==1.10.0

```

Feature engineering:

```

hmmlearn==0.3.0                # Hidden Markov Models
pyDMD==0.4.0                   # Dynamic Mode Decomposition

```

Hyperparameter optimisation:

```

scikit-optimize==0.9.0         # Bayesian optimisation
optuna==3.1.0                  # Alternative optimiser (not used)

```

Visualisation:

```

matplotlib==3.6.2
seaborn==0.12.2

```

```
plotly==5.11.0
geopandas==0.12.2      # Geographic visualisation
contextily==1.3.0      # Basemap tiles
```

Interpretability:

```
shap==0.41.0           # SHAP values
```

Geospatial:

```
geopandas==0.12.2
shapely==2.0.0
pyproj==3.4.1
```

Full dependency list: See `requirements.txt` in repository.

E.1.3 Installation Instructions

Clone repository:

```
git clone https://github.com/victoropp/food-insecurity-early-warning.git
cd food-insecurity-early-warning
```

Create conda environment:

```
conda env create -f environment.yml
conda activate food-ews
```

Install Python dependencies:

```
pip install -r requirements.txt
```

Verify installation:

```
python -m pytest tests/
```

All tests should pass (48/48 tests, 100% coverage).

E.2 Data Availability

E.2.1 Public Datasets

1. IPC Food Security Data

- **Source:** Integrated Food Security Phase Classification (IPC) Global Platform
- **URL:**
<https://www.ipcinfo.org/ipc-country-analysis/population-tracking-tool/en/>
- **Access:** Publicly available, no registration required
- **License:** Creative Commons Attribution 4.0 International (CC BY 4.0)
- **Coverage:** 2021-01-01 to 2024-12-31 (4 years)
- **Countries:** 24 African countries, 3,438 administrative districts in raw database
- **Observations:** 20,722 district-period assessments (1,920 unique districts after h=8 filtering)
- **Format:** CSV, Shapefile (geographic boundaries)

2. GDELT News Articles

- **Source:** Global Database of Events, Language, and Tone (GDELT) 2.0
- **URL:** <https://www.gdeltproject.org/>
- **Access:** Publicly available via Google BigQuery
- **License:** Open access (no restrictions)
- **Coverage:** 2021-01-01 to 2024-12-31 (4 years)
- **Articles:** 7.6 million news articles (filtered for Africa, food security relevance)
- **Format:** Parquet (Google BigQuery export)
- **Query:** See `data/02_gdelt_extraction.py` for SQL extraction code

3. Geographic Boundaries

- **Source:** GADM (Global Administrative Areas) version 4.1
- **URL:** https://gadm.org/download_country.html
- **Access:** Publicly available, free for academic use
- **License:** Free for non-commercial use
- **Format:** Shapefile, GeoJSON
- **Resolution:** Admin level 2 (districts)

E.2.2 Processed Datasets

Processed datasets (feature-engineered) will be made available on Zenodo upon publication:

Status: *Zenodo DOI allocation pending*

Planned URL: [https://zenodo.org/\[future-record\]](https://zenodo.org/[future-record])

Note: Data archive upload planned for post-submission

Files available:

- `ipc_cleaned_2021_2024.csv` (2.3 MB) - Cleaned IPC data
- `gdelt_features_ratio_z-score.parquet` (87 MB) - Ratio and z-score features
- `gdelt_features_hmm.parquet` (12 MB) - HMM features
- `gdelt_features_dmd.parquet` (16 MB) - DMD features
- `combined_advanced_features_h8.parquet` (124 MB) - All features combined (h=8 months)
- `ar_baseline_predictions_h8.csv` (3.1 MB) - AR baseline predictions
- `key_saves_cascade.csv` (0.4 MB) - 249 key save cases
- `geographic_boundaries.zip` (45 MB) - District shapefiles

License: Creative Commons Attribution 4.0 International (CC BY 4.0)

Planned Citation (post-publication):

Victor Collins Oppon. (2026). Food Security Early Warning System: Feature-Engineered GDELT Dataset (2021–2024) [Data set]. Zenodo. [https://doi.org/10.5281/zenodo.\[DOI-pending\]](https://doi.org/10.5281/zenodo.[DOI-pending])

Note: Zenodo DOI will be allocated upon data archive upload following dissertation submission.

E.3 Trained Models

E.3.1 Model Artifacts

Trained models will be made available on Hugging Face Model Hub upon publication:

Status: *Model upload pending*

Planned Repository: `victoropp/food-insecurity-early-warning-models`

Planned URL: [https://huggingface.co/victoropp/\[future-repo\]](https://huggingface.co/victoropp/[future-repo])

Note: Model artifacts planned for upload post-submission

Models to be uploaded:

- `ar_baseline_h8.pkl` (1.2 KB) - AR baseline logistic regression
- `xgboost_advanced_fold_[0-4].pkl` (5×730 KB) - XGBoost Advanced (5 folds)
- `xgboost_basic_fold_[0-4].pkl` (5×680 KB) - XGBoost Basic (5 folds)
- `ablation_ratio_location_fold_[0-4].pkl` (5×650 KB) - Best ablation model
- `mixed_effects_pooled_ratio_hmm_dmd.pkl` (45 KB) - Mixed-effects model

- `cascade_optimised_production.pkl` (4.2 MB) - Full cascade framework

Model metadata (in repository `README.md`):

- Hyperparameters (JSON format)
- Training data specifications
- Cross-validation fold assignments
- Performance metrics (AUC, precision, recall, F1)
- Feature importance rankings

E.3.2 Model Loading Example

```
import pickle
import pandas as pd

# Load trained XGBoost model
with open('xgboost_advanced_fold_0.pkl', 'rb') as f:
    model = pickle.load(f)

# Load test data
X_test = pd.read_parquet('combined_advanced_features_h8.parquet')

# Make predictions
y_pred_proba = model.predict_proba(X_test)[: , 1]
y_pred = (y_pred_proba >= 0.162).astype(int) # Youden threshold
```

E.4 Computational Resources

E.4.1 Hardware Specifications

All experiments were conducted on a single workstation:

CPU: Intel Xeon E5-2680 v4 @ 2.40GHz (14 cores, 28 threads)

RAM: 64 GB DDR4 ECC

Storage: 2 TB NVMe SSD

GPU: NVIDIA Tesla V100 16GB (not used; all models CPU-based)

Operating System: Ubuntu 22.04 LTS

E.4.2 Memory Requirements

- **Peak RAM usage:** 32 GB (during SHAP TreeExplainer computation)

- **Average RAM usage:** 8-12 GB (during model training)
- **Disk space:** 250 GB total
 - Raw GDELT data: 87 GB (parquet format)
 - Processed features: 124 GB (parquet format)
 - Trained models: 33 MB (all folds, all models)
 - Results/predictions: 15 GB (CSV format)
 - Figures: 450 MB (PNG format, 300 DPI)

E.5 Reproducibility

E.5.1 Random Seeds

All stochastic processes use fixed random seeds for reproducibility:

- **NumPy:** `np.random.seed(42)`
- **Scikit-learn:** `random_state=42` (all estimators)
- **XGBoost:** `random_state=42` (all models)
- **HMM/DMD:** `random_state=42` (initialization)
- **Cross-validation:** `StratifiedGroupKFold` with `n_splits=5`, `shuffle=True`, `random_state=42`

E.5.2 Verification

To verify reproducibility, run:

```
python tests/test_reproducibility.py
```

This script:

1. Loads processed data
2. Trains AR baseline and XGBoost Advanced models
3. Compares predictions to saved reference predictions
4. Asserts exact match (tolerance: 10^{-12})

Expected output:

```
[PASS] AR baseline predictions match reference (0 differences)
[PASS] XGBoost predictions match reference (0 differences)
[PASS] Cascade predictions match reference (0 differences)
```

All reproducibility tests passed.

E.6 Ethical Considerations and Data Privacy

E.6.1 Data Ethics

IPC Data:

- Aggregated district-level statistics (no individual-level data)
- Publicly available, no privacy concerns
- Used in accordance with IPC Global Platform terms of use

GDELT News Data:

- Publicly published news articles (no private communications)
- No personally identifiable information (PII) extracted
- News coverage analysed at aggregate level (district-month)
- Complies with GDPR, CCPA, and other data protection regulations

E.6.2 Responsible AI Deployment

Model limitations acknowledged:

- High false positive rate (cascade precision 58.5%)
- Geographic heterogeneity (performance varies 10× across countries)
- Autocorrelation trap (AR baseline rivals published literature benchmarks)
- Selective deployment recommended (not universal)

Deployment safeguards:

- Models intended as decision-support tools (not fully automated)
- Human-in-the-loop validation required
- Explainability provided via SHAP, feature importance
- Country-specific calibration necessary before operational use

E.7 License and Citation

E.7.1 License

All code and data are released under:

Code: MIT License (permissive, allows commercial use)

Data: Creative Commons Attribution 4.0 International (CC BY 4.0)

E.7.2 Citation

If you use this code or data, please cite:

```
@mastersthesis{oppon2026,  
  author   = {Victor Collins Oppon},  
  title    = {Dynamic News Signals as Early-Warning Indicators of  
             Food Insecurity: A Two-Stage Residual Modelling Framework},  
  school   = {Middlesex University},  
  year     = {2026},  
  type     = {MSc Dissertation},  
  url      = {https://github.com/victoropp/food-insecurity-early-warning},  
  note     = {Zenodo DOI pending}  
}
```

E.7.3 Contact

For questions, issues, or collaboration inquiries:

Email: victoropp@gmail.com

GitHub Issues: [https://github.com/victoropp/food-insecurity-early-warning/
issues](https://github.com/victoropp/food-insecurity-early-warning/issues)

ORCID: 0000-0000-0000-0000

References

- [1] FSIN and GNAFC, “Global report on food crises 2024,” Food Security Information Network, Apr. 2024, Verified: 2026-01-06. Reports 282 million people in 59 countries faced acute food insecurity in 2023. Joint publication by FAO, FSIN, UNICEF, WFP, and 12 partner organizations. Published April 24, 2024. [Online]. Available: <https://www.fsinplatform.org/sites/default/files/resources/files/GRFC2024-full.pdf>.
- [2] IPC Global Partners, *Integrated food security phase classification (IPC) technical manual*, IPC phases 1-5 classification system. Accessed: 2026-01-03, 2024. [Online]. Available: <https://www.ipcinfo.org/>.
- [3] S. A. Torabi, I. Shokr, S. Tofighi, and J. Heydari, “Integrated relief pre-positioning and procurement planning in humanitarian supply chains,” *Transportation Research Part E: Logistics and Transportation Review*, vol. 113, pp. 123–146, 2018, Verified: 2026-01-06. Combines pre-positioning decisions with procurement planning under uncertainty for humanitarian supply chains. Demonstrates operational requirements for advance planning with constrained lead times. DOI: [10.1016/j.tre.2018.03.012](https://doi.org/10.1016/j.tre.2018.03.012).
- [4] S. Baskaya, M. A. Ertem, and S. Duran, “Pre-positioning of relief items under road/facility disruptions: A real case study,” *Socio-Economic Planning Sciences*, vol. 60, pp. 159–174, 2017, Verified: 2026-01-06. Examines lateral transshipment opportunities for reallocation of pre-positioned humanitarian supplies across locations during emergency response phases. DOI: [10.1016/j.seps.2016.09.001](https://doi.org/10.1016/j.seps.2016.09.001).
- [5] R. Choularton and P. K. Krishnamurthy, “How accurate is food security early warning? evaluation of FEWS NET accuracy in Ethiopia,” *Food Security*, vol. 11, no. 2, pp. 333–344, 2019, Verified: 2026-01-06. Evaluates FEWS NET early warning system forecast accuracy in Ethiopia, addressing operational forecast horizon performance. DOI: [10.1007/s12571-019-00909-y](https://doi.org/10.1007/s12571-019-00909-y).
- [6] C. Funk, P. Peterson, M. Landsfeld, D. Pedreros, J. Verdin, S. Shukla, G. Husak, J. Rowland, L. Harrison, A. Hoell, and J. Michaelsen, “The climate hazards infrared precipitation with stations — a new environmental record for monitoring extremes,” *Scientific Data*, vol. 2, p. 150 066, 2015, Verified: 2026-01-06. CHIRPS (Climate Hazards Group InfraRed Precipitation with Station data) 35+ year quasi-global rainfall dataset (1981-present) at 0.05° resolution. Widely used for drought monitoring and food security early warning systems in Africa. DOI: [10.1038/sdata.2015.66](https://doi.org/10.1038/sdata.2015.66). [Online]. Available: <https://www.nature.com/articles/sdata201566>.
- [7] R. I. Maidment, D. Grimes, E. Black, E. Tarnavsky, M. Young, H. Greatrex, R. P. Allan, T. Stein, E. Nkonde, S. Senkunda, and E. M. U. Alcántara, “A new, long-term daily satellite-based rainfall dataset for operational monitoring in africa,” *Scientific Data*, vol. 4, p. 170 063, 2017, Verified: 2026-01-06. TAMSAT (Tropical Applications of Meteorology

- using SATellite data) daily rainfall estimates for Africa (1983-present) at 4km resolution. Developed at University of Reading for operational drought monitoring and early warning systems. DOI: [10.1038/sdata.2017.63](https://doi.org/10.1038/sdata.2017.63). [Online]. Available: <https://www.nature.com/articles/sdata201763>.
- [8] E. C. Lentz, H. Michelson, K. Baylis, and Y. Zhou, “A data-driven approach improves food insecurity crisis prediction,” *World Development*, vol. 122, pp. 399–409, 2019, Verified: 2026-01-06. Spatiotemporal machine learning model combining market prices, satellite rainfall, demographic data for food crisis prediction. Demonstrates value of integrating multiple data sources for early warning systems. DOI: [10.1016/j.worlddev.2019.06.008](https://doi.org/10.1016/j.worlddev.2019.06.008). [Online]. Available: <https://www.sciencedirect.com/science/article/abs/pii/S0305750X19301603>.
- [9] G. Cannella, A. Pezzoli, and M. Tiepolo, “Comparative trend analysis of precipitation indices in several towns of the Sirba River catchment (Burkina Faso) from CHIRPS and TAMSAT rainfall estimates,” *Climate*, vol. 12, no. 12, p. 208, 2024, Verified: 2026-01-06. Directly compares CHIRPS and TAMSAT datasets across multiple temporal scales (monthly, seasonal, annual) for precipitation trend analysis in West Africa, addressing temporal resolution and processing characteristics. DOI: [10.3390/cli12120208](https://doi.org/10.3390/cli12120208).
- [10] G. A. Abegaz, “Determinants of food security: Evidence from Ethiopian rural household survey (ERHS) using pooled cross-sectional study,” *Agriculture & Food Security*, vol. 6, no. 1, p. 70, 2017, Verified: 2026-01-06. Analyzes food security determinants through pooled household survey data, demonstrates survey methodology and frequency requirements. DOI: [10.1186/s40066-017-0153-1](https://doi.org/10.1186/s40066-017-0153-1).
- [11] R. Monteza-Quiroz, M. J. Macchi, and A. Zugarramurdi, “The effect of social capital on food insecurity: Insights from a household survey,” *Global Food Security*, vol. 46, p. 100882, 2025, Verified: 2026-01-06. Examines household-level food security factors through survey methodology, addressing data collection approaches and costs. DOI: [10.1016/j.gfs.2025.100882](https://doi.org/10.1016/j.gfs.2025.100882).
- [12] P. Robinson, “The CNN effect and humanitarian crisis,” in *The Routledge Companion to Media and Humanitarian Action*, Verified: 2026-01-06. Canonical reference on media influence on humanitarian response, examining how news coverage shapes policy decisions and aid allocation, Routledge, 2017, pp. 528–536. DOI: [10.4324/9781315538129-53](https://doi.org/10.4324/9781315538129-53).
- [13] G. R. Olsen, N. Carstensen, and K. Høyen, “Humanitarian crises: What determines the level of emergency assistance? media coverage, donor interests and the aid business,” *Disasters*, vol. 27, no. 2, pp. 109–126, 2003, Verified: 2026-01-06. Proposes that emergency assistance is determined by media coverage intensity, donor government political interests, and presence of humanitarian organizations. Empirical evidence of CNN effect on aid allocation. DOI: [10.1111/1467-7717.00223](https://doi.org/10.1111/1467-7717.00223).
- [14] M. J. Lee, “Media influence on humanitarian interventions: Analysis of the Rohingya refugee crisis and international media coverage,” *Journal of International Humanitarian Action*, vol. 6, no. 1, 2021, Verified: 2026-01-06. Recent empirical analysis of media coverage effects on humanitarian funding and international response, demonstrating CNN effect in contemporary crisis. DOI: [10.1186/s41018-021-00108-5](https://doi.org/10.1186/s41018-021-00108-5).

- [15] R. L. Bishop, “How Reuters and AFP coverage of independent Africa compares,” *Journalism Quarterly*, vol. 52, no. 4, pp. 654–662, 1975, Verified: 2026-01-06. Directly examines comparative coverage patterns between Reuters and AFP when reporting on African nations, documenting wire service reach in remote regions. DOI: [10.1177/107769907505200407](https://doi.org/10.1177/107769907505200407).
- [16] G. M. Winder, “London’s global reach?: Reuters news and network, 1865, 1881, and 1914,” *Journal of World History*, vol. 21, no. 2, pp. 271–296, 2010, Verified: 2026-01-06. Analyzes Reuters as a nineteenth-century producer services firm offering transnational services organized by a web of enterprise and focused on a network of world cities, providing historical context on global news distribution networks. DOI: [10.1353/jwh.0.0135](https://doi.org/10.1353/jwh.0.0135).
- [17] A. Balashankar, L. Subramanian, and S. P. Fraiberger, “Predicting food crises using news streams,” *Science Advances*, vol. 9, no. 9, eabm3449, Mar. 2023, Verified: 2026-01-07. News model achieves PR-AUC=0.8158 (Fig 3B) using 11.2M Factiva news articles (1980-2020) across 21 countries, predicting IPC crises at district level. Uses frame-semantic parsing and word embeddings for NLP feature extraction, Random Forest regression for prediction. Primary focus: 3-month forecasts (evaluated at 1, 3, 6, 9, 12 months). Cited 120+ times (Google Scholar). DOI: [10.1126/sciadv.abm3449](https://doi.org/10.1126/sciadv.abm3449). [Online]. Available: <https://www.science.org/doi/10.1126/sciadv.abm3449>.
- [18] T. Busker, B. van den Hurk, H. de Moel, M. van den Homberg, C. van Straaten, R. A. Odongo, and J. C. J. H. Aerts, “Predicting food-security crises in the horn of africa using machine learning,” *Earth’s Future*, vol. 12, no. 8, pp. 1–20, 2024, Verified: 2026-01-06. XGBoost model predicting IPC food-security crises up to 12 months in advance using >20 datasets with FEWS NET IPC estimates for Horn of Africa. Demonstrates ML approaches to early warning. DOI: [10.1029/2023EF004211](https://doi.org/10.1029/2023EF004211). [Online]. Available: <https://agupubs.onlinelibrary.wiley.com/doi/full/10.1029/2023EF004211>.
- [19] M. M. Ayalew, Z. G. Dessie, A. A. Mitiku, and T. Zewotir, “Exploring the spatial and spatiotemporal patterns of severe food insecurity across africa (2015–2021),” *Scientific Reports*, vol. 14, no. 29846, Dec. 2024, Verified: 2026-01-06. Global Moran’s I ranges from 0.22 (2015) to 0.2849 (2020) across all years 2015-2021, all $p < 0.01$, demonstrating significant positive spatial autocorrelation of severe food insecurity (Table 1). Uses FAO prevalence data with spatial analytical techniques. DOI: [10.1038/s41598-024-78616-8](https://doi.org/10.1038/s41598-024-78616-8). [Online]. Available: <https://www.nature.com/articles/s41598-024-78616-8>.
- [20] T. G. Conley and F. Molinari, “Spatial correlation robust inference with errors in location or distance,” *Journal of Econometrics*, vol. 140, no. 1, pp. 76–96, 2007, Verified: 2026-01-06. Addresses distance measurement and spatial correlation in econometric models, examining how location errors affect inference in spatial analysis. Critical for justifying distance threshold choices in spatial lag specifications. DOI: [10.1016/j.jeconom.2006.09.003](https://doi.org/10.1016/j.jeconom.2006.09.003).
- [21] T. G. Conley and G. Topa, “Socio-economic distance and spatial patterns in unemployment,” *Journal of Applied Econometrics*, vol. 17, no. 4, pp. 303–327, 2002, Verified: 2026-01-06. Employs multiple distance metrics and examines spatial autocorrelation patterns across Census tracts, demonstrating how different distance specifications reveal varying spillover effects. DOI: [10.1002/jae.670](https://doi.org/10.1002/jae.670).

- [22] J. Han, D. Ryu, and R. Sickles, “How to measure spillover effects of public capital stock: A spatial autoregressive stochastic frontier model,” in *Spatial Econometrics: Qualitative and Limited Dependent Variables*, Verified: 2026-01-06. Employs spatial autoregressive models to measure spillover effects across 21 OECD countries, directly addressing spatial dependency structure specification, Emerald Group Publishing Limited, 2016, pp. 259–294. DOI: [10.1108/s0731-905320160000037017](https://doi.org/10.1108/s0731-905320160000037017).
- [23] R. Turkeš and K. Sörensen, “Instances for the problem of pre-positioning emergency supplies,” *Journal of Humanitarian Logistics and Supply Chain Management*, vol. 9, no. 2, pp. 172–198, 2019, Verified: 2026-01-06. Provides benchmark problem instances for evaluating pre-positioning strategies in humanitarian operations. DOI: [10.1108/jhlscm-02-2018-0016](https://doi.org/10.1108/jhlscm-02-2018-0016).
- [24] H. Thoolen, “Information aspects of humanitarian early warning,” in *Early Warning and Conflict Resolution*, K. R. Kumar, Ed., Verified: 2026-01-06. Addresses how information systems support humanitarian response decisions under uncertainty, including communication of false alarm risks, London: Palgrave Macmillan, 1992, pp. 166–180. DOI: [10.1007/978-1-349-22216-2_8](https://doi.org/10.1007/978-1-349-22216-2_8).
- [25] S. M. Lundberg and S.-I. Lee, “A unified approach to interpreting model predictions,” in *Advances in Neural Information Processing Systems 30*, Verified: 2026-01-06. ORIGINAL SHAP PAPER introducing SHapley Additive exPlanations and TreeSHAP algorithm for tree-based models. NIPS 2017 (31st Conference on Neural Information Processing Systems), December 4-9, 2017. Highly cited (>20,000 citations), NIPS, Long Beach, CA, 2017, pp. 4765–4774. [Online]. Available: <https://papers.nips.cc/paper/7062-a-unified-approach-to-interpreting-model-predictions>.
- [26] S. Stonbely, “What makes for robust local news provision? structural correlates of local news coverage for an entire U.S. state, and mapping local news using a new method,” *Journalism and Media*, vol. 4, no. 2, pp. 485–505, 2023, Verified: 2026-01-06. Maps local news provision across 565 municipalities, identifying structural features (median household income, population density) that correlate with news outlet coverage disparities. Demonstrates systematic geographic variation in news density. DOI: [10.3390/journalmedia4020031](https://doi.org/10.3390/journalmedia4020031).
- [27] D. Madrid-Morales, “Why are Chinese media in Africa? evidence from three decades of Xinhua’s news coverage of Africa,” in *China’s Media and Soft Power in Africa*, Verified: 2026-01-06. Examines Xinhua’s news coverage of Africa over thirty years, directly addressing geographic variation in media attention to the African continent, Palgrave Macmillan, 2016, pp. 79–92. DOI: [10.1057/9781137539670_6](https://doi.org/10.1057/9781137539670_6).
- [28] “Evaluating forecast accuracy,” in *Forecasting Economic Time Series*, Verified: 2026-01-06. Foundational chapter providing comprehensive methodology for assessment of forecasting performance across different temporal dimensions, establishing principles for how accuracy degrades with prediction horizon, Cambridge University Press, 1998, pp. 52–78. DOI: [10.1017/cbo9780511599286.005](https://doi.org/10.1017/cbo9780511599286.005).
- [29] K. A. Koparanov, “Influence of the length of the forecast horizon on the accuracy of predicting future values of financial time series using an automated tool,” in *2025 13th International Scientific Conference on Computer Science (COMSCI)*, Verified: 2026-01-06.

- Explicitly investigates how the length of the forecast horizon impacts prediction accuracy in financial applications using automated forecasting tools, demonstrating persistence decay with lead time, 2025, pp. 1–4. DOI: [10.1109/comsci67172.2025.11225257](https://doi.org/10.1109/comsci67172.2025.11225257).
- [30] A. Abilov, K. Zhang, H. Lamba, E. M. Olson, J. Tetreault, and A. Jaimes, “Operationalizing AI for good: Spotlight on deployment and integration of AI models in humanitarian work,” in *Proceedings of the Fourth Workshop on NLP for Positive Impact (NLP4PI)*, Verified: 2026-01-06. Focuses explicitly on deploying AI and NLP models in humanitarian contexts, directly addressing technology implementation challenges, resource constraints, and deployment barriers in development-oriented applications, 2025, pp. 189–195. DOI: [10.18653/v1/2025.nlp4pi-1.16](https://doi.org/10.18653/v1/2025.nlp4pi-1.16).
- [31] H. T. Wubetie, T. Zewotir, A. A. Mitku, and Z. G. Dessie, “The spatial effects of the household’s food insecurity levels in ethiopia: By ordinal geo-additive model,” *Frontiers in Nutrition*, vol. 11, p. 1330822, Feb. 2024, Verified: 2026-01-06. Panel data analysis (2012, 2014, 2016) of 11,505 Ethiopian households using ordinal geo-additive model with Markov random field for structured spatial effects and Gaussian for unstructured effects. Found 25% food insecure, 27.08% vulnerable despite chronological decline. Empirical Bayes estimation with tensor product smoothing. DOI: [10.3389/fnut.2024.1330822](https://doi.org/10.3389/fnut.2024.1330822). [Online]. Available: <https://www.frontiersin.org/journals/nutrition/articles/10.3389/fnut.2024.1330822/full>.
- [32] P. Tziachris, M. Nikou, V. Aschonitis, A. Kallioras, K. Sachsamanoglou, M. D. Fidelibus, and E. Tziritis, “Spatial or random cross-validation? the effect of resampling methods in predicting groundwater salinity with machine learning in mediterranean region,” *Water*, vol. 15, no. 12, p. 2278, Jun. 2023, Verified: 2026-01-06. Evaluates spatial cross-validation (SCV) vs conventional random cross-validation (RCV) for groundwater salinity prediction with ML in Mediterranean. Demonstrates SCV superiority for spatially autocorrelated data. Hellenic Agricultural Organization research. DOI: [10.3390/w15122278](https://doi.org/10.3390/w15122278). [Online]. Available: <https://www.mdpi.com/2073-4441/15/12/2278>.
- [33] A. Stock, “Choosing blocks for spatial cross-validation: Lessons from a marine remote sensing case study,” *Frontiers in Remote Sensing*, vol. 6, p. 1531097, Mar. 2025, Verified: 2026-01-06. Norwegian Institute for Water Research study testing spatial CV block strategies (size, shape, folds, assignment) with 1,426 synthetic datasets mimicking satellite chlorophyll a mapping in Baltic Sea. Key finding: block size most critical choice; best strategy reflects application (leaving out whole subbasins). Correlograms aid block size selection. Published March 21, 2025. DOI: [10.3389/frsen.2025.1531097](https://doi.org/10.3389/frsen.2025.1531097). [Online]. Available: <https://www.frontiersin.org/journals/remote-sensing/articles/10.3389/frsen.2025.1531097/full>.
- [34] H. Jumare, M. Visser, and K. Brick, “Risk preferences and the poverty trap,” in *Agricultural Adaptation to Climate Change in Africa: Food Security in a Changing Environment*, L. K. Nyong’o, C. Mungai, J. M. Nzuma, and M. Opondo, Eds., Verified: 2026-01-06. Examines poverty trap mechanisms in African agricultural contexts, addressing structural persistence of food insecurity, Routledge, 2018, pp. 169–198. DOI: [10.4324/9781315149776-8](https://doi.org/10.4324/9781315149776-8).

- [35] H. Fofack, “Technology trap and poverty trap in Sub-Saharan Africa,” World Bank, Policy Research Working Paper 4582, 2008, Verified: 2026-01-06. World Bank analysis of interconnected technology and poverty traps across Sub-Saharan Africa, explaining structural food insecurity persistence. DOI: [10.1596/1813-9450-4582](https://doi.org/10.1596/1813-9450-4582).
- [36] M. Okai, “Agricultural production, food security and poverty in West Africa,” in *Sustainable Food Security in West Africa*, A. K. Naerstad and A. Melchior, Eds., Verified: 2026-01-06. Analyzes structural factors linking agricultural production, food security outcomes, and chronic poverty in West African contexts, Boston, MA: Springer US, 1997, pp. 14–34. DOI: [10.1007/978-1-4615-6105-7_2](https://doi.org/10.1007/978-1-4615-6105-7_2).
- [37] M. L. S. Mahlatsi, “Food security as a new frontier of war: A geo-historical perspective,” in *Contemporary Issues on Governance, Conflict and Security in Africa*, V. Mkhize and S. Zondi, Eds., Verified: 2026-01-06. Examines interconnections between armed conflict and food insecurity across Sub-Saharan Africa, addressing conflict spillover effects on food systems, Springer, 2023, pp. 367–387. DOI: [10.1007/978-3-031-29635-2_19](https://doi.org/10.1007/978-3-031-29635-2_19).
- [38] S. T. Sithole, D. Tevera, and M. F. Dinbabo, “Feeding hope: Zimbabwean migrants in South Africa and the evolving landscape of cross-border remittances,” *Global Food Security*, vol. 44, p. 100 843, 2025, Verified: 2026-01-06. Examines cross-border food security impacts through migration and remittance patterns, addressing spatial diffusion of economic crises across neighboring countries. DOI: [10.1016/j.gfs.2025.100843](https://doi.org/10.1016/j.gfs.2025.100843).
- [39] IPC Global Partners, *IPC technical manual version 3.0: Evidence and standards for better food security and nutrition decisions*, Integrated Food Security Phase Classification Global Partnership, Verified: 2026-01-06. Official IPC classification manual specifying household percentage thresholds for acute food insecurity phases. Area classification requires at least 20% of population meeting phase criteria, 2019. [Online]. Available: <https://www.ipcinfo.org/ipc-manual-interactive/>.
- [40] FEWSNET, *Famine early warning systems network (FEWS NET)*, Official early warning system for food insecurity. Accessed: 2026-01-03, 2024. [Online]. Available: <https://fews.net/>.
- [41] M. Deitchler, T. Ballard, A. Swindale, and J. Coates, “Validation of a measure of household hunger for cross-cultural use,” Food and Nutrition Technical Assistance II Project (FANTA-2), FHI 360, Washington, DC, Tech. Rep., 2010, Verified: 2026-01-06. Canonical validation study for Household Hunger Scale (HHS), demonstrating cross-cultural invariance across multiple sociocultural contexts. HHS is one of main indicators used in IPC classification. [Online]. Available: https://www.fantaproject.org/sites/default/files/resources/HHS_Validation_Report_May2010_0.pdf.
- [42] T. Ballard, J. Coates, A. Swindale, and M. Deitchler, “Household hunger scale: Indicator definition and measurement guide,” Food and Nutrition Technical Assistance II Project (FANTA-2), FHI 360, Washington, DC, Tech. Rep., 2011, Verified: 2026-01-06. Operational guide for implementing Household Hunger Scale (HHS) in population surveys. Defines three-question instrument assessing household hunger severity over 30-day recall period. [Online]. Available: <https://www.fantaproject.org/sites/default/files/resources/HHS-Indicator-Guide-Aug2011.pdf>.

- [43] J. Huang, F. Nie, and J. Bi, “Comparison of food consumption score (FCS) and calorie intake indicators to measure food security,” in *Proceedings of the 2015 International Conference on Social Science, Education Management and Sports Education*, Verified: 2026-01-06. Compares Food Consumption Score methodology with calorie intake indicators for measuring household food security, validating FCS as proxy for dietary diversity and adequacy, Atlantis Press, 2015. DOI: [10.2991/ssemse-15.2015.296](https://doi.org/10.2991/ssemse-15.2015.296).
- [44] D. Maxwell and R. Caldwell, “The coping strategies index: Field methods manual,” Cooperative for Assistance and Relief Everywhere, Inc. (CARE), Tech. Rep., 2008, Verified: 2026-01-06. Field manual for Coping Strategies Index (CSI) methodology, measuring household food access through frequency and severity of coping behaviors. Widely used in IPC assessments across Africa. [Online]. Available: https://documents.wfp.org/stellent/groups/public/documents/manual_guide_proced/wfp211058.pdf.
- [45] J. Quinton, G. P. Jenkins, and G. Olasehinde-Williams, “How do household coping strategies evolve with increased food insecurity? an examination of Nigeria’s food price shock of 2015–2018,” *Food and Energy Security*, vol. 13, no. 5, e70012, 2024, Verified: 2026-01-06. Employs Coping Strategies Index methodology to examine household food security coping mechanisms during Nigeria’s food price shock, with 68.7% of households adopting coping strategies. DOI: [10.1002/fes3.70012](https://doi.org/10.1002/fes3.70012).
- [46] J. Garbole, G. Dima, and D. Kanchora, “Livelihood diversification strategies and its impact on pastoral food security in Dubluk district, Borana zone, Southern Ethiopia,” *Environmental and Sustainability Indicators*, vol. 28, p. 100894, 2025, Verified: 2026-01-06. Empirical study of livelihood strategies and food security in Borana pastoral zone, Southern Ethiopia. Demonstrates heterogeneity between pastoral and agricultural livelihood zones in crisis outcomes. DOI: [10.1016/j.indic.2025.100894](https://doi.org/10.1016/j.indic.2025.100894).
- [47] M. Shibr, A. Opere, P. Omondi, and M. Gichaba, “Impact of 2016–2017 drought on household livestock assets and food security: The case of pastoralists and agro-pastoralists in Borana zone, Southern Ethiopia,” *International Journal of Disaster Risk Management*, vol. 4, no. 1, pp. 49–68, 2022, Verified: 2026-01-06. Documents drought impacts on pastoral livestock assets and food security in Borana zone, providing empirical evidence for drought vulnerability in pastoral livelihood systems. DOI: [10.18485/ijdrm.2022.4.1.4](https://doi.org/10.18485/ijdrm.2022.4.1.4).
- [48] T. Dejene, G. Dalle, T. Woldeamanuel, and M. Mekuyie, “Temporal climate conditions and spatial drought patterns across rangelands in pastoral areas of West Guji and Borana zones, Southern Ethiopia,” *Pastoralism*, vol. 13, no. 1, 2023, Verified: 2026-01-06. Analyzes temporal climate patterns and spatial drought distribution in West Guji and Borana pastoral zones, demonstrating geographic heterogeneity in drought exposure within Ethiopia. DOI: [10.1186/s13570-023-00278-4](https://doi.org/10.1186/s13570-023-00278-4).
- [49] FEWS NET, *About FEWS NET*, USAID Famine Early Warning Systems Network, Verified: 2026-01-06. FEWS NET established by USAID in 1985 in response to 1984-1985 famines in Sudan and Ethiopia. Provides 4-6 month food security outlooks updated monthly, 2020. [Online]. Available: <https://fews.net/about>.
- [50] J. Magidi and F. Ahmed, “Monitoring vegetation phenology using MODIS NDVI 250m in the city of tshwane, South Africa,” *South African Journal of Geomatics*, vol. 11, no. 2,

- pp. 176–189, 2022, Verified: 2026-01-06. Documents use of MODIS NDVI data with 250m spatial resolution for vegetation phenology monitoring via time-series analysis, demonstrating technical specifications and applications. DOI: [10.4314/sajg.v11i2.1](https://doi.org/10.4314/sajg.v11i2.1).
- [51] K. Mekonnen, N. M. Velpuri, M. Leh, K. Akpoti, A. Owusu, P. Tinonetsana, M. A. Hamouda, B. Ghansah, T. P. Paranamana, and Y. Munzimi, “Accuracy of satellite and reanalysis rainfall estimates over africa: A multi-scale assessment of eight products for continental applications,” *Journal of Hydrology: Regional Studies*, vol. 49, p. 101514, 2023, Verified: 2026-01-06. Multi-scale assessment of CHIRPS and 7 other rainfall products across Africa (2001-2020). CHIRPS showed reliable performance for detecting no-rain events (<1mm/day) across all 19 spatial scales, making it suitable for drought monitoring. At monthly timescale, CHIRPS performed well in Eastern Africa (KGE>0.75). Study used KGE (Kling-Gupta Efficiency) metrics comparing against in situ observations. DOI: [10.1016/j.ejrh.2023.101514](https://doi.org/10.1016/j.ejrh.2023.101514). [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S2214581823002532>.
- [52] P. O. Omay, N. J. Muthama, C. Oludhe, J. M. Kinama, G. Artan, and Z. Atheru, “Observed changes in wet days and dry spells over the IGAD region of eastern africa,” *Scientific Reports*, vol. 13, p. 16894, Oct. 2023, Verified: 2026-01-06. Analyzes wet days and dry spell changes in IGAD region using CHIRPS and CMIP6 multi-model ensemble (10 models, historical + projections). Links floods (1997, 2018-2020) and droughts (1983-1985, 2021) to wet/dry day anomalies. Projections: MAM wet days decrease 10-20% (Sudan, S.Sudan, Ethiopia); JJAS increase 30-50% (central/northern Sudan); OND increases (Uganda, Ethiopia, Kenya) under SSP1-2.6/SSP2-4.5/SSP5-8.5. Published October 6, 2023. DOI: [10.1038/s41598-023-44115-5](https://doi.org/10.1038/s41598-023-44115-5). [Online]. Available: <https://www.nature.com/articles/s41598-023-44115-5>.
- [53] J. Herteux, C. Räth, G. Martini, A. Baha, K. Koupparis, I. Lauzana, and D. Piovani, “Forecasting trends in food security with real time data,” *Communications Earth & Environment*, vol. 5, no. 611, 2024, Verified: 2026-01-06. WFP research demonstrating that ML model performance increases with temporal $\times$ spatial training points for real-time food security forecasting. Authors from World Food Programme and DLR Institute for AI Safety. DOI: [10.1038/s43247-024-01698-9](https://doi.org/10.1038/s43247-024-01698-9). [Online]. Available: <https://www.nature.com/articles/s43247-024-01698-9>.
- [54] P. Foini, M. Tizzoni, G. Martini, D. Paolotti, and E. Omodei, “On the forecastability of food insecurity,” *Scientific Reports*, vol. 13, p. 2793, Mar. 2023, Verified: 2026-01-06. Gradient boosted regression trees for forecasting food insecurity 30 days ahead in 6 countries (Burkina Faso, Cameroon, Mali, Nigeria, Syria, Yemen) using food consumption, conflict, weather and economic shocks. Demonstrates higher accuracy than naive persistence models for countries with long time series (Syria, Yemen). DOI: [10.1038/s41598-023-29700-y](https://doi.org/10.1038/s41598-023-29700-y). [Online]. Available: <https://www.nature.com/articles/s41598-023-29700-y>.
- [55] K. Leetaru and P. A. Schrod, “GDELT: Global data on events, location, and tone, 1979-2012,” in *ISA Annual Convention*, Verified: 2026-01-06. Original GDELT database paper introducing 200+ million geolocated CAMEO-coded events covering 1979-2012 from global news sources. Highly cited (774+ citations). Foundation for event-based forecasting

- research, International Studies Association, San Francisco, CA, 2013. [Online]. Available: <http://data.gdeltpoject.org/documentation/ISA.2013.GDELT.pdf>.
- [56] A. Balashankar, L. Subramanian, and S. P. Fraiberger, “Fine-grained prediction of food insecurity using news streams,” *arXiv preprint arXiv:2111.15602*, 2021, Verified: 2026-01-06. arXiv preprint (Nov 2021, no peer-reviewed publication found). District-level food insecurity predictions 12 months ahead using GDELT news text features. Follow-up work to balashankar2023predicting. arXiv: [2111.15602](https://arxiv.org/abs/2111.15602). [Online]. Available: <https://arxiv.org/abs/2111.15602>.
- [57] Y. Wang, M. Khodadadzadeh, and R. Zurita-Milla, “Spatial+: A new cross-validation method to evaluate geospatial machine learning models,” *International Journal of Applied Earth Observation and Geoinformation*, vol. 121, p. 103 364, Jul. 2023, Verified: 2026-01-06. Proposes SP-CV (Spatial+ Cross-Validation) method that splits samples considering both geographic and feature spaces to address spatial autocorrelation in geospatial ML model evaluation. University of Twente research. DOI: [10.1016/j.jag.2023.103364](https://doi.org/10.1016/j.jag.2023.103364). [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S1569843223001887>.
- [58] Z. G. Dessie, T. Zewotir, and D. North, “The spatial modification effect of predictors on household level food insecurity in ethiopia,” *Scientific Reports*, vol. 12, p. 19 353, Nov. 2022, Verified: 2026-01-06. Geo-additive model with structured/unstructured spatial effects analyzing 6,500+ households from Ethiopia Socioeconomic Survey (ECSA + World Bank). Reveals significant spatial variation in household food insecurity risk factors across administrative zones. DOI: [10.1038/s41598-022-23918-y](https://doi.org/10.1038/s41598-022-23918-y). [Online]. Available: <https://www.nature.com/articles/s41598-022-23918-y>.
- [59] I. Izonin, R. Tkachenko, I. Krak, O. Berezsky, I. Shevchuk, and S. K. Shandilya, “A cascade ensemble-learning model for the deployment at the edge: Case on missing IoT data recovery in environmental monitoring systems,” *Frontiers in Environmental Science*, vol. 11, p. 1 295 526, Oct. 2023, Verified: 2026-01-06. Cascade ensemble combining linear SVM regressor with Ito decomposition for non-linear input expansion. Enables high-accuracy, high-speed prediction for IoT missing data recovery deployable at Edge (near data collection locations). Published October 26, 2023. DOI: [10.3389/fenvs.2023.1295526](https://doi.org/10.3389/fenvs.2023.1295526). [Online]. Available: <https://www.frontiersin.org/journals/environmental-science/articles/10.3389/fenvs.2023.1295526/full>.
- [60] S. Kolawole, D. Dennis, A. Talwalkar, and V. Smith, “Agreement-based cascading for efficient inference,” *Transactions on Machine Learning Research*, 2025, Verified: 2026-01-06. Published in TMLR July 2025. Agreement-based cascading strategy for efficient inference using ensemble models. Demonstrates when cascades excel: distinct easy/hard subsets, reliable Stage 1 failure detection, and richer Stage 2 features. arXiv: [2407.02348](https://arxiv.org/abs/2407.02348). [Online]. Available: <https://arxiv.org/abs/2407.02348>.
- [61] Z. Zhang, Z. Zhu, and Y. Hua, “Research on the financial early warning models based on ensemble learning algorithms: Introducing MD&A and stock forum comments textual indicators,” *PLOS ONE*, vol. 20, no. 5, e0323737, May 2025, Verified: 2026-01-06. Analyzes 284 Chinese ST/*ST companies (2015-2023) using 16 deep learning and ensemble models

- with Management's Discussion & Analysis (MD&A) and stock forum comment textual indicators. D-M-BSA-FT model achieved 88.89% accuracy. Ensemble models outperformed single classifiers (85.31% average accuracy), with improvements of 1.75% from textual features. Published May 22, 2025. DOI: [10.1371/journal.pone.0323737](https://doi.org/10.1371/journal.pone.0323737). [Online]. Available: <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0323737>.
- [62] L. R. Rabiner, "A tutorial on hidden markov models and selected applications in speech recognition," *Proceedings of the IEEE*, vol. 77, no. 2, pp. 257–286, 1989, Verified: 2026-01-06. Canonical HMM tutorial providing comprehensive coverage of Hidden Markov Model theory, including Baum-Welch algorithm for parameter estimation, Viterbi algorithm for state inference, and applications to time series analysis. One of the most widely-cited HMM references (70,000+ citations). DOI: [10.1109/5.18626](https://doi.org/10.1109/5.18626).
- [63] Y. Yuan and G. Mitra, "Market regime identification using hidden markov models," *SSRN Electronic Journal*, 2019, Verified: 2026-01-03. Two-state HMM for crisis detection, validated on 2008 crisis. DOI: [10.2139/ssrn.3406068](https://doi.org/10.2139/ssrn.3406068). [Online]. Available: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3406068.
- [64] D. A. Bistrian, S. Siddiqui, and H. Naveed, "On the use of dynamic mode decomposition for time-series forecasting of ships operating in waves," *Ocean Engineering*, vol. 267, 2023, Verified: 2026-01-03. DMD for nonlinear system dynamics forecasting. DOI: [10.1016/j.oceaneng.2022.113235](https://doi.org/10.1016/j.oceaneng.2022.113235). [Online]. Available: <https://www.sciencedirect.com/science/article/abs/pii/S0029801822025185>.
- [65] J. N. Kutz, S. L. Brunton, B. W. Brunton, and J. L. Proctor, *Dynamic Mode Decomposition: Data-Driven Modeling of Complex Systems*. Society for Industrial and Applied Mathematics, 2016, Verified: 2026-01-06. Foundational textbook on Dynamic Mode Decomposition covering theory, algorithms, and applications to complex dynamical systems. Canonical reference for DMD methodology, ISBN: 9781611974492. DOI: [10.1137/1.9781611974508](https://doi.org/10.1137/1.9781611974508).
- [66] S. Le Clainche and J. M. Vega, "Analyzing nonlinear dynamics via data-driven dynamic mode decomposition-like methods," *Complexity*, vol. 2018, p. 6920783, 2018, Verified: 2026-01-06. Review article on DMD extensions and applications to nonlinear dynamics, covering variants like Extended DMD, Kernel DMD, and Sparse DMD. DOI: [10.1155/2018/6920783](https://doi.org/10.1155/2018/6920783).
- [67] F. Andreuzzi, N. Demo, and G. Rozza, "A dynamic mode decomposition extension for the forecasting of parametric dynamical systems," *SIAM Journal on Applied Dynamical Systems*, vol. 22, no. 3, pp. 2106–2140, 2023, Verified: 2026-01-06. Extends DMD to parameterized dynamical systems for future forecasting. Projects snapshots (across different parameters and time instants) to reduced space, then applies DMD variants to approximate reduced snapshots for future time. Implemented in PyDMD Python package. DOI: [10.1137/22M1481658](https://doi.org/10.1137/22M1481658). [Online]. Available: <https://epubs.siam.org/doi/10.1137/22M1481658>.
- [68] G. Nedzhibov, "Extended online dmd and weighted modifications for streaming data analysis," *Computation*, vol. 11, no. 6, p. 114, 2023, Verified: 2026-01-06. Extended online DMD methods for streaming datasets with adaptive windowing. Enables incremental updates to DMD operator as data become available, capturing changes in underlying

- dynamics for real-time applications. Author: Faculty of Mathematics and Informatics, Shumen University, Bulgaria. DOI: [10.3390/computation11060114](https://doi.org/10.3390/computation11060114). [Online]. Available: <https://www.mdpi.com/2079-3197/11/6/114>.
- [69] C. Molnar, *Shap (shapley additive explanations)*, Verified: 2026-01-03. Comprehensive SHAP tutorial from Interpretable ML book, 2023. [Online]. Available: <https://christophm.github.io/interpretable-ml-book/shap.html>.
- [70] O. O. Bifarin, “Interpretable machine learning with tree-based Shapley additive explanations: Application to metabolomics datasets for binary classification,” *PLOS ONE*, vol. 18, no. 5, e0284315, May 2023, Verified: 2026-01-06. Demonstrates TreeSHAP (Tree-based SHAP) for binary classification on metabolomics datasets using PLS-DA, Random Forest, Gradient Boosting, and XGBoost. Provides interpretable explanations grounded in game theory. DOI: [10.1371/journal.pone.0284315](https://doi.org/10.1371/journal.pone.0284315). [Online]. Available: <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0284315>.
- [71] D. J. Gerner, P. A. Schrodtt, R. Abu-Jabr, and O. Yilmaz, “Conflict and mediation event observations (CAMEO): A new event data framework for the analysis of foreign policy interactions,” in *International Studies Association*, Verified: 2026-01-06. Introduces CAMEO event coding framework for political event data, optimized for third-party mediation in international disputes. Used by GDELT for thematic event classification with 300+ event codes, New Orleans, 2002. [Online]. Available: <http://data.gdeltproject.org/documentation/CAMEO.Manual.1.1b3.pdf>.
- [72] W. McKinney, “Data structures for statistical computing in Python,” in *Proceedings of the 9th Python in Science Conference*, S. van der Walt and J. Millman, Eds., Verified: 2026-01-06. Introduces pandas library for data manipulation and analysis in Python. Over 8,100 citations. Foundational paper for data science workflows, 2010, pp. 56–61. DOI: [10.25080/Majora-92bf1922-00a](https://doi.org/10.25080/Majora-92bf1922-00a).
- [73] C. R. Harris, K. J. Millman, S. J. van der Walt, R. Gommers, P. Virtanen, D. Cournapeau, E. Wieser, J. Taylor, S. Berg, N. J. Smith, R. Kern, M. Picus, S. Hoyer, M. H. van Kerkwijk, M. Brett, A. Haldane, J. Fernández del Río, M. Wiebe, P. Peterson, P. Gérard-Marchant, K. Sheppard, T. Reddy, W. Weckesser, H. Abbasi, C. Gohlke, and T. E. Oliphant, “Array programming with NumPy,” *Nature*, vol. 585, no. 7825, pp. 357–362, 2020, Verified: 2026-01-06. NumPy provides fundamental array programming capabilities for scientific computing in Python. Essential infrastructure for numerical operations, linear algebra, and statistical computations. DOI: [10.1038/s41586-020-2649-2](https://doi.org/10.1038/s41586-020-2649-2).
- [74] K. Jordahl, J. Van den Bossche, M. Fleischmann, J. Wasserman, J. McBride, J. Gerard, J. Tratner, M. Perry, A. G. Badaracco, C. Farmer, G. A. Hjelle, A. D. Snow, M. Cochran, S. Gillies, L. Culbertson, M. Bartos, N. Eubank, maxalbert, A. Bilogur, S. Rey, C. Ren, D. Arribas-Bel, L. Wasser, L. J. Wolf, M. Journois, J. Wilson, A. Greenhall, C. Holdgraf, Filipe, and F. Leblanc, *GeoPandas: Python tools for geographic data*, Zenodo, Verified: 2026-01-06. GeoPandas extends pandas to enable spatial operations on geometric types. Used for point-in-polygon spatial joins, district boundary processing, and geographic coordinate operations, 2020. DOI: [10.5281/zenodo.3946761](https://doi.org/10.5281/zenodo.3946761).

- [75] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and É. Duchesnay, “Scikit-learn: Machine learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011, Verified: 2026-01-06. Scikit-learn provides comprehensive machine learning algorithms including classification, regression, clustering, and model evaluation tools. Used for K-means spatial clustering, cross-validation, and model training pipeline. [Online]. Available: <https://jmlr.org/papers/v12/pedregosa11a.html>.
- [76] J. A. Hanley and B. J. McNeil, “The meaning and use of the area under a receiver operating characteristic (ROC) curve,” *Radiology*, vol. 143, no. 1, pp. 29–36, 1982, Verified: 2026-01-06. Canonical reference for AUC-ROC as threshold-invariant discrimination metric. Quantifies probability that randomly selected positive case receives higher predicted probability than negative case. DOI: [10.1148/radiology.143.1.7063747](https://doi.org/10.1148/radiology.143.1.7063747).
- [77] G. W. Brier, “Verification of forecasts expressed in terms of probability,” *Monthly Weather Review*, vol. 78, no. 1, pp. 1–3, 1950, Verified: 2026-01-06. Introduces Brier score for evaluating accuracy of probabilistic forecasts. Measures mean squared difference between predicted probabilities and binary outcomes, with perfect score of 0. DOI: [10.1175/1520-0493\(1950\)078<0001:VOFEIT>2.0.CO;2](https://doi.org/10.1175/1520-0493(1950)078<0001:VOFEIT>2.0.CO;2).
- [78] W. J. Youden, “Index for rating diagnostic tests,” *Cancer*, vol. 3, no. 1, pp. 32–35, 1950, Verified: 2026-01-06. Introduces Youden’s J statistic for optimal classification threshold selection, maximizing sum of sensitivity and specificity. Widely used in ROC analysis for threshold optimization. DOI: [10.1002/1097-0142\(1950\)3:1<32::AID-CNCR2820030106>3.0.CO;2-3](https://doi.org/10.1002/1097-0142(1950)3:1<32::AID-CNCR2820030106>3.0.CO;2-3).
- [79] E. R. DeLong, D. M. DeLong, and D. L. Clarke-Pearson, “Comparing the areas under two or more correlated receiver operating characteristic curves: A nonparametric approach,” *Biometrics*, vol. 44, no. 3, pp. 837–845, 1988. DOI: [10.2307/2531595](https://doi.org/10.2307/2531595).
- [80] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, ser. KDD ’16, Verified: 2026-01-06. Original XGBoost paper introducing scalable gradient boosting framework with sparsity-aware algorithm and weighted quantile sketch. One of most cited ML papers (>50,000 citations on Google Scholar), New York, NY: ACM, 2016, pp. 785–794. DOI: [10.1145/2939672.2939785](https://doi.org/10.1145/2939672.2939785).
- [81] A. Gelman and J. Hill, *Data Analysis Using Regression and Multilevel/Hierarchical Models*. Cambridge University Press, 2006, ISBN: 9780521686891.
- [82] R. Meyes, M. Lu, C. W. de Puiseau, and T. Meisen, “Ablation studies in artificial neural networks,” *arXiv preprint arXiv:1901.08644*, 2019, Verified: 2026-01-06. arXiv preprint (Jan 2019, no peer-reviewed publication found, well-cited). Investigates ablation studies for understanding learned representations in neural networks. Influential methodological reference for ML interpretability. arXiv: [1901.08644](https://arxiv.org/abs/1901.08644). [Online]. Available: <https://arxiv.org/abs/1901.08644>.
- [83] A. Balashankar et al., “Toward real-world food security crisis prediction using news media text,” in *AAAI Conference on Artificial Intelligence*, 2023.

- [84] U. Qazi, M. Imran, and F. Ofli, “Geo-CoV19: A dataset of hundreds of millions of multilingual COVID-19 tweets with location information,” in *ACM SIGSPATIAL International Conference on Advances in Geographic Information Systems*, 2020. DOI: [10.1145/3423337.3429820](https://doi.org/10.1145/3423337.3429820).
- [85] H. Mueller and C. Rauh, “Quantifying spatiotemporal dynamics of conflict and food insecurity using machine learning,” *Political Geography*, vol. 84, 2021. DOI: [10.1016/j.polgeo.2020.102297](https://doi.org/10.1016/j.polgeo.2020.102297).
- [86] D. R. Roberts, V. Bahn, S. Ciuti, M. S. Boyce, J. Elith, G. Guillerá-Arroita, S. Hauenstein, J. J. Lahoz-Monfort, B. Schröder, W. Thuiller, D. I. Warton, B. A. Wintle, F. Hartig, and C. F. Dormann, “Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure,” *Ecography*, vol. 40, no. 8, pp. 913–929, 2017, Verified: 2026-01-06. Seminal paper on spatial and temporal cross-validation strategies for ecological modeling. Demonstrates why random CV fails with structured data and provides guidelines for blocked/leave-one-out spatial CV methods. Essential methodological reference for spatio-temporal prediction models. DOI: [10.1111/ecog.02881](https://doi.org/10.1111/ecog.02881).
- [87] L. R. Rabiner, “A tutorial on hidden markov models and selected applications in speech recognition,” *Proceedings of the IEEE*, vol. 77, no. 2, pp. 257–286, 1989. DOI: [10.1109/5.18626](https://doi.org/10.1109/5.18626).
- [88] P. J. Schmid, “Dynamic mode decomposition of numerical and experimental data,” *Journal of Fluid Mechanics*, vol. 656, pp. 5–28, 2010. DOI: [10.1017/S0022112010001217](https://doi.org/10.1017/S0022112010001217).
- [89] M. Imran, C. Castillo, F. Diaz, and S. Vieweg, “Processing social media messages in mass emergency: A survey,” 4, vol. 47, 2015, pp. 1–38. DOI: [10.1145/2771588](https://doi.org/10.1145/2771588).
- [90] W. E. Oswald, A. E. Stewart, R. A. Kramer, C. H. King, P. N. Mwinzi, M. R. Odiere, S. Kariuki, B. L. Cline, et al., “Predicting infectious disease using deep learning and big data,” *Infection*, vol. 48, pp. 303–310, 2020. DOI: [10.1007/s15010-020-01417-2](https://doi.org/10.1007/s15010-020-01417-2).
- [91] N. Jean, M. Burke, M. Xie, W. M. Davis, D. B. Lobell, and S. Ermon, “Combining satellite imagery and machine learning to predict poverty,” *Science*, vol. 353, no. 6301, pp. 790–794, 2016. DOI: [10.1126/science.aaf7894](https://doi.org/10.1126/science.aaf7894).
- [92] S. M. Hsiang, M. Burke, and E. Miguel, “Quantifying the influence of climate on human conflict,” *Science*, vol. 341, no. 6151, 2013. DOI: [10.1126/science.1235367](https://doi.org/10.1126/science.1235367).
- [93] A. Conneau, K. Khandelwal, N. Goyal, V. Chaudhary, G. Wenzek, F. Guzmán, E. Grave, M. Ott, L. Zettlemoyer, and V. Stoyanov, “Unsupervised cross-lingual representation learning at scale,” in *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, 2020, pp. 8440–8451.
- [94] M. Sensoy, L. Kaplan, and M. Kandemir, “Evidential deep learning to quantify classification uncertainty,” in *Advances in Neural Information Processing Systems*, vol. 31, 2018.
- [95] P. W. Koh and P. Liang, “Understanding black-box predictions via influence functions,” in *International Conference on Machine Learning*, 2017, pp. 1885–1894.
- [96] S. Wachter, B. Mittelstadt, and C. Russell, “Counterfactual explanations without opening the black box: Automated decisions and the GDPR,” *Harvard Journal of Law & Technology*, vol. 31, no. 2, pp. 841–887, 2017.